



Deploying Julia

Strategies, Architectures and Pitfalls

Avik Sengupta
JuliaHub, Inc.



Who am I

- Julia developer since 2012
- Production Julia code since 2016
- Products and services in Julia

What I'll talk about today

- Packaging – Manifests, Docker, Precompilation
- Sysimages, PackageCompiler
- JuliaC
- Miscellaneous Tips and tricks



What does deployment mean to you?

- Where do you deploy
 - Public or Private Cloud
 - HPC / Supercomputer
 - Desktop / Workstations
- What do you deploy
 - Batch processes (*Model training / Simulations*)
 - Long running service (*Web apps / API services*)
 - Desktop UI
 - OpenSource / Internal / Commercial



What does deployment mean to me?

- Run code on a machine that is not yours
- Code comes with implicit assumptions about the environment they will run in
 - Operating System
 - Instruction set
 - Dependent libraries and their versions
 - Configuration files



Send/Publish a file with Julia source code

- Simple 😊
- User needs Julia installed 😕
- User need to manually install packages used 😕
- User may get different versions of packages used 😕
- User may use a different Julia version 😕
- Source code might be proprietary 😕
- Not really feasible except for the very simplest of codes



Packaging – Project + Manifest

- Project.toml + Manifest.toml
 - Always check in Manifest for Applications
 - User gets the same version of packages 😊
- User needs to install Julia & instantiate 🙄
 - The correct version of Julia
 - New Juliaup functionality to pick Julia versions from project
- Same package works across Operating Systems
- Multiple files to track if all you need is a script 🙄
 - In 1.13, you can store project+manifest in source files



Packaging – Registry

- Always use a private registry for private code
 - Unless you are single developer, deploying to single machine
- A registry is only a git repository, so no major infrastructure is required
- LocalRegistry.jl can help with managing a registry
- Makes dependency versioning and version tracking possible



Considerations for HPC Clusters

- Set your depot to a shared filesystem
 - Shared filesystems are slow
- Make sure to precompile before running a distributed job
 - Don't want thousands of repeated precompiles
 - Consider running with `--compiled-modules=strict`
 - Precompile on login node if possible
- Set CPU target appropriately
 - Especially for heterogenous nodes
- Figure out correct MPI config



Packaging - Docker

- Public Julia docker images
- Instantiate a manifest during docker build
 - Internet access needed at build time
 - Precompile: `Pkg.compile()` or use a precompile workload
 - Manage CPU targets unless building on the same machines as deploying
 - `JULIA_CPU_TARGET=generic;sandybridge,-xsaveopt,clone_all;haswell,-rdrnd,base(1)`
- Use memory hints at runtime
 - `julia --heap-size-hint=4G`
- Set thread limits appropriately



Aside on Precompile workloads

- Julia only compiles the code that runs at top level
- If you have functions that take a long time to compile, but haven't been run at top level (ie at compile time) then you will compile at runtime, and pay a latency cost
- Run those functions at top level or
- Use PrecompileTools.jl



System Images (or sysimage)

- “Serialized snapshot of a julia session”,
- Contains julia objects, compiled code, code yet to be compiled, etc.
- Relatively fast to load
- Julia comes with a sysimage containing the base and (most) standard libraries



Custom Sysimages solve the compile time latency

```
pkg> add ModellingToolkit
```

```
...
```

```
186 dependencies successfully precompiled in 611 seconds
```

```
julia> @time using ModelingToolkit
```

```
3.443558 seconds (4.96 M allocations: 341.655 MiB, 7.86% gc time)
```

```
julia> @time using ModelingToolkit
```

```
32.879128 seconds (6.31 M allocations: 414.853 MiB, 1.66% gc time,  
1.48% compilation time: 84% of which was recompilation)
```





Enter PackageCompiler.jl



Custom Sysimages solve the compile time latency

```
julia> using PackageCompiler
```

```
julia> PackageCompiler.create_sysimage(["ModelingToolkit", "GLMakie"];  
    sysimage_path="mtk_makie_sys.so")
```

```
C:\> julia +1.12 --sysimage=mtk_makie_sys.so --project=.
```

```
julia> @time using ModelingToolkit  
0.000991 seconds (505 allocations: 27.688 KiB)
```



Custom sysimages with packagecompiler

- Reduce compile time latency 😊
- Needs a C compiler on the build machine
- Need to use the same version of Julia to load it 😰
- Artifacts need to be downloaded separately 😰
- Need sources to resolve new packages in environment
- Build for each OS (and no ability to cross-compile)
- Needs a loooong time to create, and results are biiiiiig 😰

```
.rwxr-xr-x 1.5G avik 29 Dec 19:07 my_custom_sys.so
```



Apps with PackageCompiler

- Package Julia+sysimage+artifacts+entrypoint

```
julia> PackageCompiler.create_app( "App", "AppCompiled")
```

```
$ tree AppCompiled/
AppCompiled/
├── bin
│   └── HelloApp
└── lib
    └── julia

      ┌── libLLVM.so.18.1jl
      ┌── libatomic.so.1.2.0
      ┌── libblastrampoline.so.5
      ┌── libdSFMT.so
      ┌── libgcc_s.so.1
```



Relocatability

- Application should not depend on particular filesystem path
- Paths are used to find different kinds of resources
 - Binaries / Libraries / Data artifacts
- Your package and ALL its dependencies must be relocatable
- Use jll packages and artifacts for native dependencies
- Don't store file paths or pointers in module globals – use `__init__` functions instead.
 - Actually, use `OncePerProcess` instead (1.12+)



JuliaC

- Create small, standalone binaries (~1.5Mib)
- All code needs to be statically inferable
- 1.12+ only
- Key innovation is ***trim*** support in the compiler
- Needs infrastructure around the compiler, including PackageCompiler



JuliaC

```
pkg> app add JuliaC
```

```
$ juliac HelloApp --output-exe hello -trim
```

```
$ ls -al hello
```

```
-rwxr-xr-x 1 avik avik 1691304 Jan  5 10:54 hello
```

```
$ ./hello
```

```
Hello, world
```



JuliaC

- JuliaC binaries depend on libjulia

```
$ ldd ./hello
    libjulia.so.1.12 => /home/avik/.julia/juliaup/julia-
1.12.3+0.x64.linux.gnu/lib/libjulia.so.1.12 (0x00007cac188a1000)
    libjulia-internal.so.1.12 => /home/avik/.julia/juliaup/julia-
1.12.3+0.x64.linux.gnu/lib/julia/libjulia-internal.so.1.12
(0x00007cac18200000)
    libunwind.so.8 => /home/avik/.julia/juliaup/julia-
1.12.3+0.x64.linux.gnu/lib/julia/libunwind.so.8 (0x00007cac17a00000)
```



Packaging with JuliaC

```
$ juliac HelloApp --output-exe hello --trim --bundle hellodir
```

```
$ tree hellodir
├── bin
│   └── hello
└── lib
    ├── julia
    │   └── libatomic.so
    .....
    └── libz.so.1.3.1
        └── libjulia.so.1.12.3
4 directories, 57 files
```



Considerations for long running services

- Latency of first request
 - Good precompile workload, or
 - Warmup requests
- Concurrency – threads or tasks or processes?
- Recommend not putting Julia web services directly on the public internet
 - Proxy via apache/nginx/caddy etc
 - Pay attention to package vulnerabilities



Considerations for cloud native usage

- Interfaces to cloud platforms - AWS.jl / Azure.jl
- Memory is the biggest cost on public cloud services
- Shared disks are slow, fast IO costs more
- Consider running Julia processes on Kubernetes
 - Kuber.jl provides complete set of entities and operations





Tell us your stories



References

- <https://github.com/GunnarFarneback/LocalRegistry.jl>
- <https://juliahpc.github.io/>
- Package compiler webinar: <https://youtu.be/J6h6Tj8IluE>
- <https://github.com/JuliaLang/JuliaC.jl/>
- Acknowledgements: Kristoffer Carlsson, Mosè Giordano, Chris Rackauckas, Jeff Bezanson

