






Julia Hackathon Challenge: High Performance Ocean Internal Wave Detection

Arun Kumar Shukla
Data Science Developer,
Atlantic International Research Centre,
Azores, Portugal

Objective

- ***Build a high-performance Julia Pipeline for Ocean Internal Wave Detection***
- *Challenge:* Improve the existing Python-dependent machine learning pipeline into a high-performance Julia implementation.
 - *Mandatory:*
 - Improves AUC-ROC from 0.89 to >0.95
 - Reduces execution time from 1 hour to under 10 minutes
 - Adds GPU acceleration
 - *Desirable:*
 - Eliminates all Python dependencies




Dataset



- 19,528 Sentinel-1 Wave(WV) Mode satellite images
- 13,668 training + 5,860 test images (50% +ve and 50% -ve samples)
- 224×224 resolution, RGB format
- **Processing Requirements**
Images → Feature Extraction → Classification → Predictions

OCEAN INTERNAL WAVE CLASSIFICATION SYSTEM

Technical Architecture Specification



The Existing Pipeline (What I Built)

- **Four Main Components:**
- **Feature Extraction** (`extract_features.jl`)
 - Hybrid Julia/Python with ONNX runtime
 - Handles GPU/CPU switching
 - Extracts 1000 features per image
- **Data Alignment** (`fix_data_alignment.jl`)
 - Solves ID mismatch problem
 - Creates translation between systems
 - Critical for correct training
- **Model Training** (`train_final.jl`)
 - XGBoost with GPU support
 - Cross-validation with AUC monitoring
 - Regularization to prevent overfitting
- **Prediction** (`predict.jl`)
 - Generates competition submissions
 - Probability clipping and validation
 - Ready for deployment

Current Performance Analysis

- **Where We Stand Today**
- **Model Performance:**
 - Cross-validation AUC: 0.8905 ± 0.0005
 - 5-fold stratified validation
 - Strong regularization applied
 - Predictions: 100% wave classification (imbalance issue)
- **Computational Performance:**
 - Feature extraction: ~1 hour (CPU)
 - Training: ~42 minutes
 - Inference: ~15 minutes
 - Memory: ~12 GB peak usage
- **Known Issues:**
 - Class imbalance not fully addressed
 - GPU acceleration partially implemented
 - Pipeline could be more efficient

The Problem Statement


- **Classify Satellite Images for Internal Waves**
 - **What:** Detect ocean internal wave patterns from satellite imagery
Why: Important for oceanography, climate studies, marine navigation
Dataset: 13,668 training images, 5,860 test images
Format: PNG images → 1000-dimensional feature vectors
 - **Starting Performance Metrics:**
 - Current AUC: 0.8905
 - Runtime: ~1 hour (CPU)
 - **Your Challenge:** Take my working code and make it better!

Available Code Base

Repository Structure:

```
├── data/                                # Features, labels, test data
├── models/                               # Pre-trained models
├── scripts/
│   ├── extract_features.jl               # Feature extraction
│   ├── fix_data_alignment.jl            # ID translation
│   ├── auc_roc_cv.jl                   # Validation
│   ├── train_final.jl                  # Model training
│   └── predict.jl                      # Prediction
└── results/                             # My results
    └── README.md                       # Setup instructions
```

Optimization Opportunities



Areas for Improvement



Model Accuracy:

Goal: Improve AUC from 0.89 to >0.98

Approaches: Address class imbalance,
better features, ensemble methods



Performance Optimization:

Goal: Reduce total runtime from 1 hour to
<10 minutes

Approaches: Better GPU utilization,
parallel processing, optimized batches

Technical Stack & Constraints



What You Can Use

Required:

Julia 1.9+ (primary language)
My existing code base as starting point




Available:

GPU access (CUDA.jl compatible)
ONNX models for feature extraction
XGBoost.jl for classification
Full dataset (13K training, 5K test)



Constraints:

Must maintain competition submission format
Need reproducible results
Should be able to run on standard hardware



Evaluation: Based on AUC improvement, runtime reduction, and code quality

Hackathon Structure

- **Monday-Friday Schedule**
- **Monday:**
 - Introduction & setup
 - Understanding the pipeline
 - Initial baseline runs
- **Tuesday-Thursday:**
 - Team work on optimizations
 - Mentor support available
 - Progress check-ins
- **Friday:**
 - Final presentations (10 min each)
 - Live Code demonstrations
 - Technical discussions
- **Submission Requirements:**
 - Optimized code
 - Performance metrics (AUC, runtime)
 - Brief presentation
 - Documentation of improvements

Learning Outcomes & Next Steps

- **What You'll Gain**
- **Technical Skills:**
 - Julia ML pipeline optimization
 - GPU acceleration techniques
 - Handling real-world data challenges
 - Production ML engineering
- **Research Skills:**
 - Debugging complex ML pipelines
 - Performance benchmarking
 - Experiment design and validation
 - Documentation and presentation
- **Community Benefits:**
 - Learn from my working implementation
 - Share improvements with others
 - Build network with Julia ML practitioners
 - Contribute to open-source ML in Julia
- **Final Goal:**
 - Everyone leaves with better Julia ML skills and potentially publishable optimizations!



Thank you



Contact email:
arun.shukla@aircentre.org