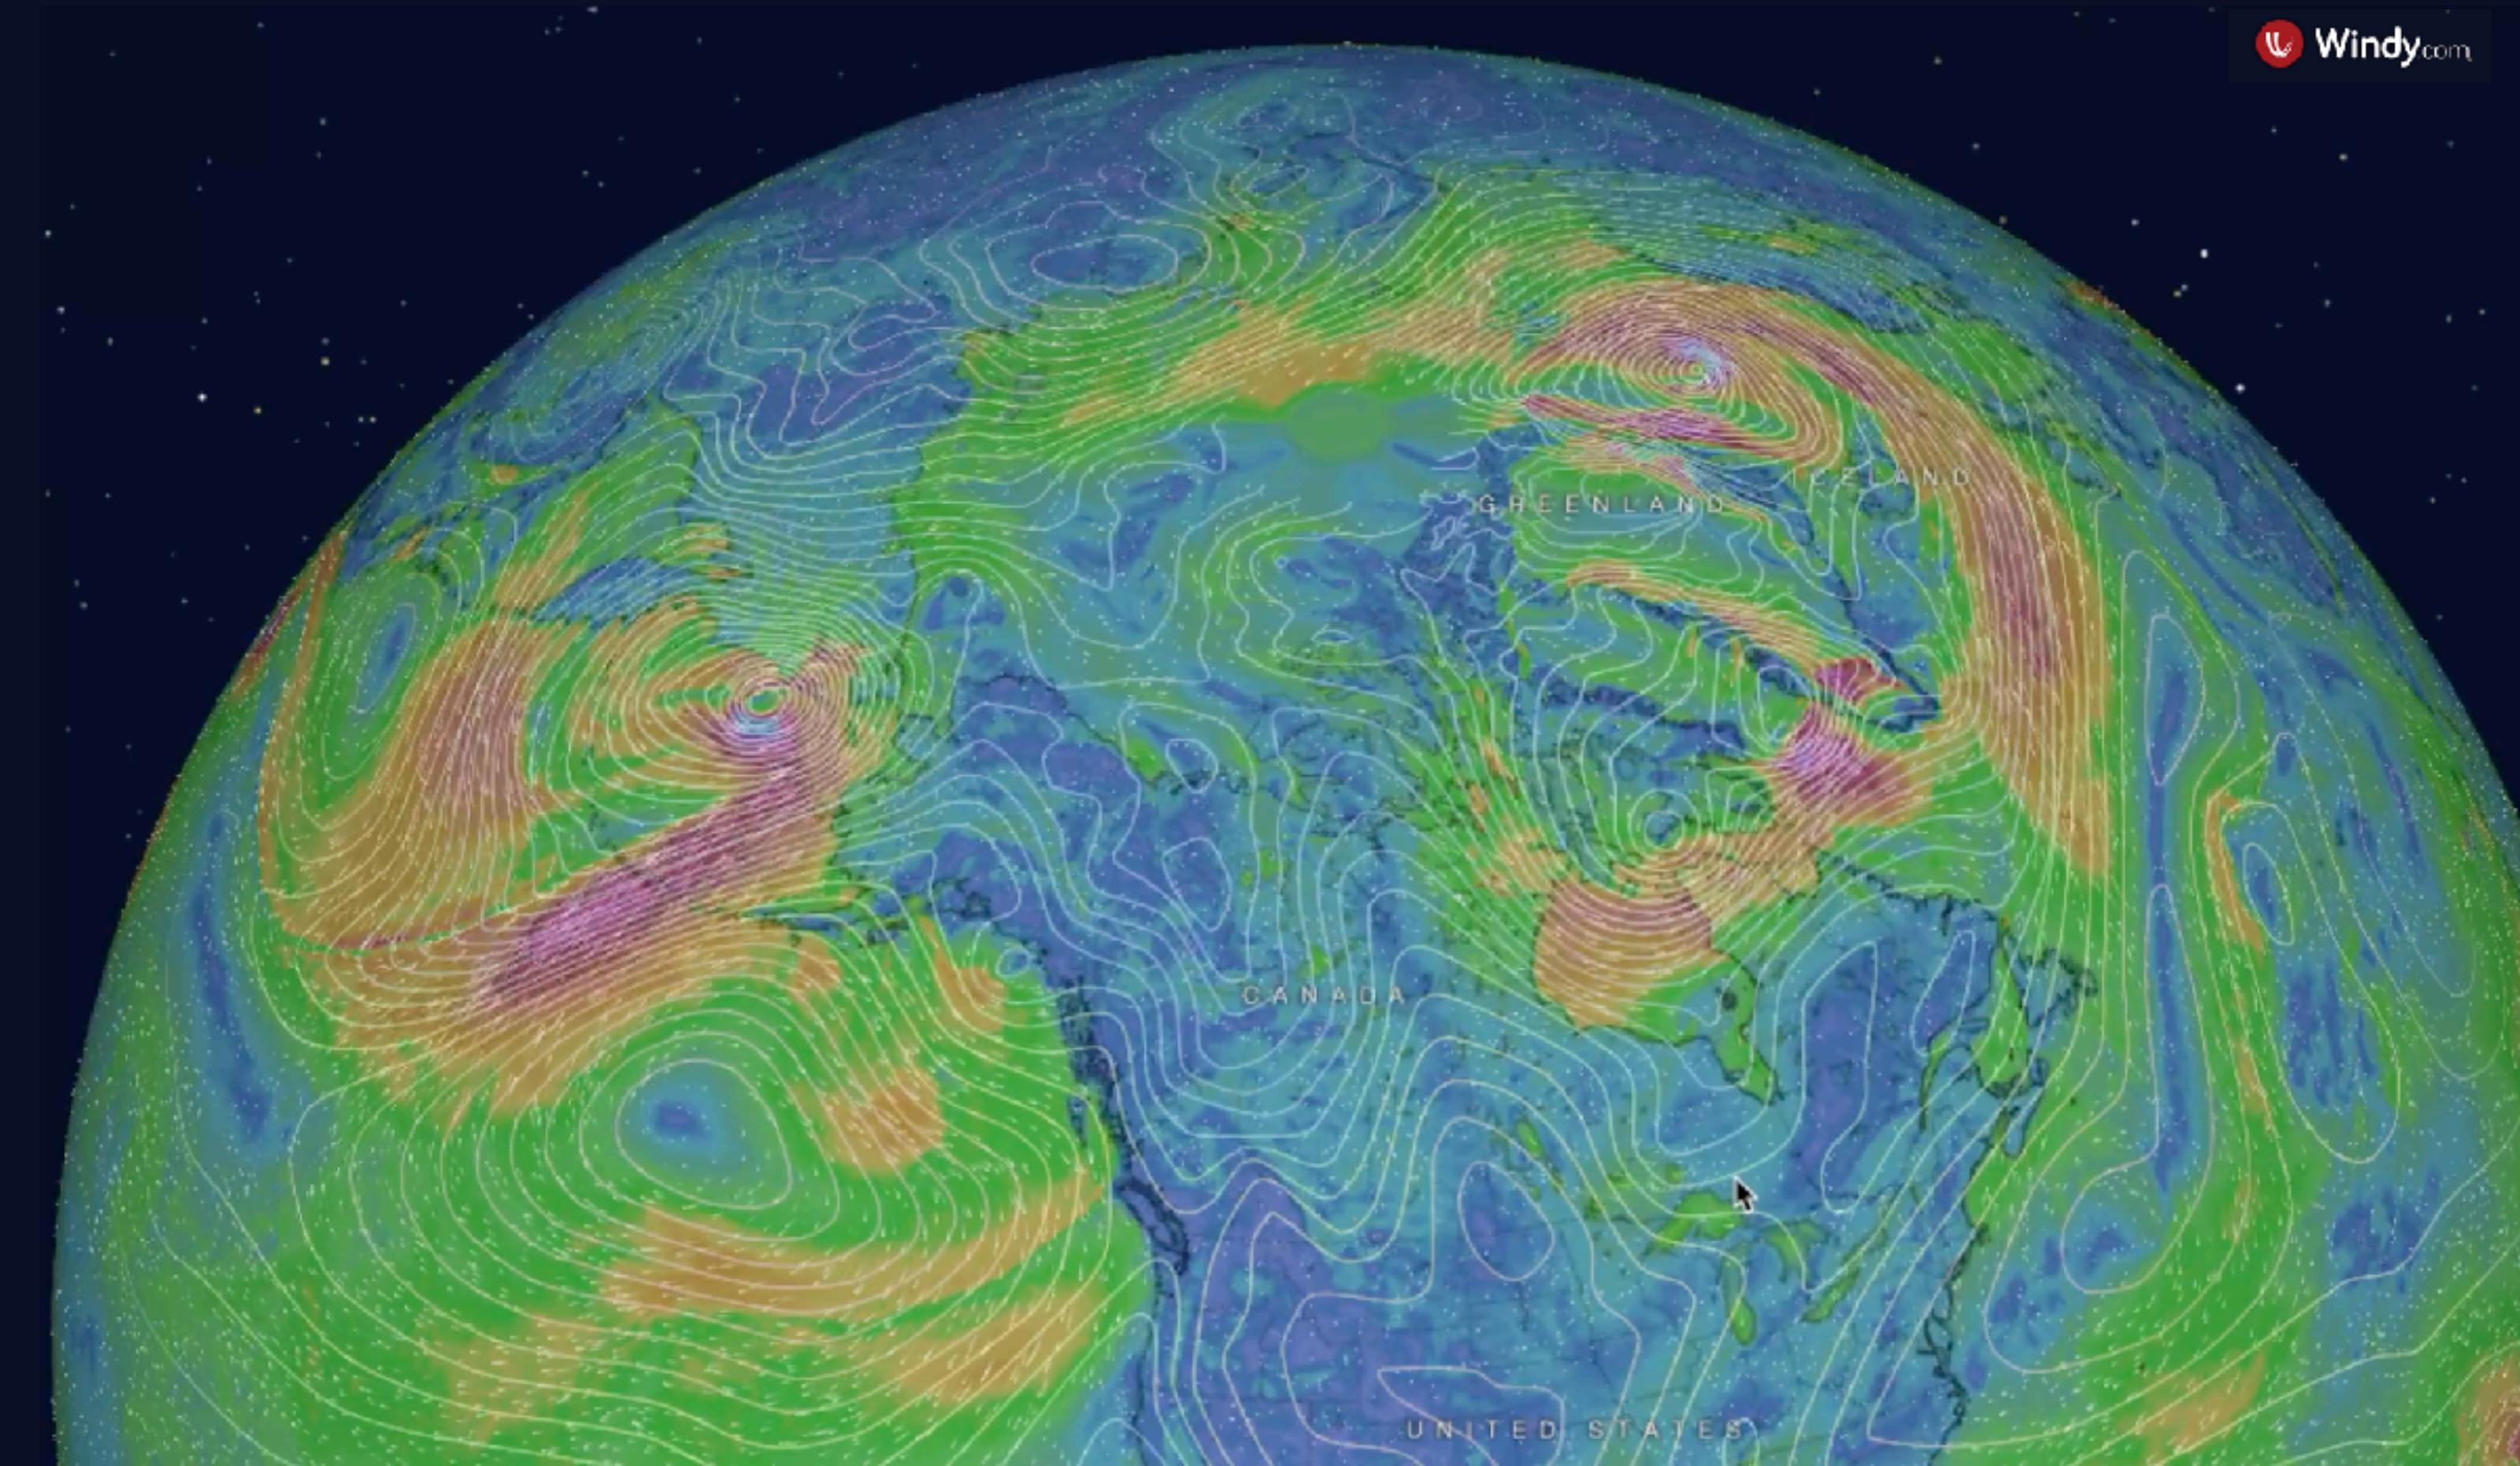


Interactive climate modelling

JuliaEO 2026, Terceira

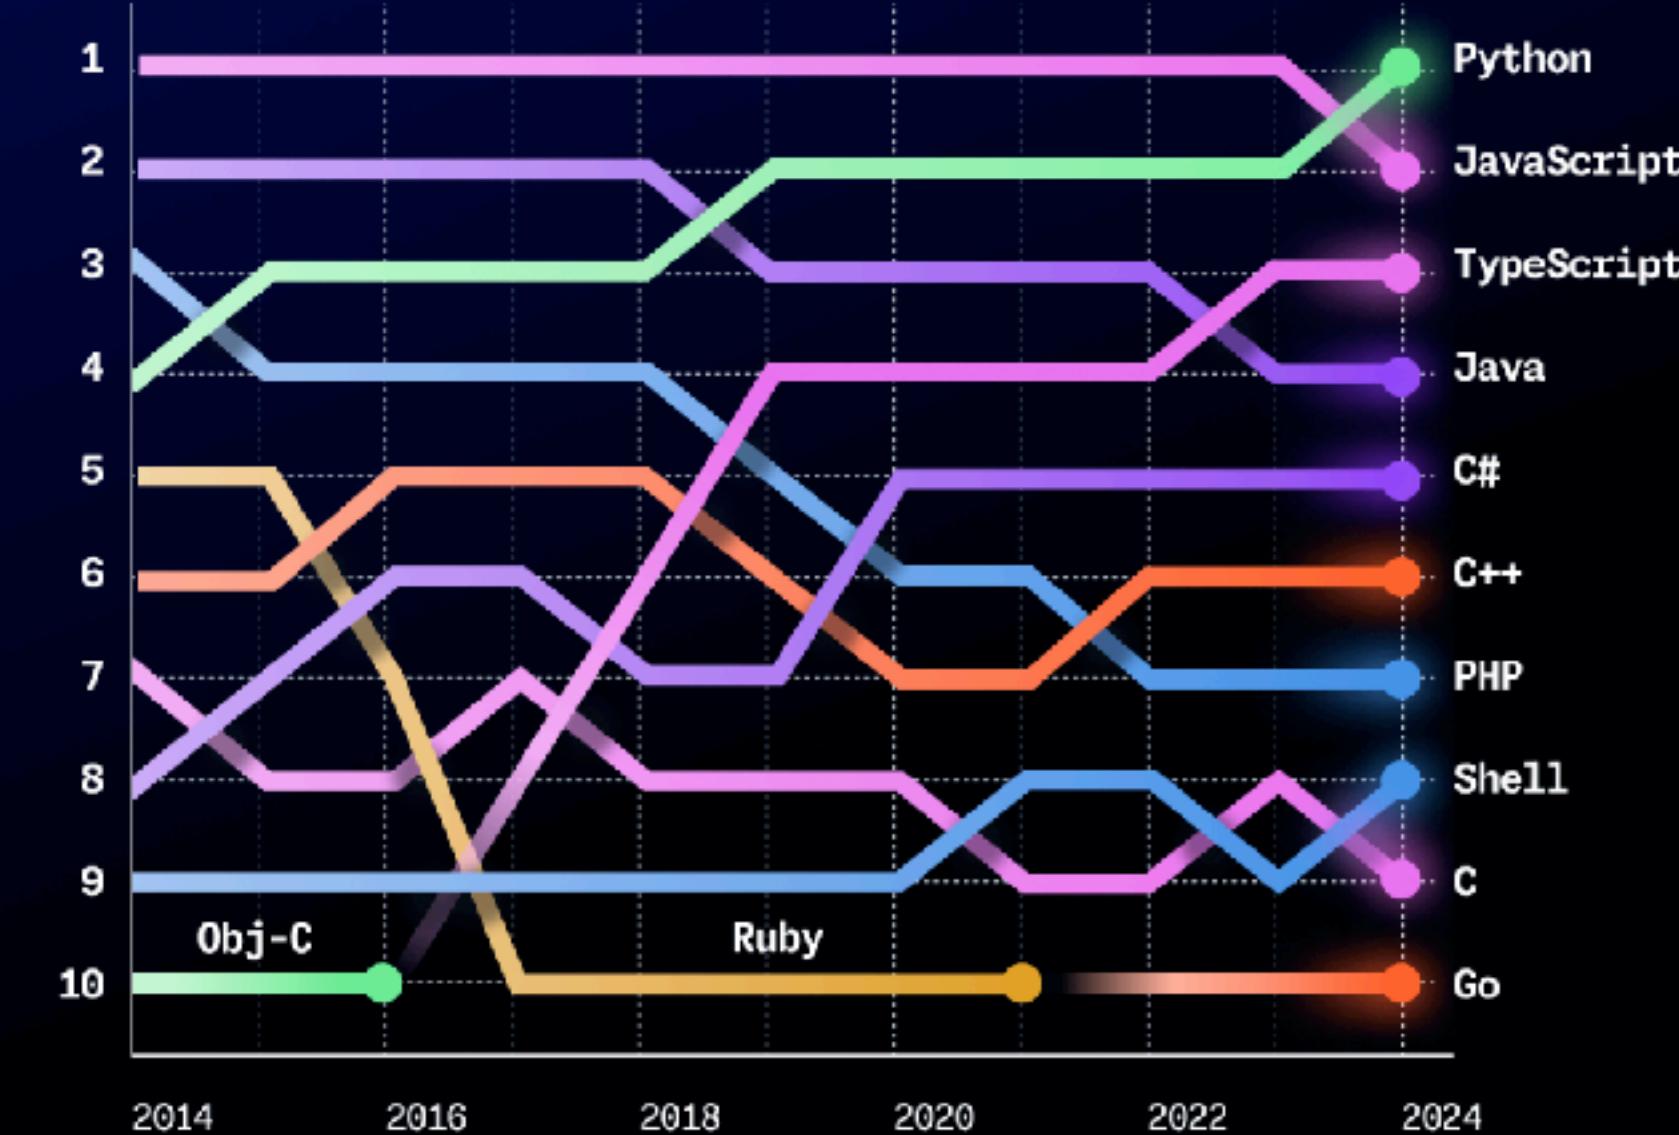
Milan Klöwer and many colleagues

University of Oxford



Top programming languages on GitHub

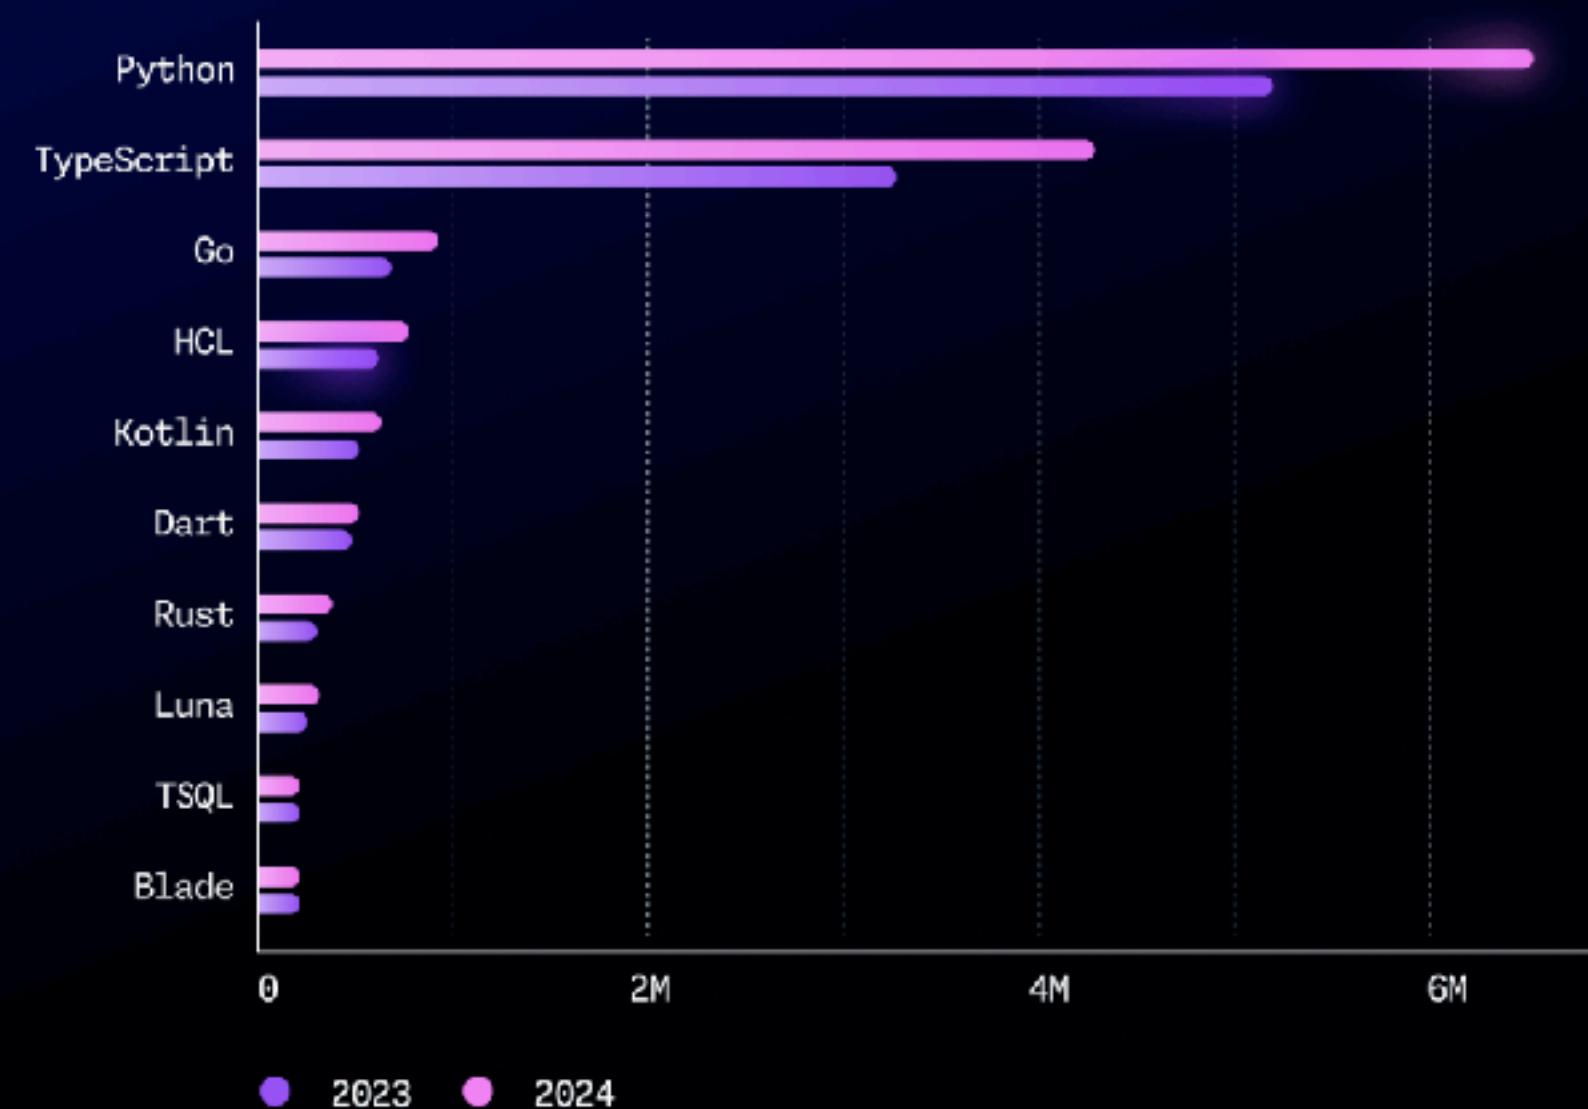
RANKED BY COUNT OF DISTINCT USERS CONTRIBUTING TO PROJECTS OF EACH LANGUAGE.



From Octoverse 2024

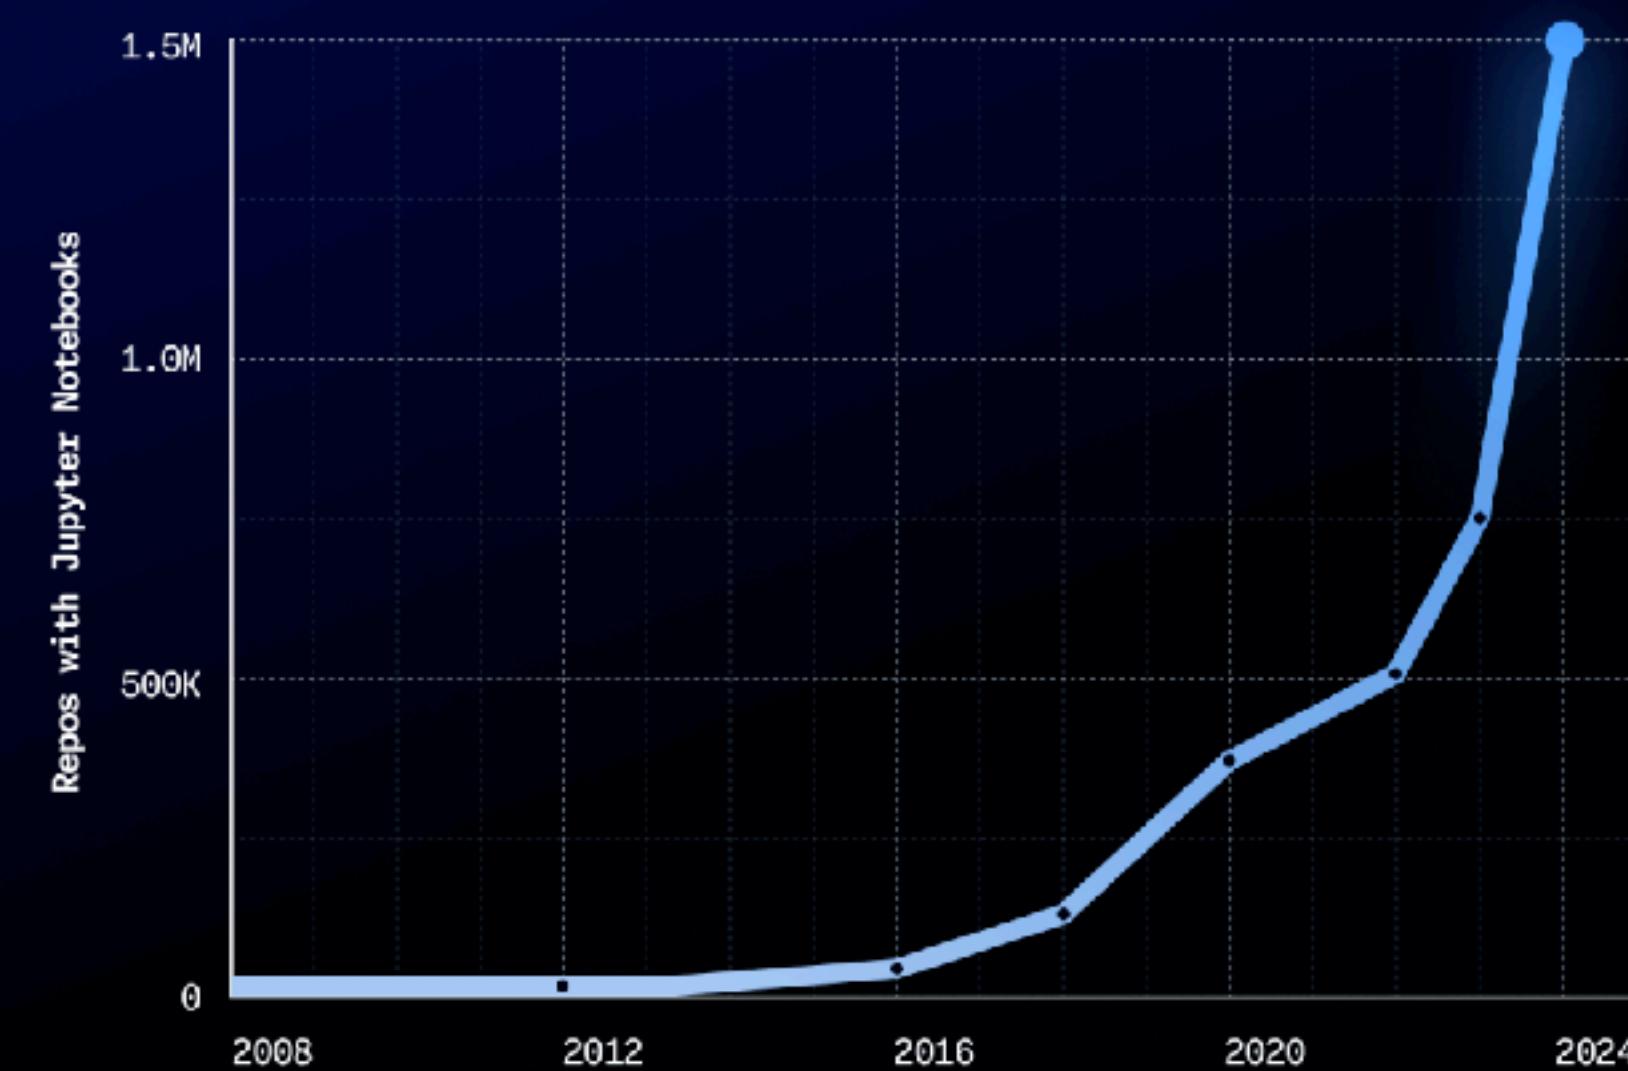
Top 10 fastest growing languages in 2024

TAKEN BY PERCENTAGE GROWTH OF CONTRIBUTORS ACROSS ALL CONTRIBUTIONS ON GITHUB.



Jupyter Notebook usage on GitHub

BY DISTINCT PUBLIC REPOSITORIES WITH AT LEAST ONE JUPYTER NOTEBOOK BY THE YEAR THAT THE REPOSITORY WAS CREATED.



Why is Python successful?

- Speed but using C
- Easy, interactive
- Cross-platform
- Extensions and software stack
- Used for open-source

Previous work on
0-dimensional climate
models by, e.g.



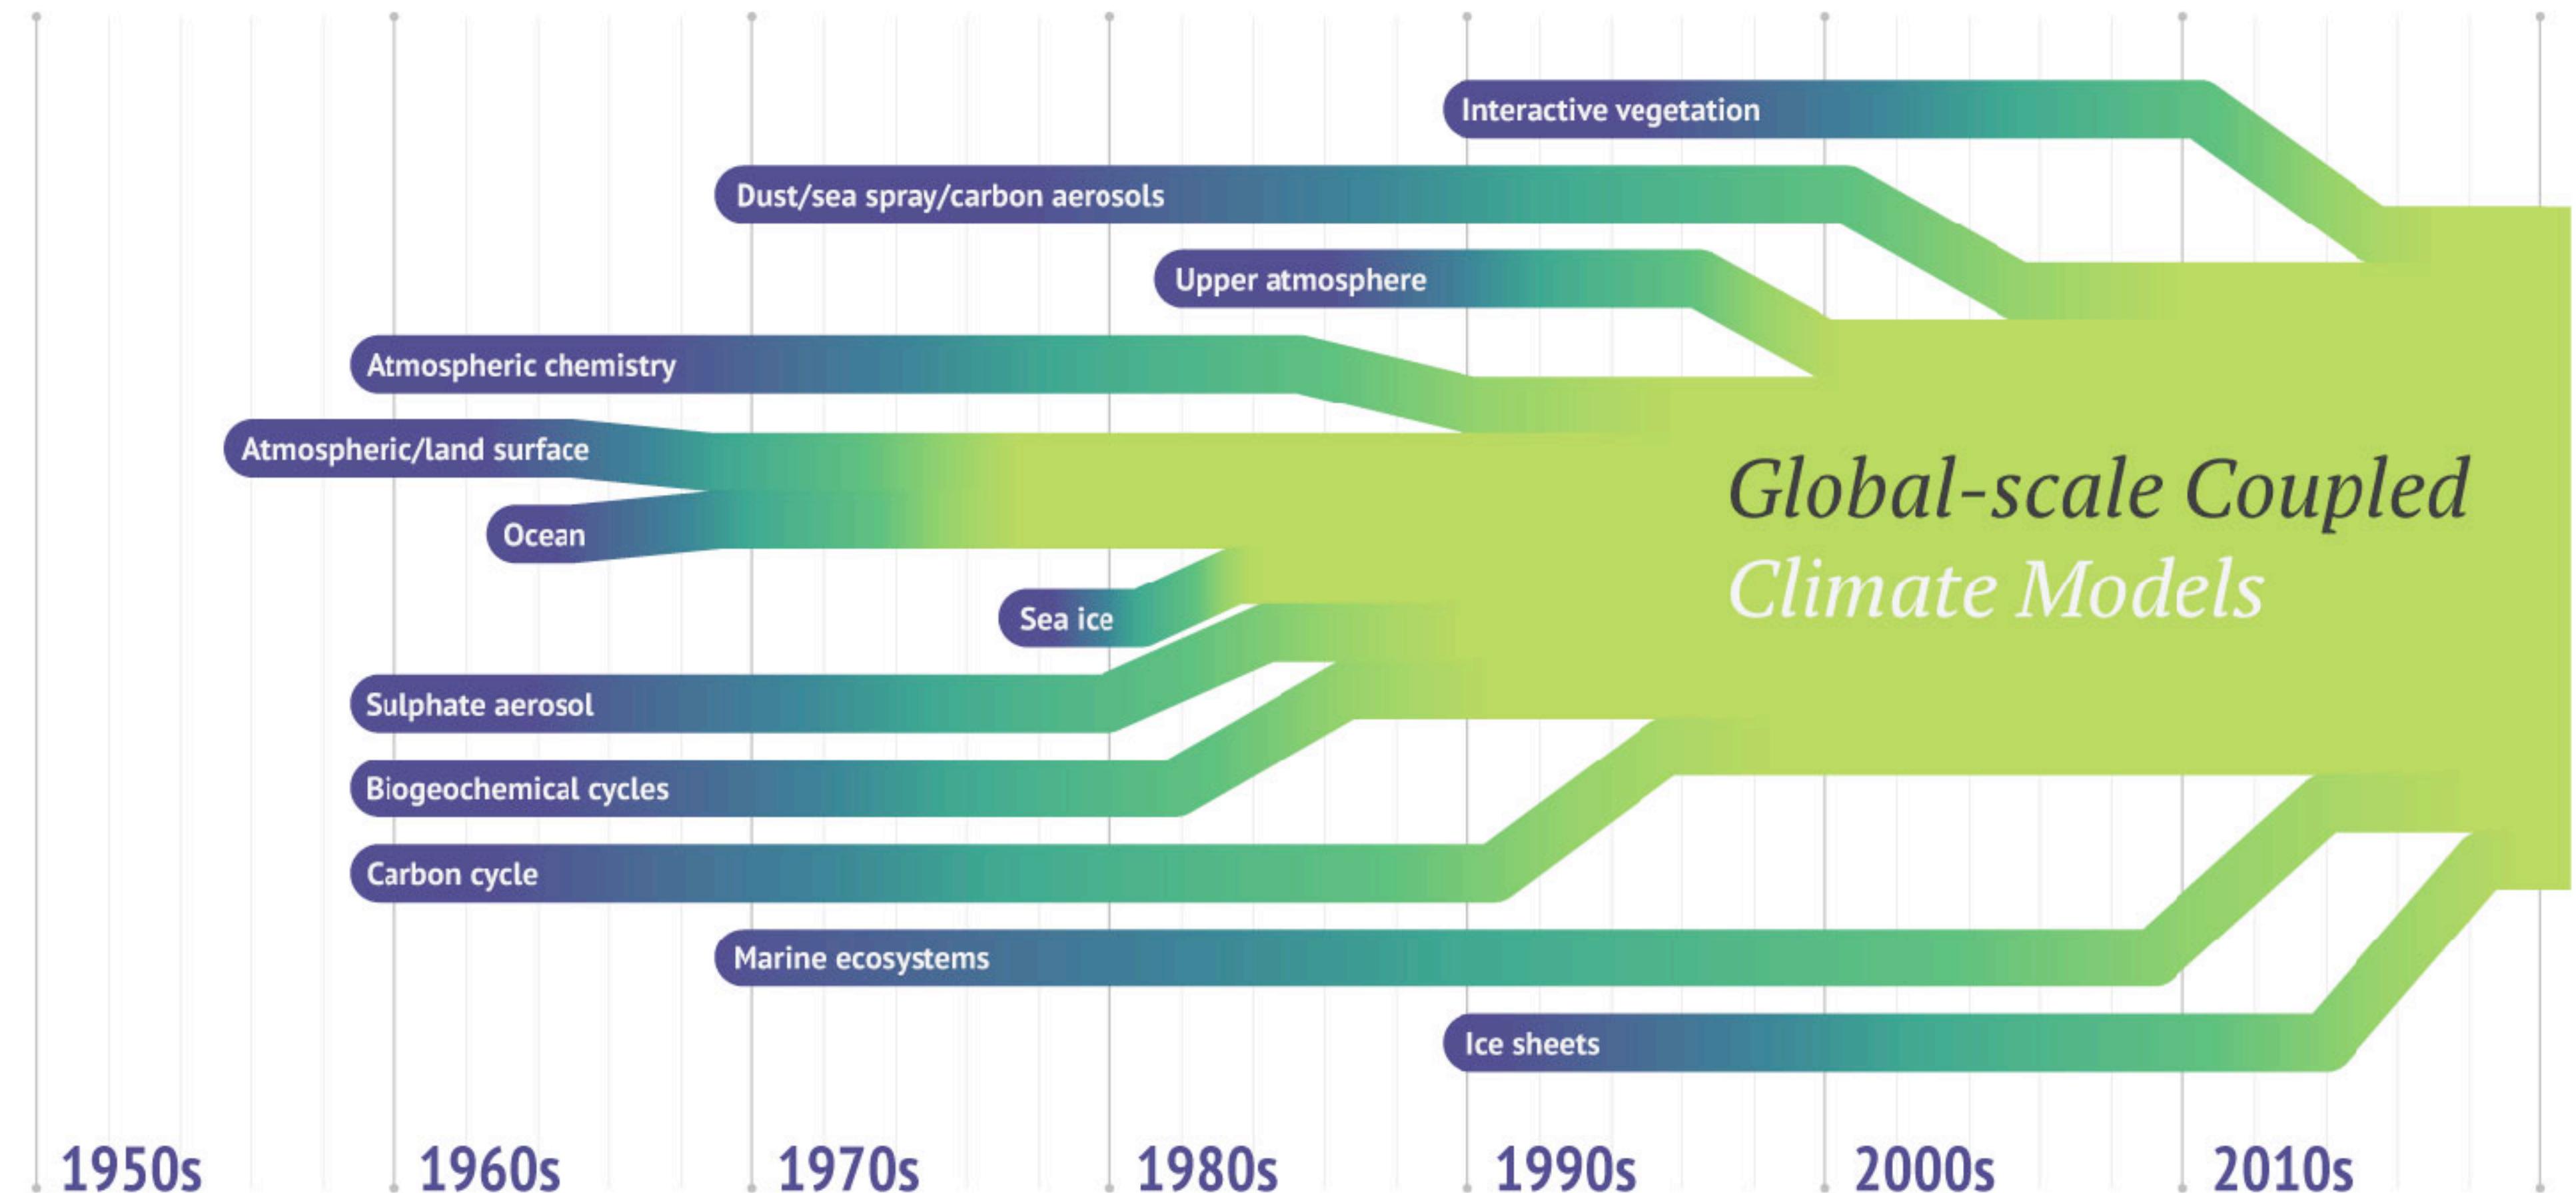
Arrhenius, 1896

*On the Influence of
Carbonic Acid in the Air
upon the Temperature of
the Ground*

$$\Delta F = \alpha \ln(C/C_0)$$

Climate models

For decades scientists have been using mathematical models to help us learn more about the Earth's climate. Known as climate models, they are driven by the fundamental physics of the atmosphere and oceans, and the cycling of chemicals between living things and their environment. Over time they have increased in complexity, as separate components have merged to form coupled systems.



Note: There were some very simplified models before the dates mentioned.

The equations: It's complicated!

$$\rho \left(\frac{\partial \mathbf{u}}{\partial t} + (\mathbf{u} \cdot \nabla) \mathbf{u} + 2\Omega \times \mathbf{u} \right) = \rho \mathbf{g} - \nabla p + \mu \nabla^2 \mathbf{u}$$
$$\frac{\partial \rho}{\partial t} + \nabla \cdot (\rho \mathbf{u}) = 0$$
$$p = \rho R T$$
$$\frac{\partial T}{\partial t} + (\mathbf{u} \cdot \nabla) T - \frac{RT}{c_p} \frac{D \ln p}{Dt} = \frac{Q}{c_p}$$

Momentum advection

Coriolis

Gravity

Pressure gradient

Viscous forces

Conservation of mass (compressible)

Temperature advection

Adiabatic conversion

Ideal gas law, equation of state

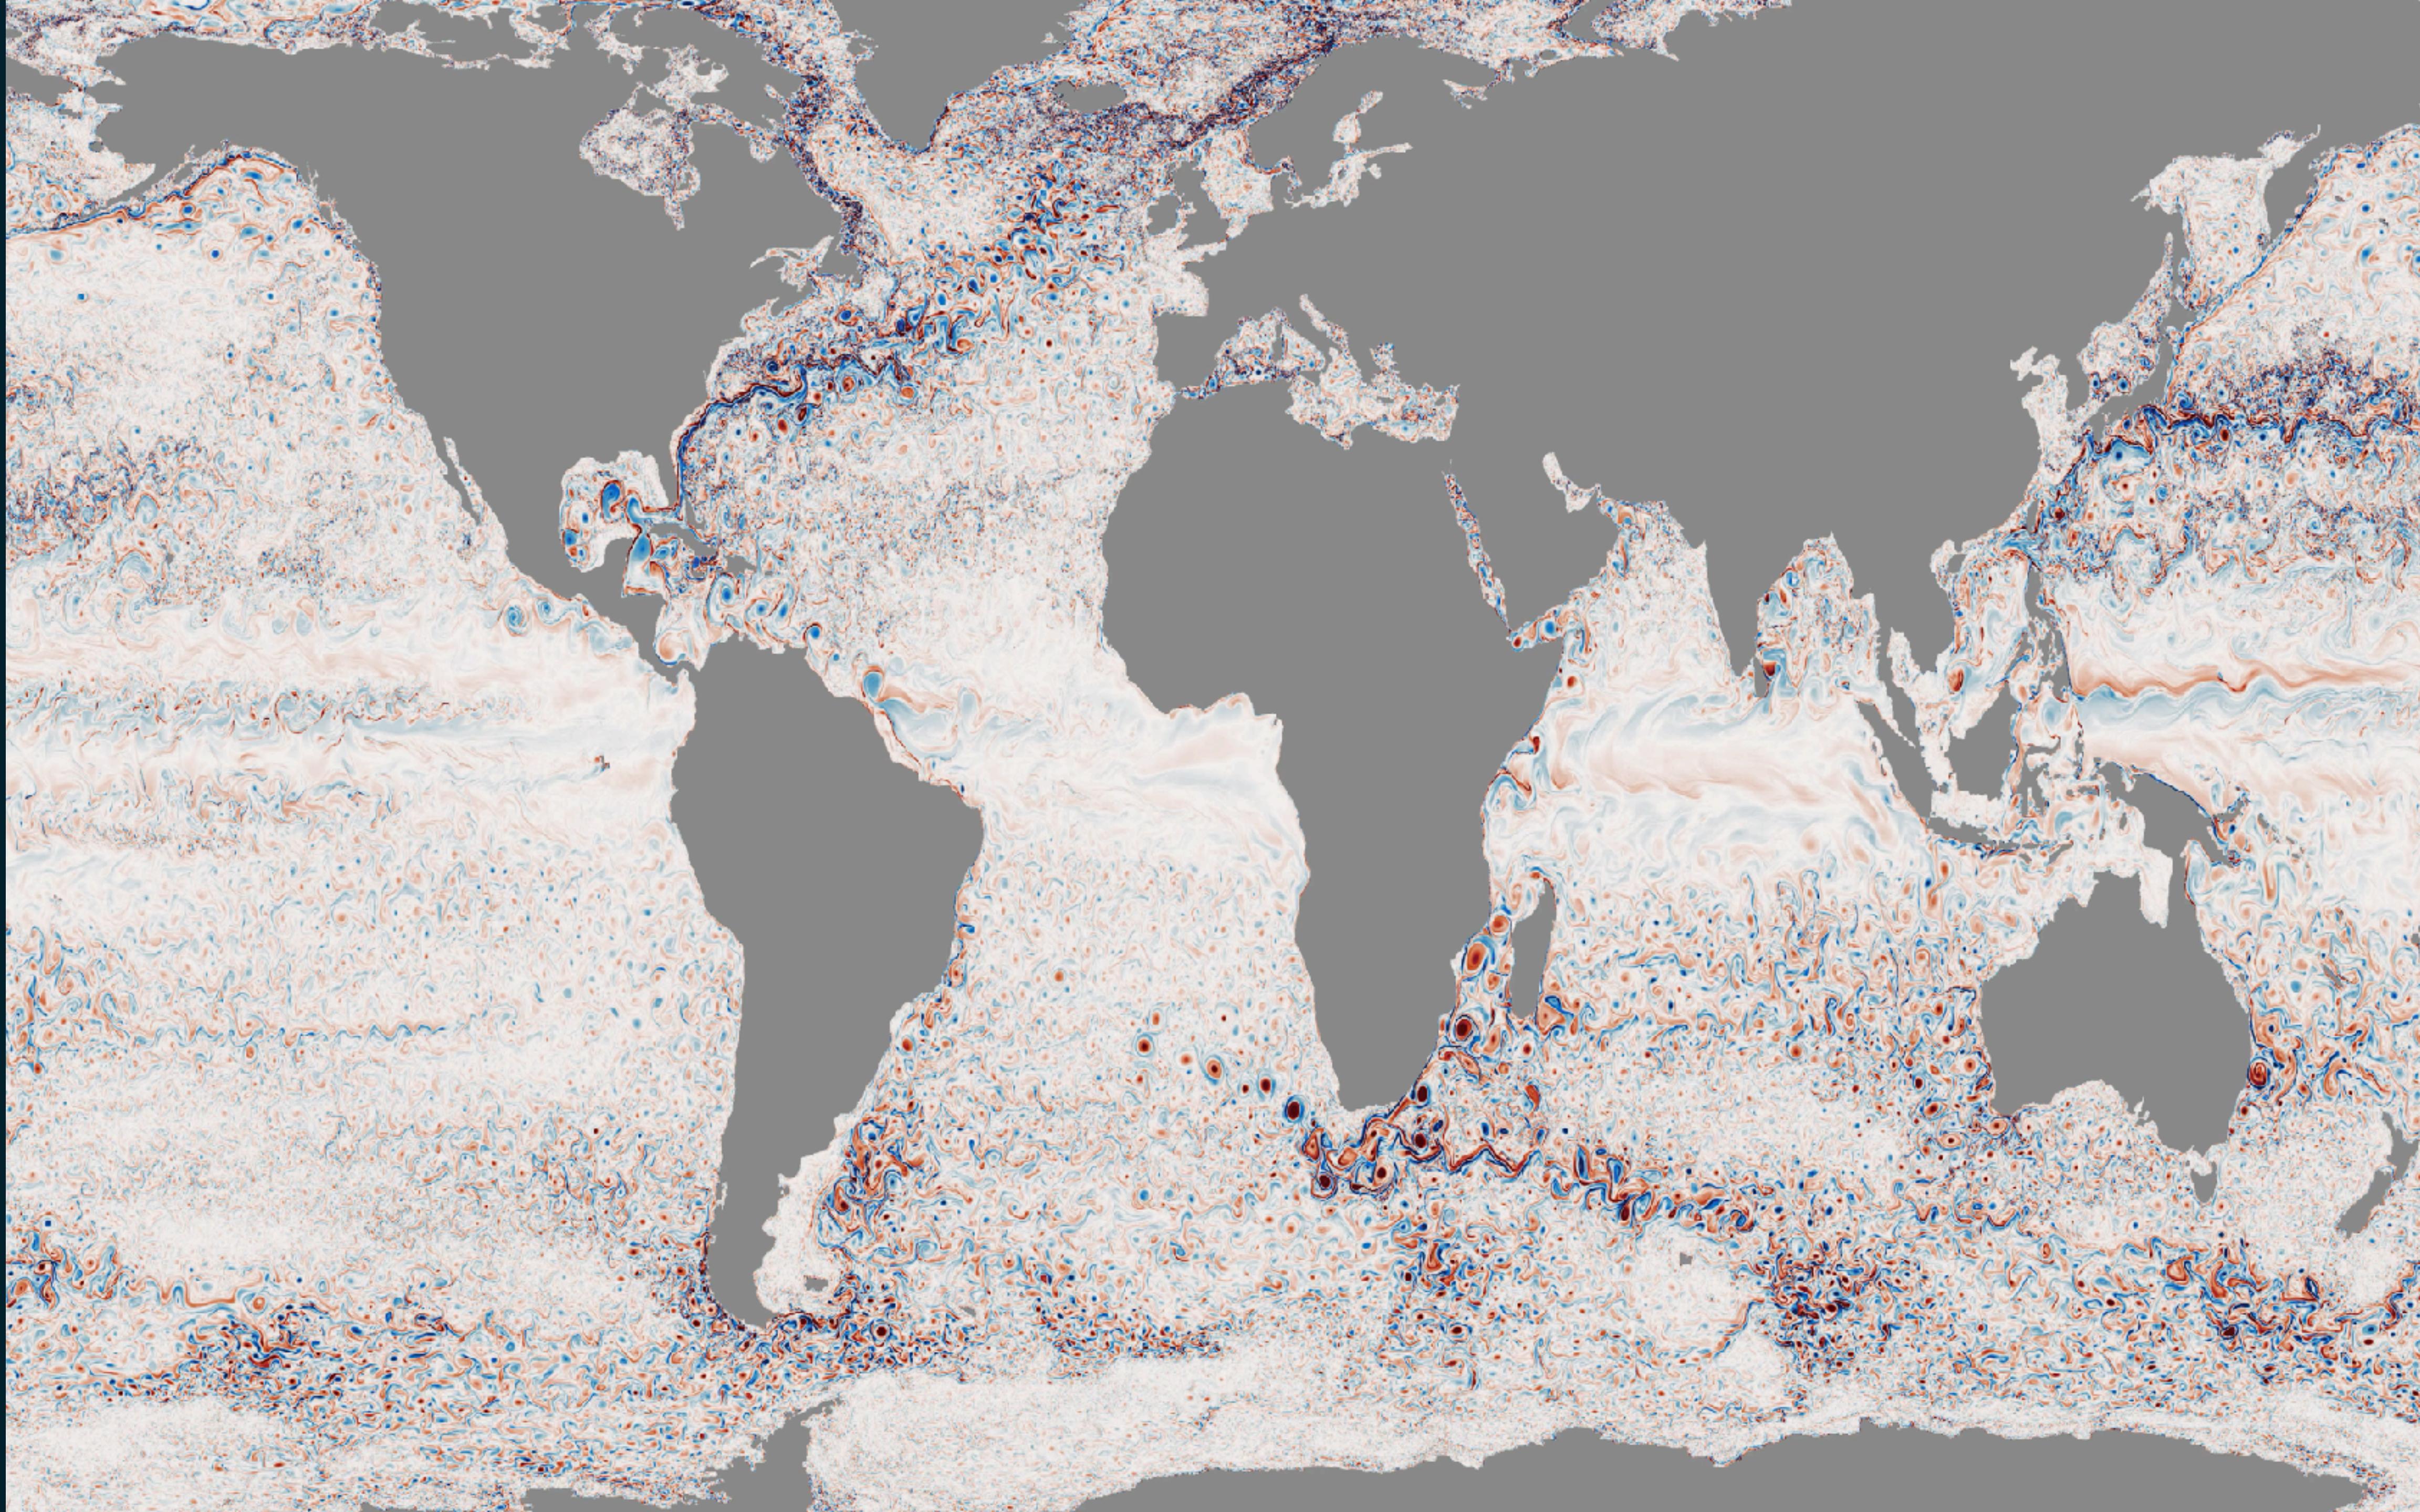
Heat sources

```
graph LR; A["\rho \left( \frac{\partial \mathbf{u}}{\partial t} + (\mathbf{u} \cdot \nabla) \mathbf{u} + 2\Omega \times \mathbf{u} \right) = \rho \mathbf{g} - \nabla p + \mu \nabla^2 \mathbf{u}"] -- "Momentum advection" --> B["\frac{\partial \rho}{\partial t} + \nabla \cdot (\rho \mathbf{u}) = 0"]; A -- "Coriolis" --> C["p = \rho R T"]; A -- "Gravity" --> D["\frac{\partial T}{\partial t} + (\mathbf{u} \cdot \nabla) T - \frac{RT}{c_p} \frac{D \ln p}{Dt} = \frac{Q}{c_p}"]; A -- "Pressure gradient" --> E["p = \rho R T"]; A -- "Viscous forces" --> F["\frac{\partial T}{\partial t} + (\mathbf{u} \cdot \nabla) T - \frac{RT}{c_p} \frac{D \ln p}{Dt} = \frac{Q}{c_p}"]; B -- "Conservation of mass (compressible)" --> G["\frac{\partial \rho}{\partial t} + \nabla \cdot (\rho \mathbf{u}) = 0"]; B -- "Ideal gas law, equation of state" --> H["p = \rho R T"]; C -- "Temperature advection" --> I["\frac{\partial T}{\partial t} + (\mathbf{u} \cdot \nabla) T - \frac{RT}{c_p} \frac{D \ln p}{Dt} = \frac{Q}{c_p}"]; C -- "Heat sources" --> J["\frac{\partial T}{\partial t} + (\mathbf{u} \cdot \nabla) T - \frac{RT}{c_p} \frac{D \ln p}{Dt} = \frac{Q}{c_p}"]; D -- "Adiabatic conversion" --> K["\frac{\partial T}{\partial t} + (\mathbf{u} \cdot \nabla) T - \frac{RT}{c_p} \frac{D \ln p}{Dt} = \frac{Q}{c_p}"]
```

What's missing? Humidity, clouds, precipitation, snow, radiation, surface fluxes, ocean, land, ...

Global Ocean simulated at 10km

Silvestri et al. 2023 *arXiv:2309.06662*



Parameterizations

Representing the unresolved processes

Atmosphere

- Radiation
- Clouds
- Precipitation
- Convection
- Surface fluxes
- Turbulence
- Gravity waves

Ocean

- Vertical diffusion and mixing
- Eddies and turbulence
- Cross-thermocline transport
- Interaction with bathymetry
- Surface fluxes (atmosphere and ice)

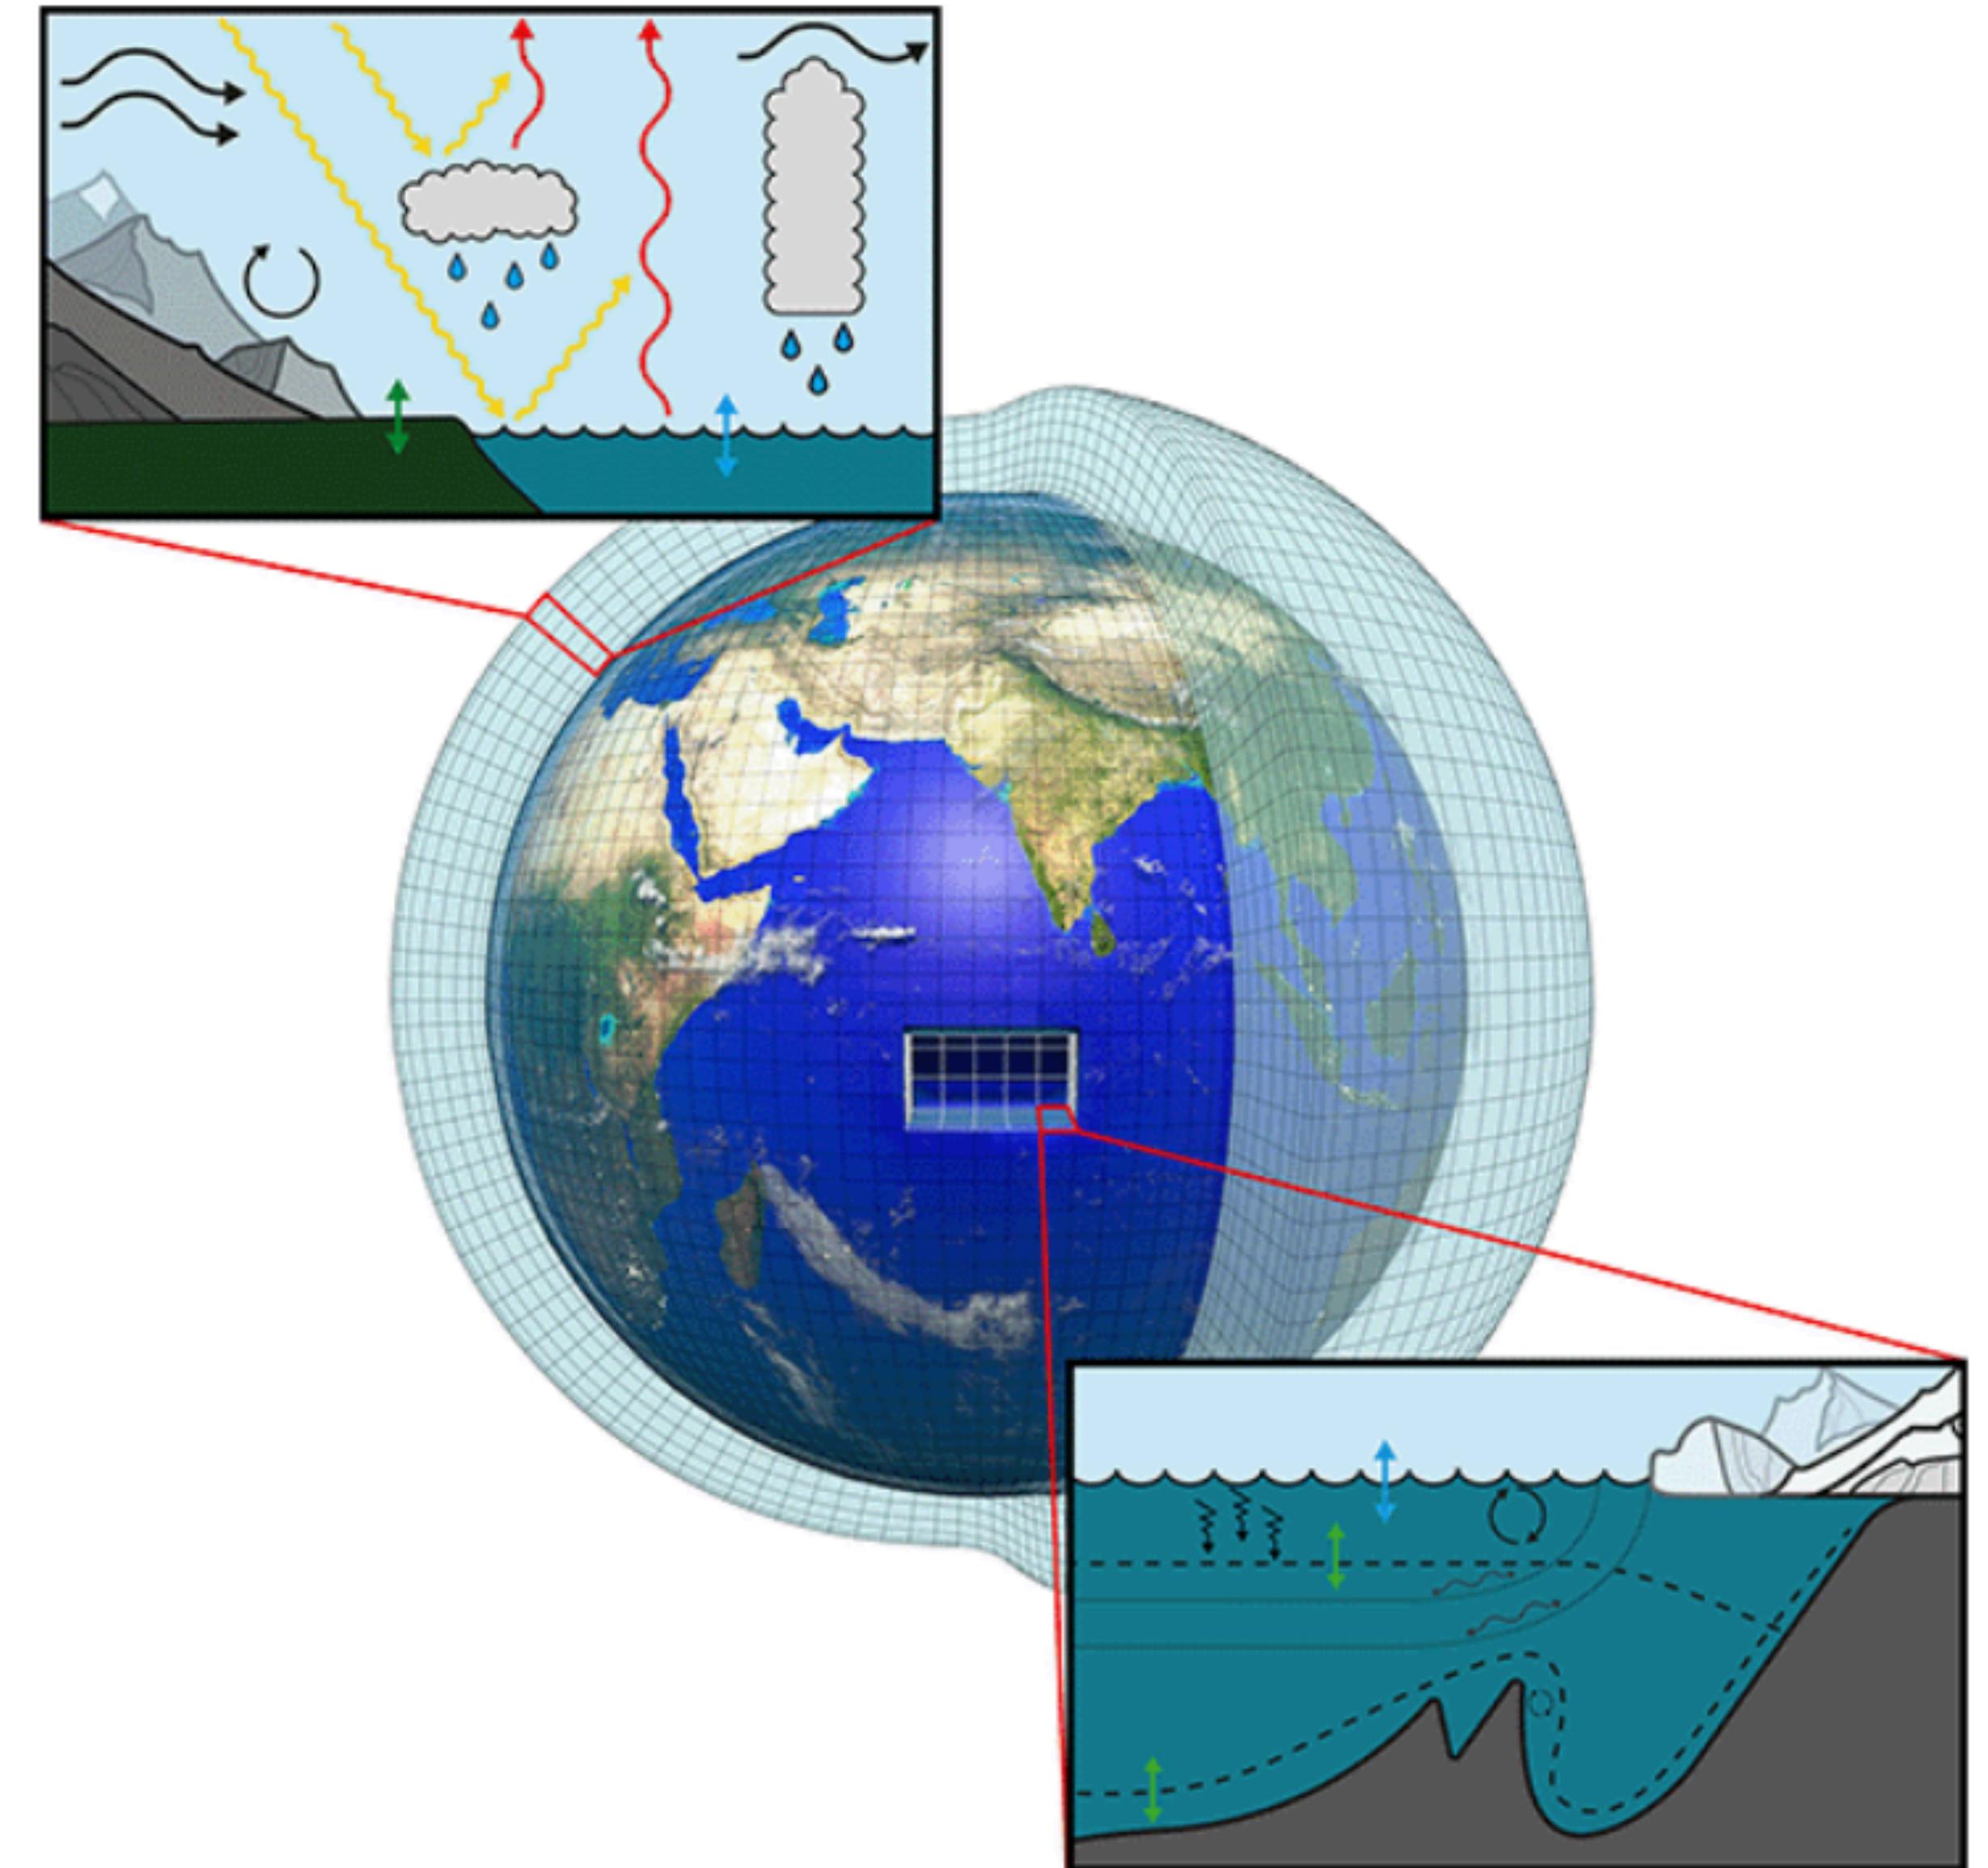


Figure 2. Atmospheric and oceanic parametrization schemes. Within an atmospheric column, typical parametrization schemes include moist convection, clouds and cloud microphysics, radiation, turbulence, turbulent exchange with the land and ocean, and both orographic and non-orographic gravity wave drag. Within an oceanic column, typical parametrization schemes include vertical diffusion, sub-mesoscale and mesoscale eddies, turbulent exchange with the atmosphere, and cross-thermocline transports.



*High-performance computing is cool
but how to accelerate climate research?*

Oldschool Fortran namelists: A *monolithic* interface?

You can only change what the developers wants you to change

- Speed ✓
- Cross-platform ✗
- Extensions ✗
- Open-source ✗
- Easy, interactive ✗
- Physical accuracy ✓

The NAMELIST input data file specifies the following:

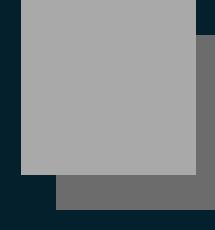
NFIRST=1	start from scratch
NLAST=42	stop after timestep 42
NNERGY=20	print fields every 20 timesteps
NTSI=1	print single line of information every timestep
NMIX=10	do a mixing timestep every 10 timesteps
AMF=1.E9	horizontal mixing of momentum
AHF=2.E7	horizontal mixing of heat, salinity, tracers
FKPMF=1.0	vertical mixing of momentum
FKPHF=1.0	vertical mixing of heat, salinity, tracers
DTTSF=172800.	length of timestep on temp., sal., tracers
DTUVF=7200.	length of timestep on internal mode velocities
DTSFF=7200.	length of timestep on stream function
NXSCAN=25	cut off relaxation at 25 scans if not converged
SORF=1.50	coefficient of over-relaxation
CRITF=4.E9	criterion for convergence of relaxation
ACORF=0.5	weight forward and backward timestep equally in the semi-implicit Coriolis term
TINITF=...	initial values of temperature, salinity, tracers (note.. for purposes of precision, salinity is carried in the model in units of parts per part, with .035 subtracted off, however TINITF is ex- pected in parts per part; .035 is subtracted after read-in)
ISIS=13,...	the island and its perimeter points are included in a box between I=13-16 and J=9-12

Monolithic interface

(Spend days or weeks
to compile the model)



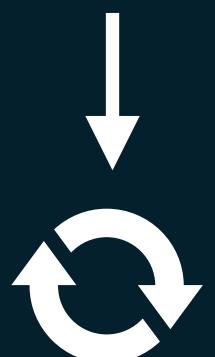
User



Namelist



Press play



Simulation

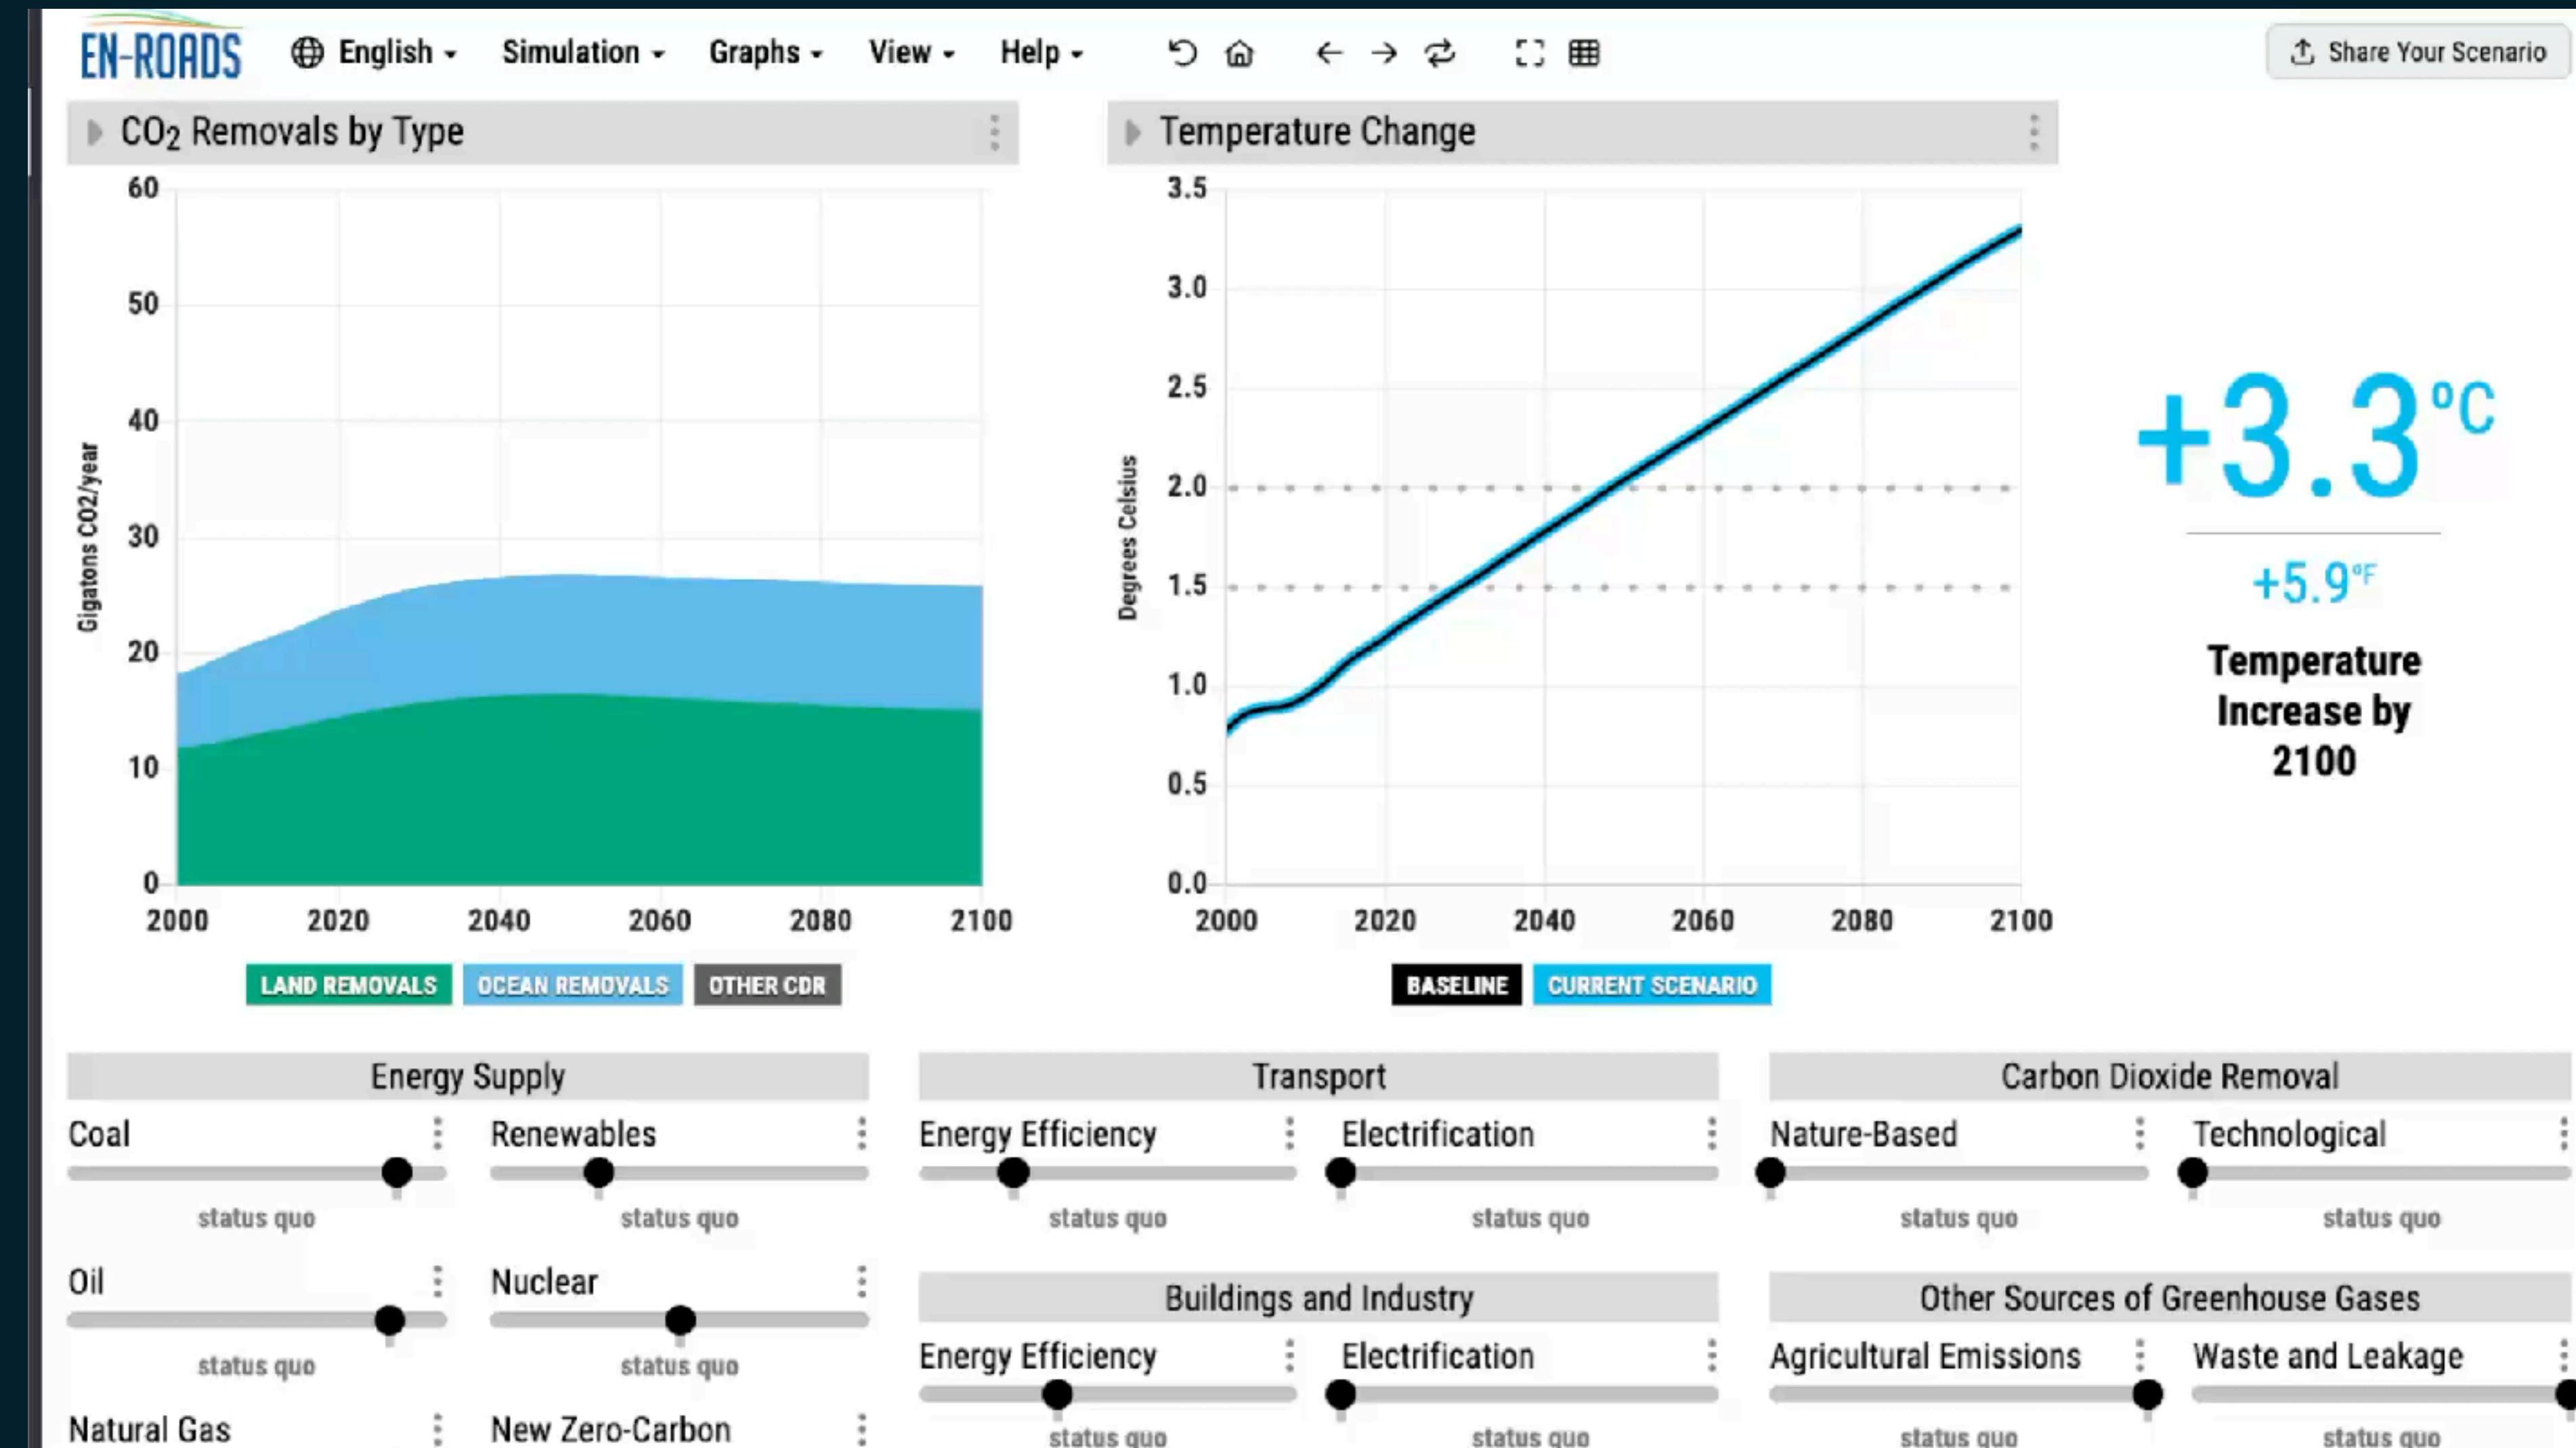


You can steer but you can't change the plane!

Interactive but also a monolithic interface

en-roads.climateinteractive.org

*Cool, but you still can
only change what the
developers wants you
to change*



Climate models are used for

(*not* weather forecast models)

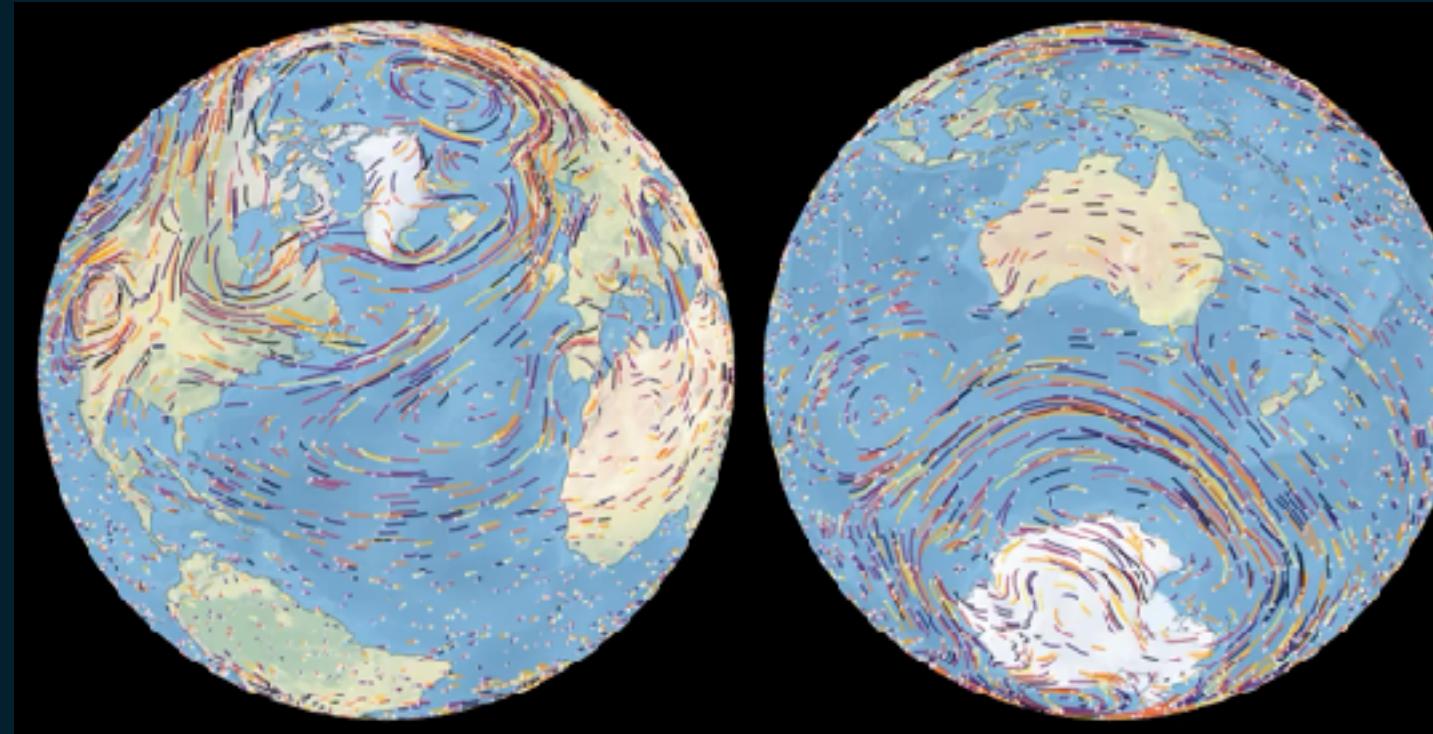
- Repeated pre-defined scenarios (CMIP etc)
- High-performance computing benchmarks
- Scientific one-off experiments
- Counterfactual worlds
- Software development
- Exoplanets
- Teaching (users become developers)
- ...

Many are not standard/repeated workflows!

Advantages of a monolithic interface

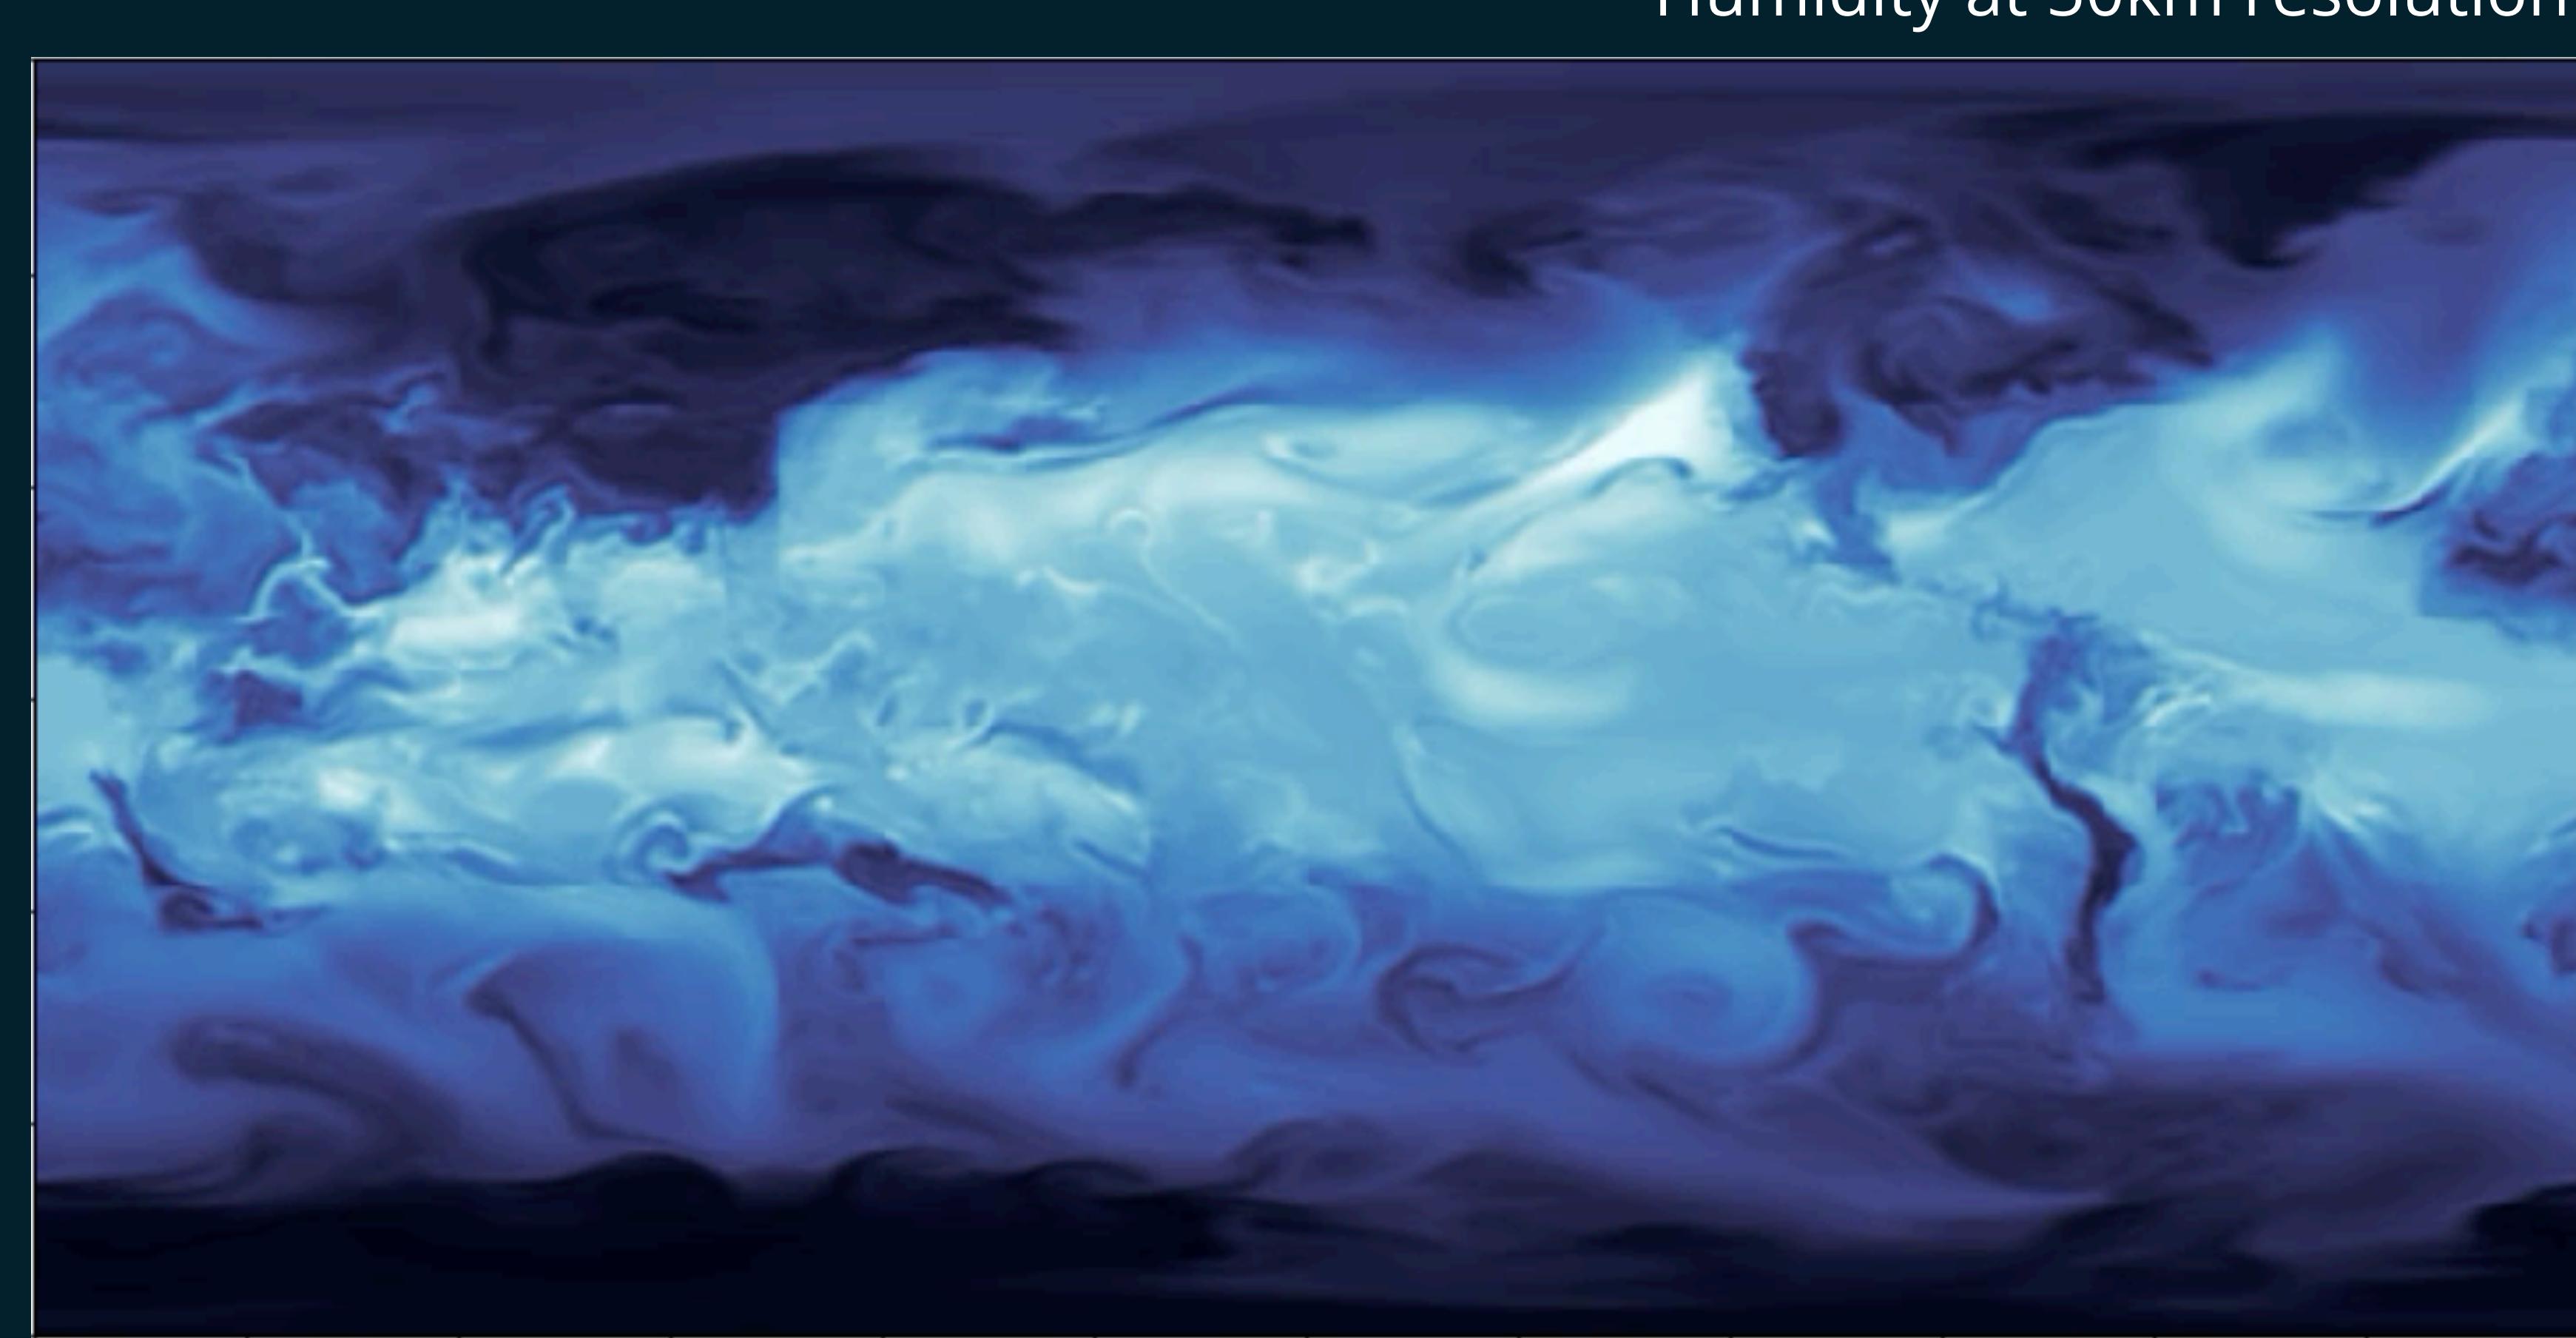
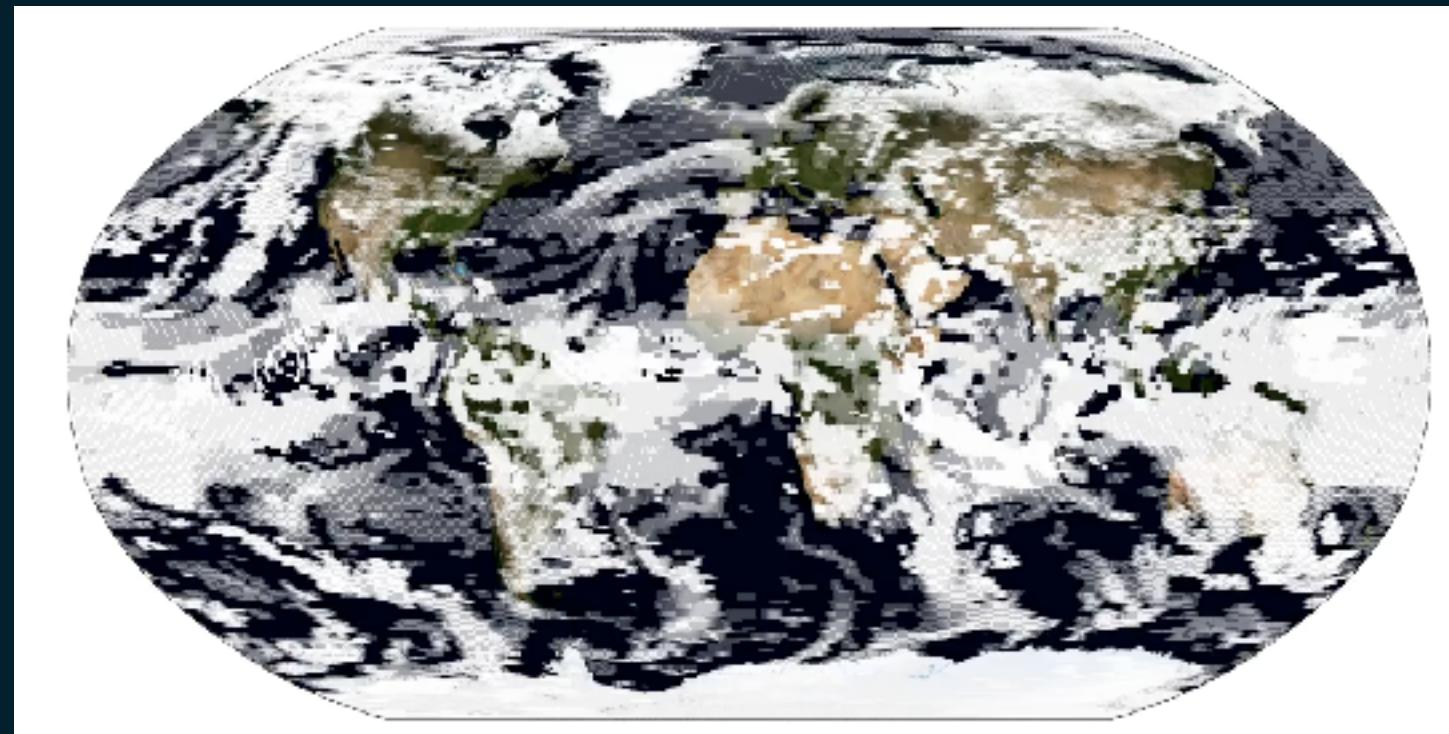
- Simpler interface
- Easier to design
- Easier to anticipate user experience
- Safety (harder to break things)
- Easier to maintain
- Developer/user separation

SpeedyWeather: Play climate modelling like it's LEGO



Particle advection

Cloud cover



Humidity at 50km resolution

Pier Luigi Vidale: "*unthinkable on such a small laptop ... like playing a video game*"

SpeedyWeather: Play climate modelling like it's LEGO

Honest (?) self assessment

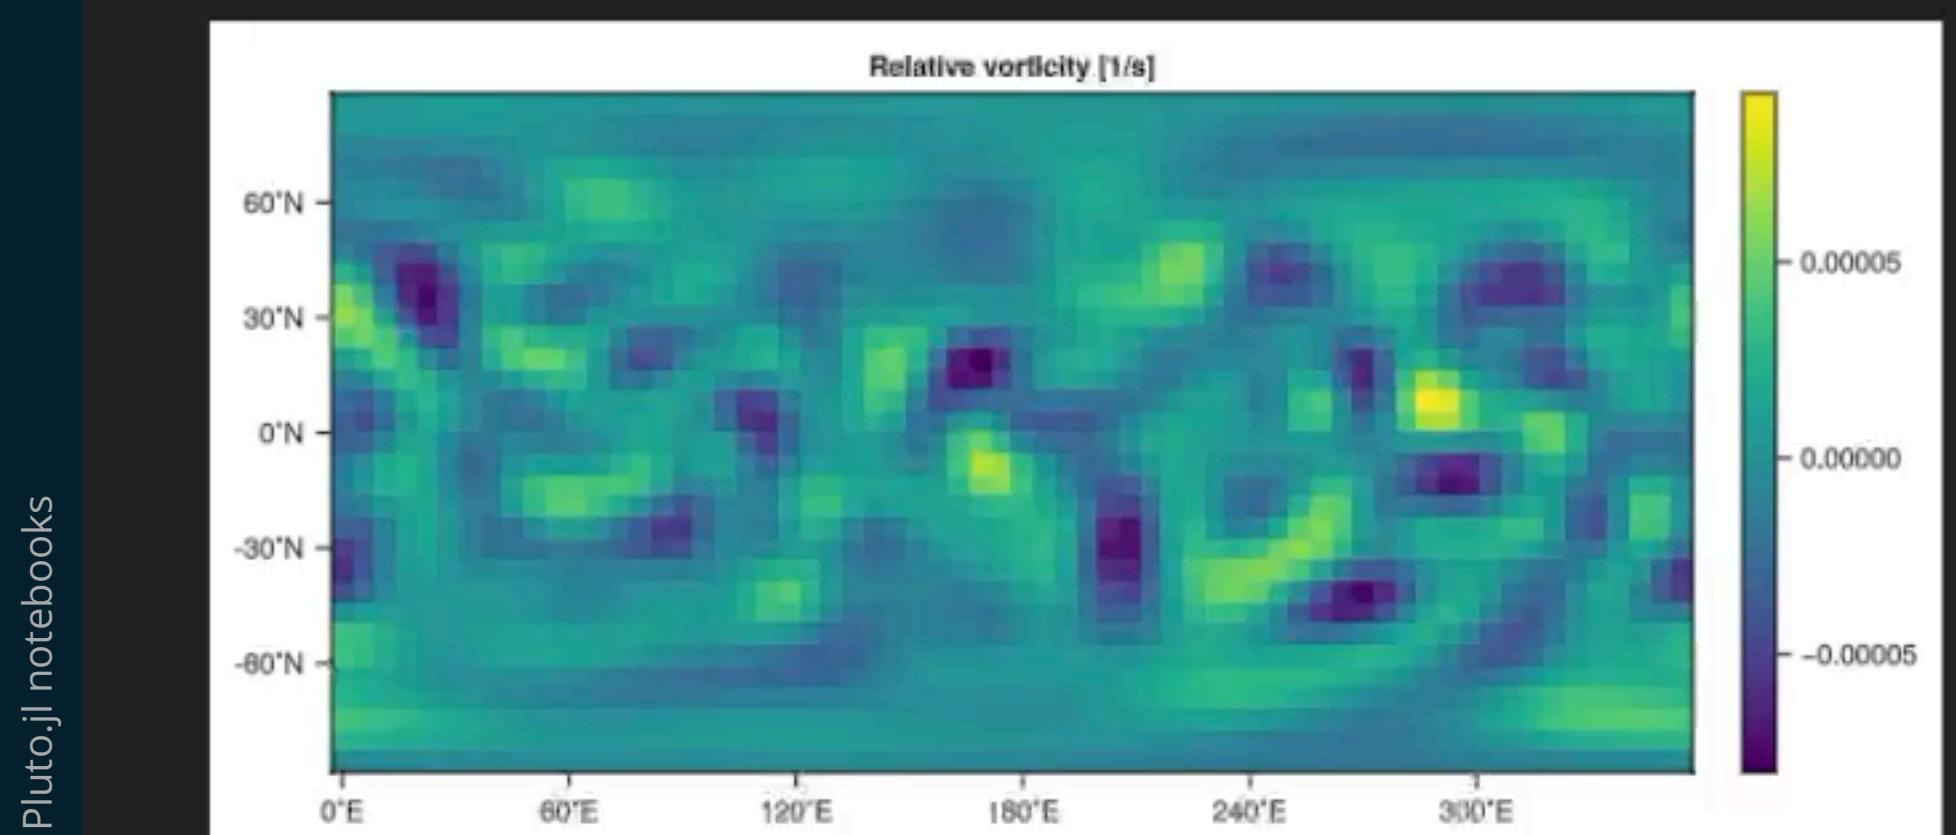
- Speed (incl GPUs)
- Cross-platform
- Extensions
- Open-source
- Easy, interactive
- Physical accuracy ?

SpeedyWeather demo

```
using SpeedyWeather ✓, PlutoUI ✓  
  
import GLMakie ✓  
  
Choose resolution  
+ ⚡  @bind resolution Slider([21, 31, 42, 63, 85])  
+ 33.4 ms  
  
Choose a forcing parameter  
+ ⚡  @bind m Slider(2:2:8)  
+ 33.9 ms  
  
Hack a quick extension to SpeedyWeather
```

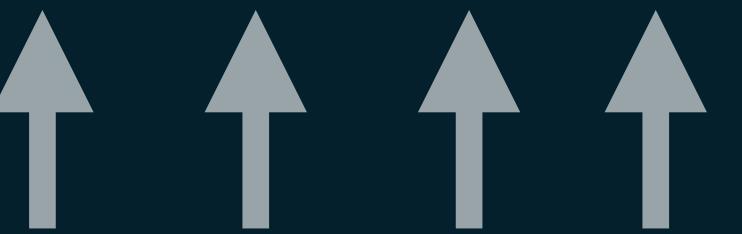
```
begin  
    struct MyKolmogorovFlow <: SpeedyWeather.AbstractForcing end  
    function SpeedyWeather.Forcing!(d, p, ::MyKolmogorovFlow, lf, model)  
        Fu = d.tendencies.u_tend_grid  
        set!(Fu, (λ, φ, σ) -> 3e-12*sind(m*φ), model.geometry)  
    end  
end
```

```
run!(simulation, period=Day(3));
```

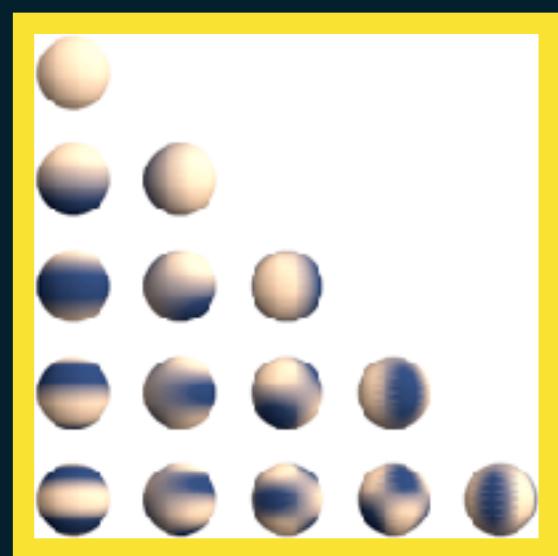


SpeedyWeather's monorepo

SpeedyWeather.jl



SpeedyTransforms.jl



LowerTriangularArrays.jl

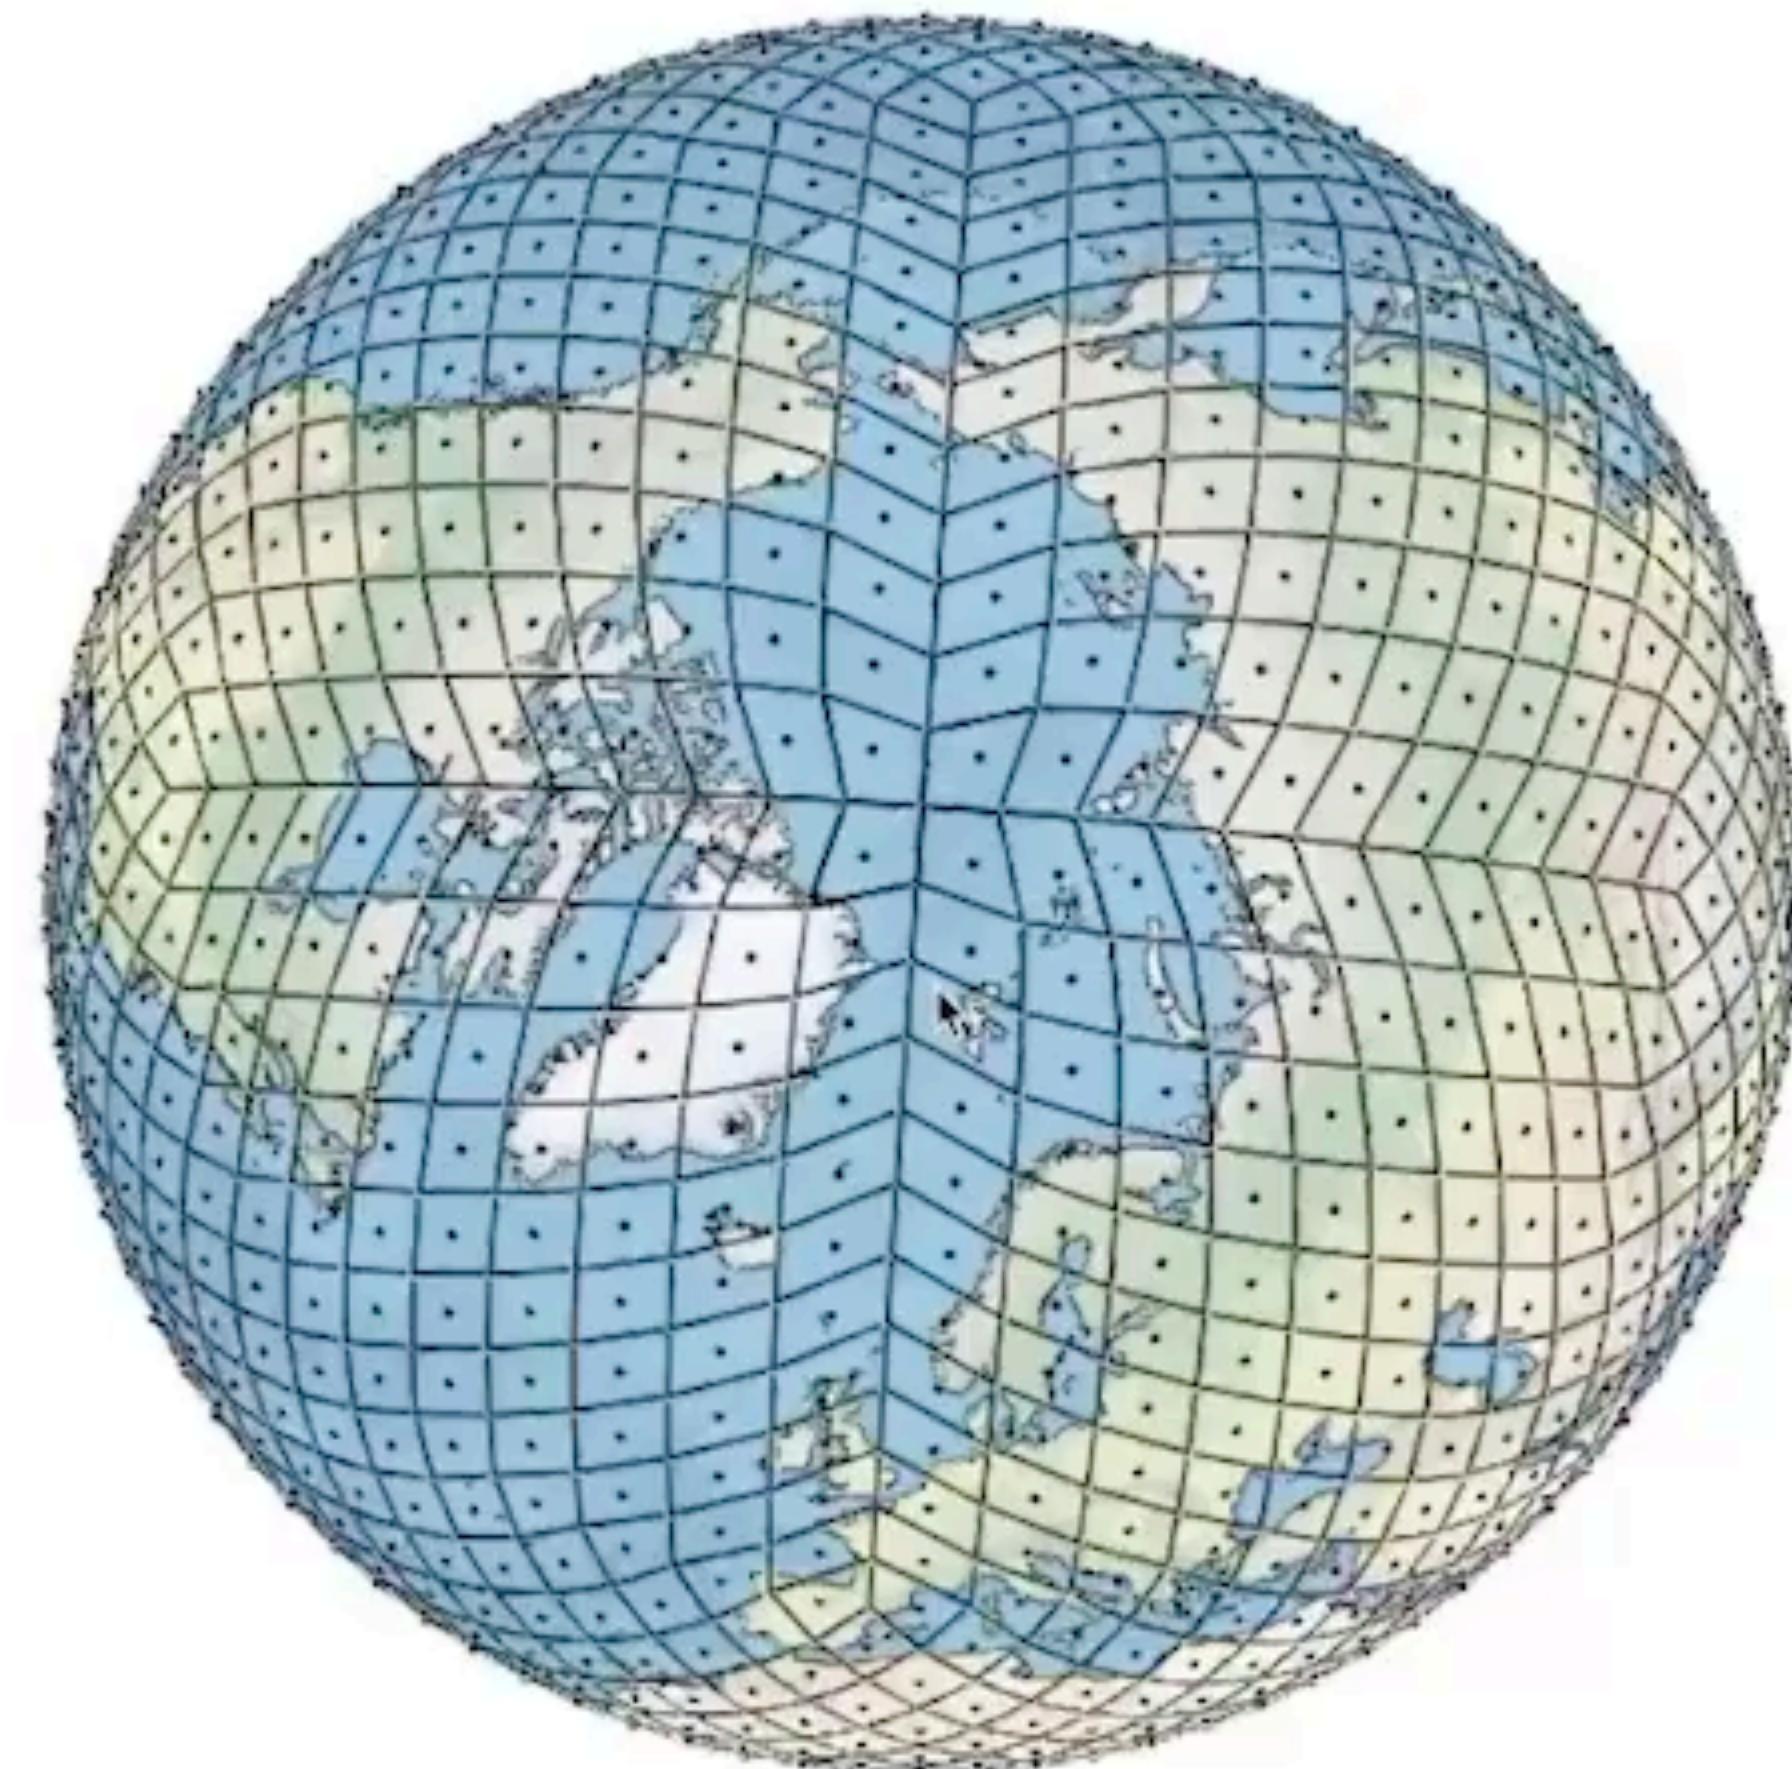


RingGrids.jl



SpeedyWeatherInternals.jl

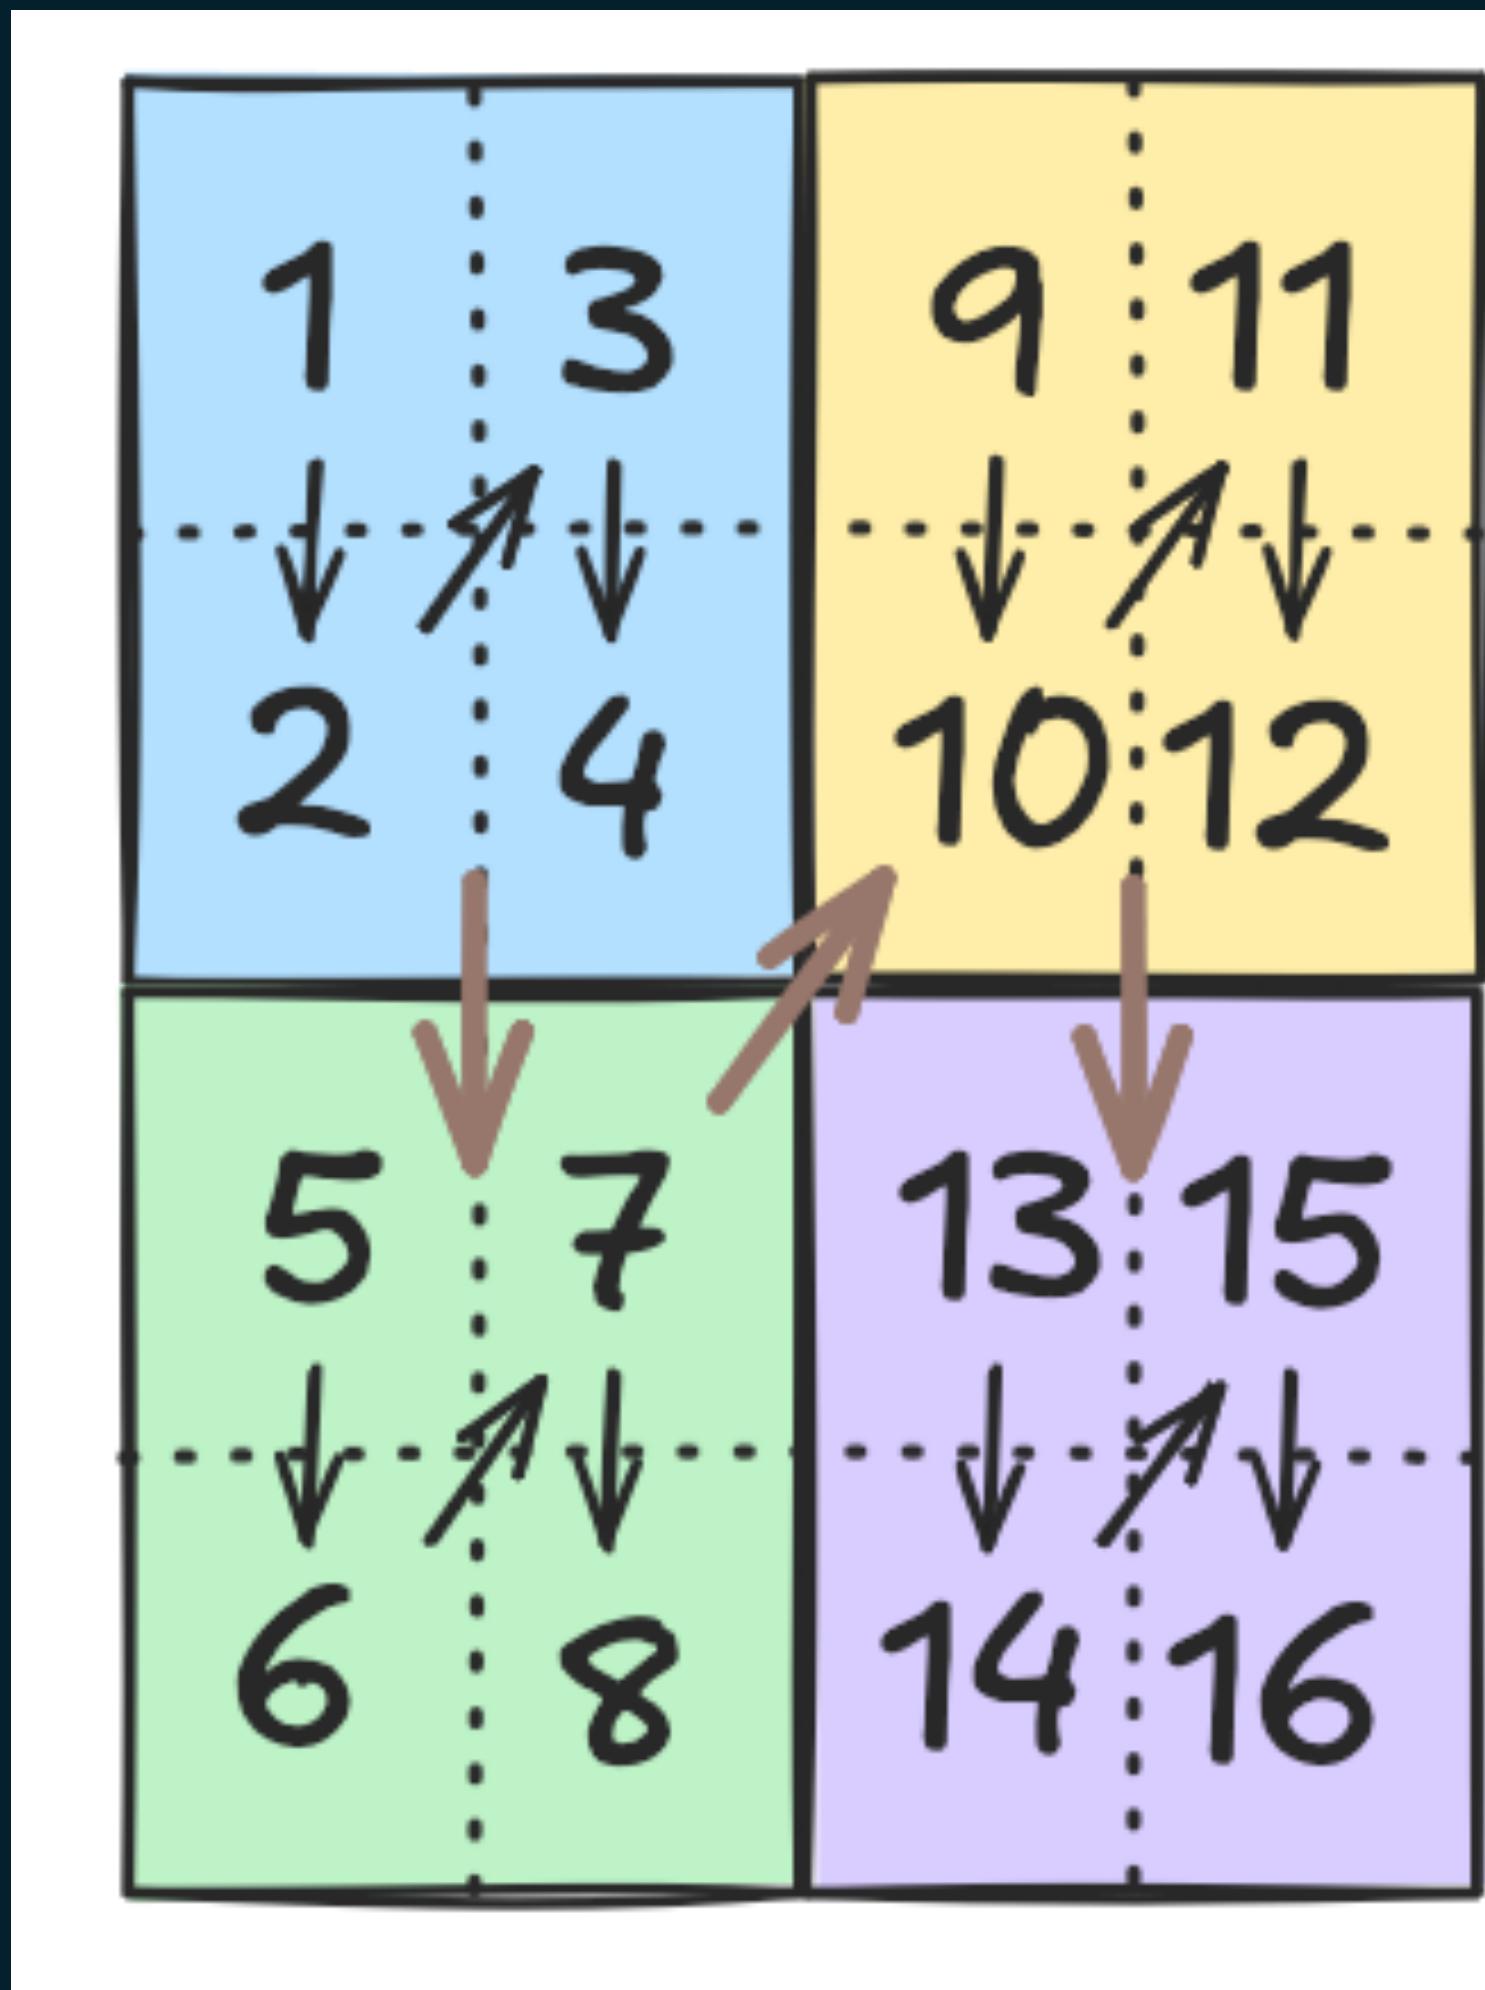




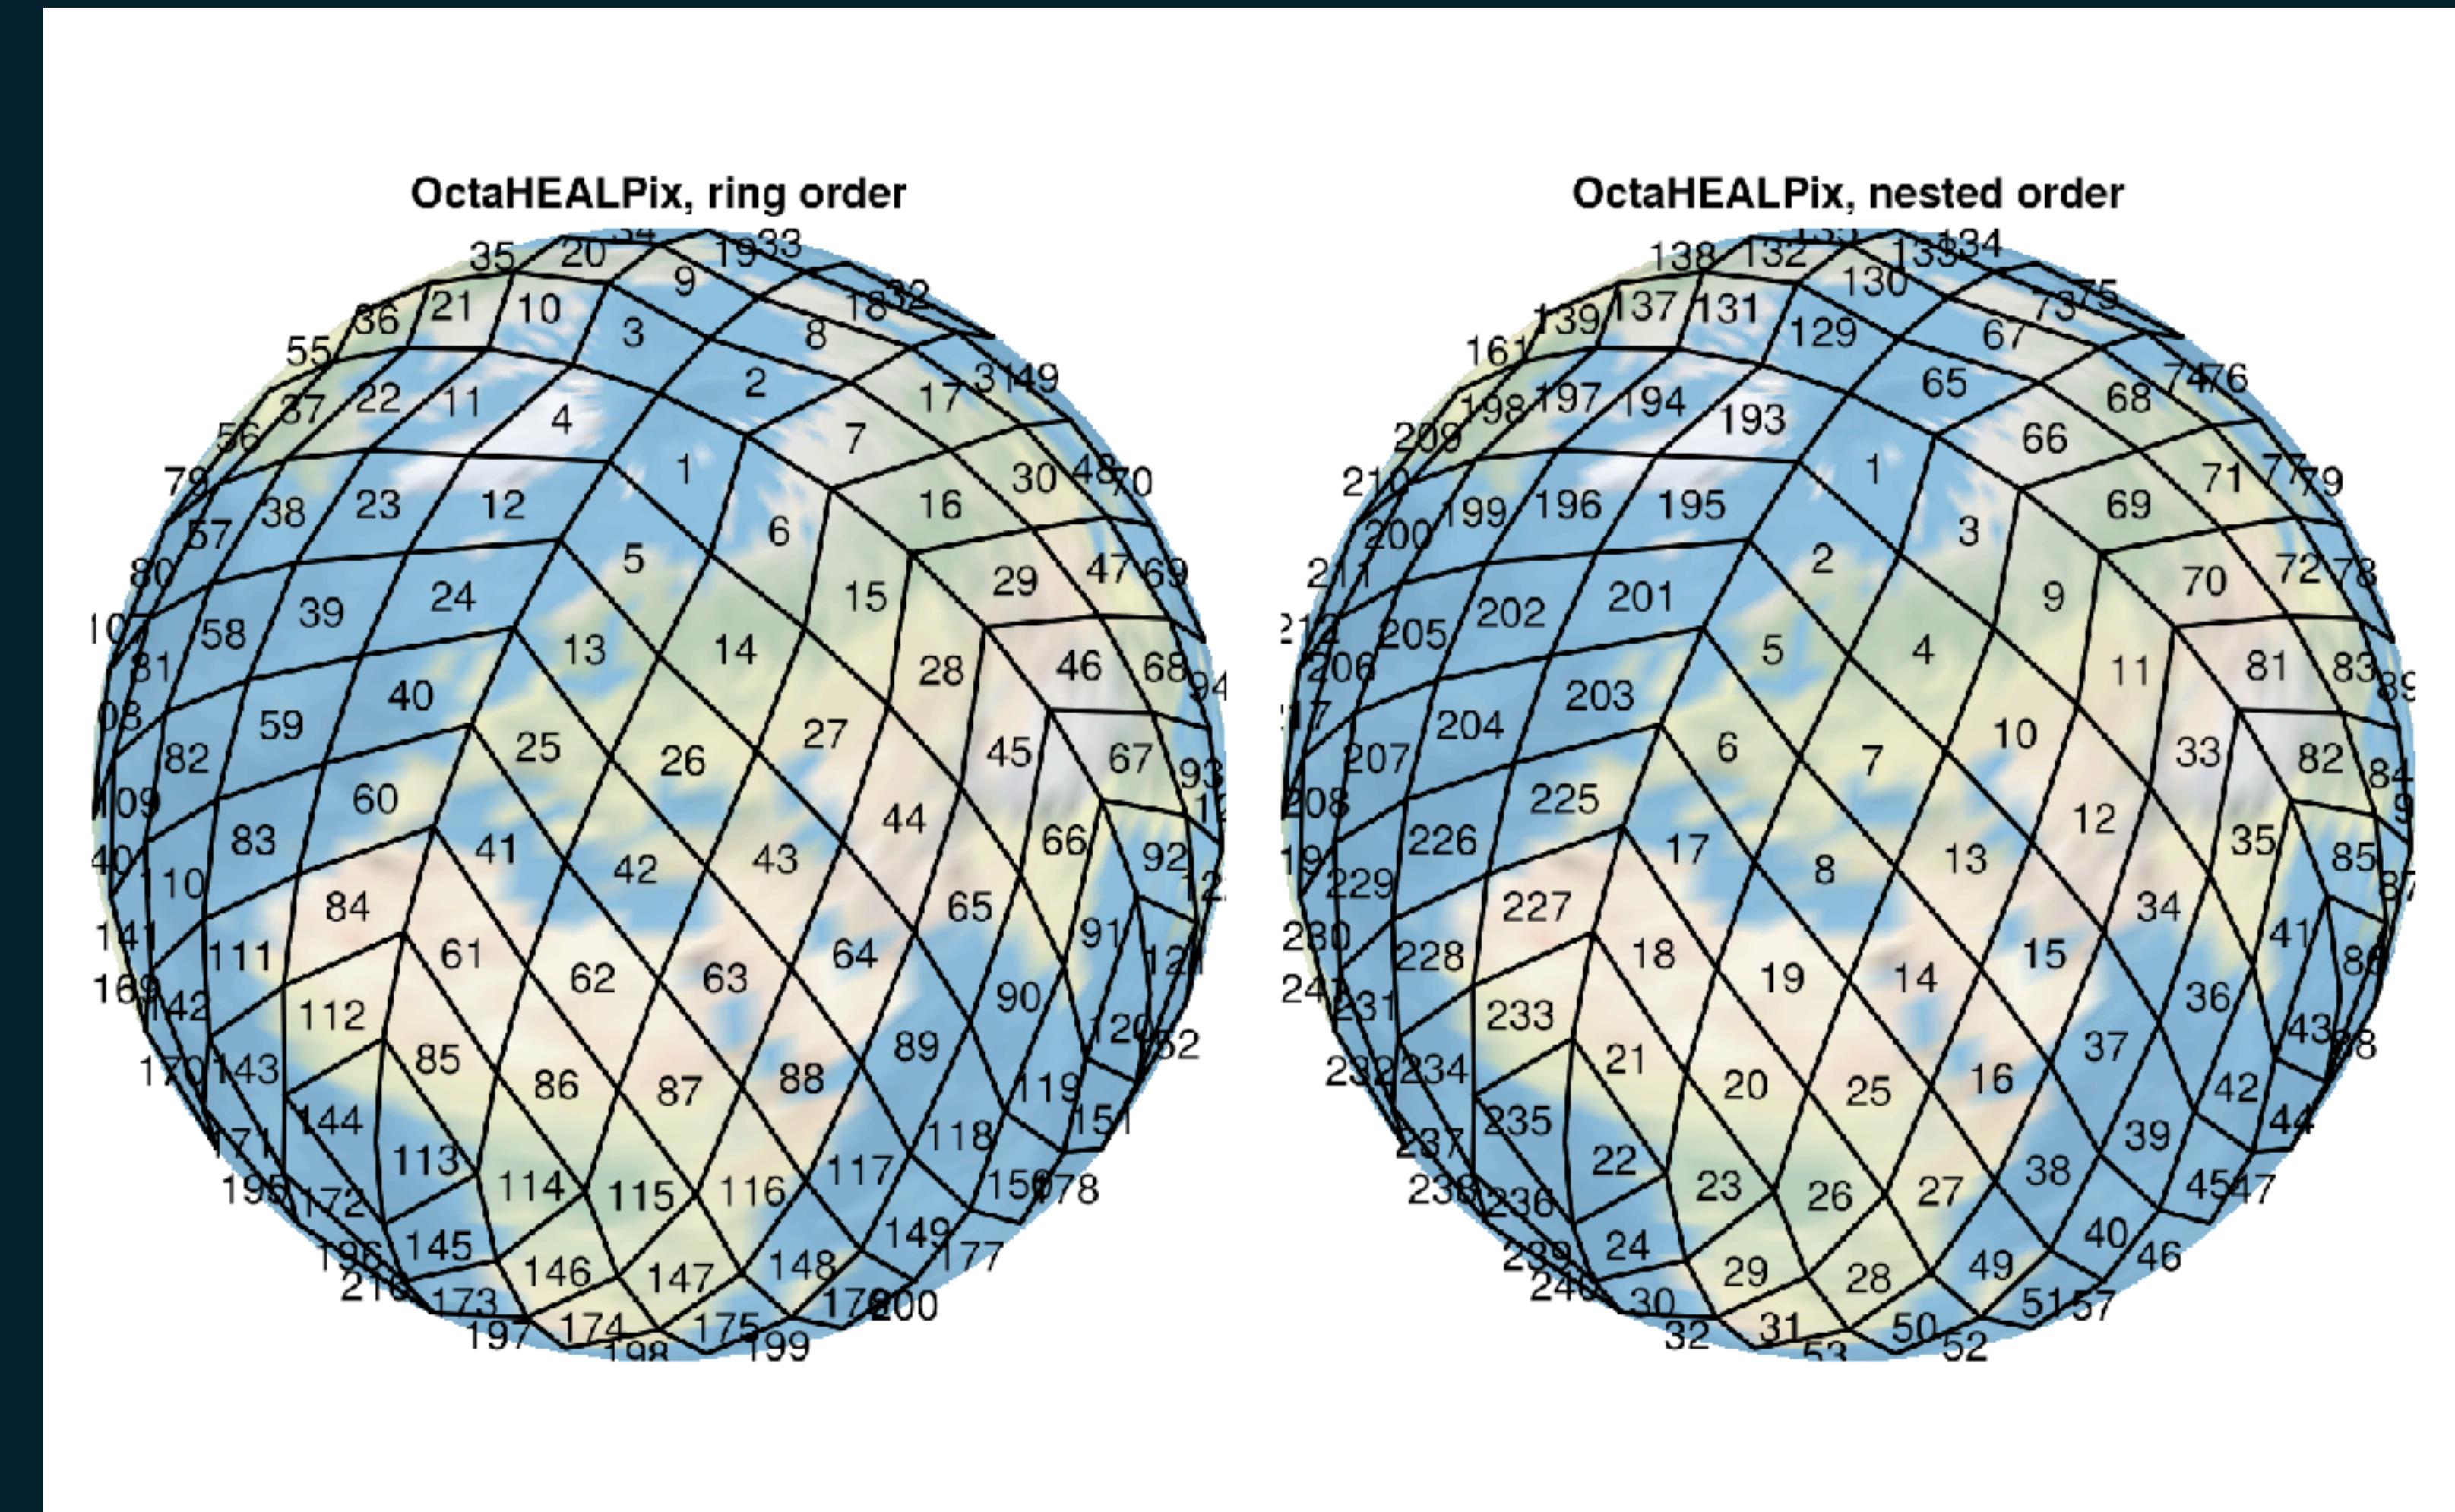
HEALPix grid visualised with GLMakie+GeoMakie

- Gaussian and octahedral grids
- HEALPix
- OctaHEALPix (new!)
- N-dimensional fields
- Array-agnostic
- GPU ready

RingGrids.jl — hierarchical HEALPix grids



*Hierarchical yields trivial
resolution changes*

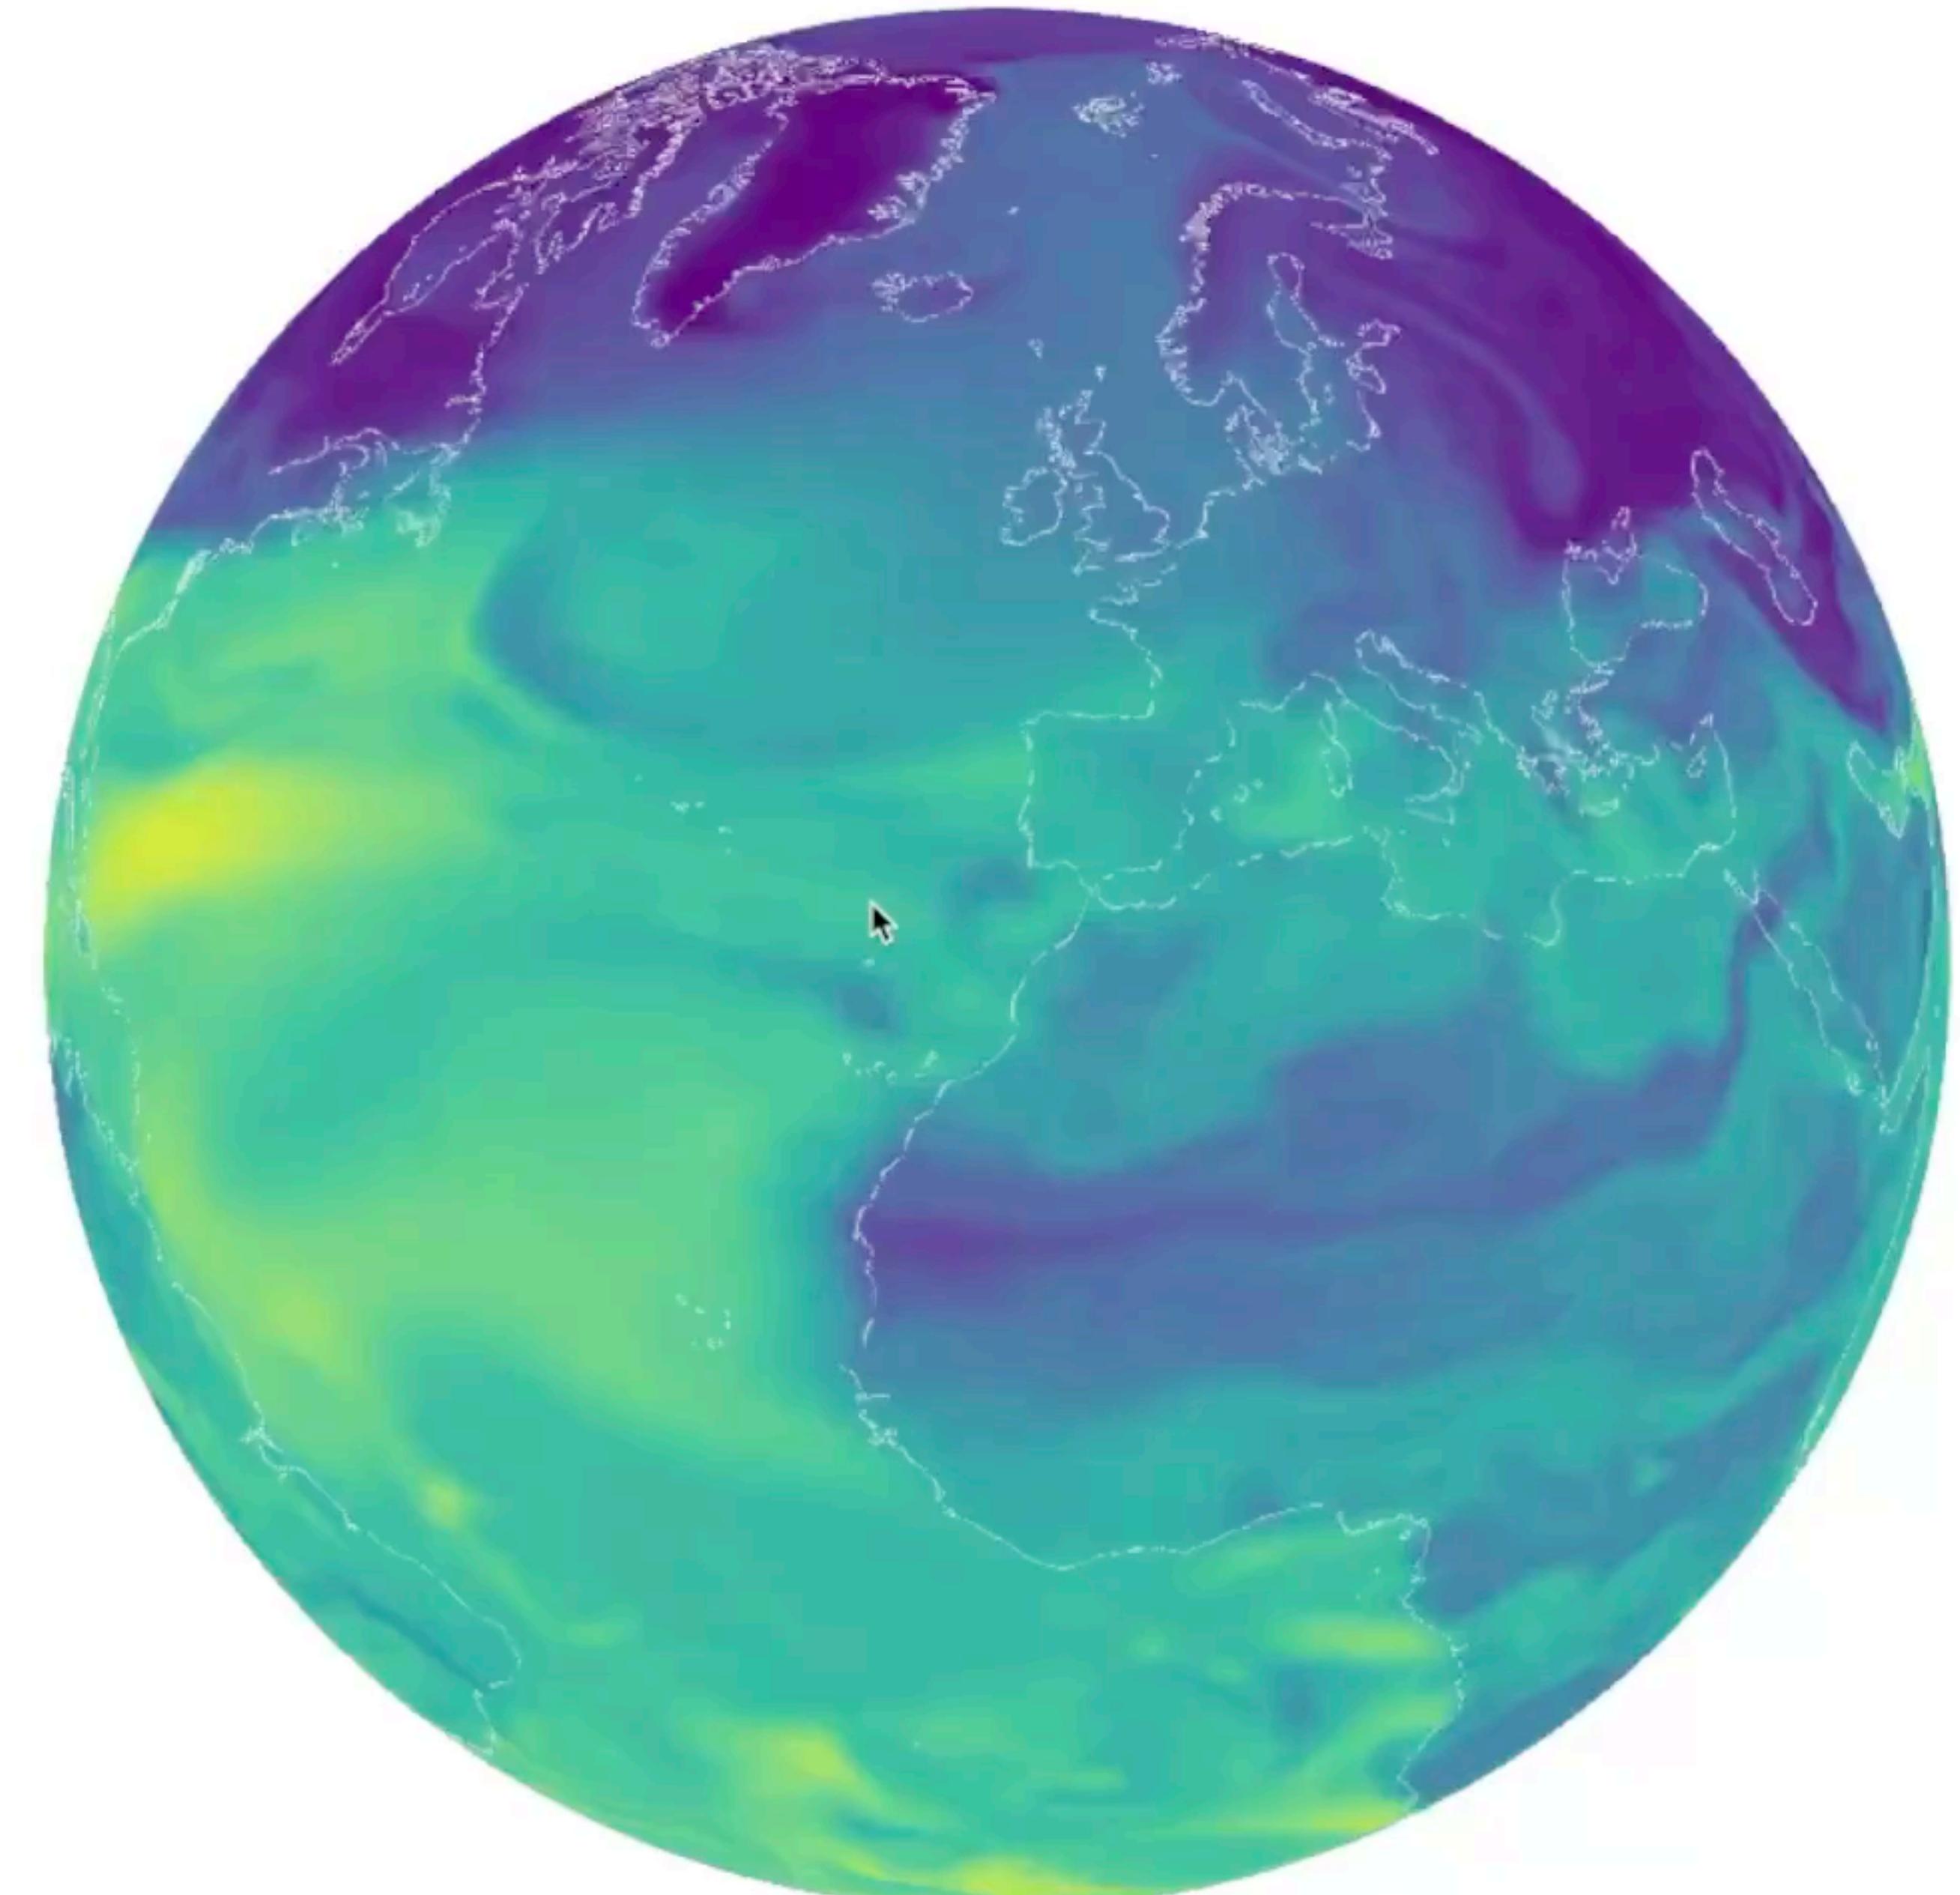
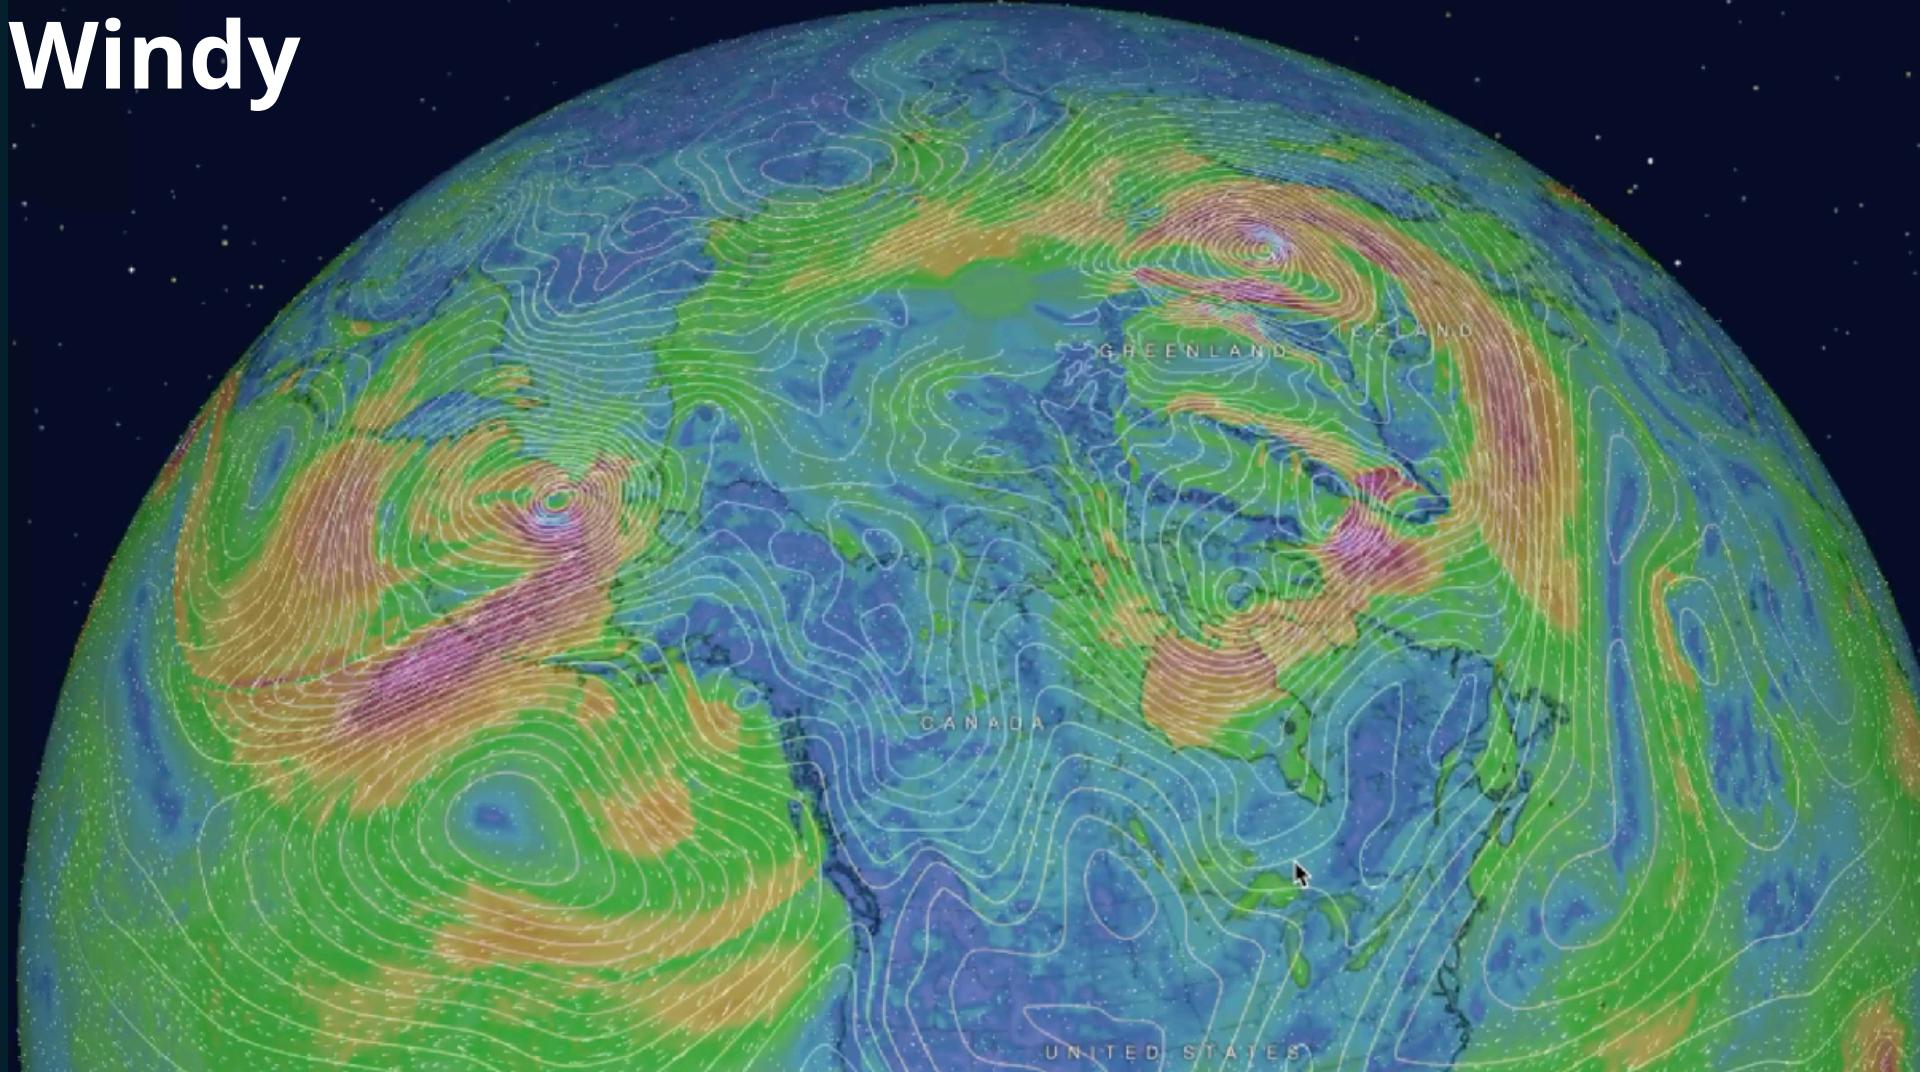


*OctaHEALPix: Equal-area, efficient spectral transform,
hierarchical, and just one big square matrix*

Climate modelling like playing with windy?

Fields plotted as polygons of the underlying grid

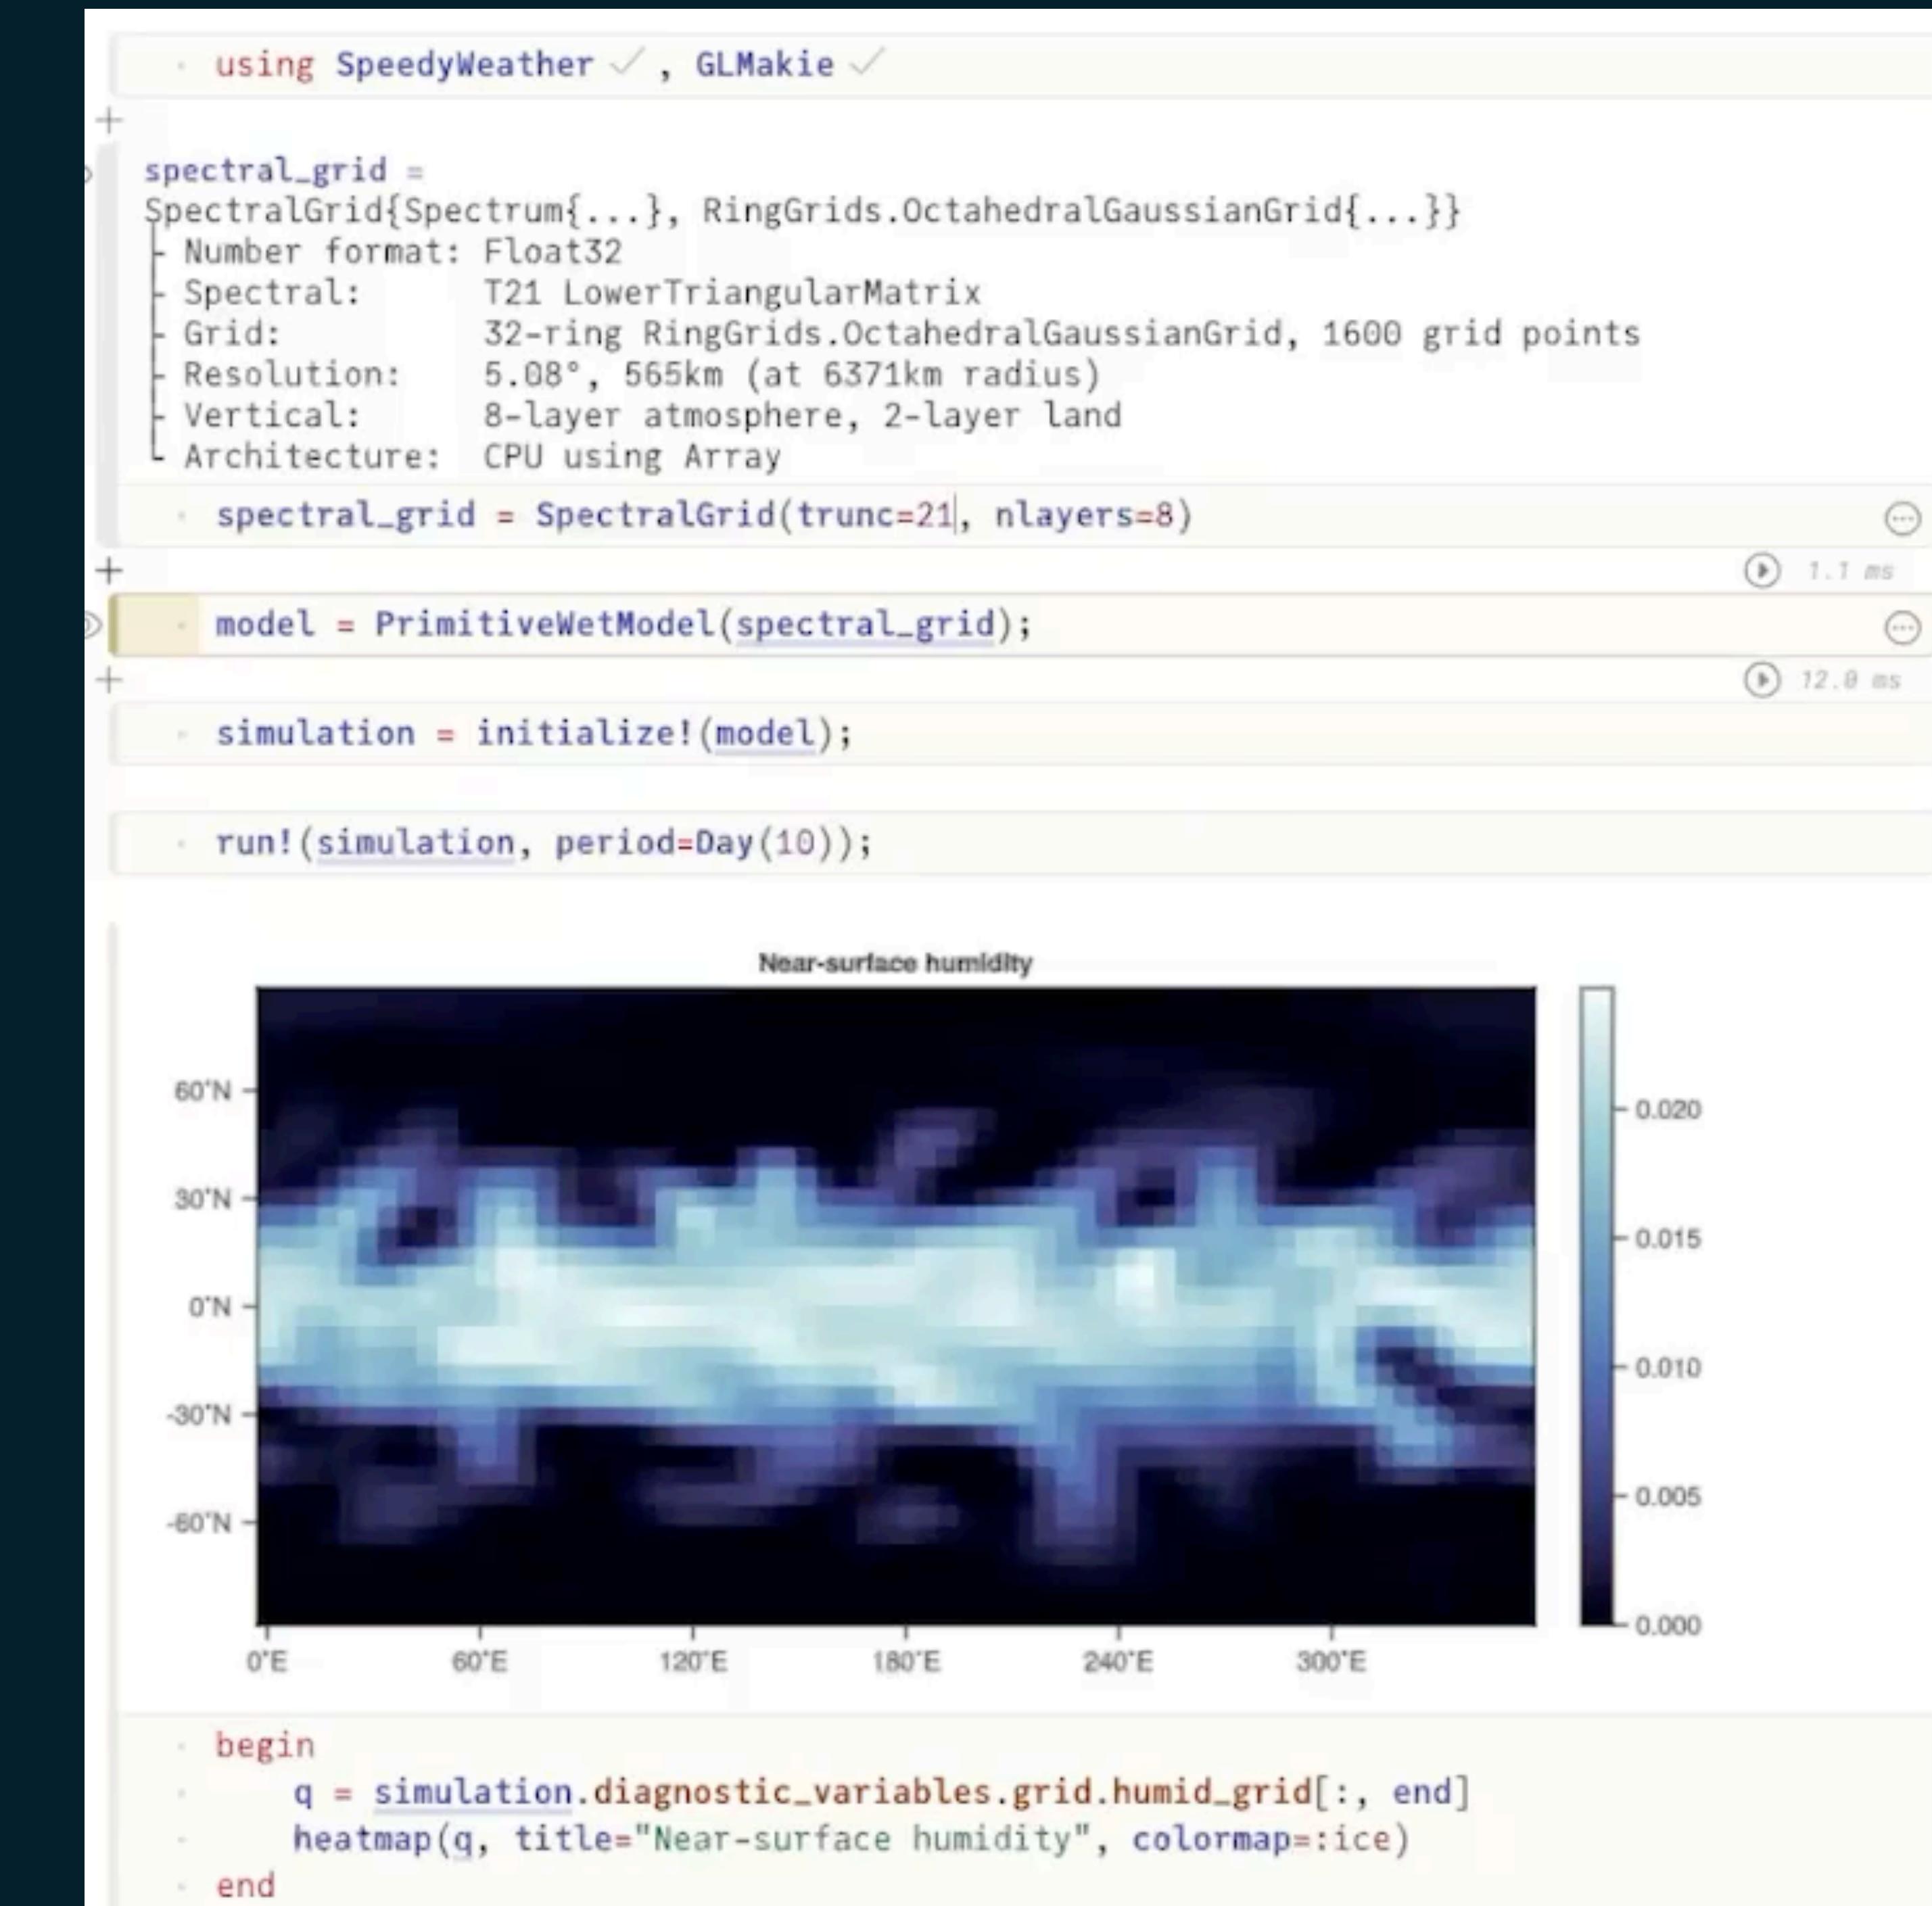
- Add a play button?
- Switch fields of a simulation?



SpeedyWeather+GeoMakie+GLMakie

SpeedyWeather: Play climate modelling like it's LEGO

- Modularity
- Flexible
- Hackable
- Easy to use
- Visualisation support
- GPU (almost) ready
- Differentiable via Enzyme



Learning surface roughness

Roughness length =

f(

vegetation,
sea ice,
snow,
orography,
surface wind

)



Learned offline in pytorch,

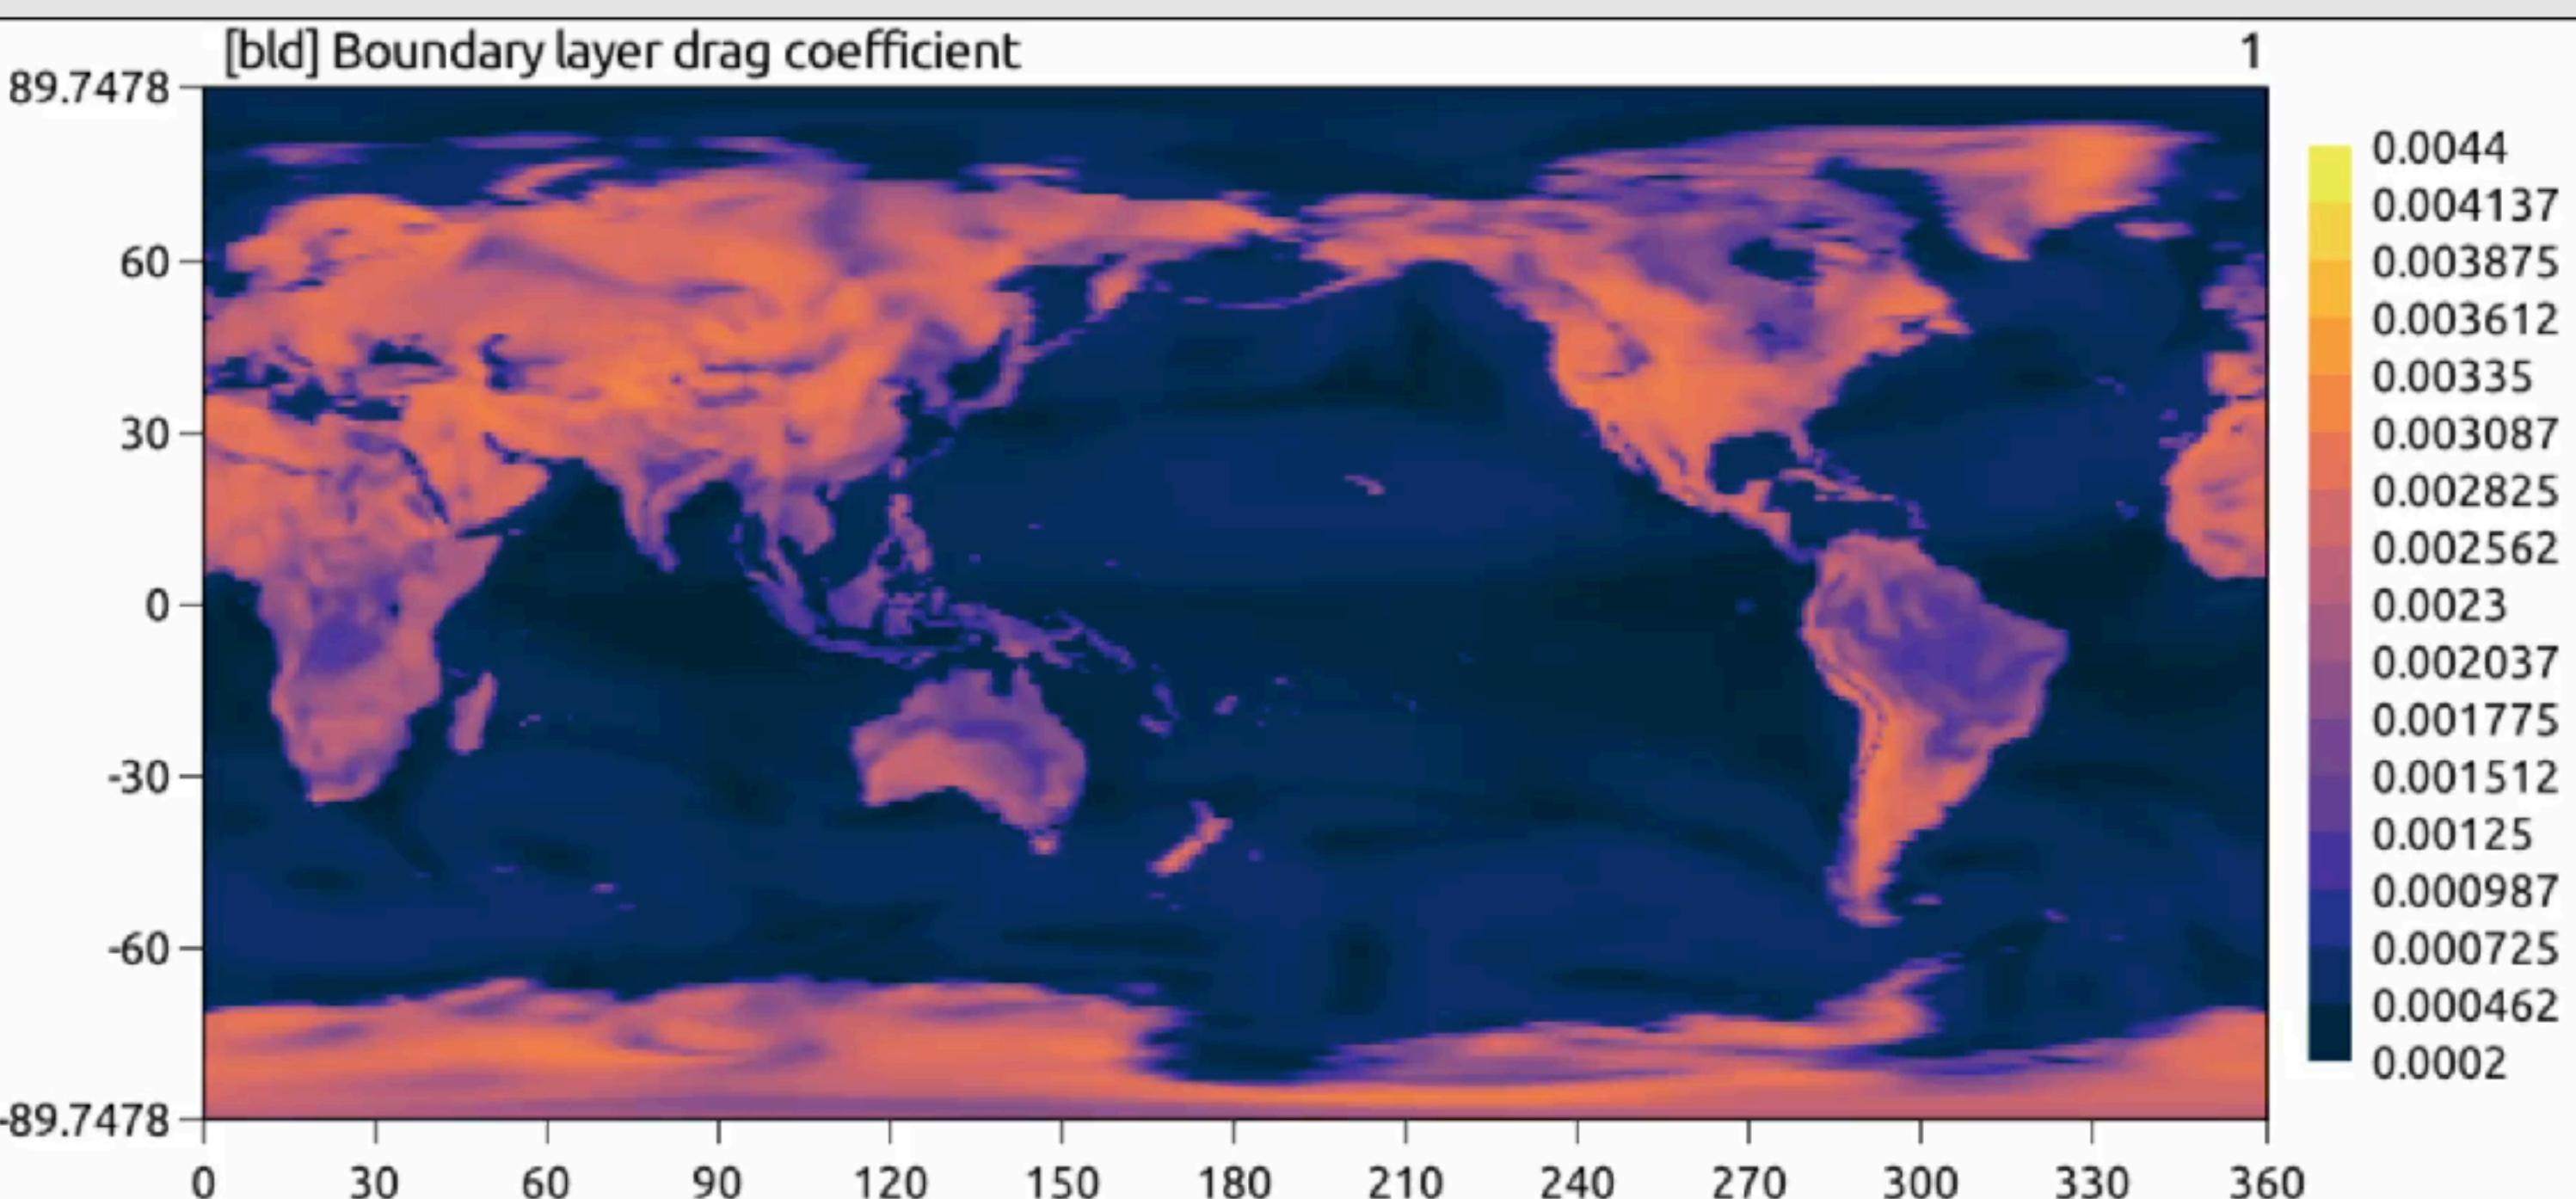
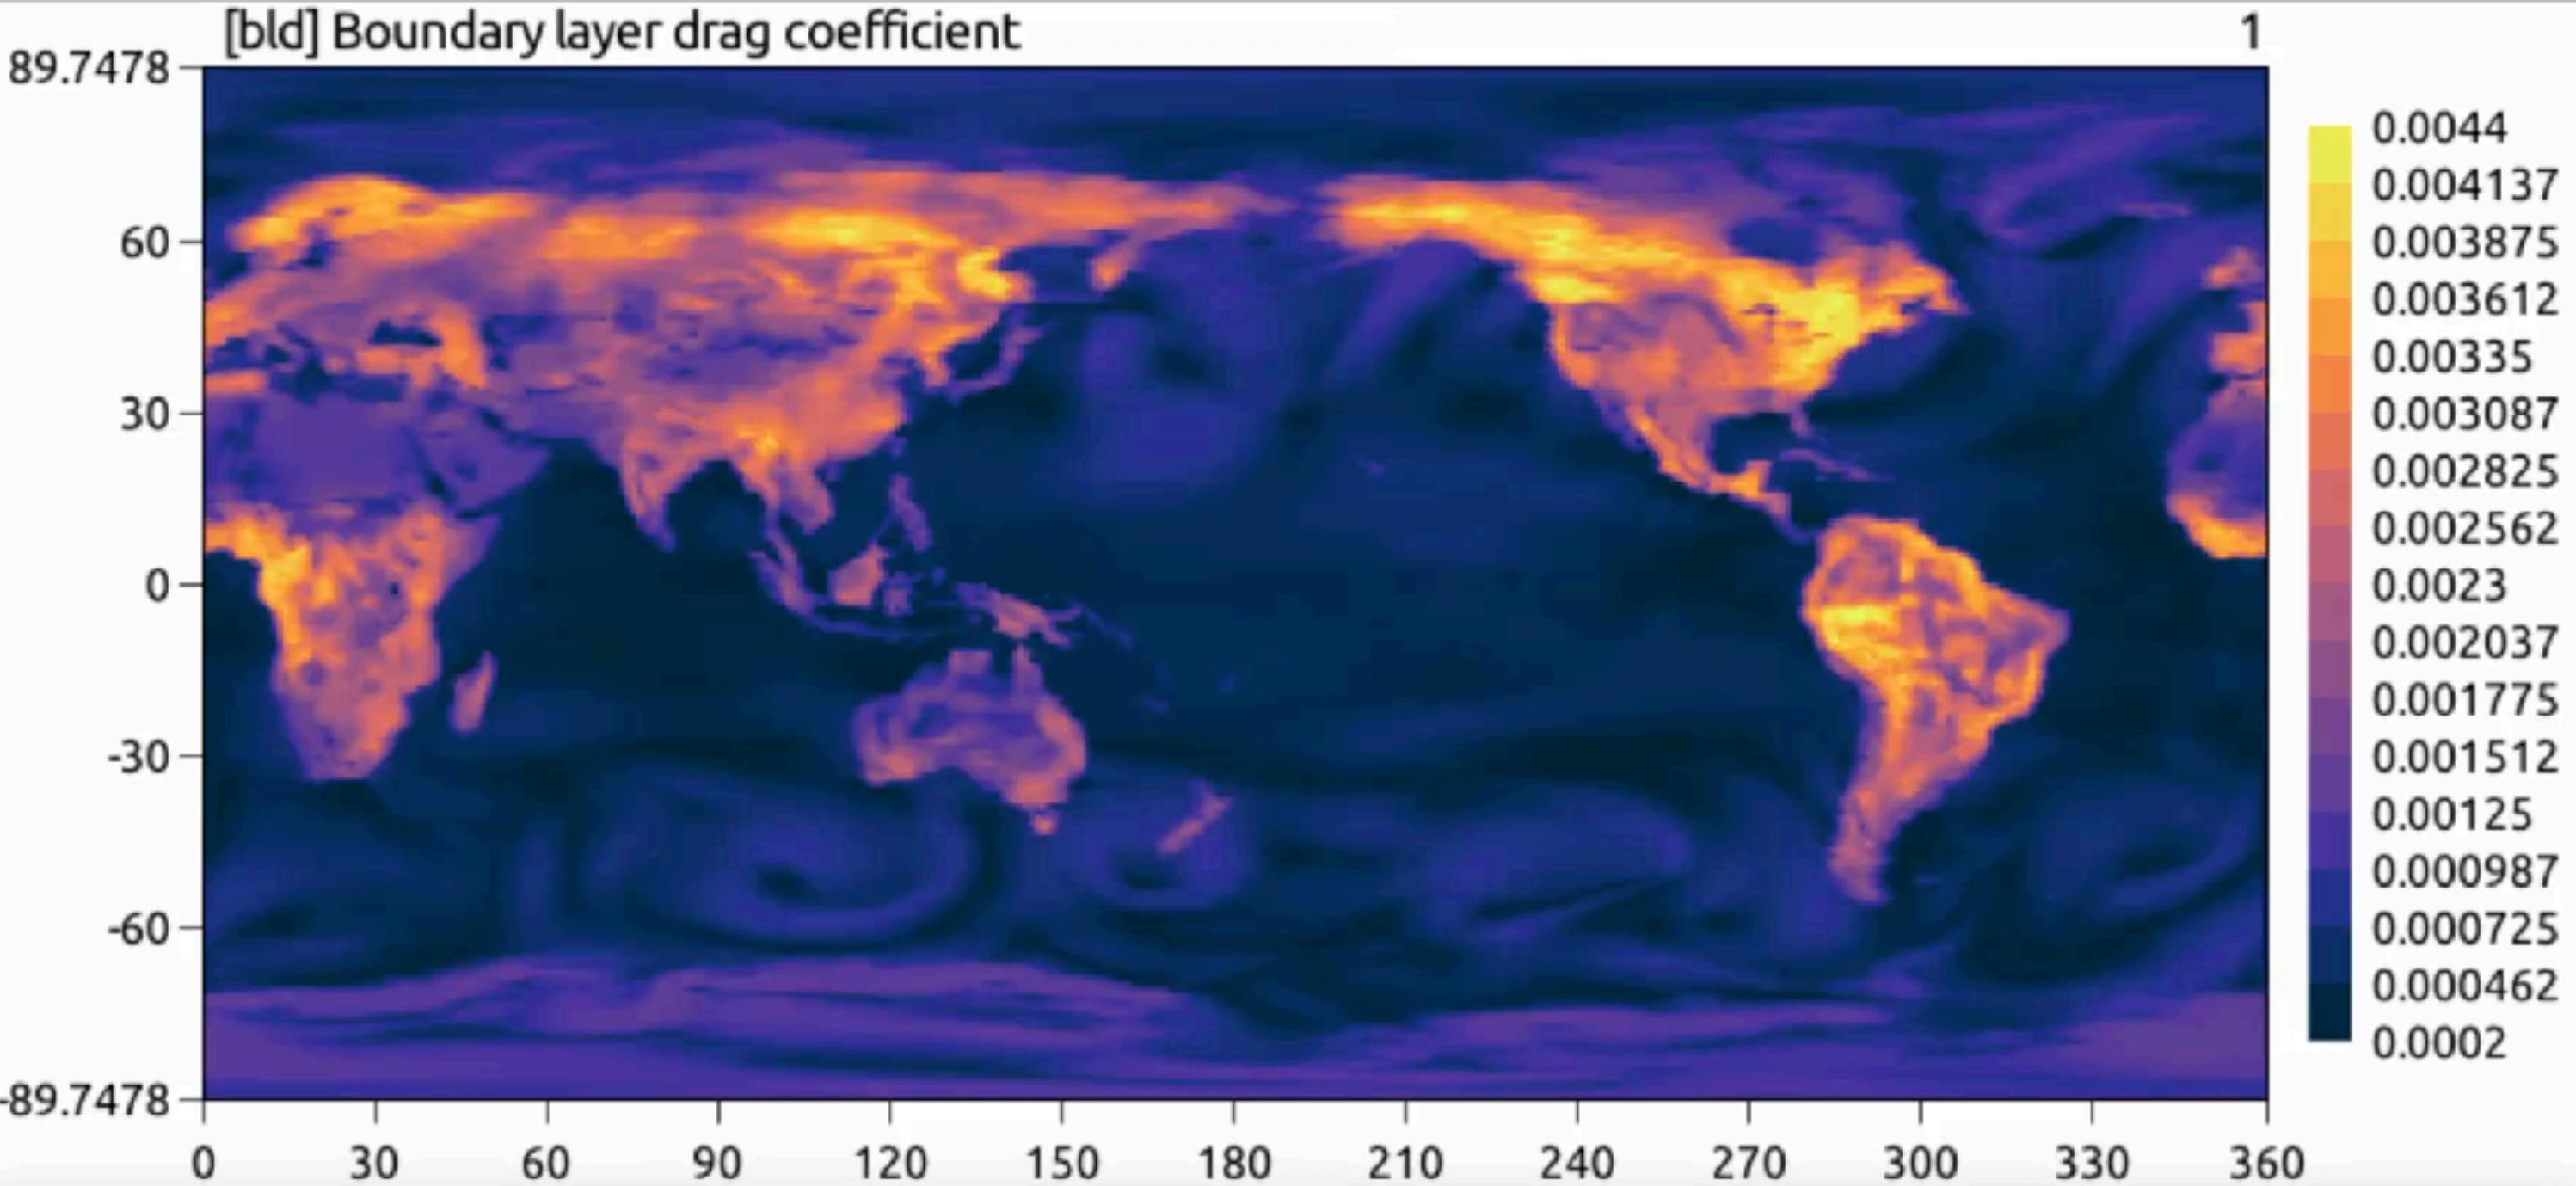
Inference done on every time step in every cell

0.5m

1cm

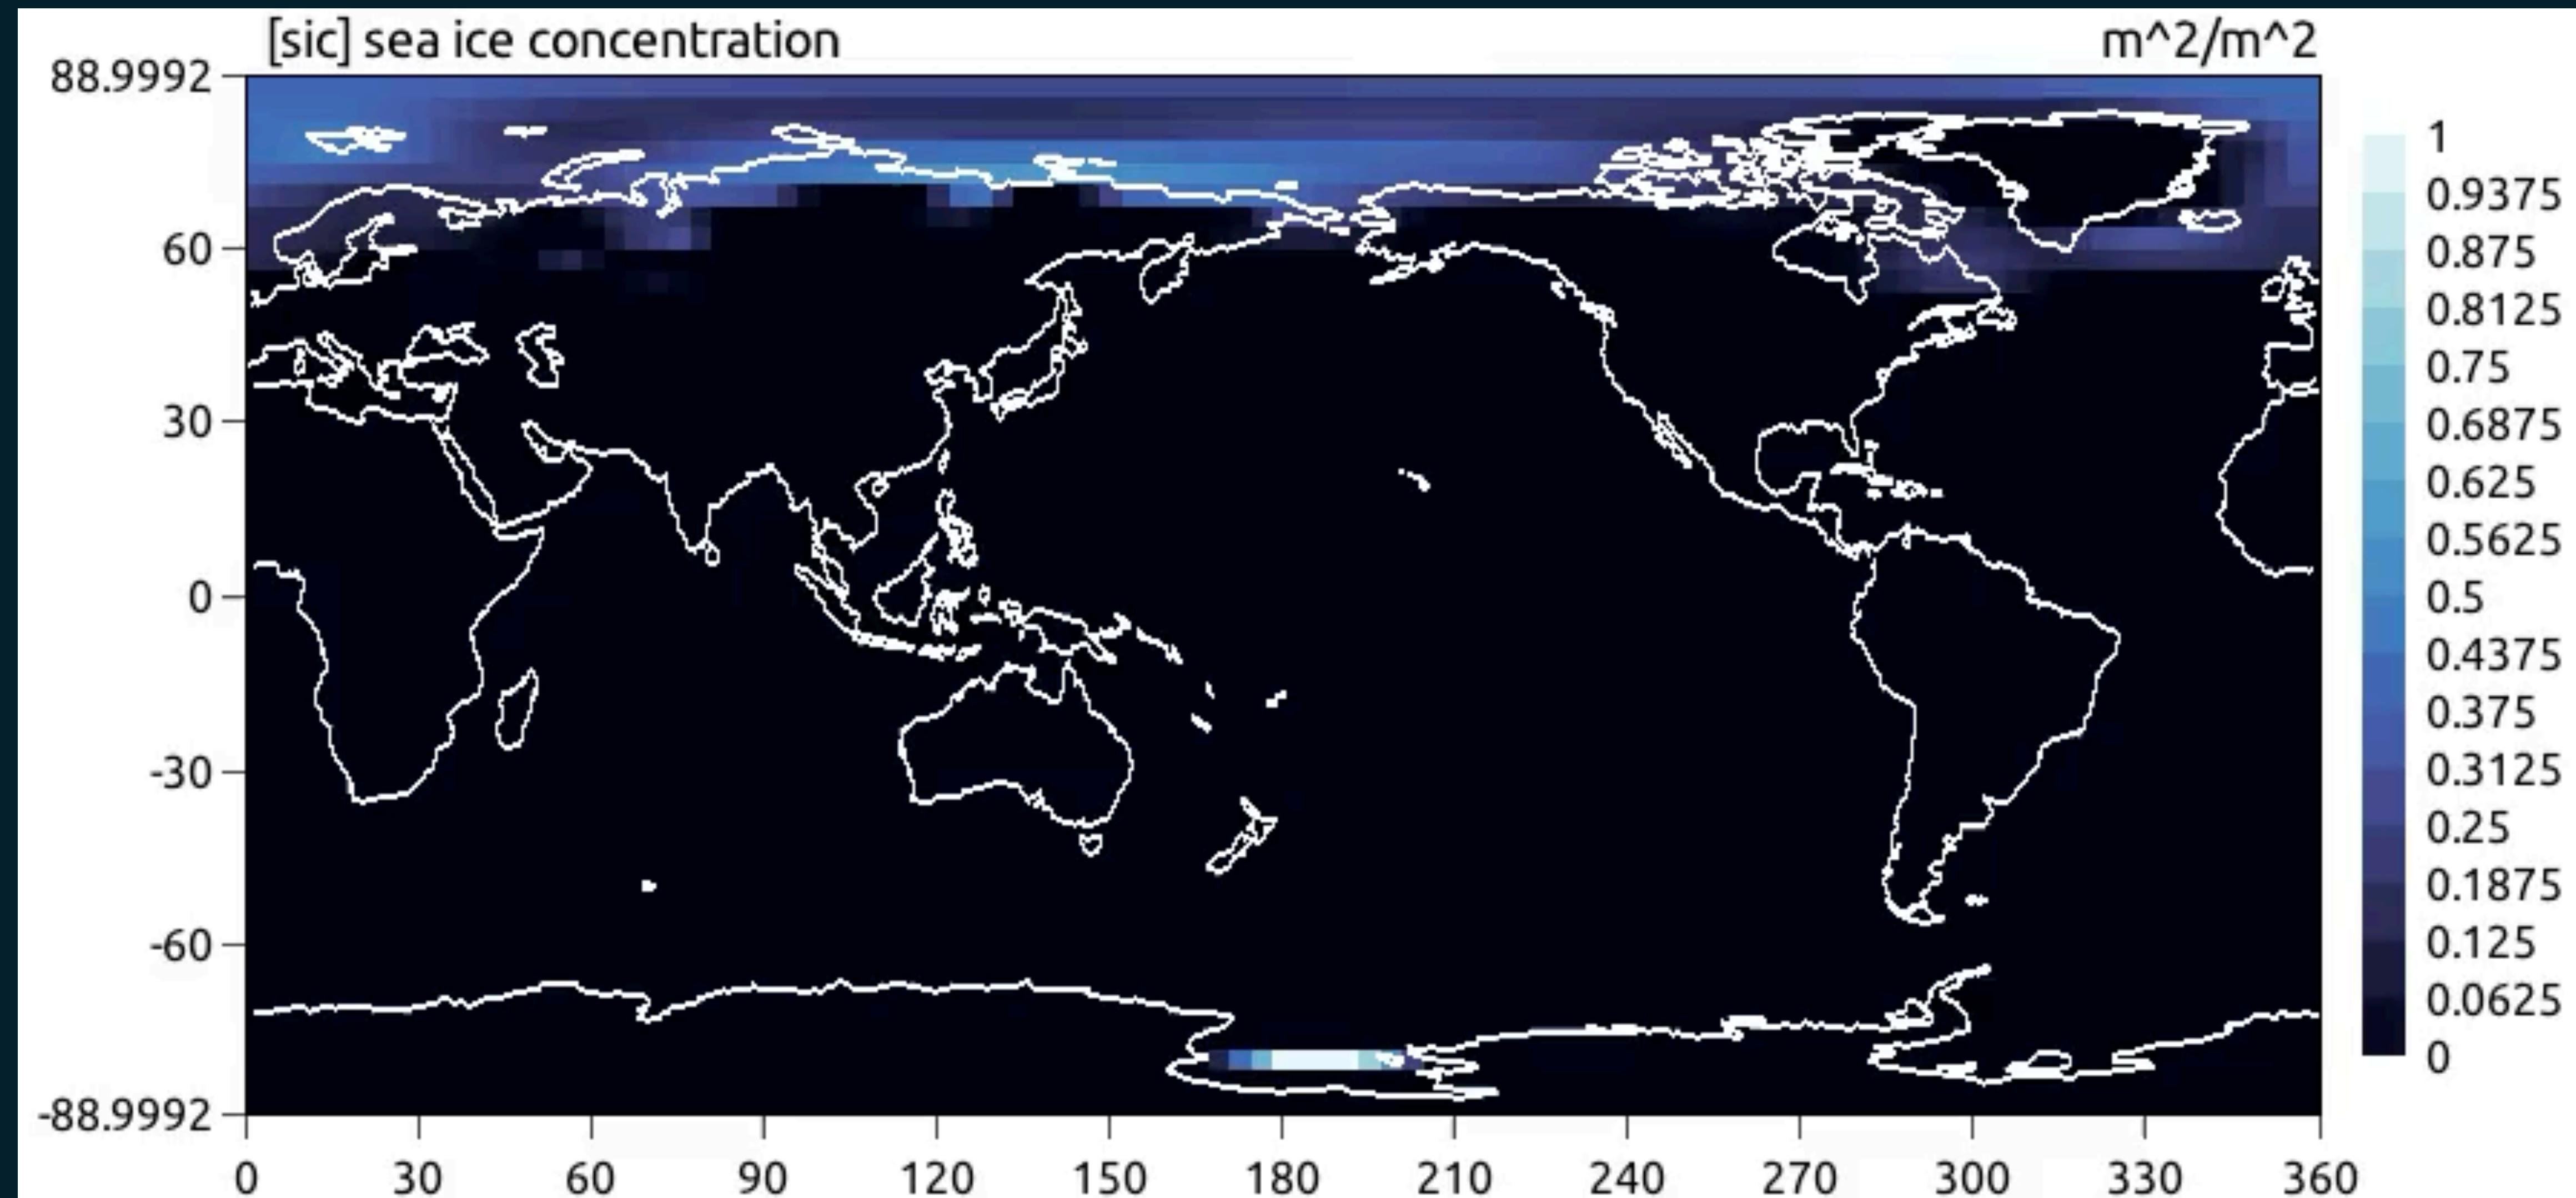
Constant over land / ocean

with Greg Munday



Speedy...Climate?

- Slab ocean model (zero velocity)
- Thermodynamic sea ice model
- Land bucket model (temperature, humidity, snow)



SpeedyWeather on the GPU

```
## OceanSeaIceAlbedo
export OceanSeaIceAlbedo

"""Albedo that scales linearly between ocean and ice albedo depending on sea ice concentration.
Fields are $(TYPEDFIELDS)"""
@kwdef struct OceanSeaIceAlbedo{NF} <: AbstractAlbedo
    "[OPTION] Albedo over open ocean [1]"
    albedo_ocean::NF = 0.06

    "[OPTION] Albedo over sea ice at concentration=1 [1]"
    albedo_ice::NF = 0.6
end

Adapt.@adapt_structure OceanSeaIceAlbedo
OceanSeaIceAlbedo(SG::SpectralGrid; kwargs...) = OceanSeaIceAlbedo{SG.NF}(;kwargs...)
initialize!(::OceanSeaIceAlbedo, ::PrimitiveEquation) = nothing
@propagate_inbounds function albedo!(ij, diagn, progn, albedo::OceanSeaIceAlbedo, model)
    (; albedo_ocean, albedo_ice) = albedo
    κ = haskey(progn.ocean, :sea_ice_concentration) ? progn.ocean.sea_ice_concentration[ij] : zero(eltype(diagn.albedo))

    # set ocean albedo linearly between ocean and ice depending on sea ice concentration
    diagn.albedo[ij] = albedo_ocean + κ * (albedo_ice - albedo_ocean)
end
```

1 line of “GPU code”



In contrast to: Gordon Bell Prize for climate modelling 2025

Not our idea of “separation of concerns”?

5.2 Separation of Concerns: From Fortran to GPU

To achieve performance and scalability on today's heterogeneous systems, ICON's code base has evolved to support a wide range of configurations: from pragmas enabling OpenMP and OpenACC parallelization, to Intel, Cray, GCC, and NEC vector annotations

dynamical core, the most compute-intensive part of the most expensive atmosphere component, contains 2728 non-empty Fortran source lines of code, of which less than 50 % are actually describing the computation. The remaining lines of code are optimizations

The diagram shows a snippet of Fortran code with various pragmas and annotations. Arrows point from specific annotations to their descriptions:

- An arrow points from the first two OpenACC pragmas to the text "OpenACC GPU pragmas".
- An arrow points from the `#ifdef __LOOP_EXCHANGE` block to the text "CPU vectorization pragma".
- An arrow points from the `!DIR$ IVDEP` directive to the text "duplicate loop header code".
- An arrow points from the NEC-specific pragma `!$NEC outerloop_unroll(4)` to the text "NEC-specific pragma".
- An arrow points from the `DO jk = 1, nlev` and `DO jc = i_startidx, i_endidx` loops to the text "duplicate loop header code".
- An arrow points from the OpenMP pragmas `ENDDO` and `!$ACC END PARALLEL` to the text "some lines down, more OpenMP pragmas".

```
!$ACC PARALLEL DEFAULT(PRESENT) ASYNC(1)
 !$ACC LOOP GANG VECTOR TILE(32, 4)
 #ifdef __LOOP_EXCHANGE
   DO jc = i_startidx, i_endidx
 !DIR$ IVDEP
   DO jk = 1, nlev
     z_ekinh(jk,jc,jb) =  &
   #else
     !$NEC outerloop_unroll(4)
       DO jk = 1, nlev
         DO jc = i_startidx, i_endidx
           z_ekinh(jc,jk,jb) =  &
   #endif
     p_int%e_bln_c_s(jc,1,jb)*z_kin_hor_e(ieidx(jc,jb,1),jk,ieblk(jc,jb,1)) + &
     p_int%e_bln_c_s(jc,2,jb)*z_kin_hor_e(ieidx(jc,jb,2),jk,ieblk(jc,jb,2)) + &
     p_int%e_bln_c_s(jc,3,jb)*z_kin_hor_e(ieidx(jc,jb,3),jk,ieblk(jc,jb,3))

   ENDDO
 ENDDO
 !$ACC END PARALLEL
```

Intertwined lines for scientist and “performance engineer”

TravellingSailorProblem.jl

Fly Christmas presents (=particles) to children (=destinations) around the globe following the wind. More points for longer distances particles fly and negative points if they miss it.

- SpeedyWeather+GLMakie+GeoMakie
- Challenge evaluated with GitHub actions
- Used for PhD-fresher teaching
- Optimisation problem
- Wither interactive Makie visualisation

Live demo on Friday!

with Anshul

