

The Bitter Lesson for the Bitter Lesson:

Where Human
Mathematicians and
Engineers go in the
Age of AI

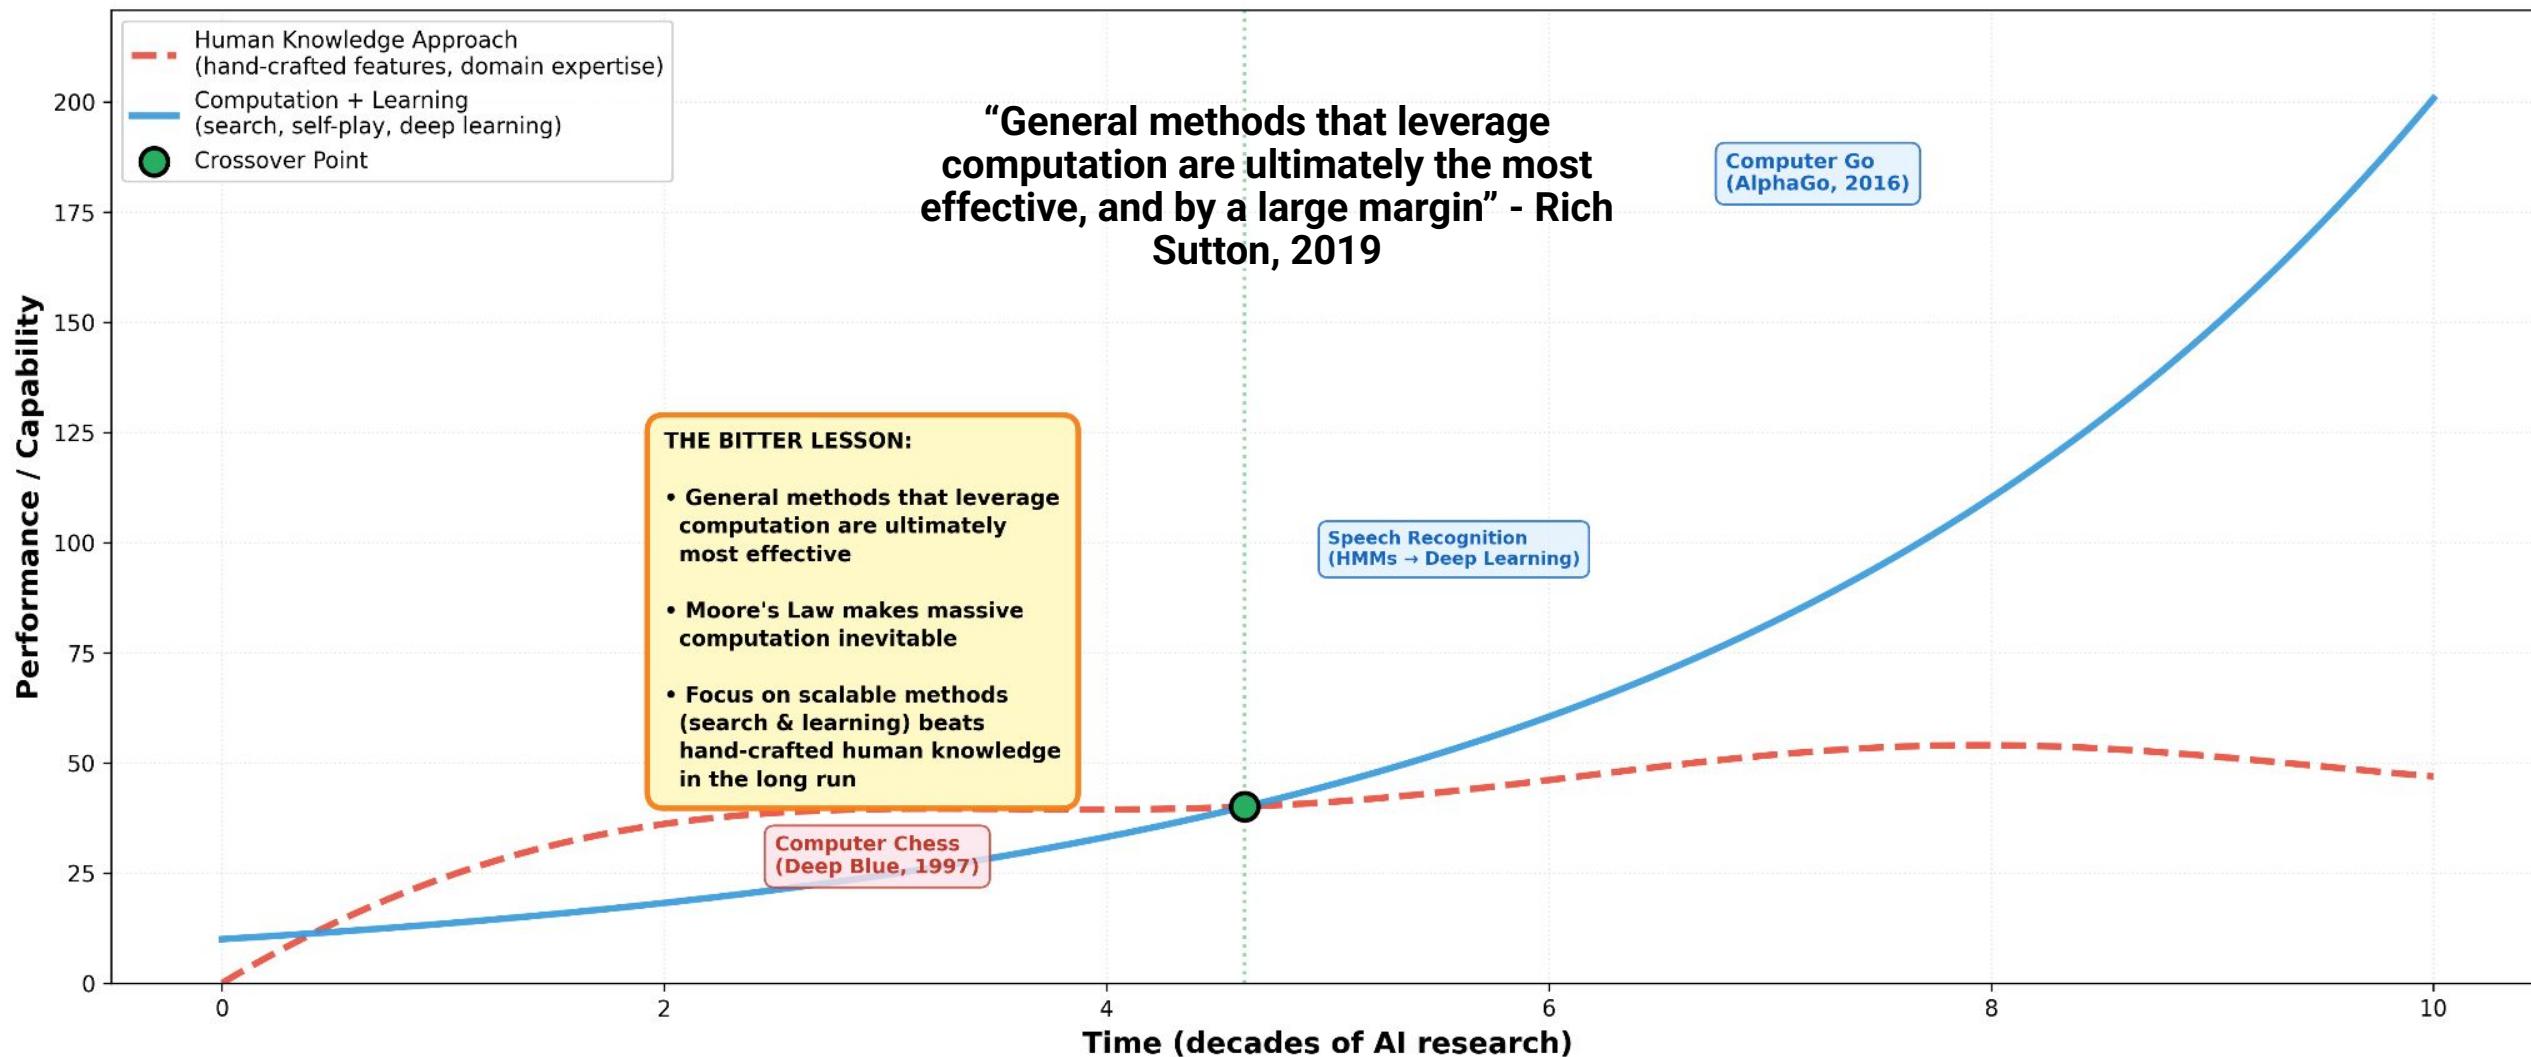
Chris Rackauckas

VP of Modeling and Simulation,
JuliaHub

Research Affiliate, Co-PI of Julia Lab,
Massachusetts Institute of Technology,
CSAIL

Director of Scientific Research,
Pumas-AI

The Bitter Lesson of AI (Rich Sutton, 2019)



Looking more closely at the “AI successes” may suggest the bitter lesson needs to learn a lesson.

Potential Problem: Out of Sample Data (Generalization Issues)



Image credit: Jordan Nelson via ViralHog

Traffic light, or moon?

**“It’s only an edge case
that happens on
planets with moons”**

Poten

sample Data (Generalization Issues)

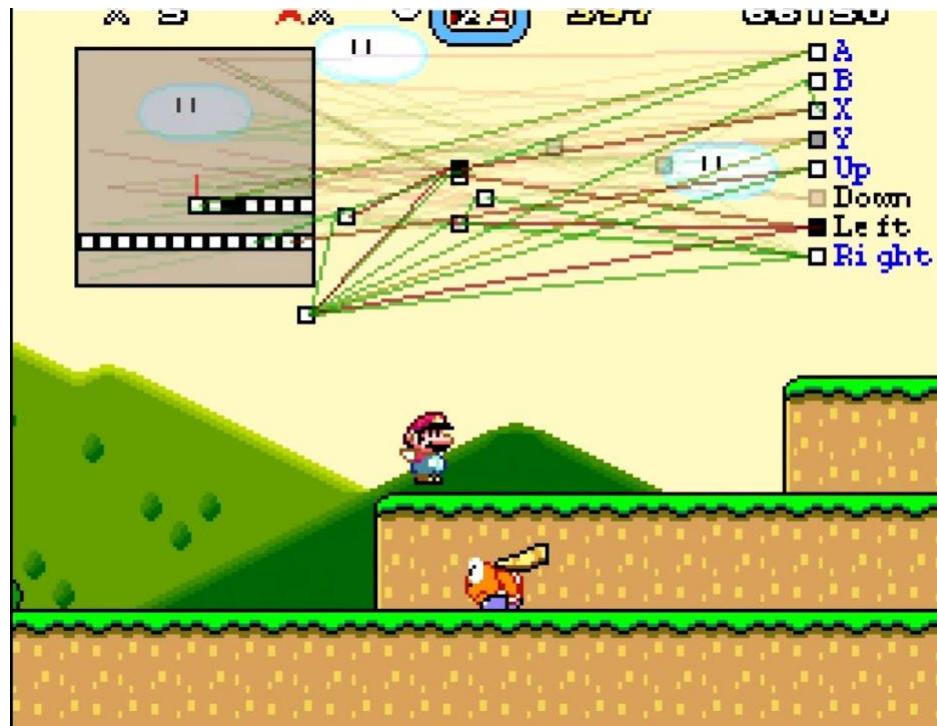


**Traffic lights,
Or truck carrying
traffic lights?**

ML/AI has issues with generalization beyond the dataset.

Anything beyond the training dataset needs some form of human intervention.

Reinforcement Learning For Game AI: Outperforming All Humans?



Game AI
(OpenAI, DeepMind, ...)

**If data is the issue,
Why not make the data
infinite?**

**Won't machine learning
be perfect there?**

=

Reinforcement Learning

Reinforcement Learning Demonstrates a Overfitting Issue

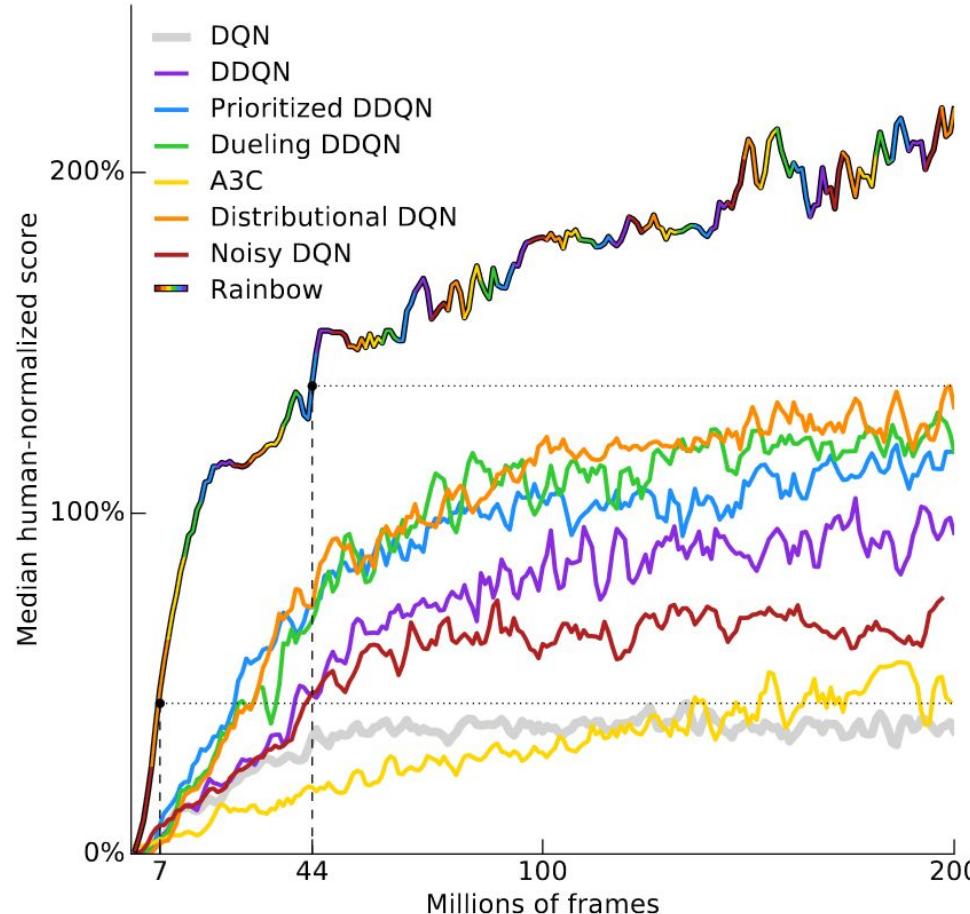
- A coworker is teaching an agent to navigate a room. The episode terminates if the agent walks out of bounds. He didn't add any penalty if the episode terminates this way. The final policy learned to be suicidal, because negative reward was plentiful, positive reward was too hard to achieve, and a quick death ending in 0 reward was preferable to a long life that risked negative reward.
 - A friend is training a simulated robot arm to reach towards a point above a table. It turns out the point was defined *with respect to the table*, and the table wasn't anchored to anything. The policy learned to slam the table really hard, making the table fall over, which moved the target point too. The target point *just so happened* to fall next to the end of the arm.
- A researcher gives a talk about using RL to train a simulated robot hand to pick up a hammer and hammer in a nail. Initially, the reward was defined by how far the nail was pushed into the hole. Instead of picking up the hammer, the robot used its own limbs to punch the nail in. So, they added a reward term to encourage picking up the hammer, and retrained the policy. They got the policy to pick up the hammer...but then it threw the hammer at the nail instead of actually using it.

ML results are good for humor.

Deep Reinforcement Learning Doesn't Work Yet

<https://www.alexirpan.com/2018/02/14/rl-hard.html>

Reinforcement Learning Demonstrates a Major Training Data Issue



Other con of reinforcement learning:

Requires an accurate simulator so it can generate hundreds of millions of training points!

The y-axis is “median human-normalized score”. This is computed by training 57 DQNs, one for each Atari game, normalizing the score of each agent such that human performance is 100%, then plotting the median performance across the 57 games. RainbowDQN passes the 100% threshold at about *18 million* frames. This corresponds to about 83 hours of play experience, plus however long it takes to train the model. A lot of time, for an Atari game that most humans pick up within a few minutes.

Deep Reinforcement Learning Doesn't Work Yet

<https://www.alexirpan.com/2018/02/14/rl-hard.html>

Reinforcement Learning Demonstrates a Major Cost Issue



AlphaStar: Closed world reinforcement learning agent takes \$1 billion of research and training! (Still needed human intervention, beat at later AI events by a university team)

Cost:

**AlphaGo – learns to play Go
\$35 million**

**GPT-3: Natural Language Processing
\$10 million**

**AlphaStar: Starcraft Bot
\$12 million for the final
>\$1 billion in research tuning**

GPT-5: \$500 million to train, total human cost in tens of billions

Claude Opus: “Billions of dollars range”, total cost high tens of billions

THE COST OF AI

Billions spent chasing artificial intelligence

Pennsylvania's dormant Three Mile Island nuclear plant would be brought back to life to feed the voracious energy needs of Microsoft under an unprecedented deal in which the tech giant would buy 100 percent of its power for 20 years.

— *Washington Post, Sept 2024*

Training a single AI model can emit as much carbon as five cars in their lifetimes. The training of GPT-3 consumed millions of dollars worth of energy.

— *Team-GPT Analysis*

OpenAI CEO Sam Altman vaguely pegged the training cost at 'more than \$100 million' for GPT-4. By 2025 we may have a \$10 billion model.

— *Anthropic CEO Dario Amodei, 2024*

Microsoft spent \$34.9 billion in just one quarter while Meta plans to spend up to \$72 billion this year on AI infrastructure.

— *CNBC, Oct 2025*

The Stargate Project intends to invest \$500 billion over the next four years building new AI infrastructure for OpenAI in the United States.

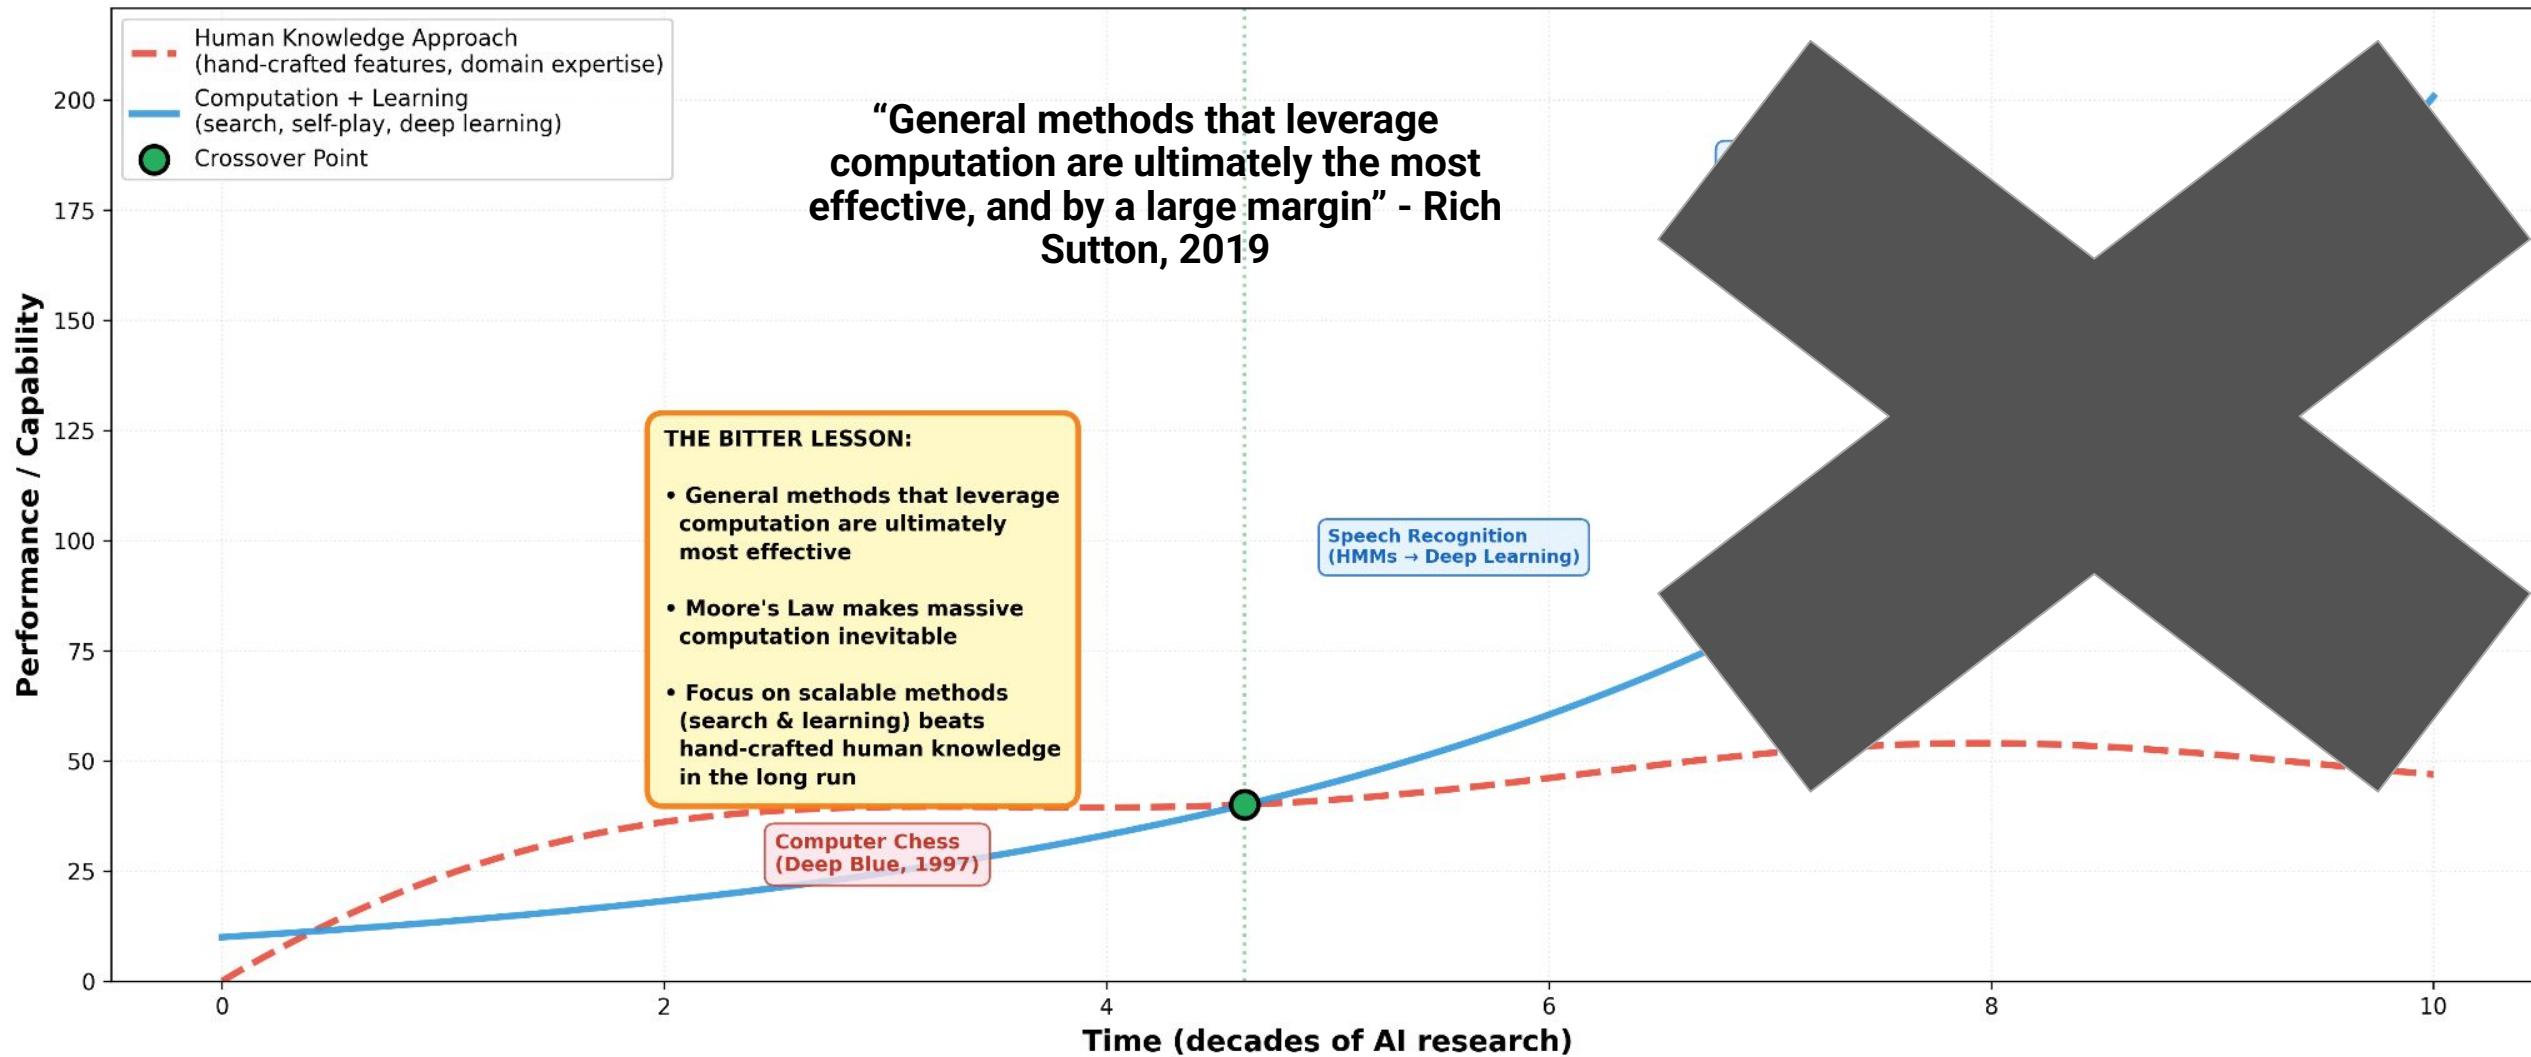
— *OpenAI, Jan 2025*

A single ChatGPT query consumes almost ten times the electricity of a standard Google search. Data centers could consume 10% of all US electricity by 2030.

— *Goldman Sachs & IEA*

Tech giants are reshaping global power infrastructure to fuel the AI race

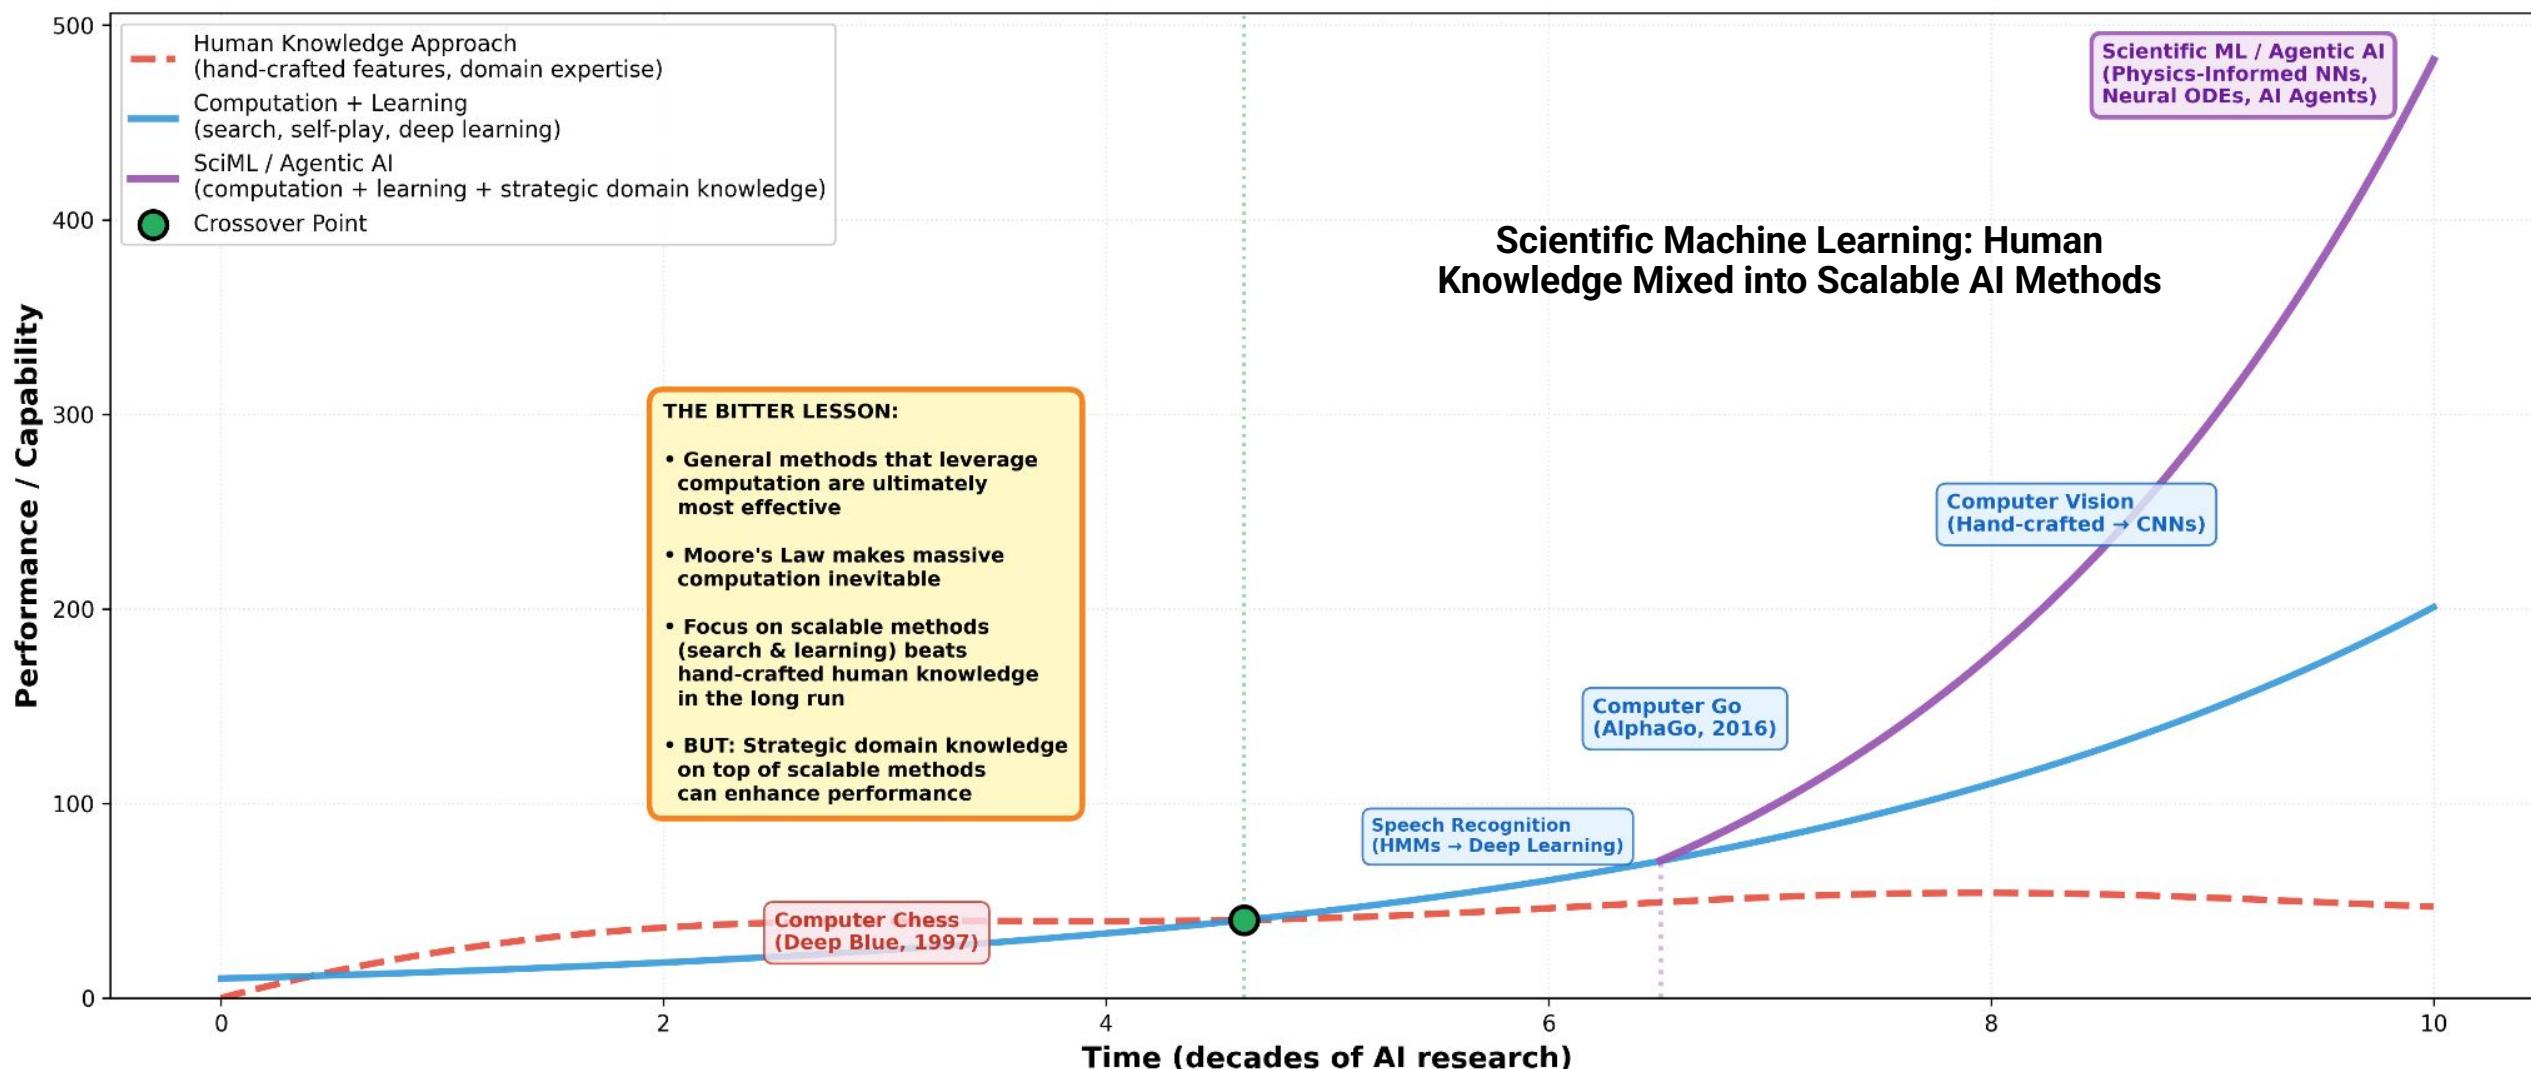
The Bitter Lesson of AI (Rich Sutton, 2019)



ML/AI scaling costs are growing exponentially

Is there a way around these issues?

The Bitter Lesson of AI + The SciML/Agentic AI Synthesis (Rich Sutton, 2019 + Modern Extensions)



Two Solutions to The Bitter Lesson

Complementary Approaches to Overcoming AI's Limitations

Scientific Machine Learning

Adding Scientific Knowledge to Scalable AI Methods



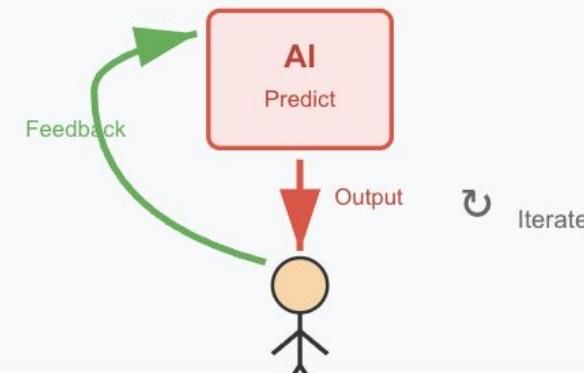
$$\begin{aligned}\partial u / \partial t &= f(u, t) \\ F &= ma \\ \nabla^2 \phi &= \rho \\ E &= mc^2\end{aligned}$$

Key Concept:

Embed physical laws, conservation principles, and domain knowledge directly into neural network architectures and training processes

Agentic AI

Agentic Loops with Humans in the Loop

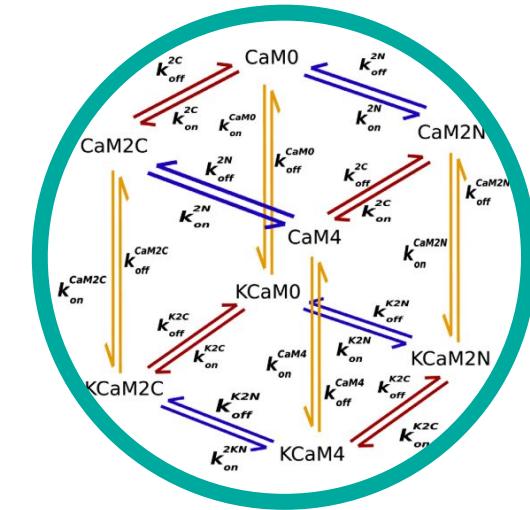
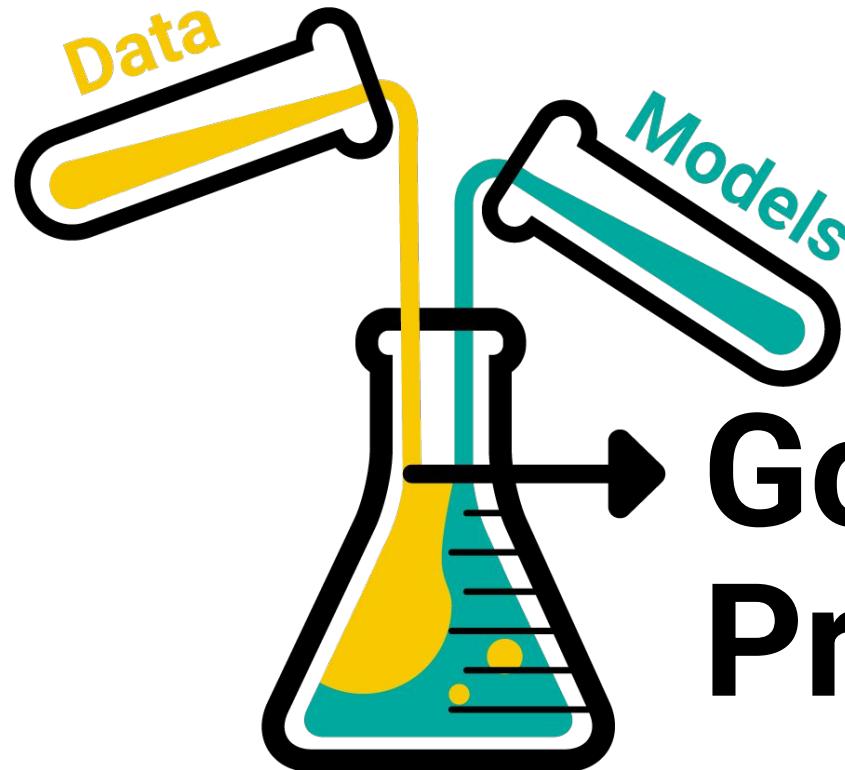
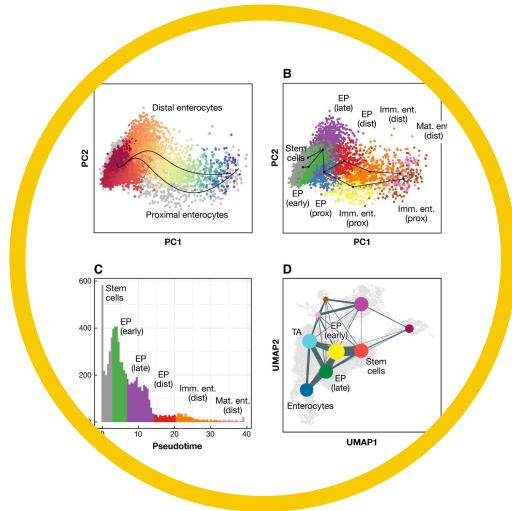


Key Concept:

Use AI as a tool in an iterative workflow where human expertise validates, corrects, and guides predictions in a continuous improvement loop

Scientific Machine Learning is model-based data-efficient machine learning

How do we simultaneously use both sources of knowledge?



Good
Predictions

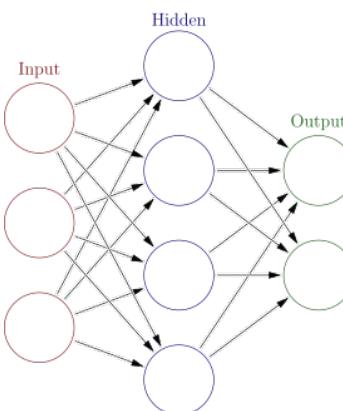
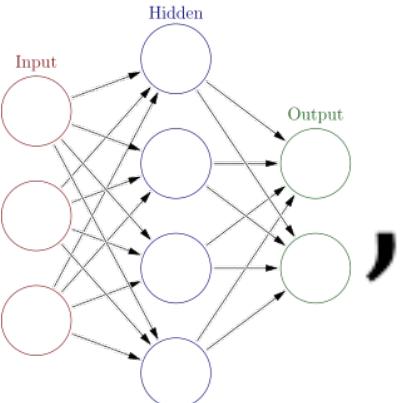
Core Point: the more structure that is encoded into the model, the less there is to learn and the more data efficient the algorithm is!

Universal (Approximator) Differential Equations

$$u' = f(u, , t)$$

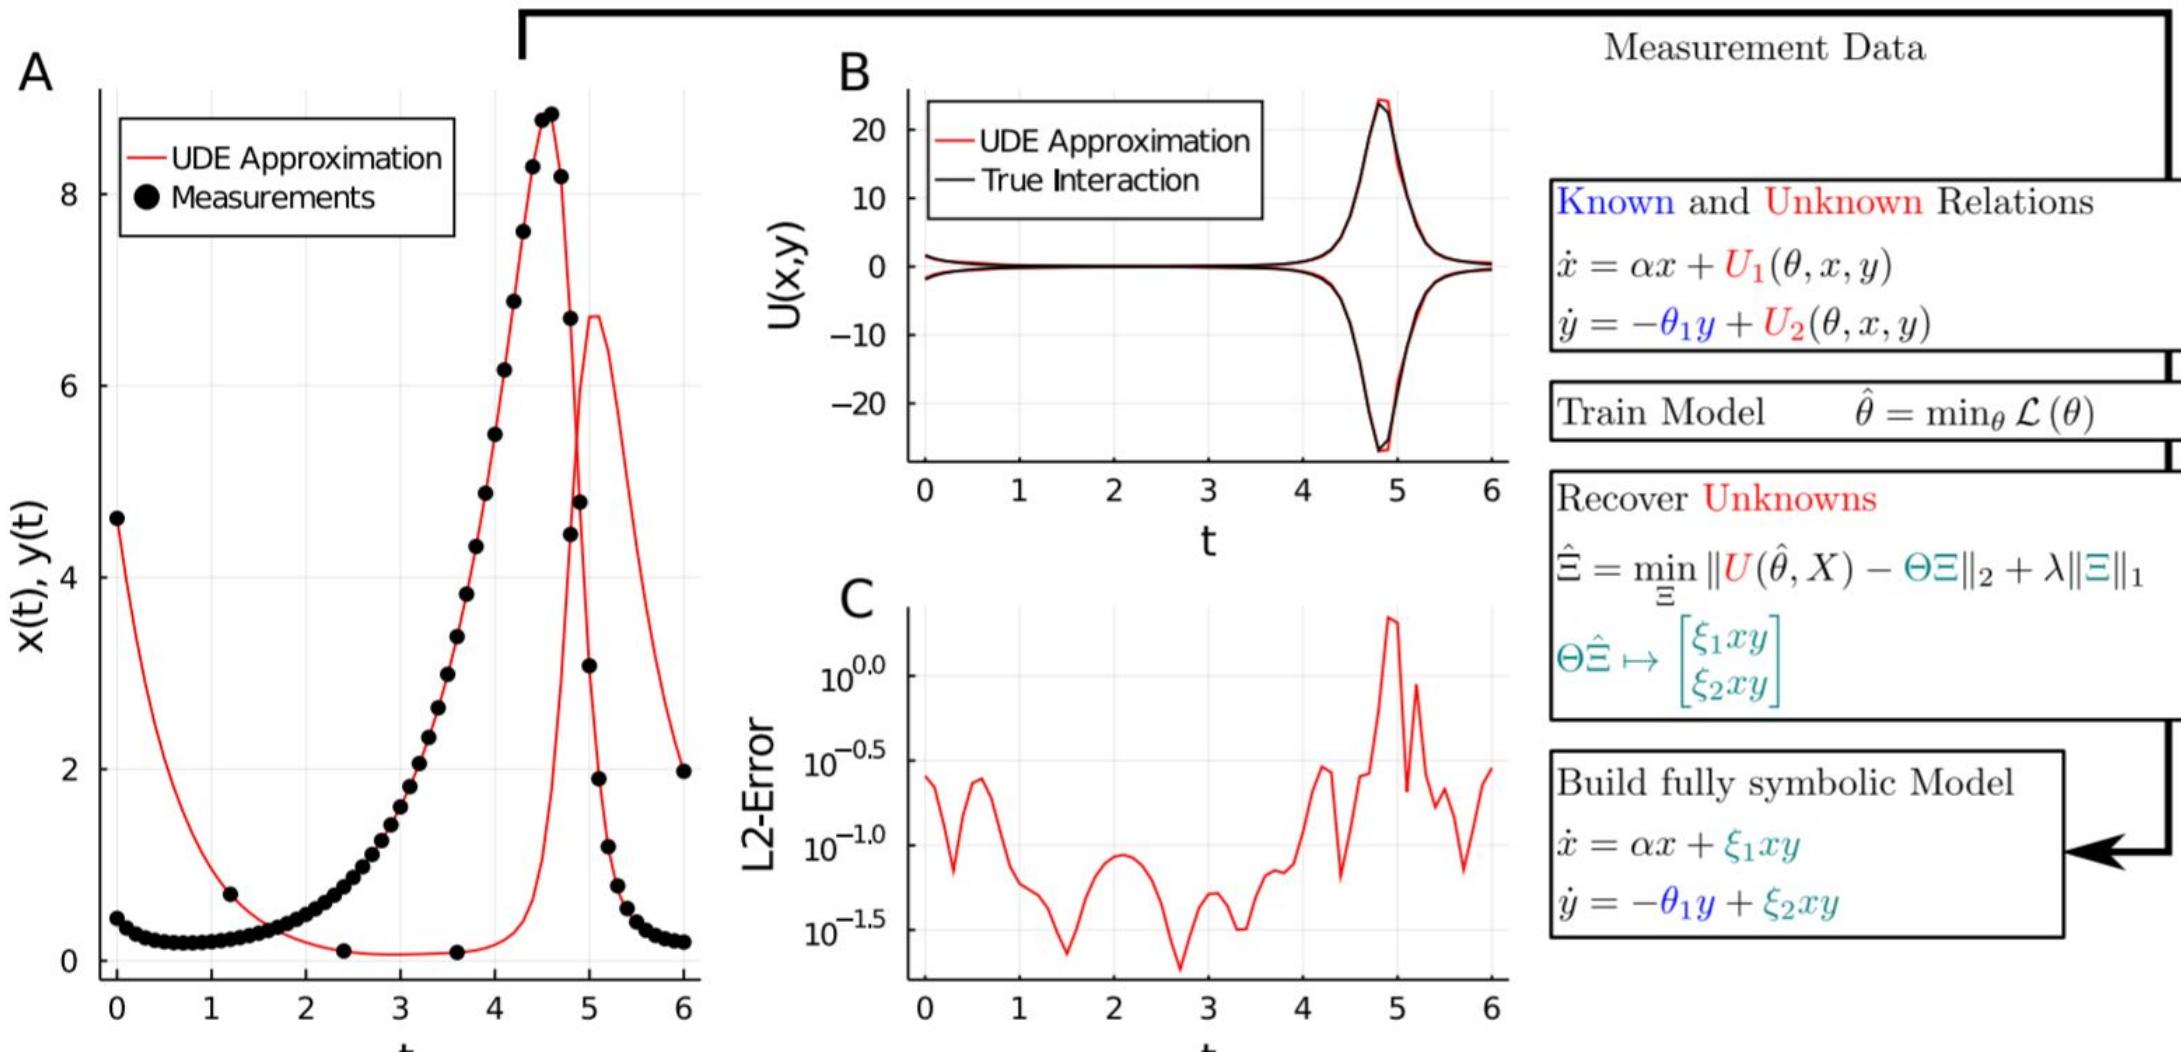


$$\begin{aligned}x' &= \alpha x + \\y' &= -\beta y +\end{aligned}$$



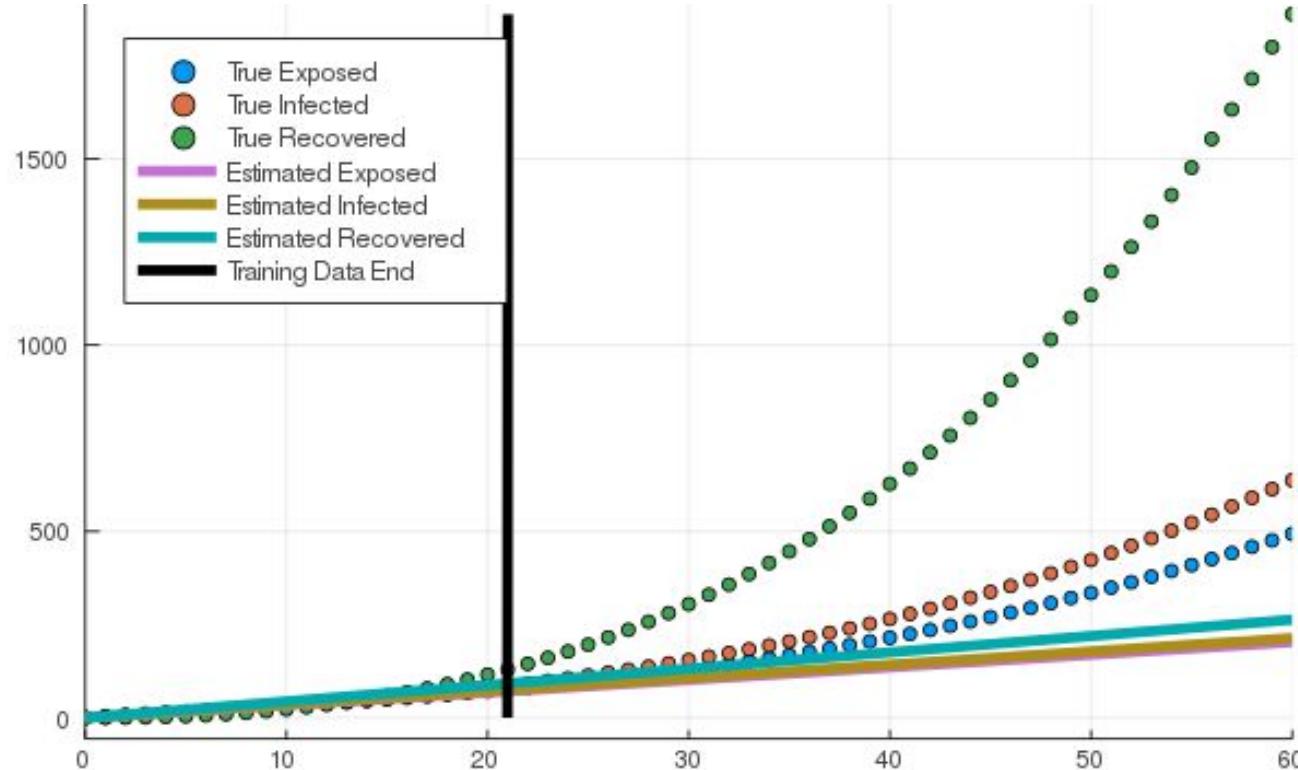
Rackauckas, Christopher, Yingbo Ma, Julius Martensen, Collin Warner, Kirill Zubov, Rohit Supekar, Dominic Skinner, Ali Ramadhan, and Alan Edelman. "Universal differential equations for scientific machine learning." *arXiv preprint arXiv:2001.04385* (2020).

Universal (Approximator) Differential Equations



Let's dive in a bit!

Standard Machine Learning: Learn the whole model



$u' = \text{NN}(u)$ trained on 21 days of data

Can fit, but not enough information
to accurately extrapolate

**Does not have the correct
asymptotic behavior**

More examples of this issue:

Ridderbusch et al. "Learning ODE Models with Qualitative Structure Using Gaussian Processes."

Universal ODE

$$S' = -\frac{\beta_0 S F}{N} - \text{Replace Unknown Portion} - \mu S,$$

$$E' = \frac{\beta_0 S F}{N} + \text{Replace Unknown Portion} - (\sigma + \mu) E,$$

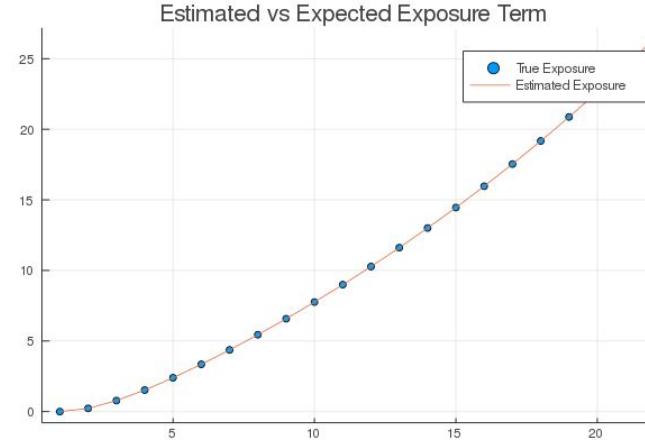
$$I' = \sigma E - (\gamma + \mu) I,$$

$$R' = \gamma I - \mu R,$$

$$N' = -\mu N,$$

$$D' = d \gamma I - \lambda D,$$

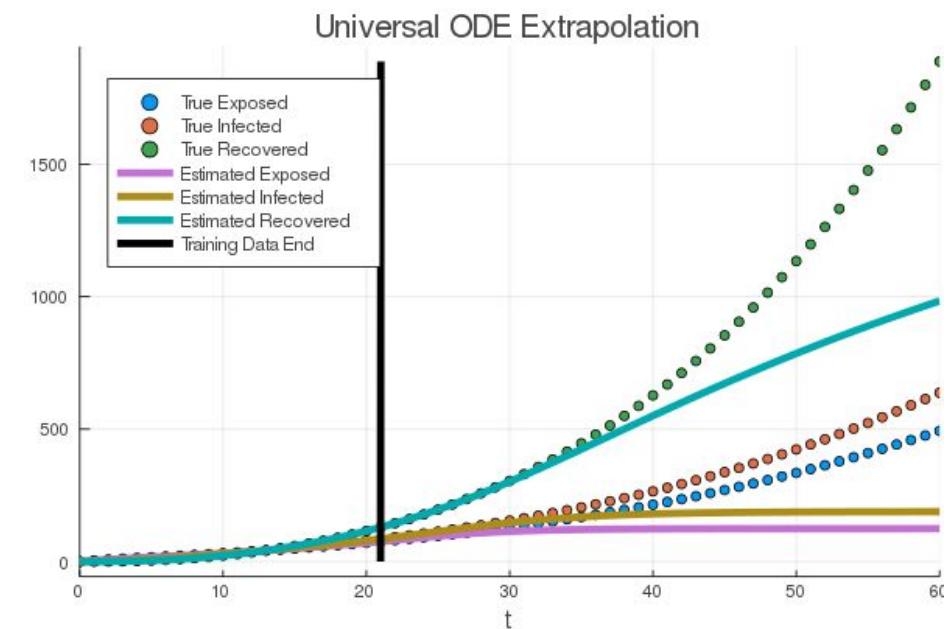
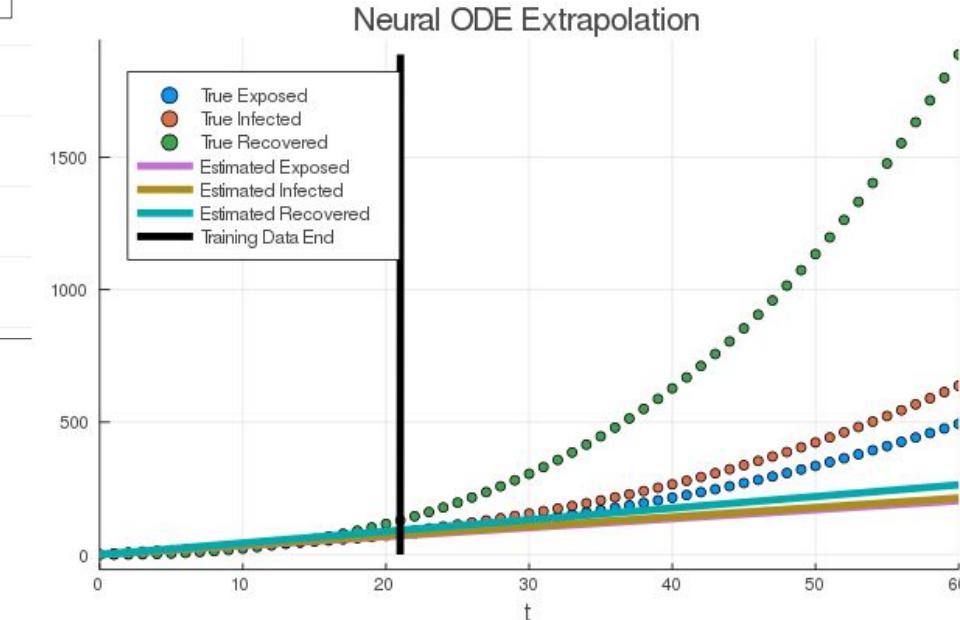
$$C' = \sigma E,$$



Exposure: Unknown

Infection rates: known
From disease quantities

Percentage of cases
known to be severe,
can be estimated



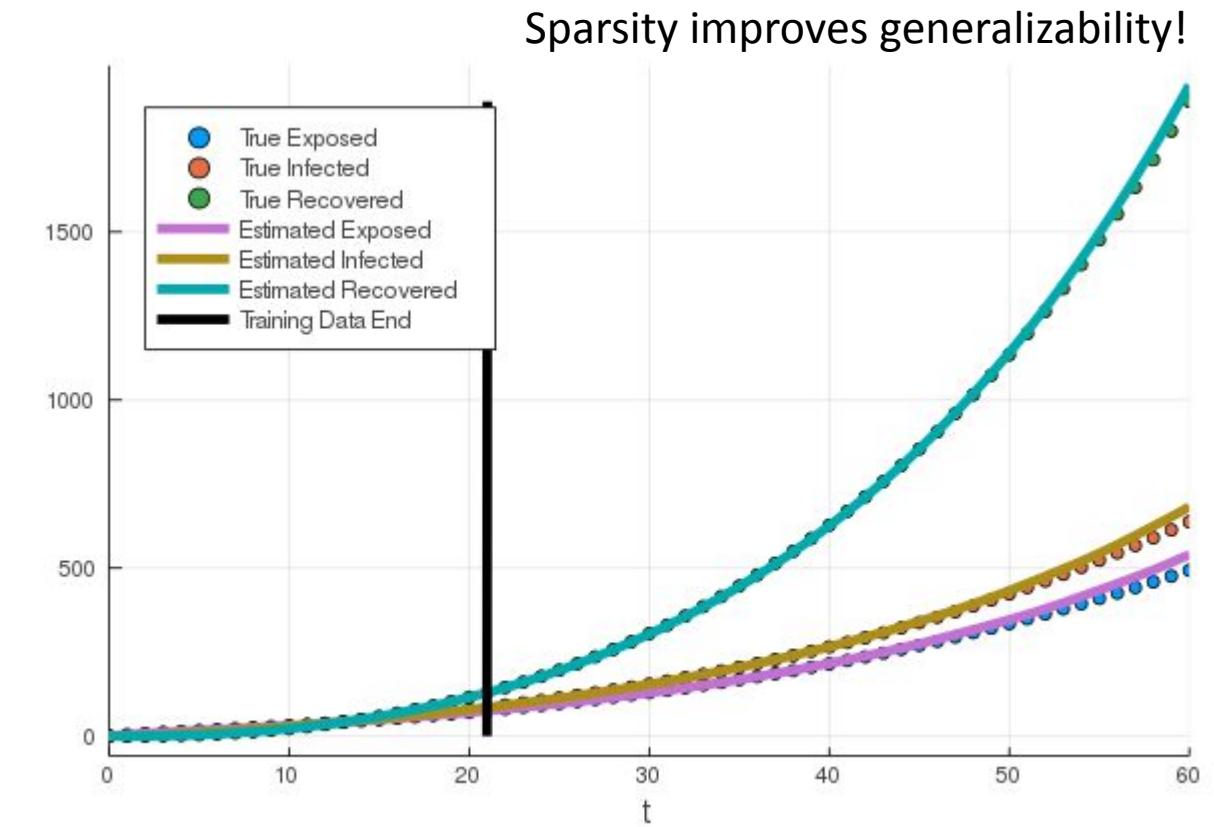
Universal ODE -> Internal Sparse Regression

Sparse Identification on only the missing term:

$$I * 0.10234428543435758 + S/N * I * 0.11371750552005416 + (S/N)^2 * I * 0.12635459799855597$$

$$\begin{aligned} S' &= -\frac{\beta_0 SF}{N} - \mu S, \\ E' &= \frac{\beta_0 SF}{N} + (\sigma + \mu)E, \\ I' &= \sigma E - (\gamma + \mu)I, \\ R' &= \gamma I - \mu R, \\ N' &= -\mu N, \\ D' &= d \gamma I - \lambda D, \quad \text{and} \\ C' &= \sigma E, \end{aligned}$$

Replace Unknown Portion



Accurate Model Extrapolation Mixing in Physical Knowledge

Upon denoting $\mathbf{x} = (\phi, \chi, p, e)$, we propose the following family of UDEs to describe the two-body relativistic dynamics:

$$\dot{\phi} = \frac{(1 + e \cos(\chi))^2}{Mp^{3/2}} (1 + \mathcal{F}_1(\cos(\chi), p, e)), \quad (5a)$$

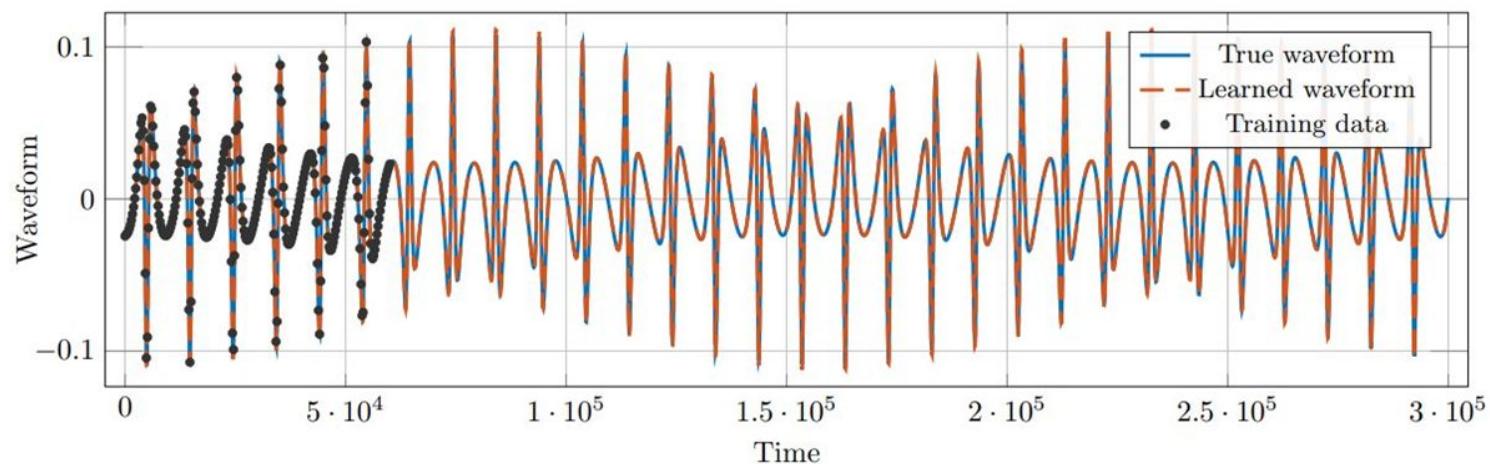
$$\dot{\chi} = \frac{(1 + e \cos(\chi))^2}{Mp^{3/2}} (1 + \mathcal{F}_2(\cos(\chi), p, e)), \quad (5b)$$

$$\dot{p} = \mathcal{F}_3(p, e), \quad (5c)$$

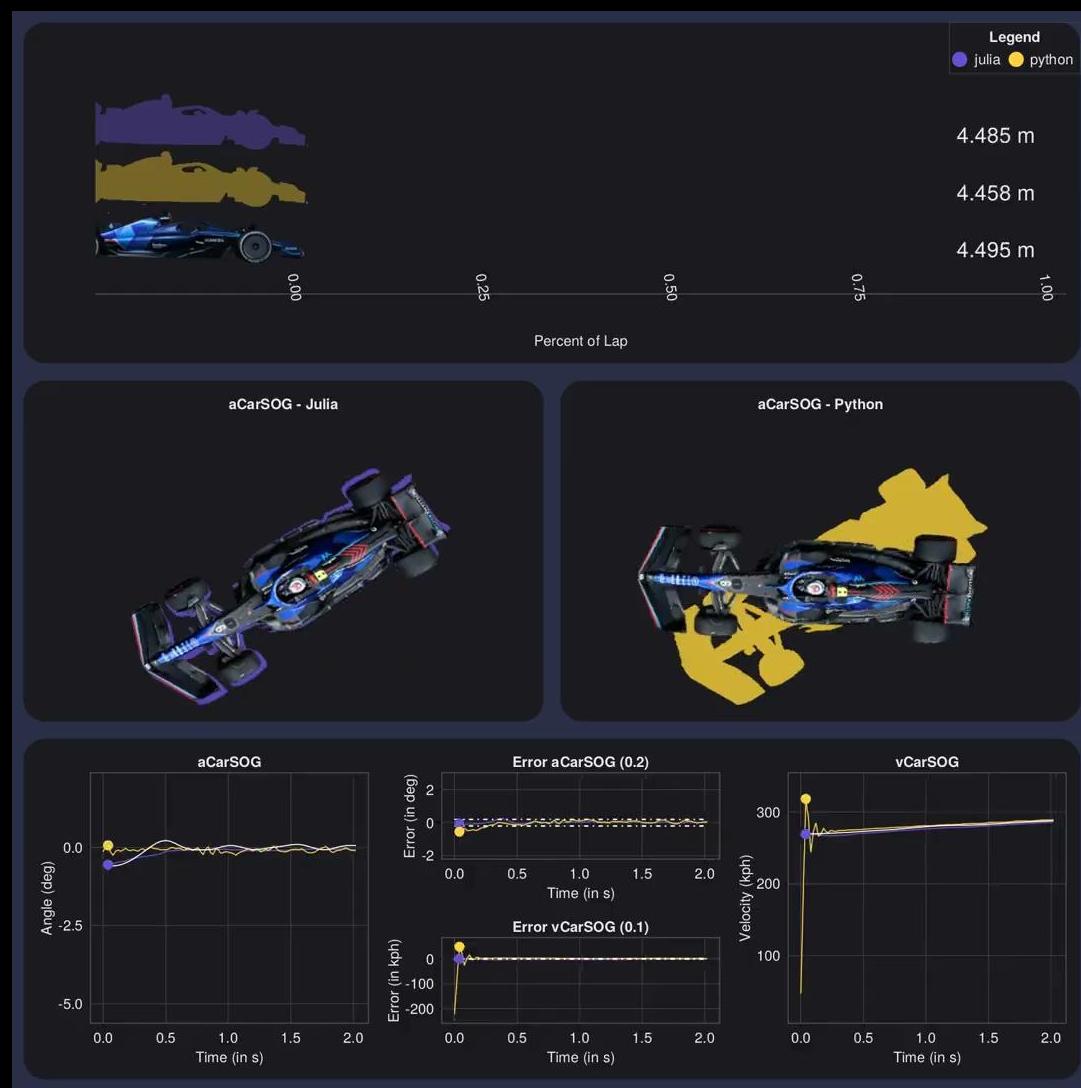
$$\dot{e} = \mathcal{F}_4(p, e), \quad (5d)$$

Automated discovery of geodesic equations from LIGO black hole data: run the code yourself!

https://github.com/Astroinformatics/ScientificMachineLearning/blob/main/neuralode_gw.ipynb

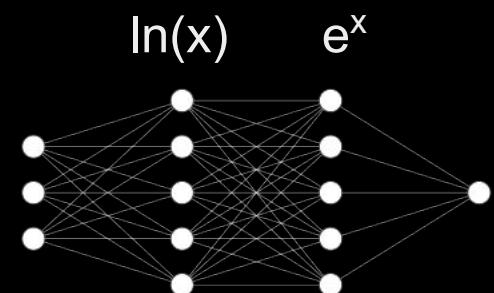


Keith, B., Khadse, A., & Field, S. E. (2021). Learning orbital dynamics of binary black hole systems from gravitational wave measurements. *Physical Review Research*, 3(4), 043101.



Physically-Informed Machine Learning

$$\dot{\beta} \approx \frac{ay}{u} - \frac{ax}{u} - r$$



Using knowledge of the physical forms as part of the design of the neural networks.

Formulation: Koopman + SIREN + Physics

Smoothen, more accurate results



Result:



Math is cool!

SciML Shows how to build Earthquake-Safe Buildings

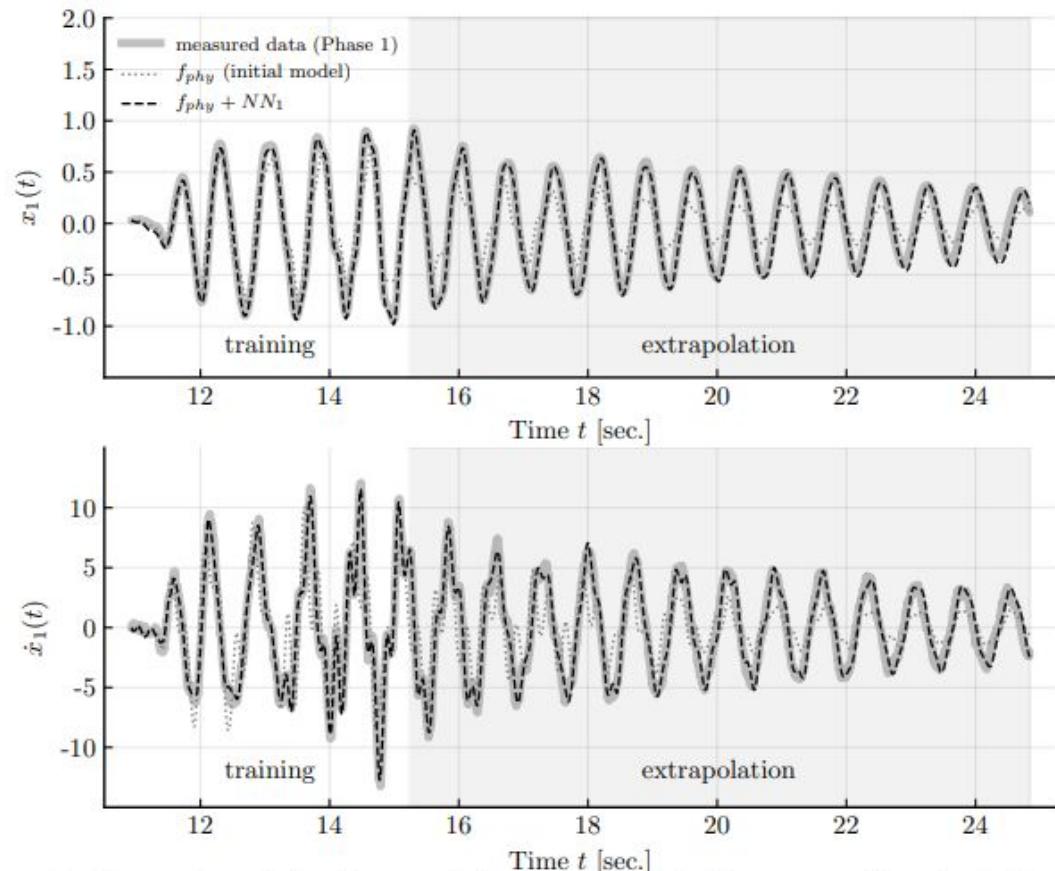


Figure 12: Comparison of time history of the response for displacement $x_1(t)$ and velocity $\dot{x}_1(t)$ for the NSD experiment (Phase 1).

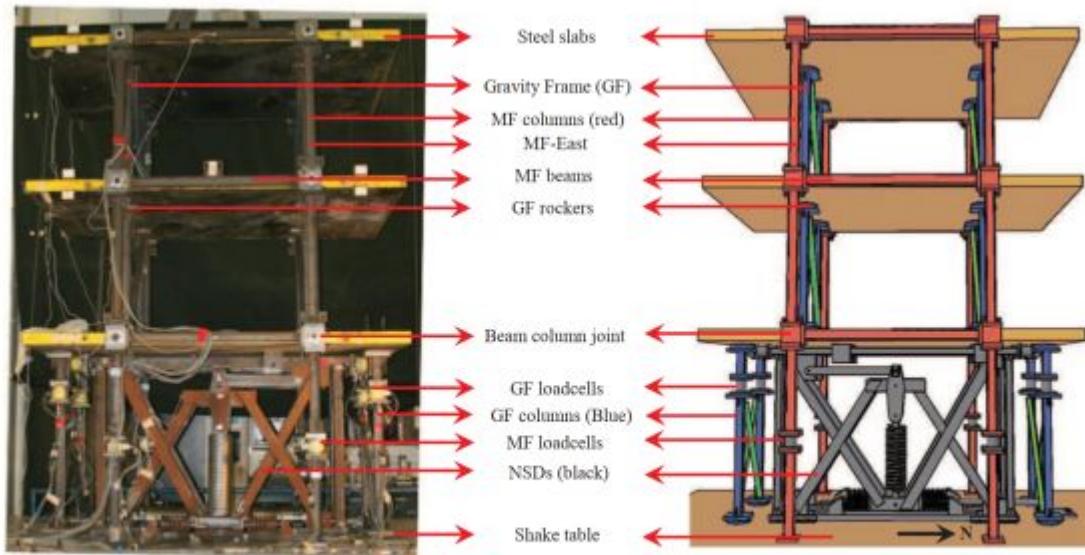


Figure 10: The structural system equipped with a negative stiffness device in between the first floor and the shake table.

Structural identification with physics-informed neural ordinary differential equations

Lai, Zhilu, Mylonas, Charilaos, Nagarajaiah, Satish, Chatzi, Eleni

For a detailed walkthrough of UDEs and applications watch on Youtube:

Chris Rackauckas: Accurate and Efficient Physics-Informed Learning Through Differentiable Simulation

Fundamental of Clinical Pharmacology: Nonlinear Mixed Effects (NLME) Modeling

Goal: Learn to predict patient behavior (**dynamics**) from simple data (**covariates**)

$$Z_i = \begin{bmatrix} wt_i, \\ sex_i, \end{bmatrix}$$

Covariates

$$g_i = \begin{bmatrix} Ka \\ CL \\ V \end{bmatrix} = \begin{bmatrix} \theta_1 e^{\eta_{i,1} \kappa_{i,k,1}}, \\ \theta_2 \left(\frac{wt_i}{70} \right)^{0.75} \theta_4^{sex_i} e^{\eta_{i,2}}, \\ \theta_3 e^{\eta_{i,3}}, \end{bmatrix}$$

Structural Model (pre)

Intuition: η (the random effects) are a fudge factor

Find θ (the fixed effect, or average effect) such that you can predict new patient dynamics as good as possible

Math: Find (θ, η) such that $E[\eta] = 0$
Requires special fitting procedures (Pumas)

$$\begin{aligned} \frac{d[\text{Depot}]}{dt} &= -Ka[\text{Depot}], \\ \frac{d[\text{Central}]}{dt} &= Ka[\text{Depot}] - \frac{CL}{V}[\text{Central}]. \end{aligned}$$

Dynamics

The Impact of Pumas (PharmacUtical Modeling And Simulation)

“

We have been using Pumas software for our pharmacometric needs to support our development decisions and regulatory submissions.

Pumas software has surpassed our expectations on its accuracy and ease of use. We are encouraged by its capability of supporting different types of pharmacometric analyses within one software. **Pumas has emerged as our "go-to" tool for most of our analyses in recent months.** We also work with Pumas-AI on drug development consulting. We are impressed by the quality and breadth of the experience of Pumas-AI scientists in collaborating with us on modeling and simulation projects across our pipeline spanning investigational therapeutics and vaccines at various stages of clinical development

Husain A. PhD (2020)

Director, Head of Clinical Pharmacology and Pharmacometrics,
Moderna Therapeutics, Inc

”

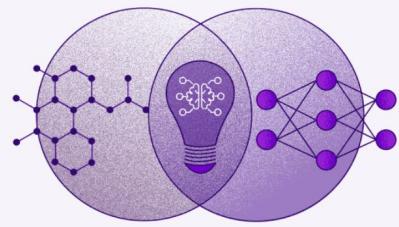
Built on SciML



moderna™

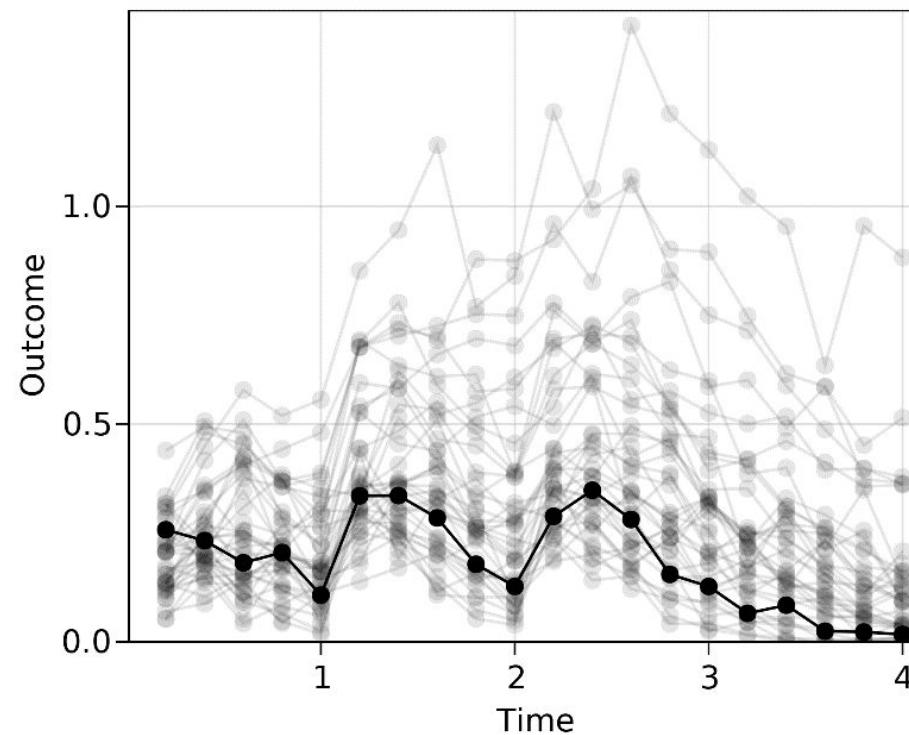
messenger therapeutics





DeepPumas: Deep Learning Nonlinear Mixed Effects Software

Method and apparatus for automating models for individualized administration of medicaments, Christopher Vincent Rackauckas, Vijay Ivaturi, US63/136,719



Typical values

$$\theta \in \mathbb{R}_+^3$$

$$\Omega \in \mathbb{R}_+^3$$

Patient data



Random effects

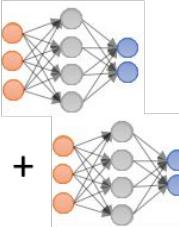
$$\eta \sim \text{MvNormal}(\Omega)$$

Individual parameters

$$Ka_i = \theta_1 \cdot e^{\eta_{i,1}} + c_1 \cdot Age_i +$$

$$CL_i = \theta_2 \cdot e^{\eta_{i,2}}$$

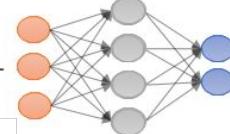
$$V_i = \theta_3 \cdot e^{\eta_{i,3}} + c_2 \cdot Weight_1^{c_3} +$$



Dynamics

$$\frac{d[\text{Depot}]}{dt} = -Ka[\text{Depot}],$$

$$\frac{d[\text{Central}]}{dt} = Ka[\text{Depot}] -$$



Error model

$$\text{Outcome} \sim \text{Normal} \left(\text{Central}, \sqrt{\text{Central}} \cdot \sigma \right)$$

Two Solutions to The Bitter Lesson

Complementary Approaches to Overcoming AI's Limitations

Scientific Machine Learning

Adding Scientific Knowledge to Scalable AI Methods



$$\begin{aligned}\partial u / \partial t &= f(u, t) \\ F &= ma \\ \nabla^2 \phi &= Q \\ E &= mc^2\end{aligned}$$

Key Concept:

Embed physical laws, conservation principles, and domain knowledge directly into neural network architectures and training processes

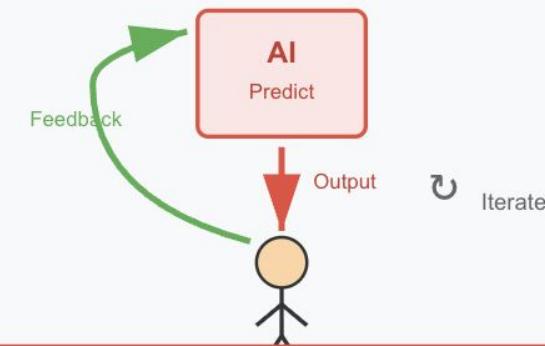
Benefits:

- ✓ Better generalization with less data
- ✓ Physical consistency guaranteed
- ✓ Interpretable predictions
- ✓ Efficient for scientific applications

Example: Universal Differential Equations

Agentic AI

Agentic Loops with Humans in the Loop



Key Concept:

Use AI as a tool in an iterative workflow where human expertise validates, corrects, and guides predictions in a continuous improvement loop

Benefits:

- ✓ Catches and corrects AI errors quickly
- ✓ Leverages human domain expertise
- ✓ Improves with each iteration
- ✓ Builds trust through transparency

Example: AI-Assisted Engineering Design

Agentic AI in science isn't tomorrow.

It's today.

How do you get started today?

AI for Software Engineering is Progressing

SWE-bench	Model	% Resolved	Avg. \$	Trajs	Org	Date	Release
	Gemini 3 Pro Preview (2025-11-18)	74.20	\$0.46	🔗	⭐	2025-11-18	1.15.0
Leaderboards	Claude 4.5 Sonnet (20250929)	70.60	\$0.56	🔗	AV	2025-09-29	1.13.3
BENCHMARKS	Claude 4 Opus (20250514)	67.60	\$1.13	🔗	AV	2025-08-02	1.0.0
SWE-bench	GPT-5 (2025-08-07) (medium reasoning)	65.00	\$0.28	🔗	🔗	2025-08-07	1.7.0
SWE-bench Verified	Claude 4 Sonnet (20250514)	64.93	\$0.37	🔗	AV	2025-07-26	1.0.0
SWE-bench Bash Only	GPT-5 mini (2025-08-07) (medium reasoning)	59.80	\$0.04	🔗	🔗	2025-08-07	1.7.0
SWE-bench Multilingual	o3 (2025-04-16)	58.40	\$0.33	🔗	🔗	2025-07-26	1.0.0
SWE-bench Multimodal	OpenB-Coder 480B/A35B Instruct	55.40	\$0.25	🔗	🔗	2025-08-02	1.0.0
SWE-bench Lite							



WHY AGENTIC AI ACTUALLY WORKS

It's not about LLM accuracy. It's about the feedback loop.

THE MYTH

"LLMs are getting accurate enough to generate correct code on the first try."

THE REALITY

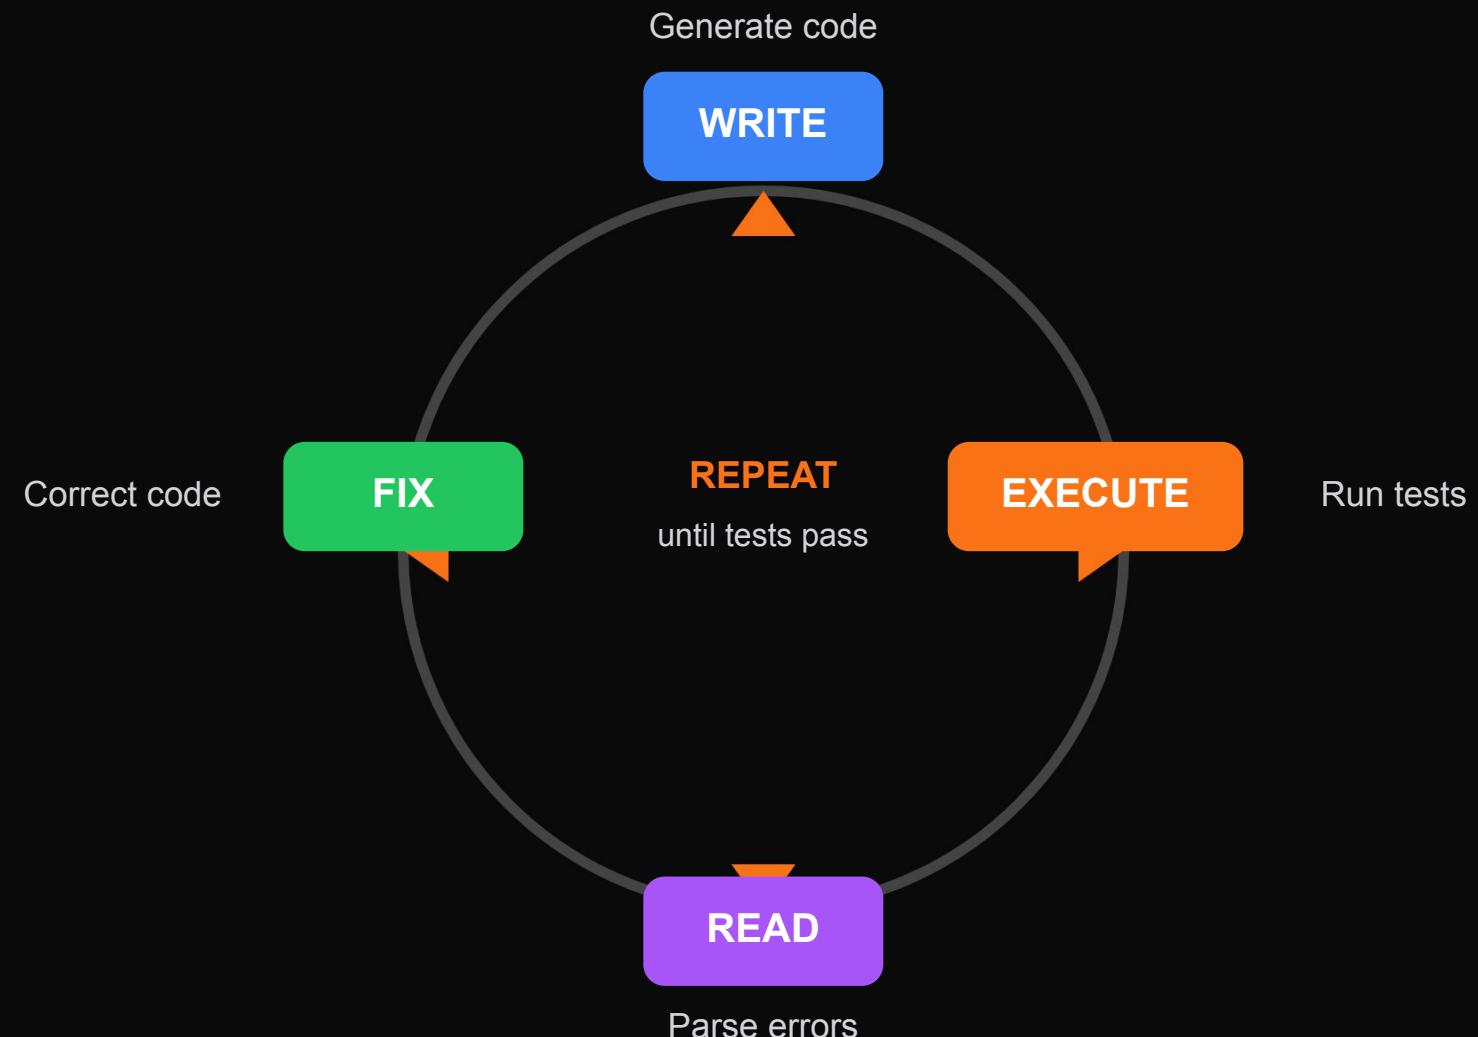
Agentic AI iterates until correct. The compiler is the oracle.

~65%

First try

95%+

After loop



The compiler provides ground truth. The loop provides convergence.

AI CLI CODING AGENTS: HEAD-TO-HEAD

Terminal-based AI assistants for software development (2025)

Claude Code

ANTHROPIC

✓ PROS

- Best code quality and accuracy
- Superior agentic navigation
- Rich customization (CLAUDE.md)
- Parallel task execution

✗ CONS

- \$20-200/month pricing

BEST FOR

Enterprise & Complex Projects

Codex CLI

OPENAI

✓ PROS

- GitHub Actions integration
- ChatGPT subscription included

✗ CONS

- Complex setup process
- Verbose code output
- Less informative CLI UX
- Fewer advanced features

BEST FOR

CI/CD & Collaborative Teams

Gemini CLI

GOOGLE

✓ PROS

- Generous free tier
- 1M token context window
- GCP/Google ecosystem
- Built-in web search grounding

✗ CONS

- Requires more guidance
- GCP setup for Workspace
- Less polished output

BEST FOR

Budget & Large Codebases

Step 1: Get Claude Code Running

Doodling ...

Built for > hackers

Work with Claude directly in your codebase. Build, debug, and ship from your terminal, IDE, Slack, or the web. Describe what you need, and Claude handles the rest.

Get Claude Code ▾

`curl -fsSL https://claude.ai/install.sh | bash` ↗

Or read the [documentation](#)

Demonstration:

1. Setting up Claude Code
2. Local orchestration in VS Code
3. Sandboxing and “agentic freedom”
 - a. Claude Code CLI
 - b. gh bot authentication
4. Multi-agent bot swarm organization

With the fundamental accuracy challenges of LLMs, what can be done to improve Agentic AI for engineering?

WHY AGENTIC AI ACTUALLY WORKS

It's not about LLM accuracy. It's about the feedback loop.

THE MYTH

"LLMs are getting accurate enough to generate correct code on the first try."

THE REALITY

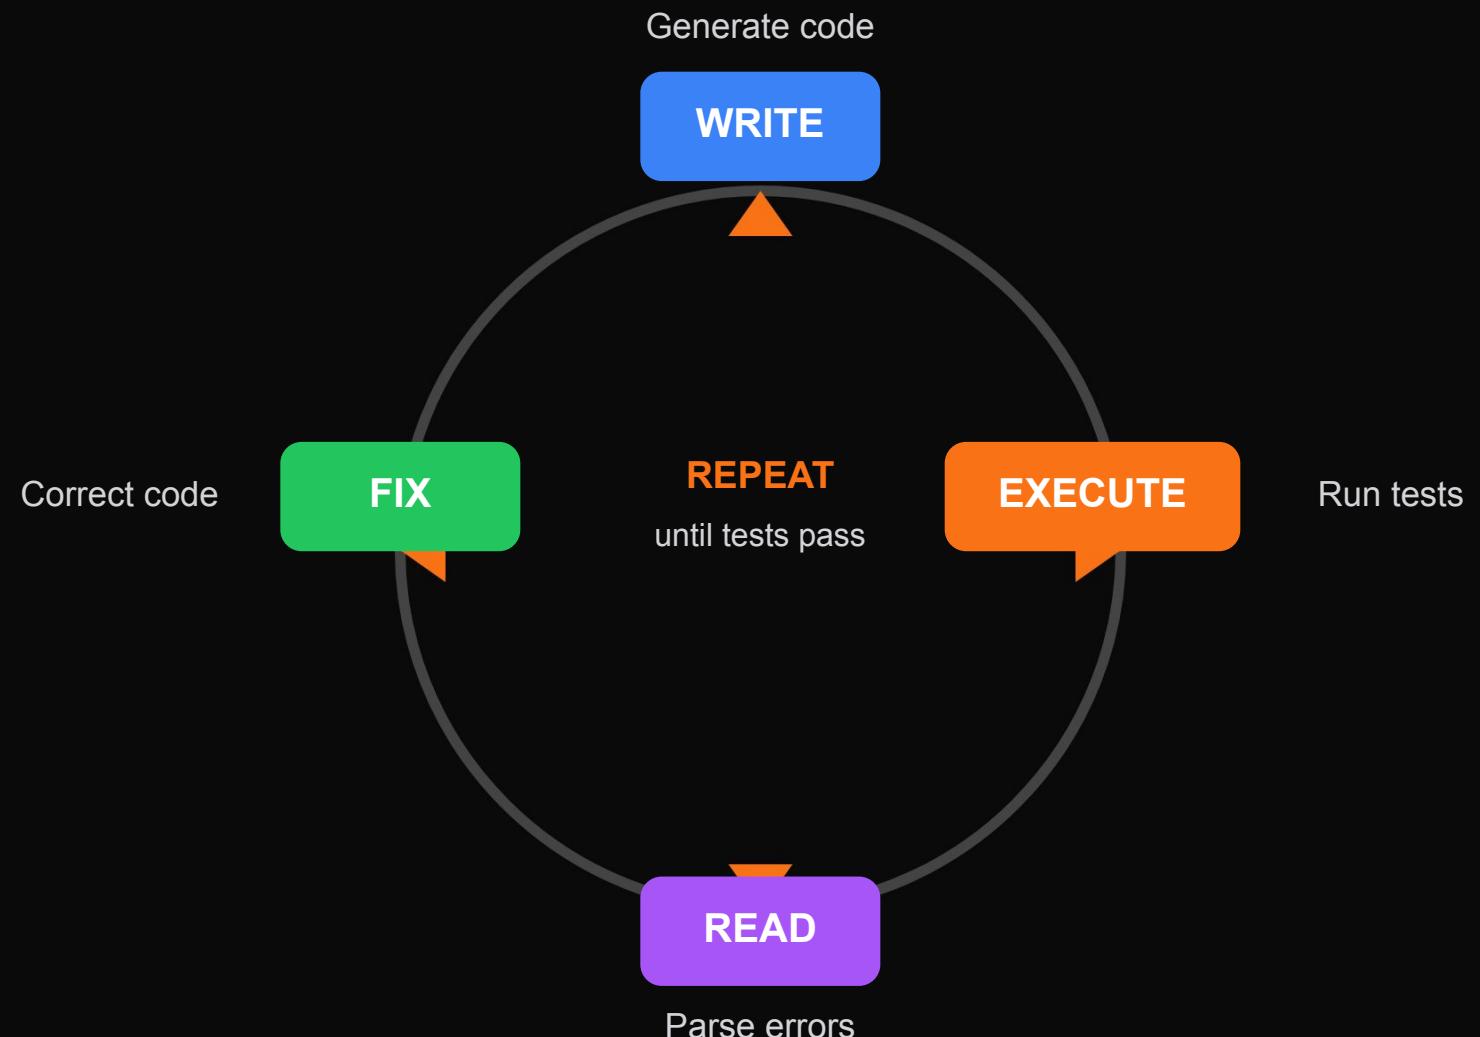
Agentic AI iterates until correct. The compiler is the oracle.

~65%

First try

95%+

After loop



The compiler provides ground truth. The loop provides convergence.

NVIDIA Research: Small Language Models are Sufficient?

Small Language Models are the Future of Agentic AI

Peter Belcak¹ Greg Heinrich¹ Shizhe Diao¹ Yonggan Fu¹ Xin Dong¹

Saurav Muralidharan¹ Yingyan Celine Lin^{1,2} Pavlo Molchanov¹

¹NVIDIA Research ²Georgia Institute of Technology

agents-research@nvidia.com

Here we lay out the position that small language models (SLMs) are *sufficiently powerful, inherently more suitable, and necessarily more economical for many invocations in agentic systems, and are therefore the future of agentic AI*. Our argumentation is grounded in the current level of capabilities exhibited by SLMs, the common architectures of agentic systems, and the economy of LM deployment. We further argue that in situations where general-purpose conversational abilities are essential, heterogeneous agentic systems (i.e., agents invoking multiple different models) are the natural choice. We discuss the potential barriers for the adoption of SLMs in agentic systems and outline a general LLM-to-SLM agent conversion algorithm.

Can this idea be scaled up to a complete language? A complete modeling system?

Dyad: Modeling and Simulation Designed for Agentic AI

Let AI write models, but let codegen make it fast and accurate

Simpler code for the AI to write and debug

Static checks for accuracy and agentic debugging

```
# Sample component
component MyComponent
    # A component can have one or more extends clauses, like this one...
    extends BaseComponent
    # A normal parameter only results in a parametric change, not a structural change,
    parameter y::Real
    # A structural parameter is a parameter changes the number of equations or variables.
    structural parameter N::Integer = 10
    # Variables and parameters can have a number of different attributes
    variable x::Real(min = 0, max = 10, guess = 5, units = "m")
    # Subcomponents are components nested inside other components forming hierarchical models
    subcomponent = SomeComponent()
    # Connectors represent points of interaction between the components
    p = Pin()
relations
    # One type of relation is an equation, like this one
    der(x) = y
    # Another type of relation is a connection
    connect(subcomponent.p, p)
end
```

Supports composition of physical models into components

Future: Allow for variables to depend on attributes that specify properties of the variable (e.g. min, max, units)

Acausal modeling framework for capturing algebraic-differential relationships

Dyad



Advancing Julia into Industrial Modeling and Simulation

1:1 GUI and Code

Physics-based modeling, controls and embedded software in one tool

Generative AI and SciML

Access to the full Julia ecosystem

Dyad accelerates breakthrough engineering with Software-Defined Machines



CUSTOMERS

Accelerating Critical Industries

Enabling aerospace, automotive, utility and high-tech manufacturing companies to innovate faster

PRODUCT

AI-Native Modeling Platform

Dyad is the next-gen platform for systems modeling, simulation, and digital twins

TECHNOLOGY

Built on Open Source

Powered by Julia and SciML. 100M+ downloads across 10,000 organizations and 1,500 universities.

Physical Modeling with Dyad



What is Dyad?

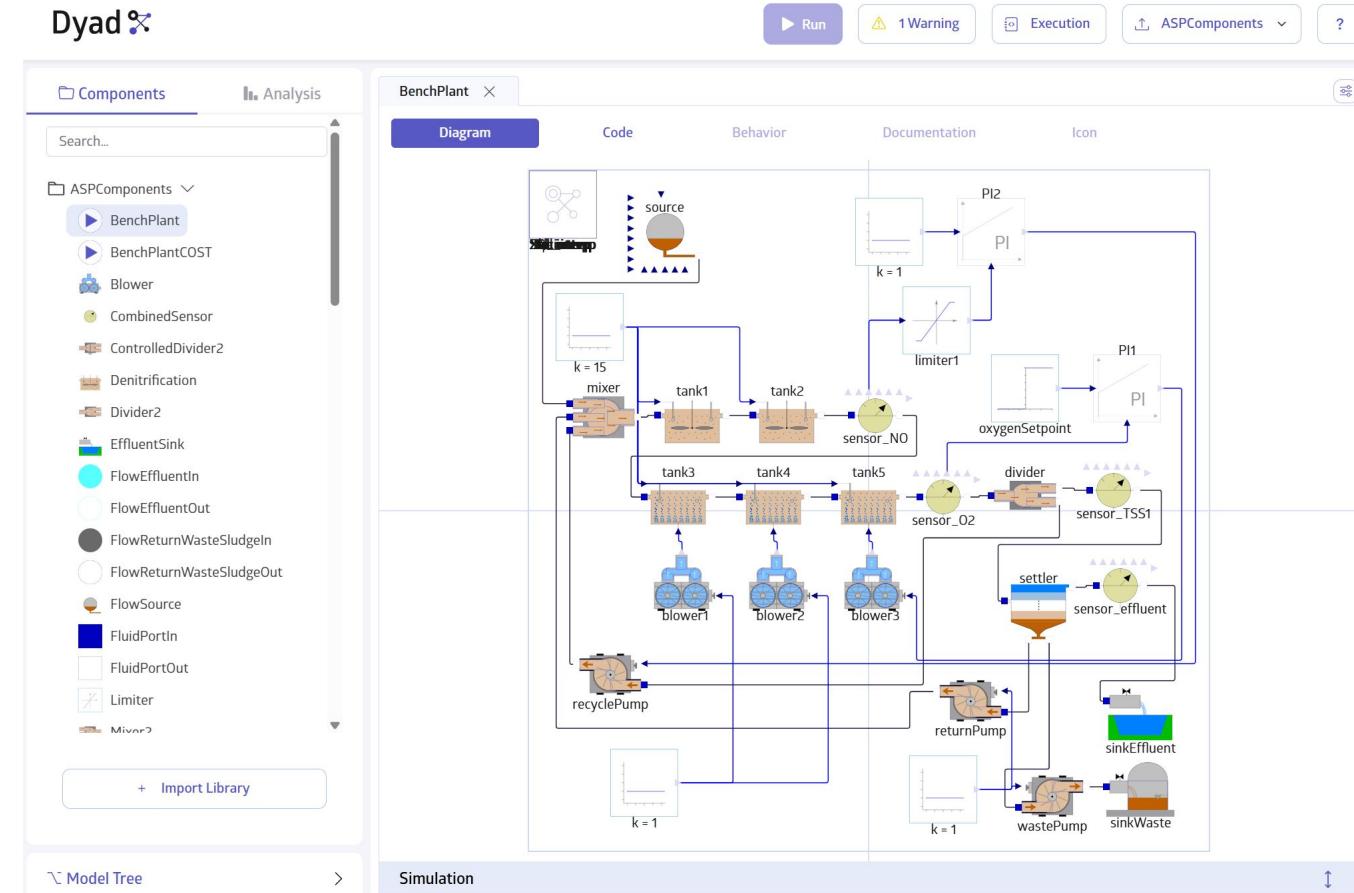
- Acausal plant modeling (0D, 1D, 2D)
- Causal synchronous modeling with state machines
- Agentic AI for engineer productivity
- Scientific AI for analysis models
- Control systems synthesis and code generation
- Modern development workflows: git, CI/CD

Dyad AI and SciML

- Surrogate modeling
- Agentic workflows
- Design space optimization through ensembles, UQ, Sensitivity analysis

Go To Market

- Partnership with Ansys to include Dyad in TwinAI
- Already demonstrated customer success



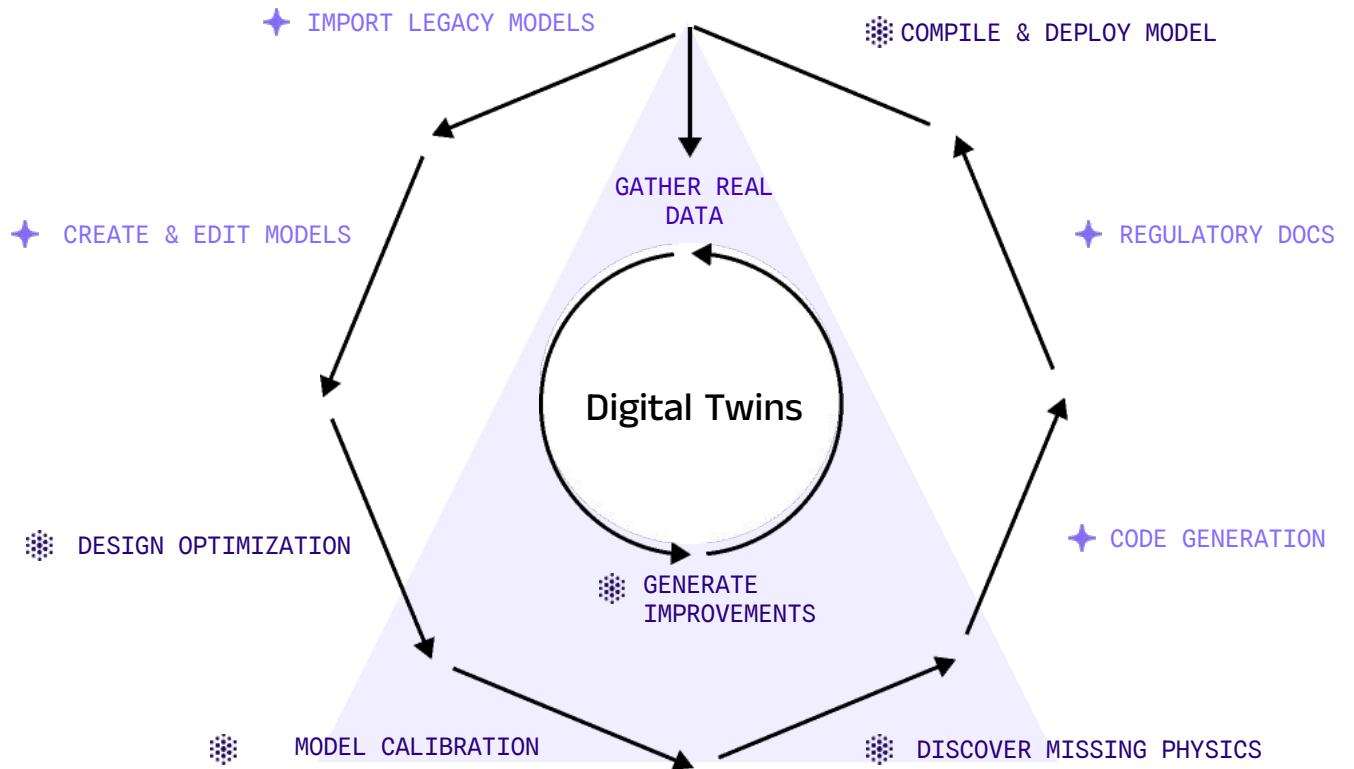
Dyad delivers a powerful combination of Scientific and Generative AI

SCIENTIFIC AI

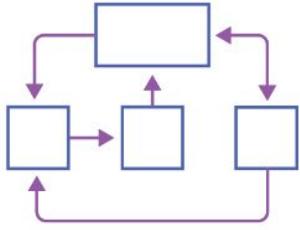
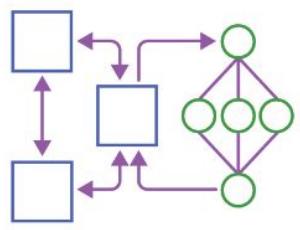
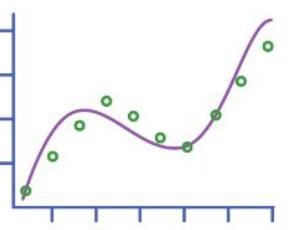
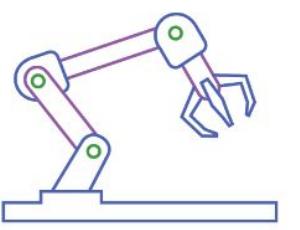
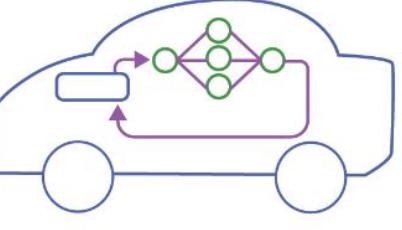
Through physics-constrained and data-driven modeling, Scientific AI drives increased productivity for engineers, and faster time to market

GENERATIVE AI

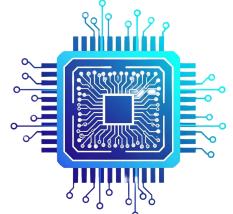
Generative AI makes it possible to break the walled gardens of legacy software which opens the door for the best compiler to win the market



Dyad at a Glance

Design	Discover	Calibrate	Control	Surrogate
<ul style="list-style-type: none">• Build realistic physical models with minimal code• Run simulations 100x faster	<ul style="list-style-type: none">• Use Machine Learning to autocomplete models• Discover missing physics	<ul style="list-style-type: none">• Turn models into Digital Twins• Robust nonlinear fitting with automatic differentiation	<ul style="list-style-type: none">• Build robust nonlinear controls• Deploy Model-Predictive Controllers (MPC)	<ul style="list-style-type: none">• Train neural networks to match models• Accelerate fast simulations by another 100x
				

Dyad Standard Libraries



ElectricalComponents

Analog/digital components as well as electrical machines



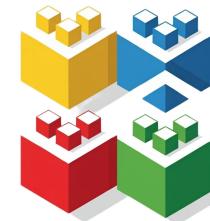
RotationalComponents

Rotational mechanical components used to make clutches, gears



ThermalComponents

Conductive and radiative components, heat transfer



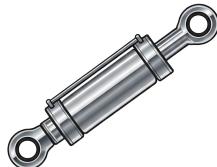
PrimitiveComponents

Mathematical operations for embedded programming



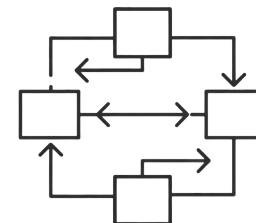
TranslationalComponents

Flanges, masses, springs, dampers, forces, translational mechanics components used to construct drive shaft and simple vehicle models



HydraulicComponents

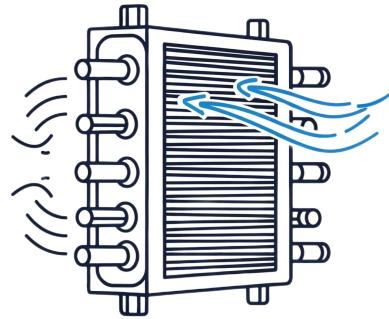
Valves, pipes, pistons



BlockComponents

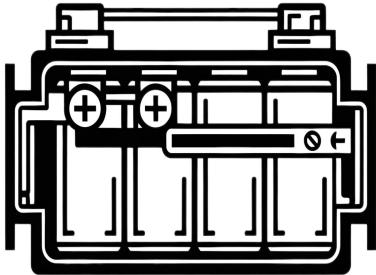
Numeric and logical components, and source blocks

Domain-Specific In-House Libraries



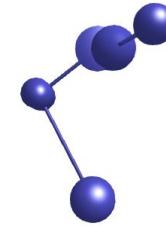
HVAC

- Single phase and two-phase flow
- Standard library of pipes, compressors, heat exchangers, etc.
- Proprietary refrigerant and media libraries with specialized splines for high accuracy and performance in two-phase scenarios



Batteries

- Equivalent Circuit Models (ECMs) (low complexity approximations)
- Single particle Models (SPM) and enhancements (medium complexity approximations)
- Doyle-Fuller-Newman Model (DFN) (high-complexity and high-accuracy)
- Battery degradation models
- Validated chemistry libraries (NiCd, Li-ion, etc.)



Multibody

- Automated generation of animations and visualizations
- Kinematic loops
- Robotics
- Vehicle Dynamics
- Trajectory planning
- Aerial, Ground and Sea Vehicles such as drones, submarines

Dyad HVAC

Pre-built components:

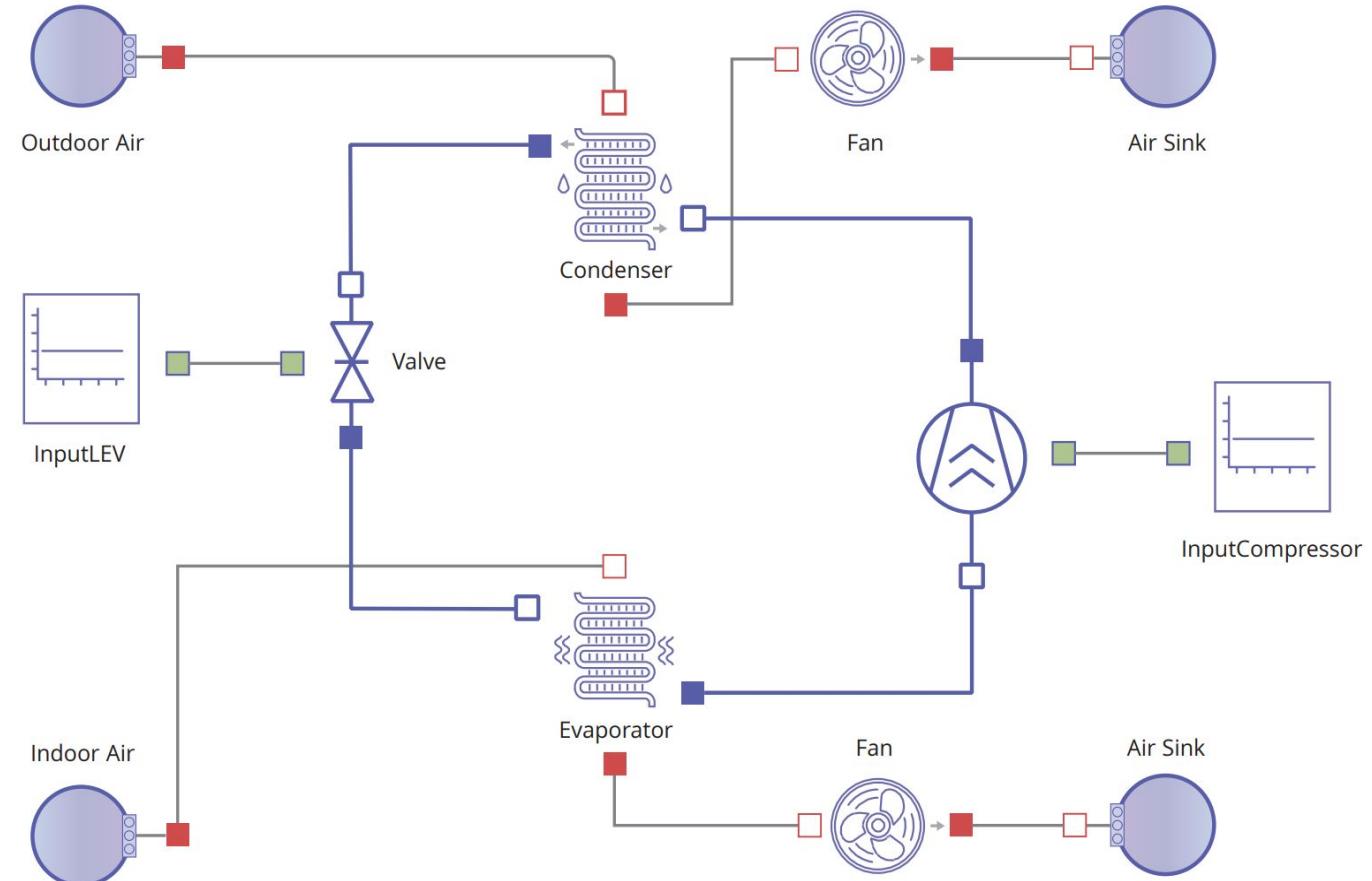
- Tube Fin, Plate heat exchangers
- Compressors, Valves, Fans, Conditioned Spaces, Pipes

Thermodynamic Property Models:

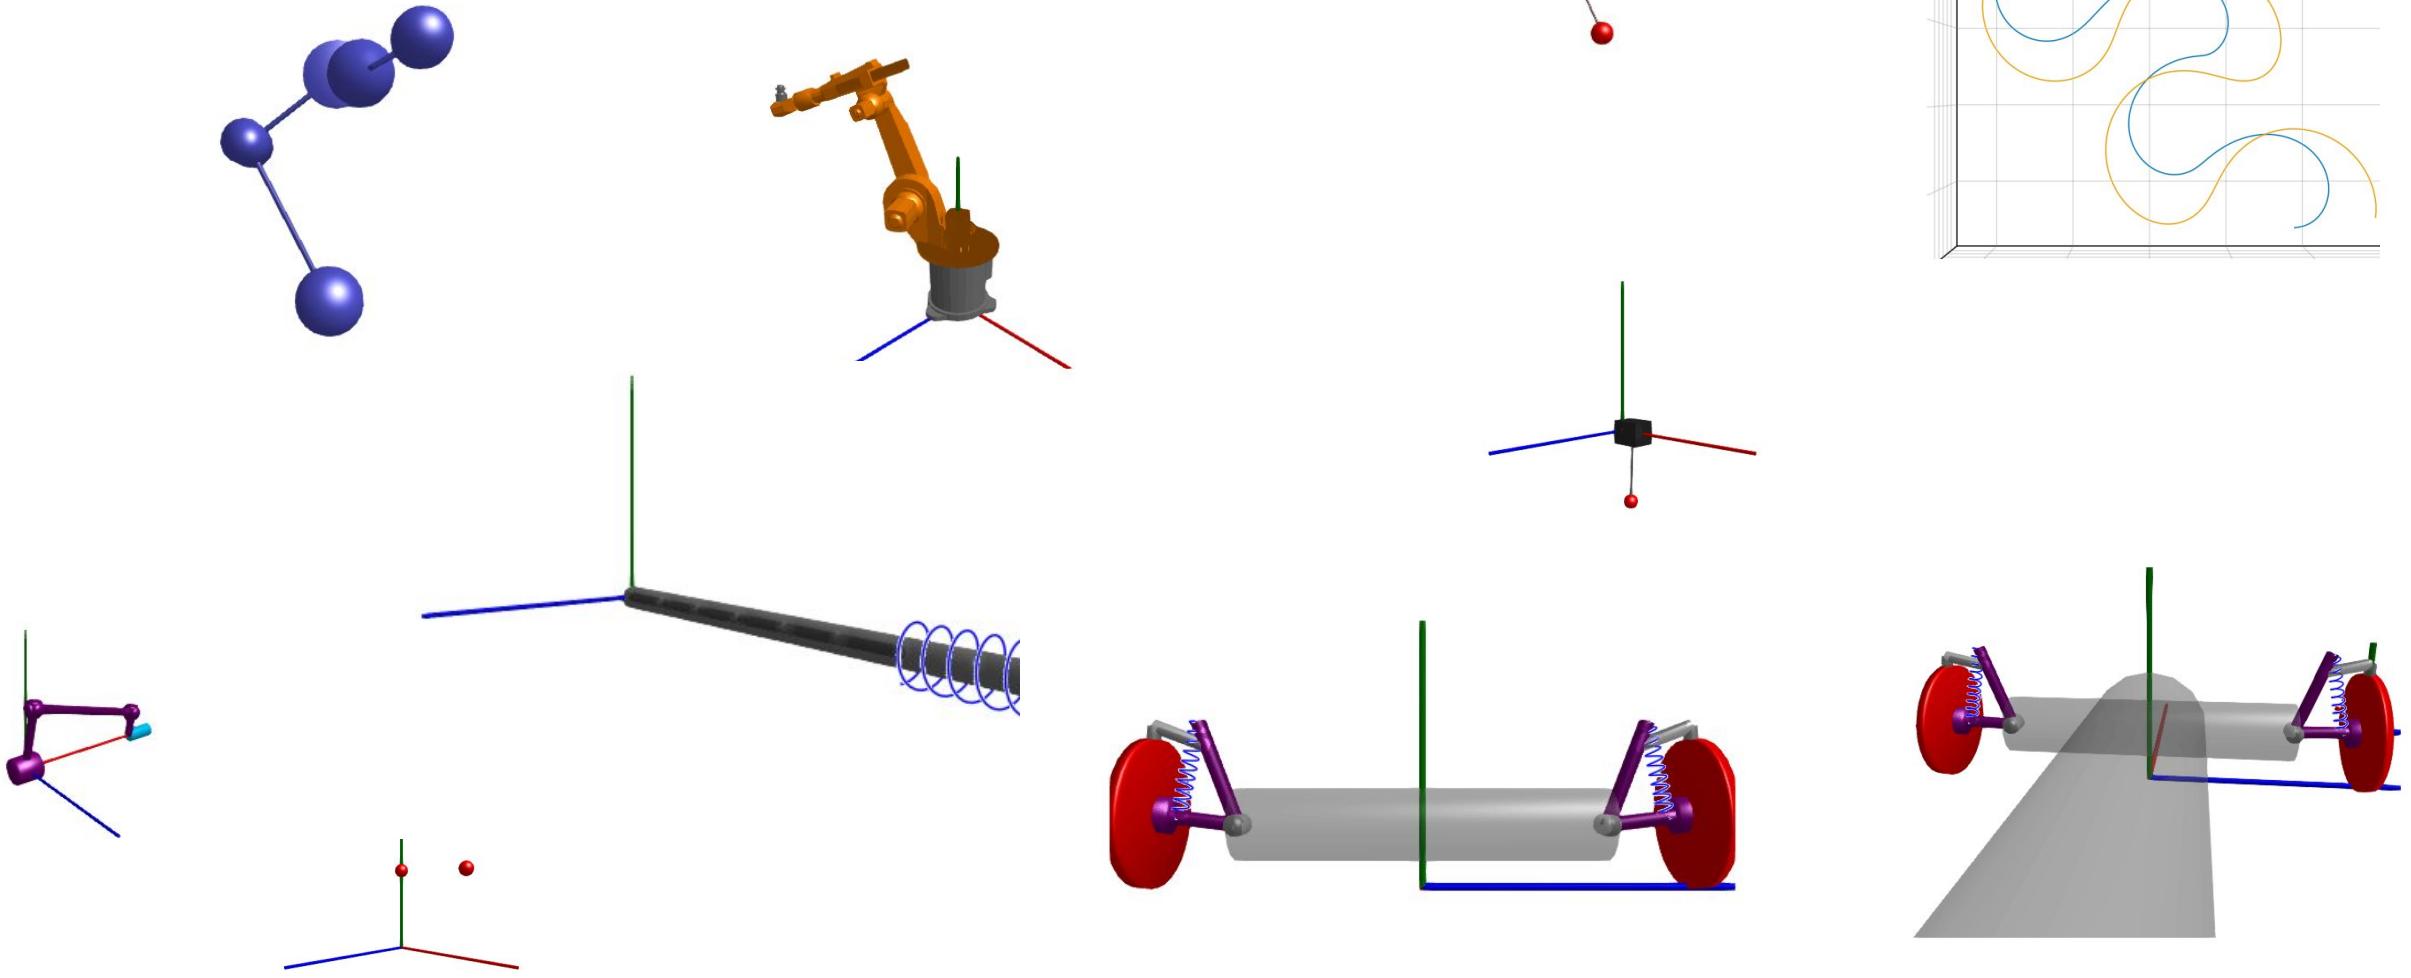
- Spline-based thermodynamic property models (R32, R1234YF, R290, R152a, R134a, R410A, R717)
- Dry and Moist air, Media models

Ecosystem:

- Numerical Solvers
- Calibration, Machine Learning, Control



Dyad Multibody

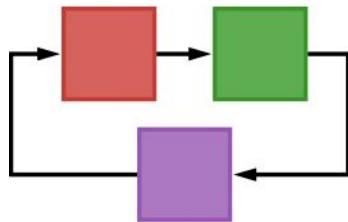


Dyad Analysis Libraries



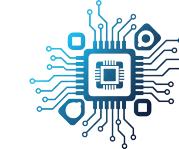
Scientific Machine Learning (SciML) Analysis Libraries

- Model Optimizer
 - Model calibration
 - Design optimization
 - Global Optimization
- Model Discovery
 - Neural Differential Equations and Universal Differential Equations
 - Symbolic Regression model generation from data
- Neural Surrogate Generation



Control Analysis Libraries

- System Identification
- Model Predictive Control
- Optimal Control and Trajectory Optimization
- Linear Controls
 - Generation and curation of analysis points
 - Model Order Reduction
 - Structured controller tuning
- Robust Controls
 - H_{Inf}
 - Robust MPC



Code Generation and Deployment Analysis Libraries

- FMU Import and Export
- Target Binary Generation
 - Generation of C code for embedded devices
 - Regulatory support (DO178, ARP26262, etc.)

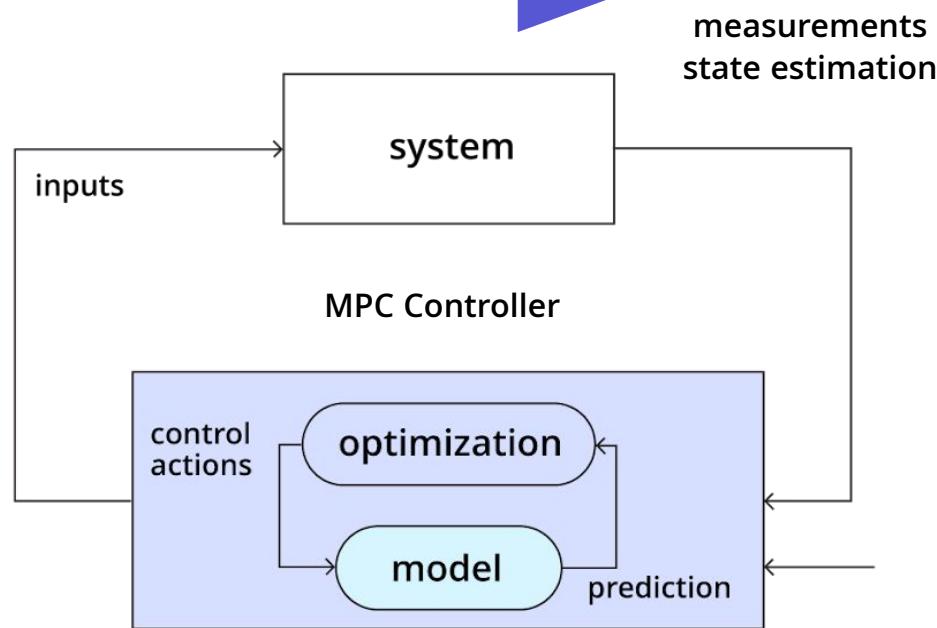
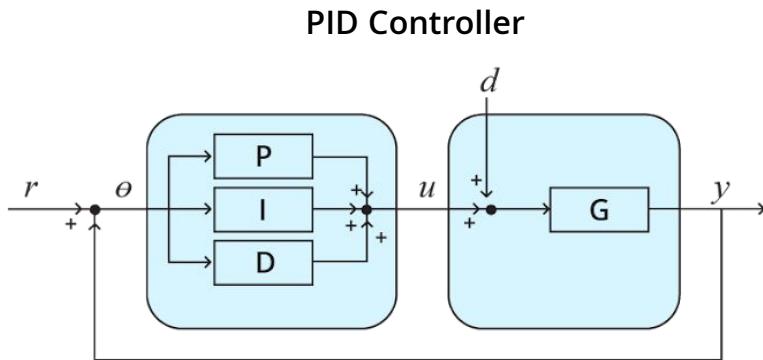
Dyad Control Systems

Take control of your hardware, explore your nonlinearities

Simple

Whole spectrum

Advanced



Linear control
analysis and
synthesis

Construct and tune
robust
**Model-predictive
controllers**

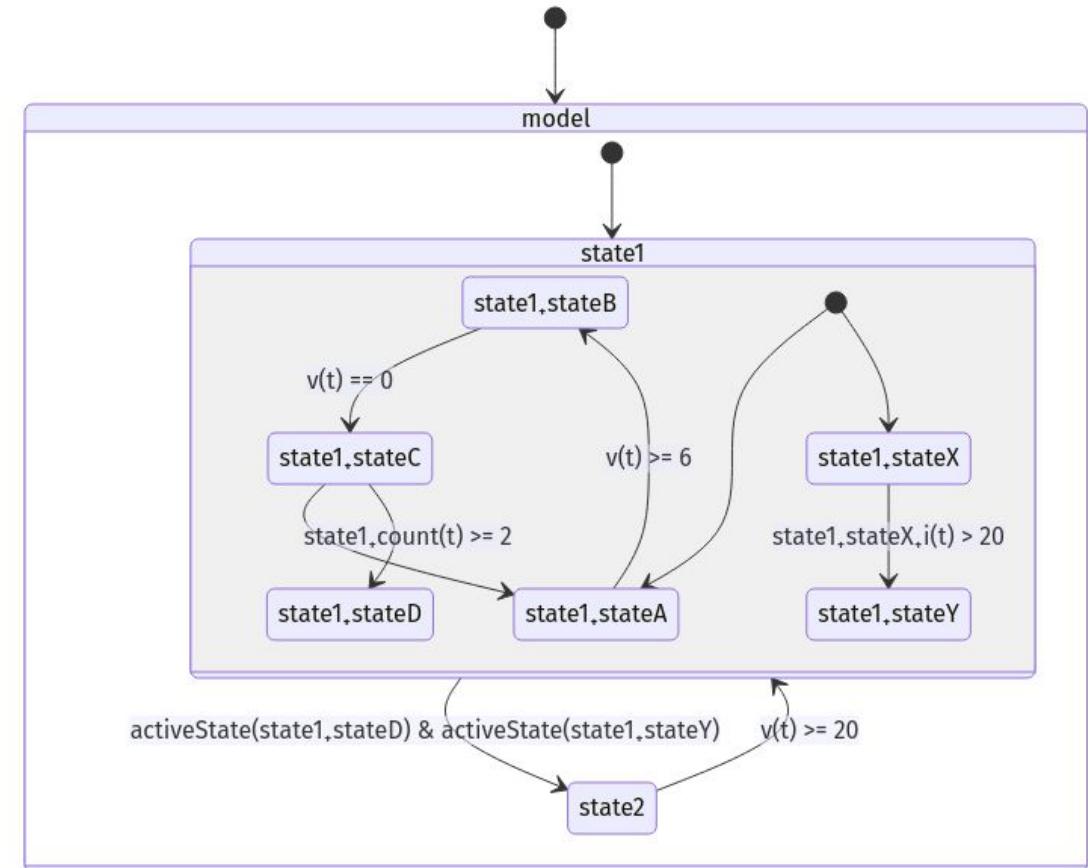
Dyad State Machines and Synchronous Programs

Features in development:

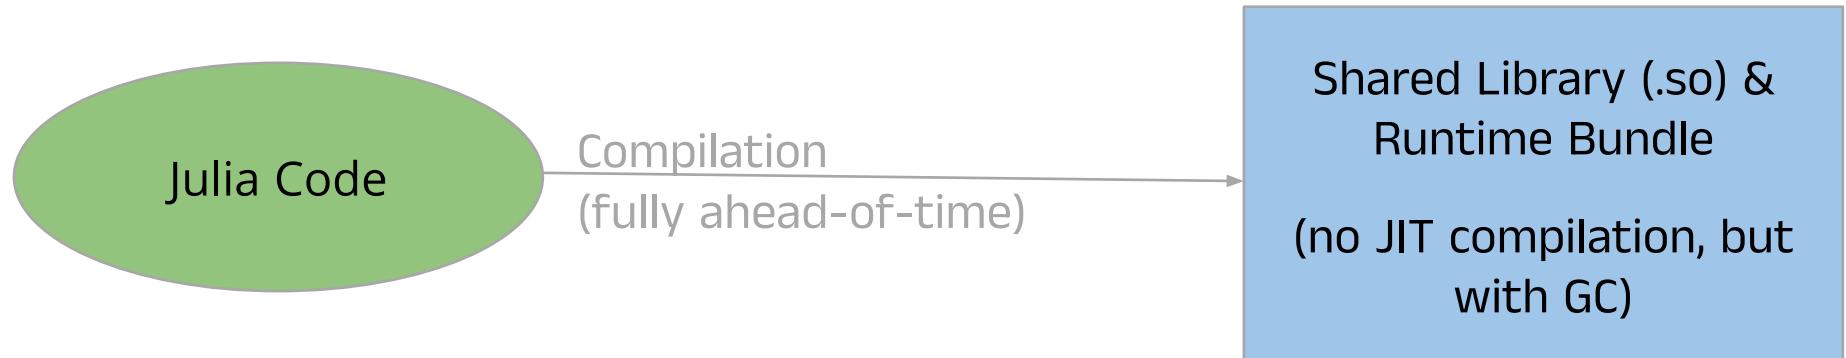
- Fundamental operations such as entry, exit, transitions, repeated operations
- Clock-based discrete time models
- Hybrid discrete-continuous systems

General Availability timeline:

- Q1 2026 for hybrid systems
- Q2 2026 for state machines



Compiling Julia down to Hardware Targets



Supported targets: ARM v8, x86-64, x86
Supported OS: Linux, macOS, Windows

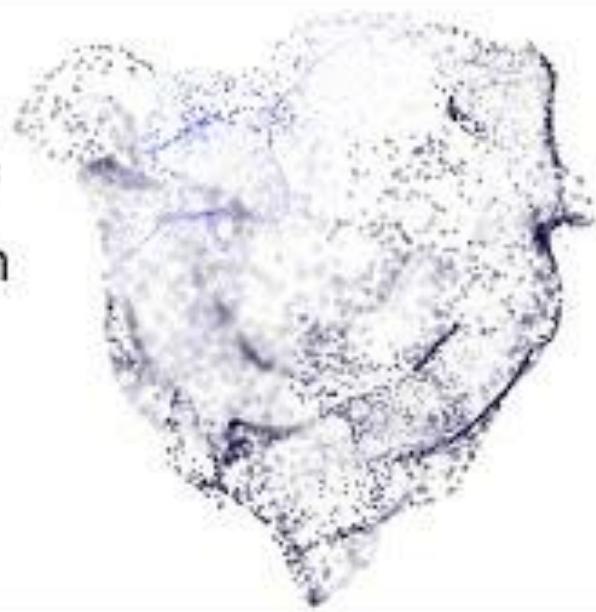


Planned OS/Target support including SoCs
Will support Traceability and Safety-critical software

Agentic Model Based Design From Scratch

Agentic modeling of
a friction brake with
Dyad

Automated & Validated



Thermo-mechanical model of a braking system built
from scratch using Dyad

Dyad has developed an AI agent for building complex system models

Through simple queries, the agent derives the physical equations and parametrization for user feedback, after which it implements the model in Dyad.

Update parameterization of model and generate test cases

Blog post:
<https://juliahub.com/blog/agentic-ai-dyad>

Agentic Translation from Modelica to Dyad

```
File Edit Selection View Go Run Terminal Help < > DyadModelTranslator (Workspace) [SSH: trantor]
translator.jl A × pipeline.jl M DyadModelTranslator.jl > translator.jl > ...
Julia REPL (v1.11.6) - DyadModelTranslator.jl + · · · | · · · ×
File Edit Selection View Go Run Terminal Help < > DyadModelTranslator (Workspace) [SSH: trantor]
translator.jl A × pipeline.jl M DyadModelTranslator.jl > translator.jl > ...
Julia REPL (v1.11.6) - DyadModelTranslator.jl + · · · | · · · ×
1 using Crayons
2 using JSON3
3 using LLMAgents
4 using ModelTranslator
5 using ModelTranslator: create_generated_lib, get_library_map,
6   modeling_agent_system_prompt, get_library_map_flattened,
7   generate_translation_state
8
9 target_lib::String = "HeaterSystemComponents"
10 target_library_path = "/home/ven-k/.julia/dev/$target_lib"
11 asset_path = "$target_library_path/assets/ModelicaDyadComponentMap"
12
13 # Load asset libraries
14 library_map = get_library_map_flattened(asset_path)
15 library_of_functions = merge(
16   JSON3.read(joinpath(asset_path, args: "Functions.json"), Dict{String, String}),
17   JSON3.read(joinpath(asset_path, "$target_lib.json"), Dict{String, String})
18 )
19 library_of_constants = JSON3.read(joinpath(asset_path, args: "Constants.json"))
20
21 # Declare agents
22 ANTHROPIC_MODEL_NAME::String = "claude-3-5-sonnet-latest"
23 API_VERSION::String = "2023-06-01"
24
25 llm_instance::LLMAgents.AnthropicModel = AnthropicModel(
26   ANTHROPIC_MODEL_NAME,
27   API_VERSION,
28   url: "https://api.anthropic.com/v1/messages",
29 );
30
31 include(x: "tools.jl");
32
33 dyad_agent::LLMAgents.Agent = Agent(llm_instance, modeling_agent_system_prompt()
34
35 flash_model::String = "gemini-2.5-flash-preview-05-20"
36 flash_agent = ModelTranslator.GeminiAgent(flash_model)
```

Translator decomposes model into a hierarchical tree and searches for Dyad equivalents to each node

Can easily convert converts Modelica equations to Dyad

Runs translated model through Dyad compiler and revises model based on error messages

By Andrew

Satisfied customer review

[Read more](#)

1,000+ reviews



1,000+

Reviews

1,000+

Photos

1,000+

Videos

Satisfied customer review

1,000+

Reviews

1,000+

Photos



Add Analysis

Type

New dimension

Find the most popular dimension



Top dimension



New



Dimension



Customer review



Customer service



Customer support



Customer feedback



Customer loyalty



Customer satisfaction



Customer retention



Customer value



Next



— 1 —

Volume 8 Number 1

- Ward, who had been working for the county with the San Joaquin River, from 1965 to 1970, told the Senate Natural Resources Committee he was asked to leave because he was critical of the agency's environmental impact statement.

100

- High school students (14, 15, 16) are given an optional 10-15 minute pretest. They get three (3) extra points of bonus credit on their final test, and excellent extra credit on their final assignment. This motivates them to do well.

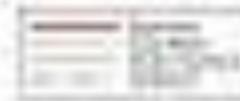
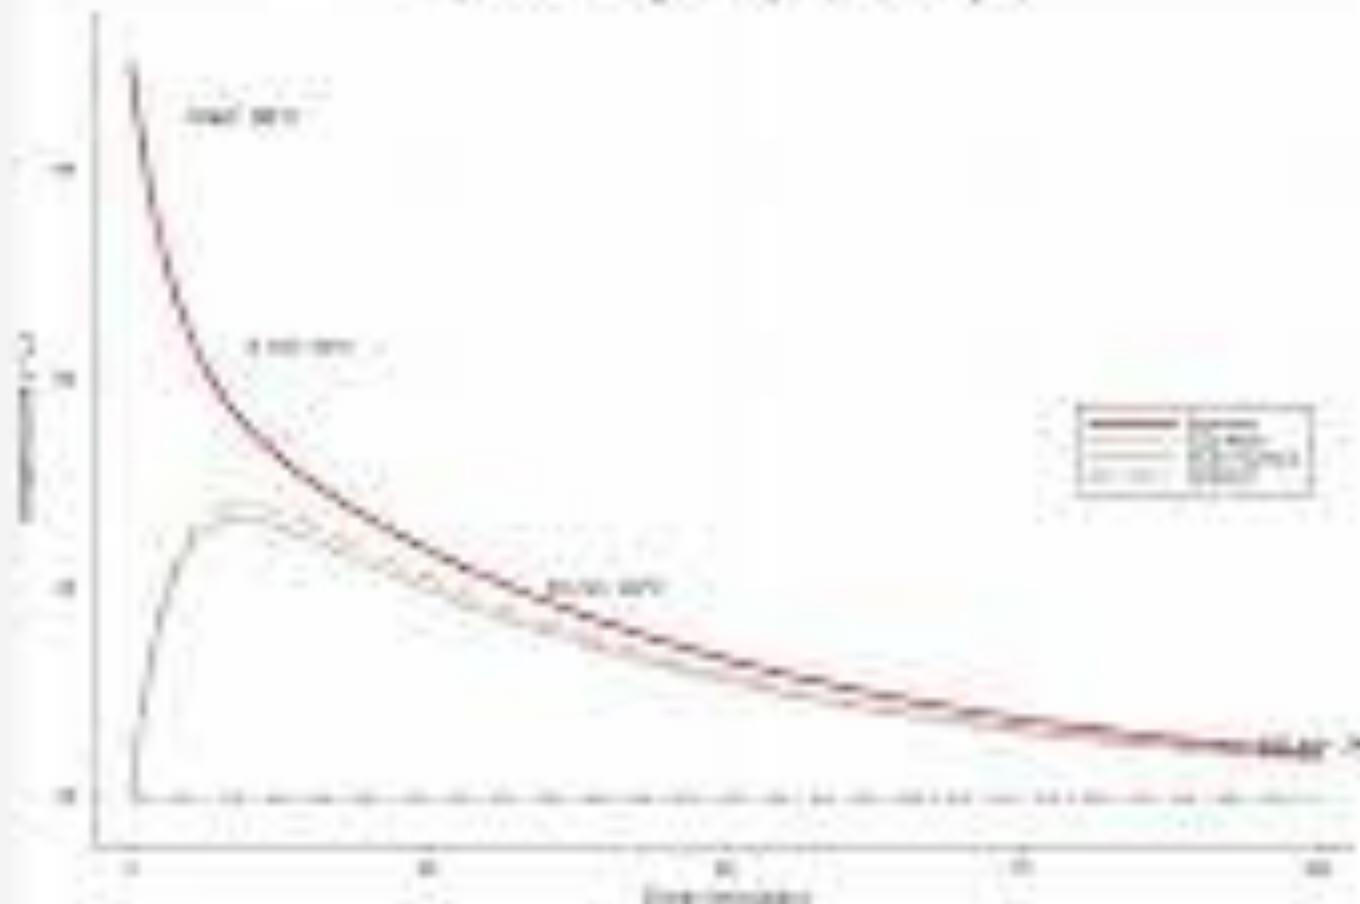
一門多學科

- Бюджетное управление включает в себя бюджетирование, бюджетный контроль, бюджетную отчетность и бюджетное планирование.

Все это ведет к тому, что в будущем придется искать новые способы извлечения, обогащения золота, а также

• 100 •

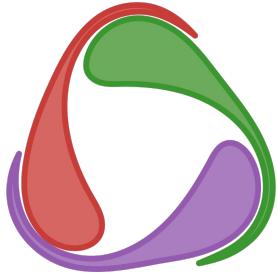
Espresso Cooling - Heavy Duty EC = 1000 RPM



••JuliaHub

JULIASIM RECREATION OF THE NASA HL20 LIFTING BODY

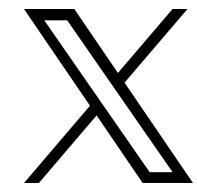




Connect with SciML



Join our community chatrooms
julialang.org/slack



Follow on Twitter
[@ChrisRackauckas](https://twitter.com/ChrisRackauckas)



Follow on LinkedIn
<https://www.linkedin.com/in/chrisrackauckas/>



Star the GitHub repositories
<https://github.com/SciML>



Dr. Chris Rackauckas
VP of Modeling and Simulation
JuliaHub
Research Affiliate
MIT CSAIL
Director of Scientific Research
Pumas-AI

SciML Open Source Software Organization

sciml.ai

- DifferentialEquations.jl: 2x-10x Sundials, Hairer, ...
- DiffEqFlux.jl: adjoints outperforming Sundials and PETSc-TS
- ModelingToolkit.jl: 15,000x Simulink
- Catalyst.jl: >100x SimBiology, gillespy, Copasi
- DataDrivenDiffEq.jl: >10x pySindy
- NeuralPDE.jl: ~2x DeepXDE* (more optimizations to be done)
- NeuralOperators.jl: ~3x original papers (more optimizations required)
- ReservoirComputing.jl: 2x-10x pytorch-esn, ReservoirPy, PyRCN
- SimpleChains.jl: 5x PyTorch GPU with CPU, 10x Jax (small only!)
- DiffEqGPU.jl: Some wild GPU ODE solve speedups coming soon

And 100 more libraries to mention...

If you work in SciML and think optimized and maintained implementations of your method would be valuable, please let us know and we can add it to the queue.

Democratizing SciML via pedantic code optimization
Because we believe full-scale open benchmarks matter

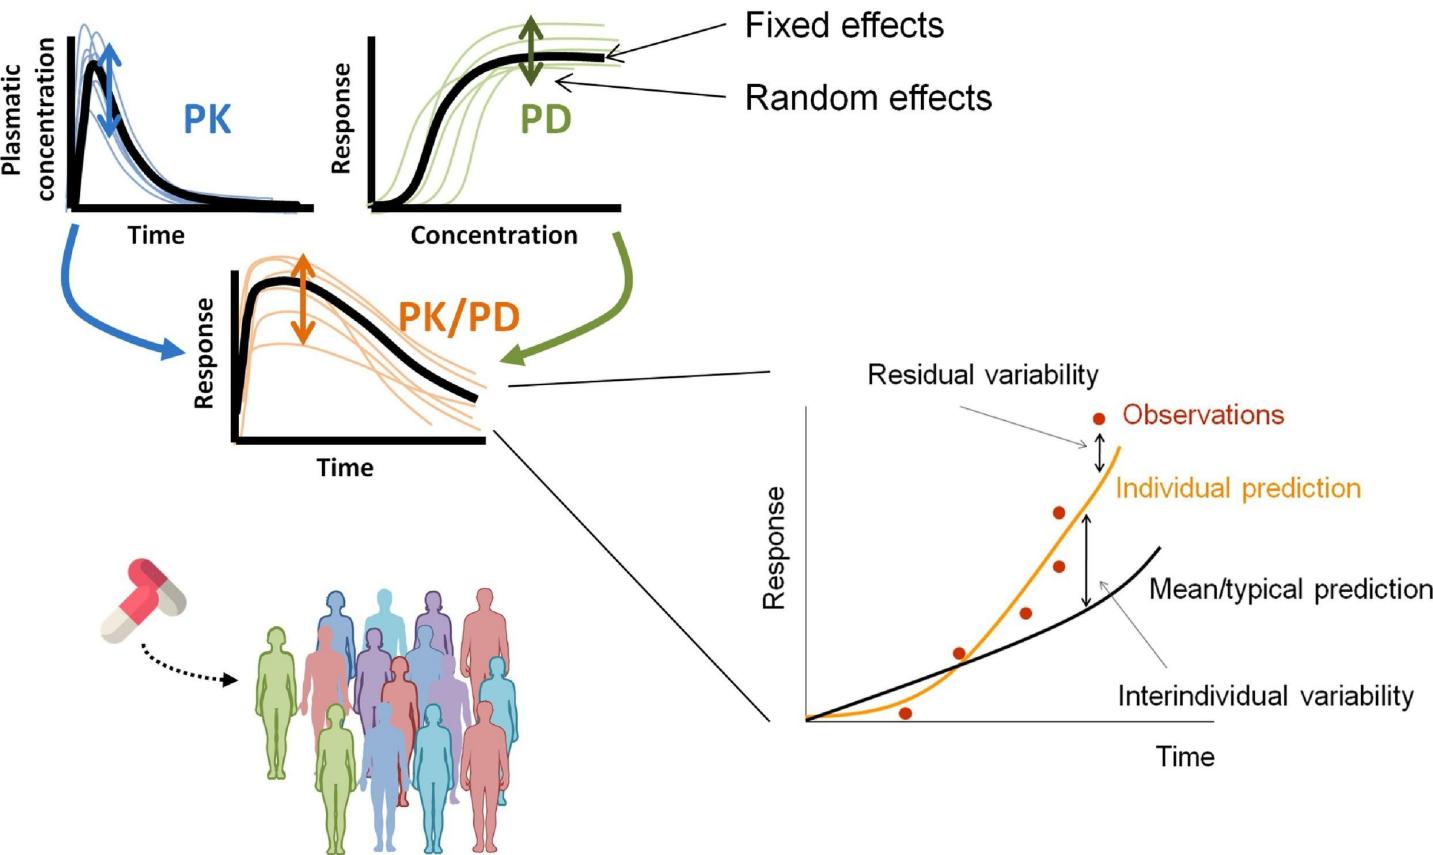


DeepNLME: Integrate neural networks into traditional NLME modeling

DeepNLME is SciML-enhanced modeling for clinical trials

DeepNLME is SciML-enhanced modeling for clinical trials

Mixed-effects modeling



- Automate the discovery of predictive covariates and their relationship to dynamics
- Automatically discover dynamical models and assess the fit
- Incorporate big data sources, such as genomics and images, as predictive covariates

From Dynamics to Nonlinear Mixed Effects (NLME) Modeling

Goal: Learn to predict patient behavior (dynamics) from simple data (covariates)

$$Z_i = \begin{bmatrix} wt_i, \\ sex_i, \end{bmatrix}$$

Covariates

$$g_i = \begin{bmatrix} Ka \\ CL \\ V \end{bmatrix} = \begin{bmatrix} \theta_1 e^{\eta_{i,1} \kappa_{i,k,1}}, \\ \theta_2 \left(\frac{wt_i}{70}\right)^{0.75} \theta_4^{sex_i} e^{\eta_{i,2}}, \\ \theta_3 e^{\eta_{i,3}}, \end{bmatrix}$$

Structural Model (pre)

Intuition: η (the random effects) are a fudge factor

Math: Find (θ, η) such that $E[\eta] = 0$
Requires special fitting procedures (Pumas)

Find θ (the fixed effect, or average effect) such that you can predict new patient dynamics as good as possible

$$\frac{d[\text{Depot}]}{dt} = -Ka[\text{Depot}],$$
$$\frac{d[\text{Central}]}{dt} = Ka[\text{Depot}] - \frac{CL}{V}[\text{Central}].$$

Dynamics

The Impact of Pumas (PharmacUtical Modeling And Simulation)

“

We have been using Pumas software for our pharmacometric needs to support our development decisions and regulatory submissions.

Pumas software has surpassed our expectations on its accuracy and ease of use. We are encouraged by its capability of supporting different types of pharmacometric analyses within one software. **Pumas has emerged as our "go-to" tool for most of our analyses in recent months.** We also work with Pumas-AI on drug development consulting. We are impressed by the quality and breadth of the experience of Pumas-AI scientists in collaborating with us on modeling and simulation projects across our pipeline spanning investigational therapeutics and vaccines at various stages of clinical development

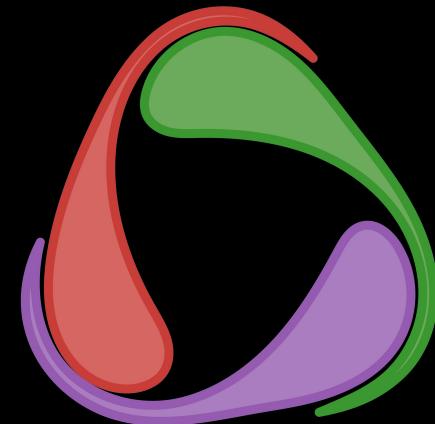
Husain A. PhD (2020)

Director, Head of Clinical Pharmacology and Pharmacometrics,
Moderna Therapeutics, Inc

moderna™

messenger therapeutics

Built on SciML



From Dynamics to Nonlinear Mixed Effects (NLME) Modeling

Goal: Learn to predict patient behavior (dynamics) from simple data (covariates)

$$Z_i = \begin{bmatrix} wt_i, \\ sex_i, \end{bmatrix}$$

Covariates

$$\text{Covariates} \rightarrow g_i = \begin{bmatrix} Ka \\ CL \\ V \end{bmatrix} = \begin{bmatrix} \theta_1 e^{\eta_{i,1} \kappa_{i,k,1}}, \\ \theta_2 \left(\frac{wt_i}{70}\right)^{0.75} \theta_4^{sex_i} e^{\eta_{i,2}}, \\ \theta_3 e^{\eta_{i,3}}, \end{bmatrix}$$

Structural Model (pre)

Intuition: η (the random effects) are a fudge factor

Math: Find (θ, η) such that $E[\eta] = 0$

How can we find these models?

Find θ (the fixed effect, or average effect) such that you can predict new patient dynamics as good as possible

$$\begin{aligned} \frac{d[\text{Depot}]}{dt} &= -Ka[\text{Depot}], \\ \frac{d[\text{Central}]}{dt} &= Ka[\text{Depot}] - \frac{CL}{V} [\text{Central}]. \end{aligned}$$

Dynamics

From Dynamics to Nonlinear Mixed Effects (NLME) Modeling

Goal: Learn to predict patient behavior (dynamics) from simple data (covariates)

$$Z_i = \begin{bmatrix} wt_i, \\ sex_i, \end{bmatrix}$$

Covariates

$$g_i = \begin{bmatrix} K_a \\ CL \\ V \end{bmatrix}$$

Structural Model (pre)

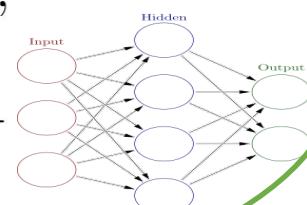
Math: Find (θ, η) such that $E[\eta] = 0$

Idea: Parameterize the model such that the models can be neural networks, where the weights of the neural networks are fixed effects!

Indirect learning of unknown functions!

How can we find these models?

$$\frac{d[\text{Depot}]}{dt} = -K_a[\text{Depot}],$$
$$\frac{d[\text{Central}]}{dt} = K_a[\text{Depot}] -$$



Dynamics

How are the networks trained?

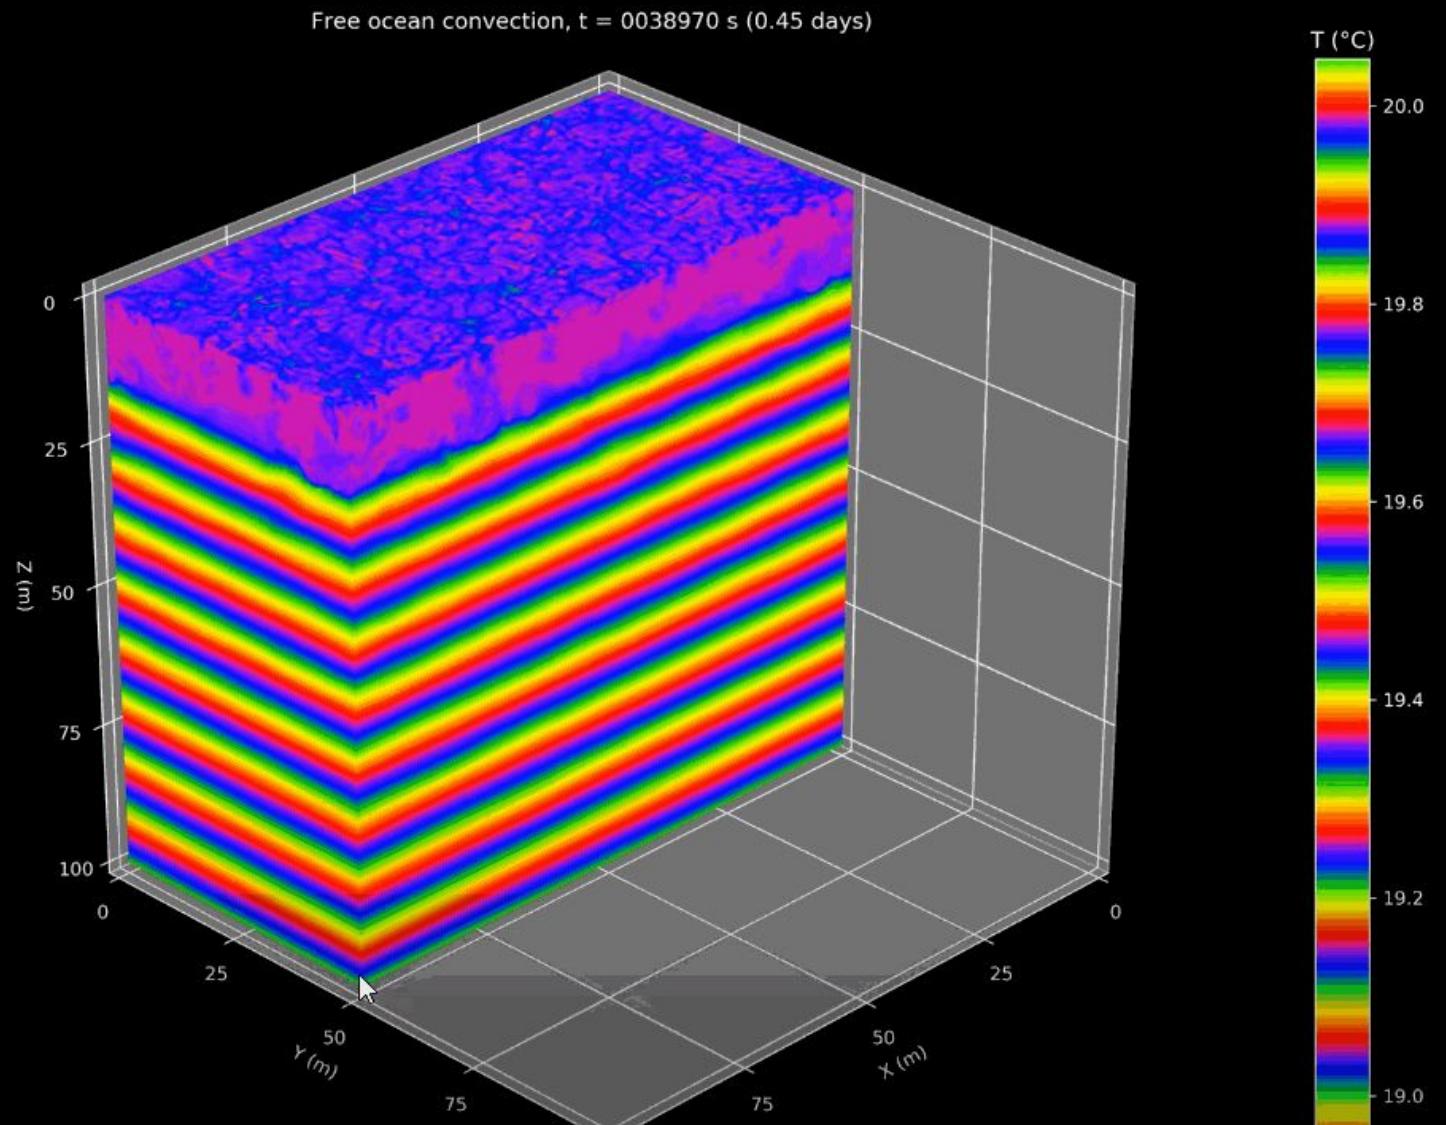
High fidelity surrogates of ocean columns for climate models

3D simulations are high resolution but too expensive.

Can we learn faster models?

Ramadhan, Ali, John Marshall, Andre Souza, Gregory LeClaire Wagner, Manvitha Ponnappati, and Christopher Rackauckas. "Capturing missing physics in climate model parameterizations using neural differential equations." *arXiv preprint arXiv:2010.12559* (2020).

(Update going online in the next month)



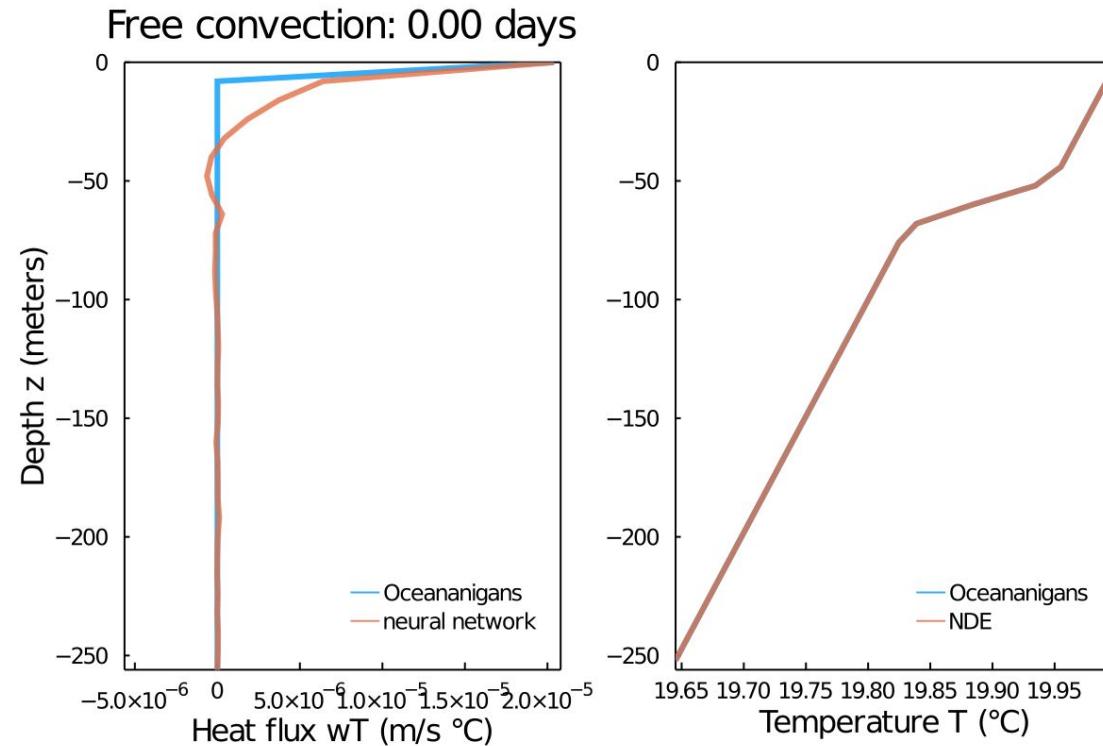
Neural Networks Infused into Known Partial Differential Equations

Derive a 1D approximation to the 3D model

$$\frac{\partial T}{\partial t} = -\frac{\partial}{\partial z} \left(\underbrace{\text{Input} \rightarrow \text{Hidden} \rightarrow \text{Output}}_{w' T'} - K \frac{\partial T}{\partial z} \right)$$

Incorporate the “convective adjustment”

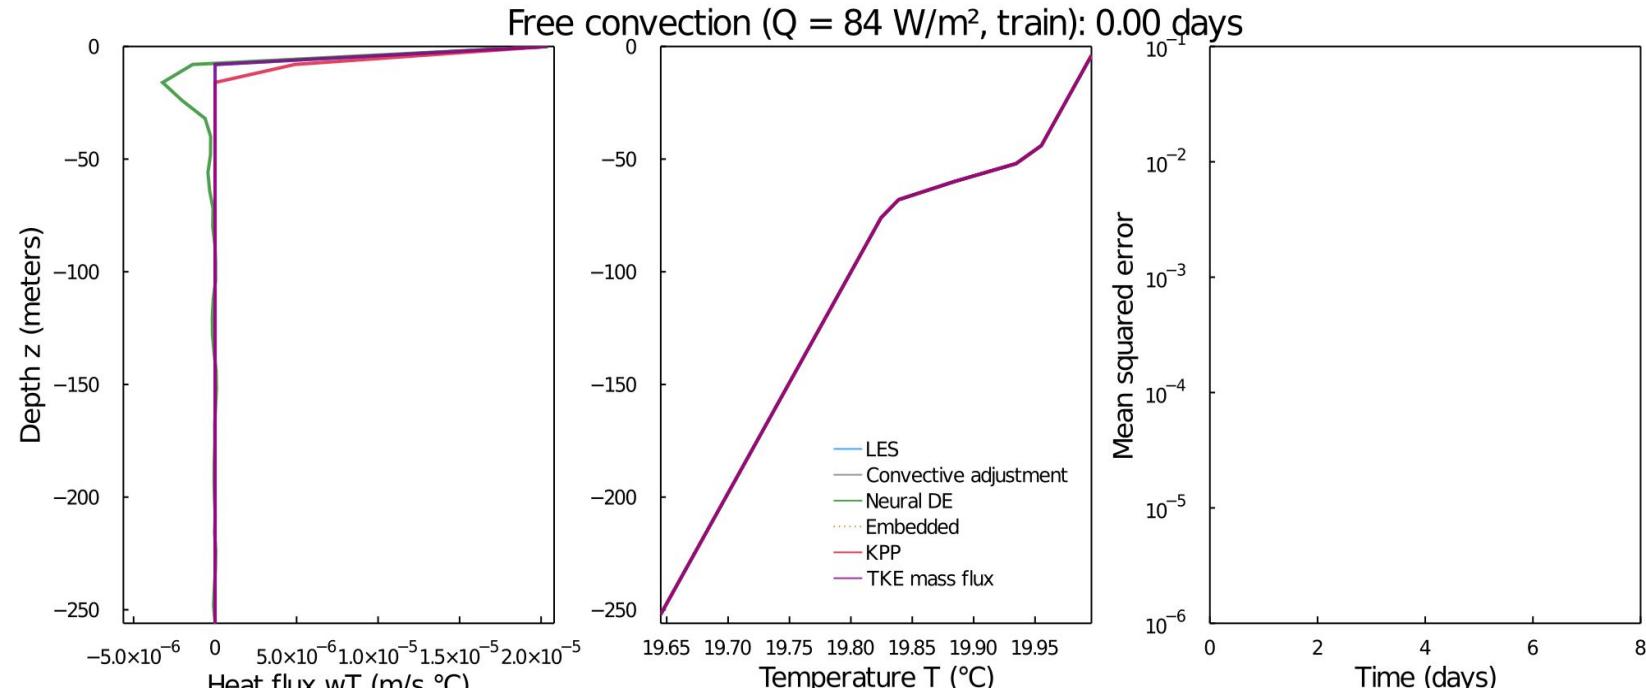
$$K = \begin{cases} 0 & \text{if } \partial_z T > 0 \\ 100 \text{ m}^2/\text{s} & \text{if } \partial_z T < 0 \end{cases}$$



$$\text{loss}(T, wT) = |NN(T) - wT|^2$$

Only okay, but why?

Good Engineering Principles: Integral Control!

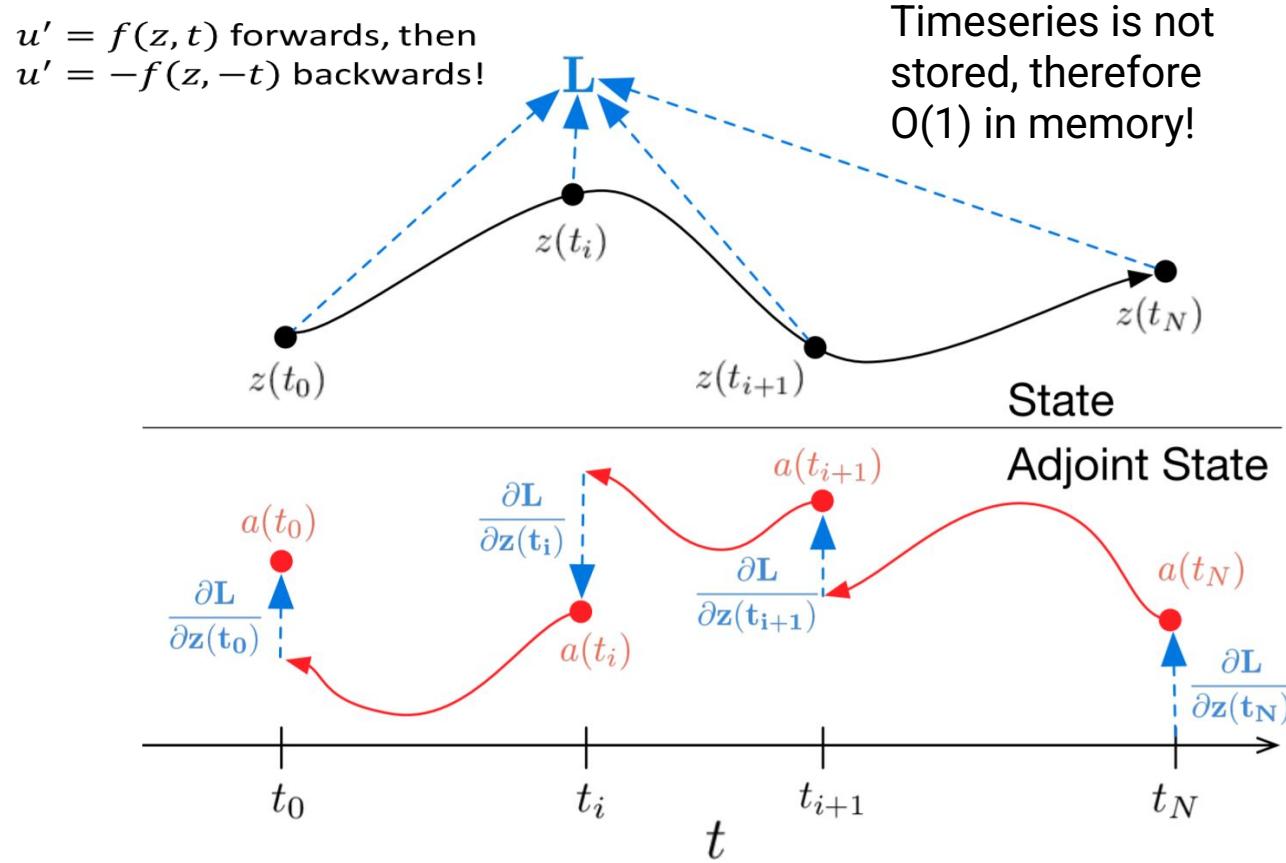


$$\frac{\partial T}{\partial t} = -\frac{\partial}{\partial z} \left(\underbrace{\left(\text{Input} \rightarrow \text{Hidden} \rightarrow \text{Output} \right)}_{w' T'} - K \frac{\partial T}{\partial z} \right)$$

$$loss(T_{NN}, T) = |T_{NN}(z, t) - T(z, t)|^2$$

**But differentiation is just automatic
differentiation... right?**

Machine Learning Neural Ordinary Differential Equations



The adjoint equation is an ODE!

$$\frac{d\mathbf{a}(t)}{dt} = -\mathbf{a}(t)^T \frac{\partial f(\mathbf{z}(t), t, \theta)}{\partial \mathbf{z}}$$

How do you get $\mathbf{z}(t)$? One suggestion:
Reverse the ODE

$$\frac{d\mathbf{a}_{aug}(t)}{dt} = - [\mathbf{a}(t) \quad \mathbf{a}_\theta(t) \quad \mathbf{a}_t(t)] \frac{\partial f_{aug}}{\partial [\mathbf{z}, \theta, t]}(t)$$

“Adjoints by reversing” also is unconditionally unstable on some problems!

Advection Equation:

$$\frac{\partial u}{\partial t} + \frac{a(\partial u)}{\partial x} = 0$$

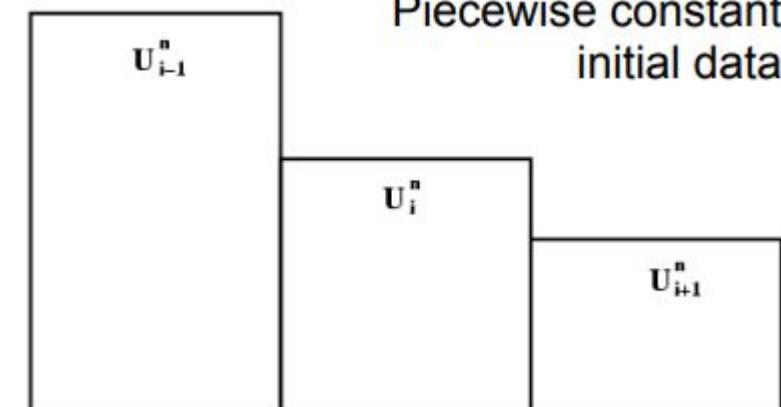
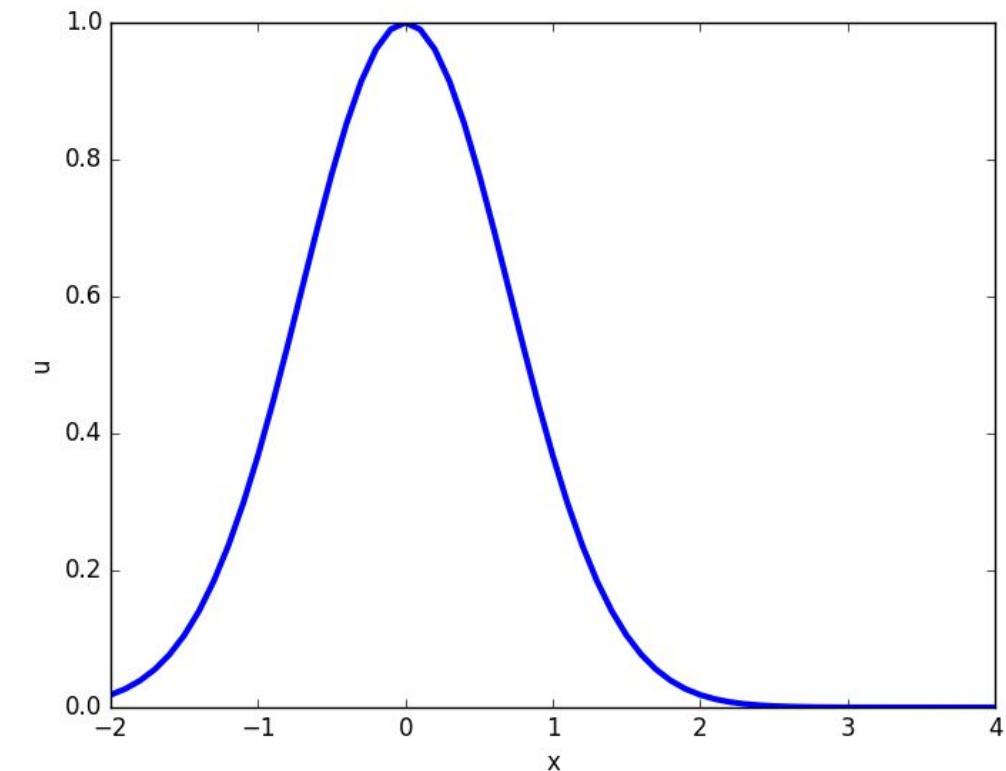
Approximating the derivative in x has two choices: forwards or backwards

$$u'_i = -\frac{a(u_i - u_{i-1})}{\Delta x} \text{ or } u'_i = -\frac{a(u_{i+1} - u_i)}{\Delta x}?$$

If you discretize in the wrong direction you get **unconditional instability**

You need to understand the engineering principles and the numerical simulation properties of domain to make ML stable on it.

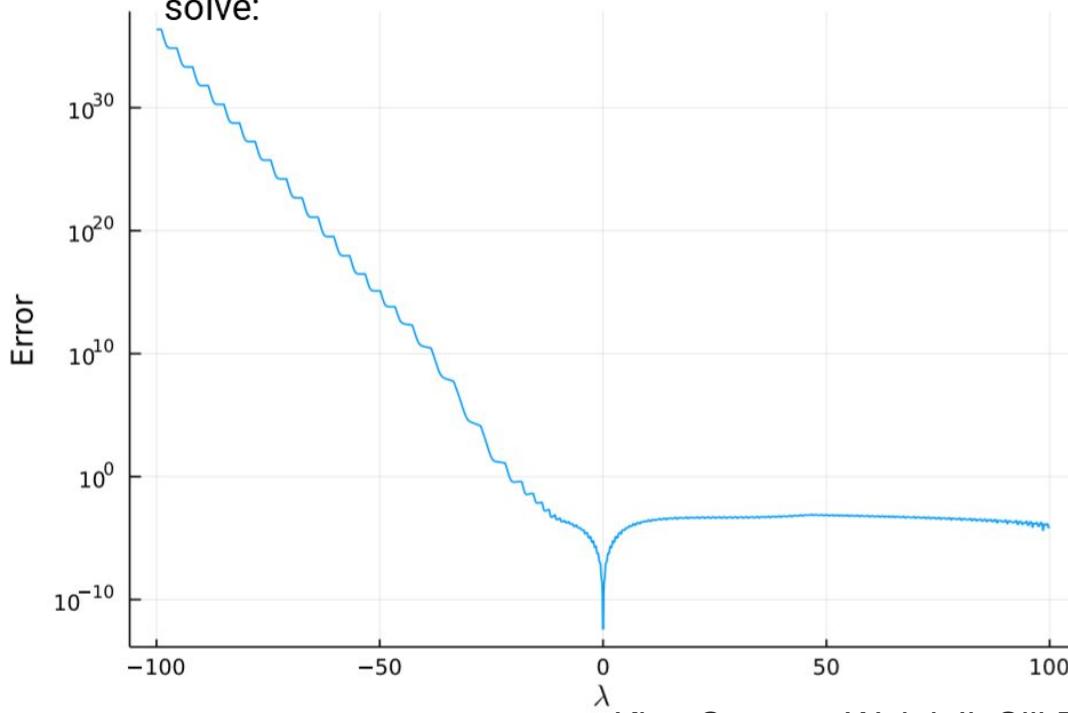
Result: the method in the PyTorch library is unconditionally unstable!!!



Problems With Naïve Adjoint Approaches On Stiff Equations

Error grows exponentially...

$u'(t) = \lambda u(t)$, plot the error in the reverse solve:

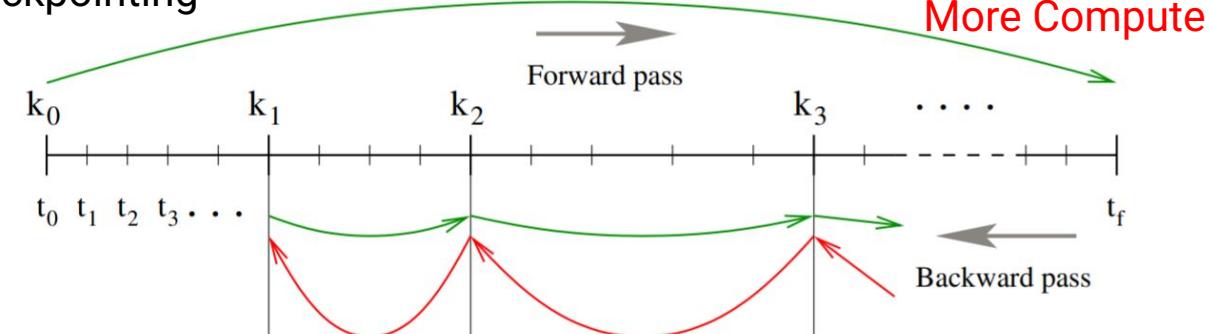


How do you get $u(t)$ while solving backwards?
3 options!

1. $u' = f(z, t)$ forwards, then
 $u' = -f(z, -t)$ backwards! **Unstable**

2. Store $u(t)$ while solving forwards (dense output) **High memory**

3. Checkpointing

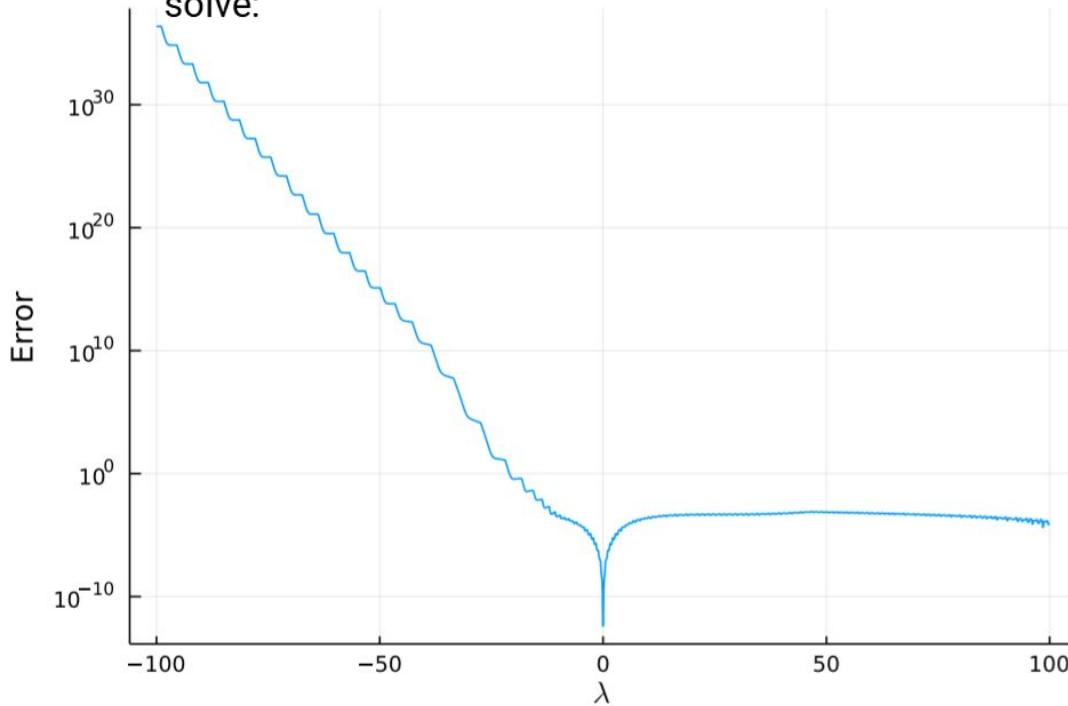


Each choice has an engineering trade-off!

Problems With Naïve Adjoint Approaches On Stiff Equations

Error grows exponentially...

$u'(t) = \lambda u(t)$, plot the error in the reverse solve:



Compute cost is cubic with parameter size when stiff

Size of reverse ODE system is:

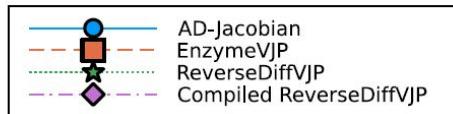
2states + parameters

Linear solves inside of stiff ODE solvers, ~cubic

Thus, adjoint cost:

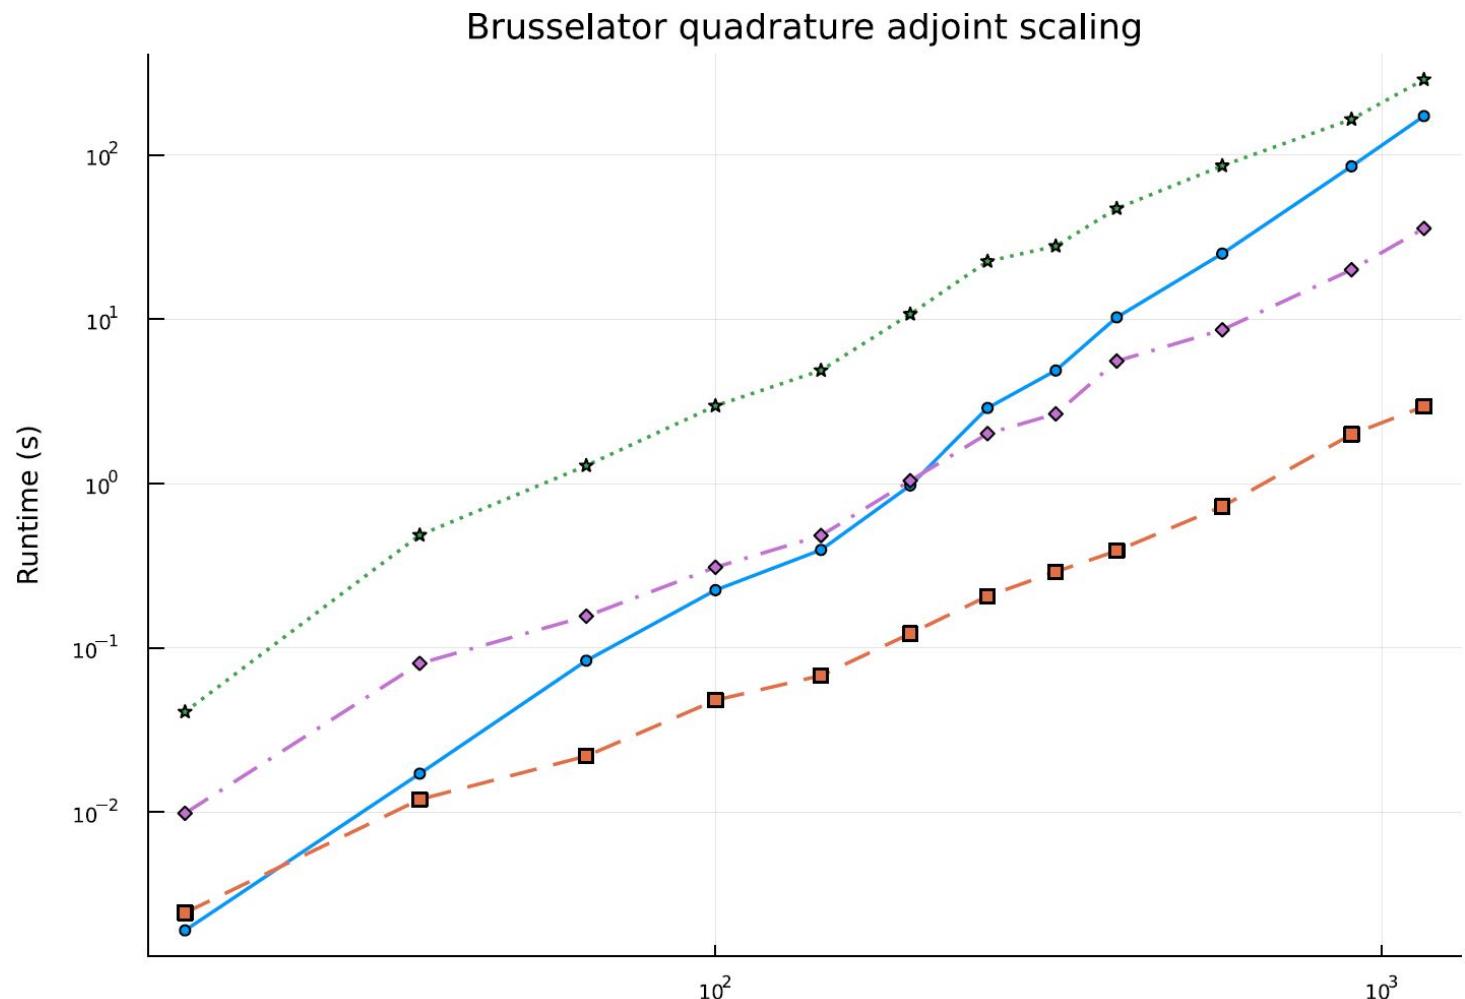
$$O((\text{states} + \text{parameters})^3)$$

Gives about a 10x performance improvement for adjoint calculations on PDEs



Ma, Yingbo, Vaibhav Dixit, Michael J. Innes, Xingjian Guo, and Chris Rackauckas. "A comparison of automatic differentiation and continuous sensitivity analysis for derivatives of differential equation solutions." In *2021 IEEE High Performance Extreme Computing Conference (HPEC)*, pp. 1-9. IEEE, 2021.

"Adjoint" can mean many
many different things...



The Test Problem

```
using OrdinaryDiffEq, ForwardDiff, Zygote, SciMLSensitivity

function f!(du, u, p, t)
    du[1] = p[1] * (u[1] - u[2])
    du[2] = p[2] * (u[2] - u[1])
end
p = [1.0, 1.0];
u0 = [1.0; 1.0];
prob = ODEProblem(f!, u0, (0.0, 5.0), p)
sol = solve(prob, Tsit5(), saveat = 0.1)

function cost(u0)
    _prob = remake(prob, u0 = u0)
    solve(_prob, Tsit5(), reltol = 1e-12, abstol = 1e-12, saveat = 0.1)[1, :]
end
```

Julia Derivative Results

```
function cost(u0)
    _prob = remake(prob, u0 = u0)
    solve(_prob, Tsit5(), reltol = 1e-12,
          abstol = 1e-12, saveat = 0.1)[1,
    end
    dx = ForwardDiff.jacobian(cost, u0)[end, :]
#=

2-element Vector{Float64}:
 11013.732897407262
 -11012.732897407262
=#
dx2 = Zygote.jacobian(cost, u0)[1][end, :]
#=

2-element Vector{Float64}:
 11013.732897407262
 -11012.732897407262
=#

```

All methods compute the same derivative to many digits

```
function cost(u0)
    _prob = remake(prob, u0 = u0)
    solve(_prob, Tsit5(), reltol = 1e-6,
          abstol = 1e-6, saveat = 0.1)[1, :]
end
dx = ForwardDiff.jacobian(cost, u0)[end, :]
#=

2-element Vector{Float64}:
 11013.73051284873
 -11012.73051284873
=#

dx2 = Zygote.jacobian(cost, u0)[1][end, :]
#=

2-element Vector{Float64}:
 11013.73051284873
 -11012.73051284873
=#

```

Jax Implementation of the Test Problem

```
import jax
import jax.numpy as jnp
from diffraz import diffeqsolve, ODETerm, Tsit5,
                    PIDController, BacksolveAdjoint,
                    RecursiveCheckpointAdjoint, DirectAdjoint

from jax import config
config.update("jax_enable_x64", True)

def vector_field(t, u, args):
    x, y = u
    a, b = args
    dx = a * (x - y)
    dy = b * (y - x)
    return dx, dy
```

Jax requires a lot more code

```
def run(y0, adjoint = RecursiveCheckpointAdjoint(), tol = 1e-12):
    term = ODETerm(vector_field)
    solver = Tsit5(scan_kind="bounded")
    stepsize_controller = PIDController(rtol=tol, atol=tol)
    t0 = 0
    t1 = 5.0
    dt0 = 0.1
    p0 = (1.0, 1.0)
    sol = diffeqsolve(term, solver, t0, t1, dt0, y0,
                      adjoint = adjoint,
                      stepsize_controller = stepsize_controller,
                      args=p0)
    ((x,), _) = sol.y
    return x
```

Jax Derivative Results: BacksolveAdjoint

```
J = jax.jacrev(lambda y0: run(y0, BacksolveAdjoint()))(y0)
# (Array(11013.73289742, dtype=float64, weak_type=True),
# Array(-11012.73289742, dtype=float64, weak_type=True))

J = jax.jacrev(lambda y0: run(y0, BacksolveAdjoint(), tol = 1e-3))(y0)
# (Array(10869.87401012, dtype=float64, weak_type=True),
# Array(-10868.87401012, dtype=float64, weak_type=True))

y1 = (jnp.array(1.000001), jnp.array(1.0))
y0 = (jnp.array(1.0), jnp.array(1.0))
(run(y1) - run(y0)) / .000001
# Array(11013.73156938, dtype=float64)
```

BacksolveAdjoint works here,
but it's the unstable one!

Jax Derivative Results: RecursiveCheckpointingAdjoint and DirectAdjoint

```
y0 = (jnp.array(1.0), jnp.array(1.0))
J = jax.jacrev(run)(y0)
# (Array(3755.79674193, dtype=float64, weak_type=True),
# Array(-3754.79674193, dtype=float64, weak_type=True))

J = jax.jacrev(lambda y0: run(y0, RecursiveCheckpointAdjoint(), tol = 1e-3))(y0)
# (Array(3755.79674193, dtype=float64, weak_type=True)
# Array(-3754.79674193, dtype=float64, weak_type=True))

J = jax.jacfwd(lambda y0: run(y0, DirectAdjoint()))(y0)
# (Array(3755.79674193, dtype=float64), Array(-3754.79674193, dtype=float64))

J = jax.jacfwd(lambda y0: run(y0, DirectAdjoint(), tol = 1e-3))(y0)
# (Array(3755.79674193, dtype=float64), Array(-3754.79674193, dtype=float64))
```

Jax's other methods don't converge, >60% error!

Conclusion:

The PyTorch library uses a numerically unstable method

The Jax library uses a non-convergent method.

The Julia library is numerically stable and convergent.

Calculating derivatives is hard.

But... any solver you can differentiate can do UDE things

Improving Coverage of Automatic Differentiation over Solvers

LinearSolve.jl: Unified Linear Solver Interface

$$A(n)x = l$$

NonlinearSolve.jl: Unified Nonlinear Solver Interface

$$f(u, p) = 0$$

DifferentialEquations.jl: Unified Interface for all
Differential Equations

$$u' = f(u, p)$$

$$du = f(u, p, t)dt + g(u, p, t)dx$$

⋮

Optimization.jl: Unified Optimization Interface

$$\text{minimize } f(u, p)$$

$$\text{subject to } g(u, p) \leq 0, h(u, p)$$

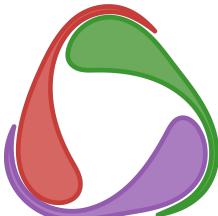
Integrals.jl: Unified Quadrature Interface

$$\int_{lb}^{ub} f(t, p) dt$$

Unified Partial Differential Equation Interface

$$u_t = u_{xx} + f(1)$$

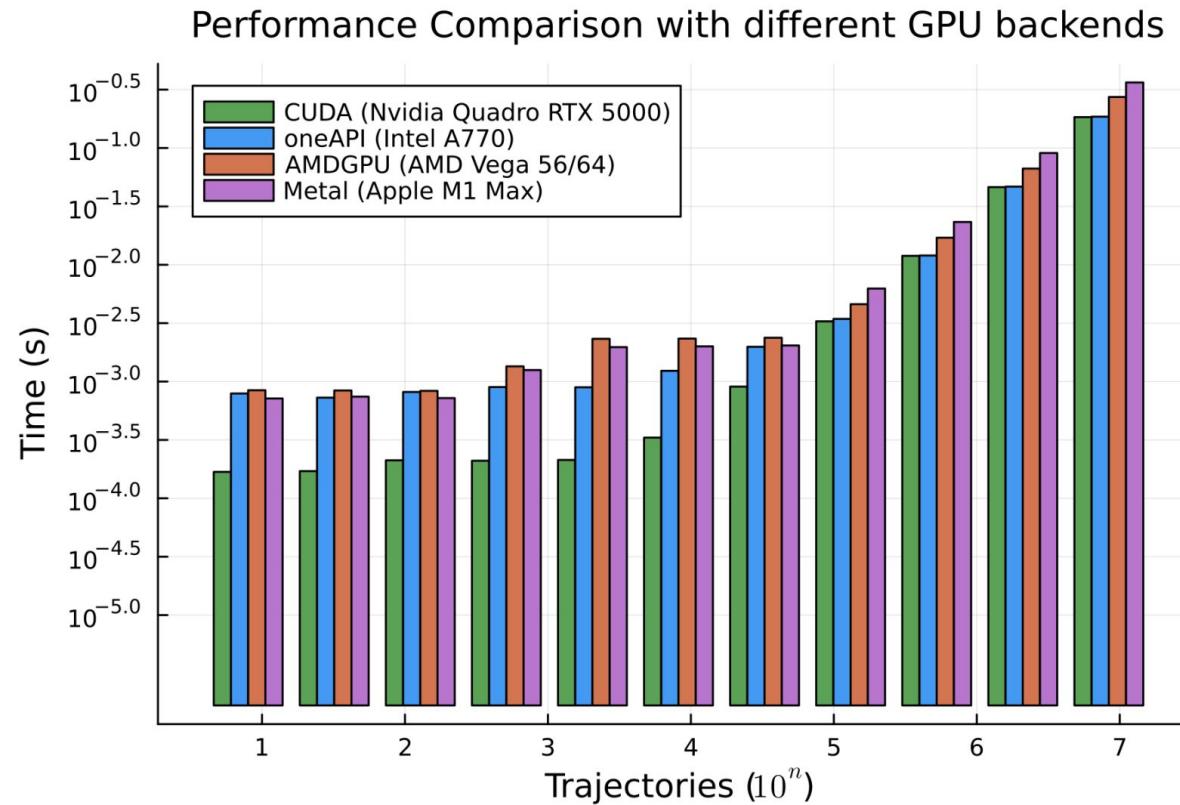
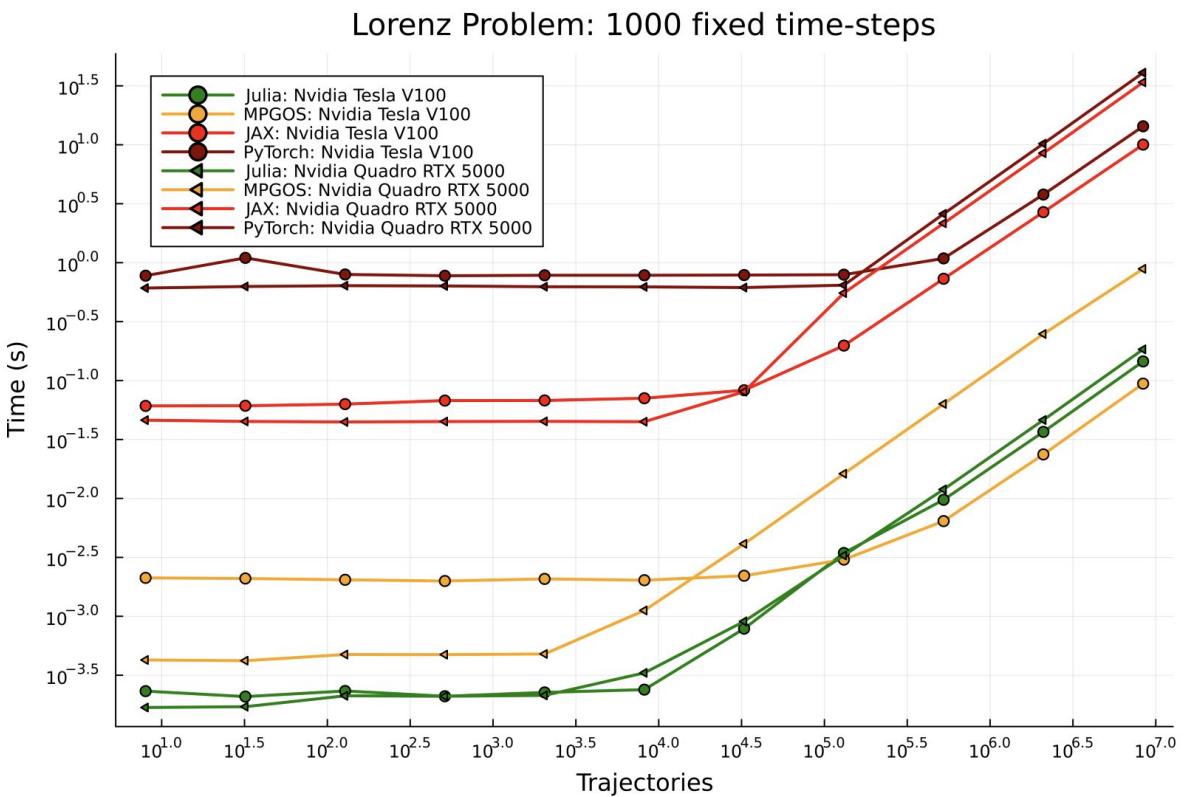
$$u_{tt} = u_{xx} + f(2)$$



The SciML Common Interface for Julia Equation Solvers

<https://scimlbase.sciml.ai/dev/>

New Parallelized GPU ODE Parallelism: 20x-100x Faster than Jax and PyTorch



Matches State of the Art on CUDA, but also works with AMD, Intel, and Apple GPUs

Utkarsh, U., Churavy, V., Ma, Y., Besard, T., Srisuma, P., Gymnich, T., ... & Rackauckas, C. (2024). Automated translation and accelerated solving of differential equations on multiple GPU platforms. *Computer Methods in Applied Mechanics and Engineering*, 419, 116591.

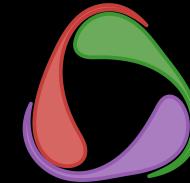
Foundation: Fast Differential Equation Solvers

1. Speed
2. Stability
3. Stochasticity
4. Adjoint and Inference
5. Parallelism



DifferentialEquations.jl is generally:

- 50x faster than SciPy
- 50x faster than MATLAB
- 100x faster than R's deSolve



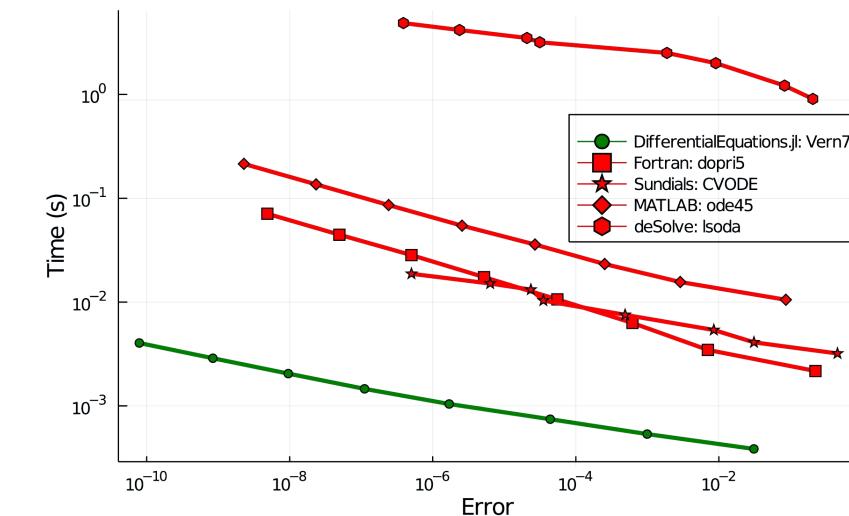
<https://github.com/SciML/SciMLBenchmarks.jl>

Rackauckas, Christopher, and Qing Nie. "DifferentialEquations.jl—a performant and feature-rich ecosystem for solving differential equations in julia." *Journal of Open Research Software* 5.1 (2017).

Rackauckas, Christopher, and Qing Nie. "Confederated modular differential equation APIs for accelerated algorithm development and benchmarking." *Advances in Engineering Software* 132 (2019): 1-6.

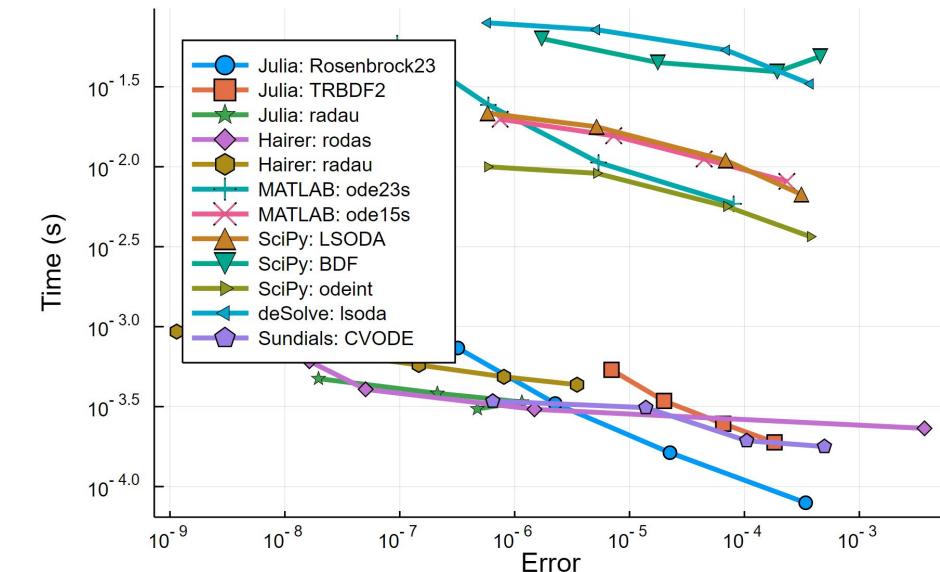
Non-Stiff ODE: Rigid Body System

Cross-Language ODE Solver Benchmark

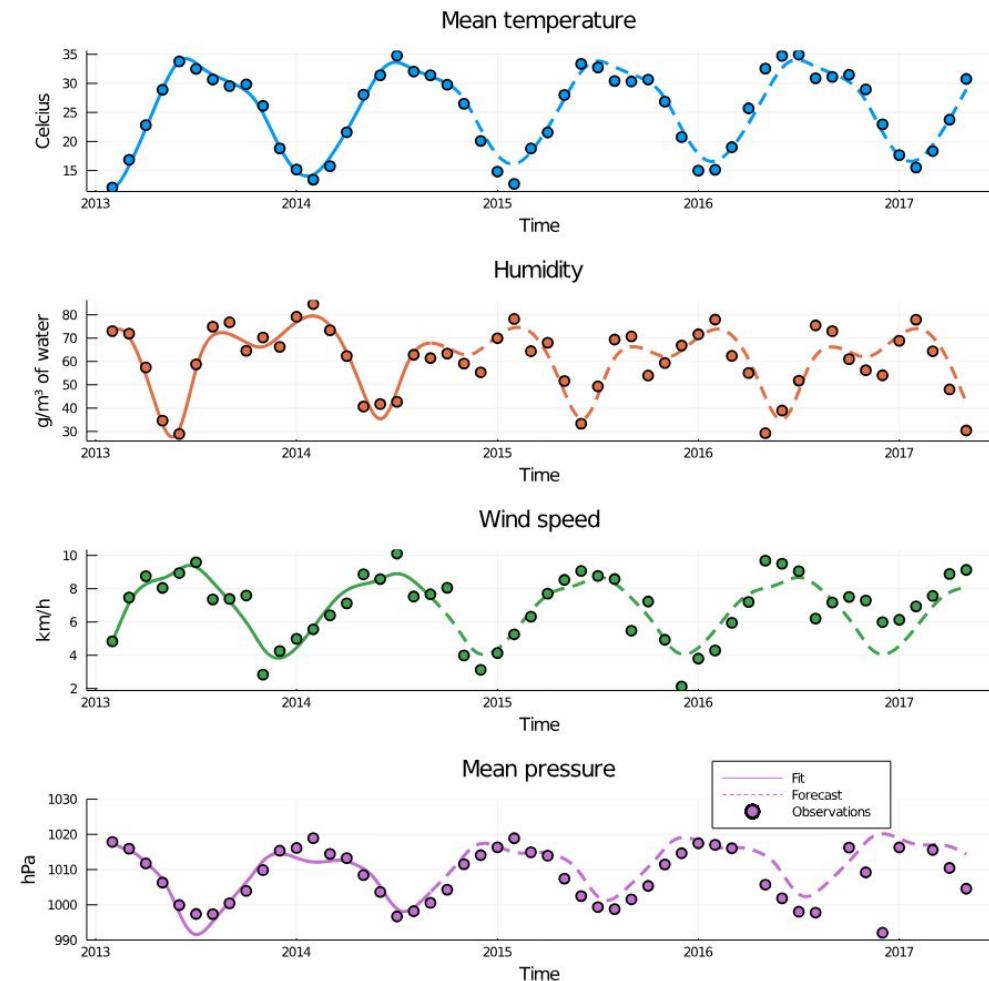
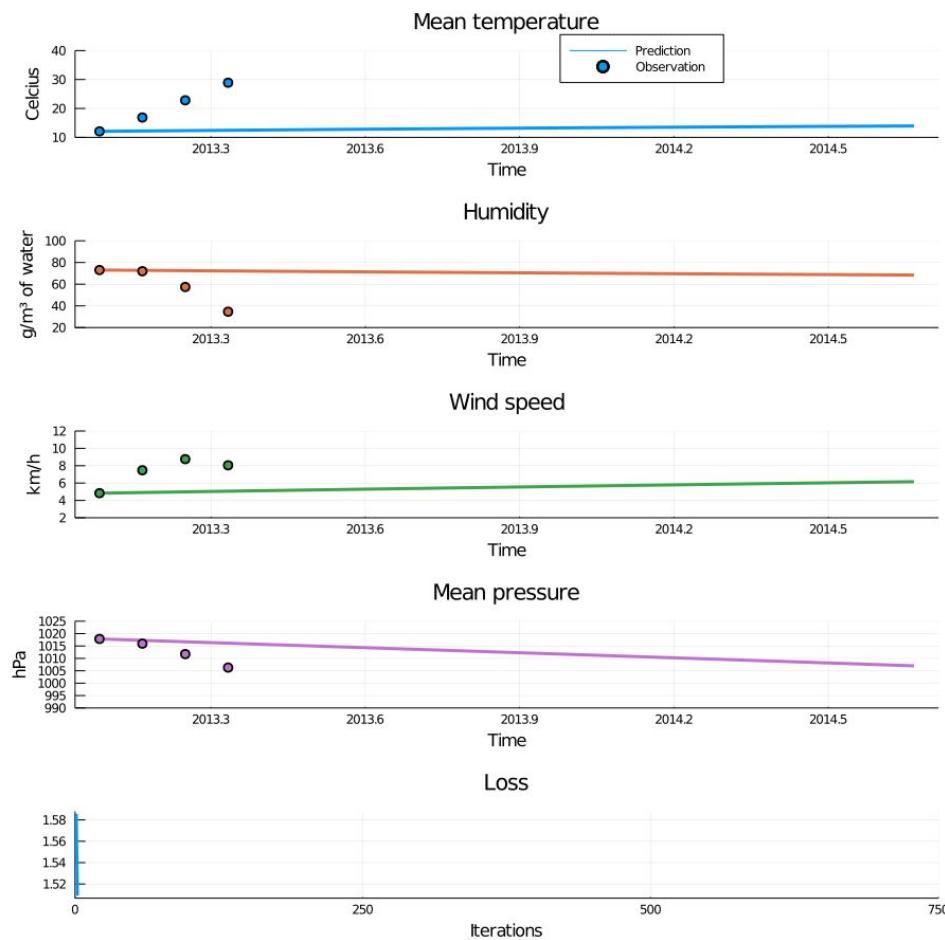


8 Stiff ODEs: Hires Chemical Reaction Network

Stiff 2: Hires

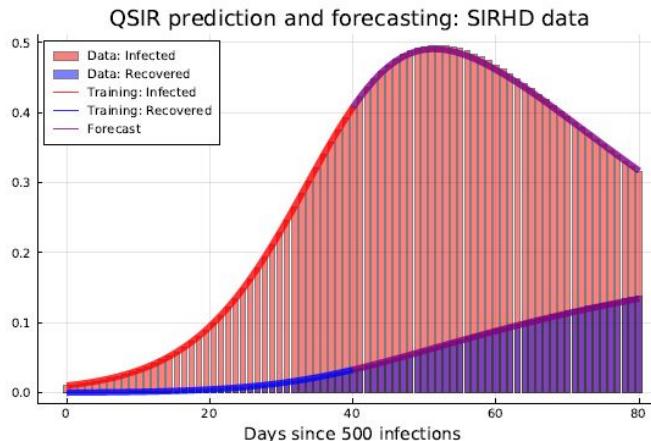


Implicit NNs are Good at Extrapolation

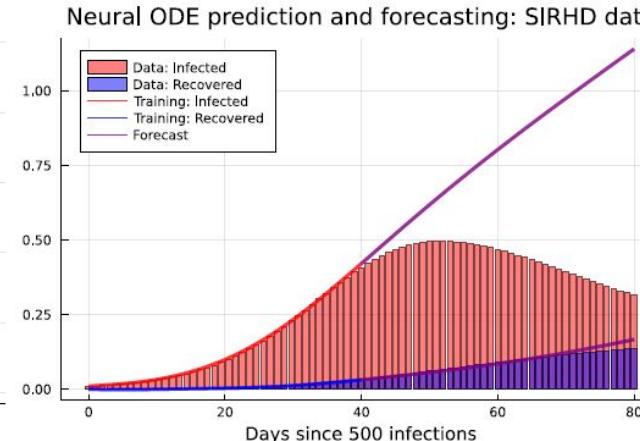


Graphics taken from <https://sebastiancallh.github.io/post/neural-ode-weather-forecast/> that was written using our open source softwares [Flux.jl](#), [DiffEqFlux.jl](#), & [SciMLSensitivity.jl](#)

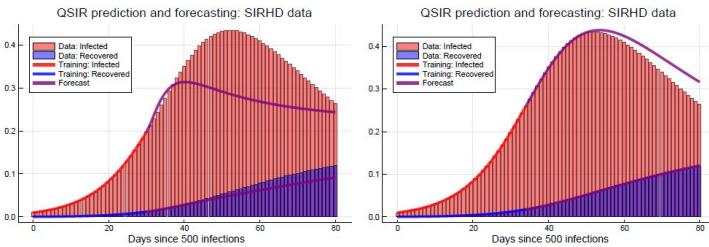
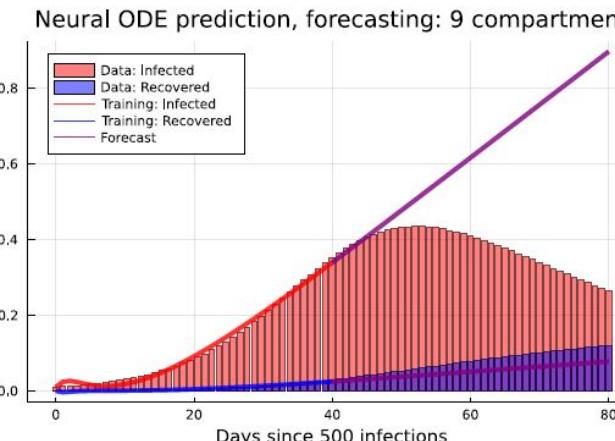
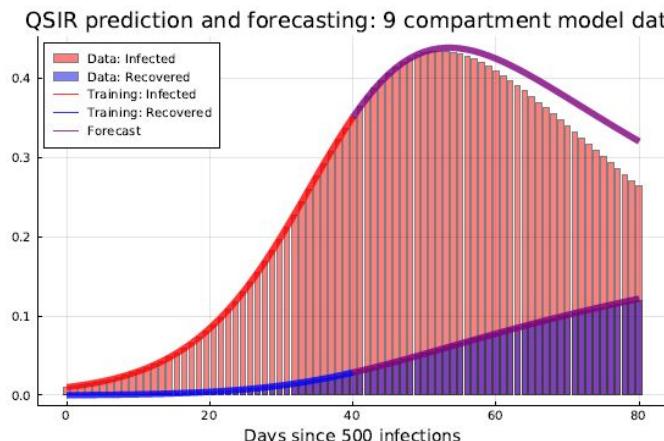
Scientific Machine Learning: Improving Predictions with Less Data



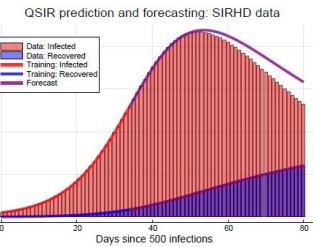
(a)



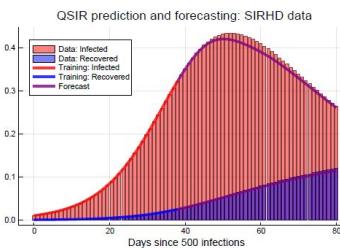
(b)



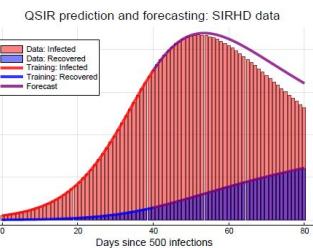
(a) train = 30 days



(b) train = 34 days



(c) train = 38 days



(d) train = 40 days

Dandekar, R., Rackauckas, C., & Barbastathis, G. (2020). A machine learning-aided global diagnostic and comparative tool to assess effect of quarantine control in COVID-19 spread. *Patterns*, 1(9), 100145.

Acuesta, E., Portone, T., Dandekar, R., Rackauckas, C., Bandy, R., & Huerta, J. (2022). *Model-Form Epistemic Uncertainty Quantification for Modeling with Differential Equations: Application to Epidemiology* (No. SAND2022-12823). Sandia National Lab.(SNL-NM), Albuquerque, NM (United States).

The UDE formulation fairly generally allows for imposing prior known structure

ODEs Effectively Recover Nonlinearities of Epidemic Models

The baseline case:

$$\frac{dS(t)}{dt} = -\frac{\tau_{SI} S(t) I(t)}{N}$$

$$\frac{dI(t)}{dt} = \frac{\tau_{SI} S(t) I(t)}{N} - \tau_{IR} I(t) - \tau_{ID} I(t)$$

$$\frac{dR(t)}{dt} = \tau_{IR} I(t)$$

$$\frac{dD(t)}{dt} = \tau_{ID} I(t).$$

Replacement of all terms with neural networks:

$$\frac{dS(t)}{dt} = -NN_{SI}$$

$$\frac{dI(t)}{dt} = NN_{SI} - NN_{IR} - NN_{ID}$$

$$\frac{dR(t)}{dt} = NN_{IR}$$

$$\frac{dD(t)}{dt} = NN_{ID}$$

Use SciML knowledge to constrain the interaction graph, but learn the nonlinearities!

Actual Equations	SINDY Active terms	SINDY Equations	Minimum AICC
------------------	--------------------	-----------------	--------------

NN_{SI}	0.85 S I	1: SI	0.74 S I	14
NN_{IR}	0.1 I	1: I	0.097 I	19
NN_{ID}	0.05 I	1: I	0.049 I	21

Table 4: SIRD: SINDY Recovered terms

Universal Differential Equations Predict Chemical Processes

$$\frac{\partial c}{\partial t^*} = -\frac{1-\varepsilon}{\varepsilon} \text{ANN}(q, q^*, \theta) - \frac{\partial c}{\partial x^*} + \frac{1}{Pe} \frac{\partial c^2}{\partial x^{*2}},$$

$$\frac{\partial q}{\partial t^*} = \text{ANN}(q, q^*, \theta),$$

$$\frac{\partial c(x^* = 1, \forall t)}{\partial x^*} = 0,$$

$$\frac{\partial c(x^* = 0, \forall t)}{\partial x^*} = Pe(c - c_{inlet}),$$

$$c(x^* \in (0, 1), t^* = 0) = c_0,$$

$$q(x^* \in (0, 1), t^* = 0) = q^*(c_0),$$

$$q^* = f(c, p),$$

Santana, V. V., Costa, E., Rebello, C. M., Ribeiro, A. M., Rackauckas, C., & Nogueira, I. B. (2023). Efficient hybrid modeling and sorption model discovery for non-linear advection-diffusion-sorption systems: A systematic scientific machine learning approach. *Chemical Engineering Science*

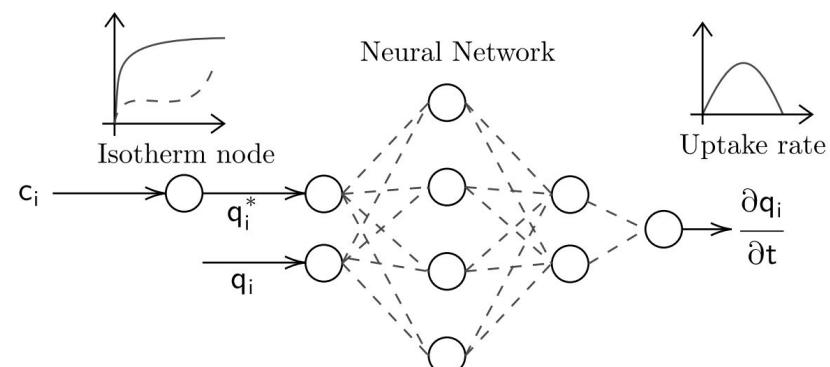
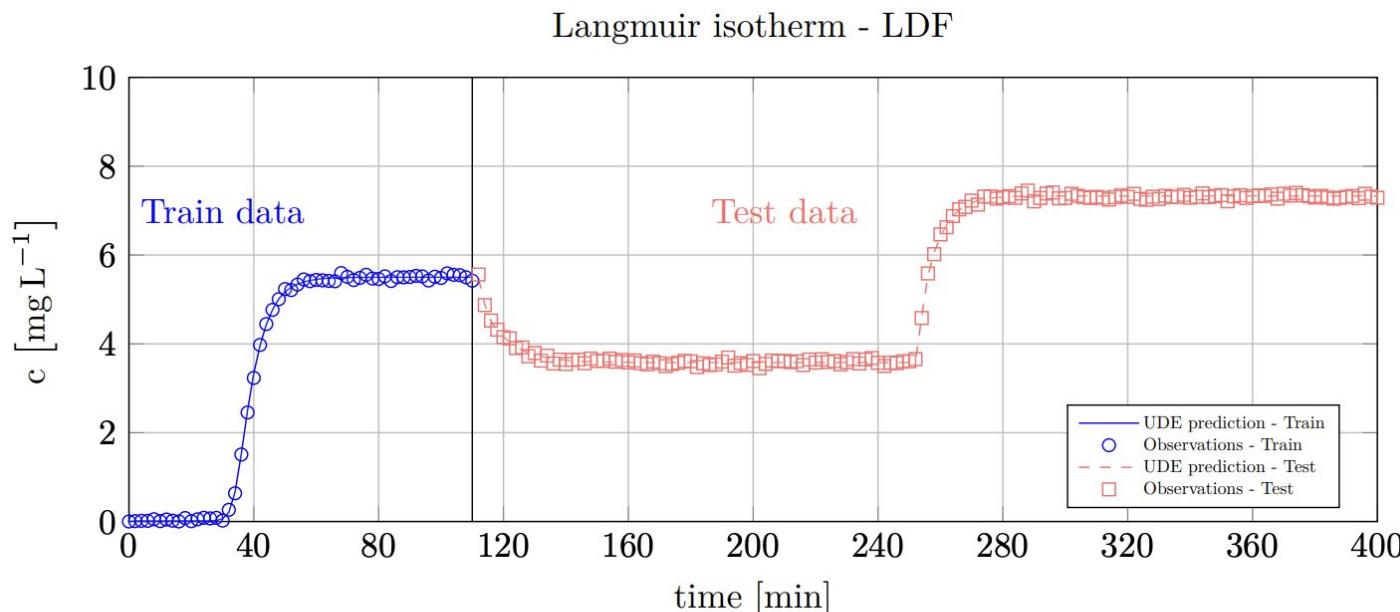


Figure 2: Schematic representation of the proposed hybrid model.

Universal Differential Equations Predict Chemical Processes

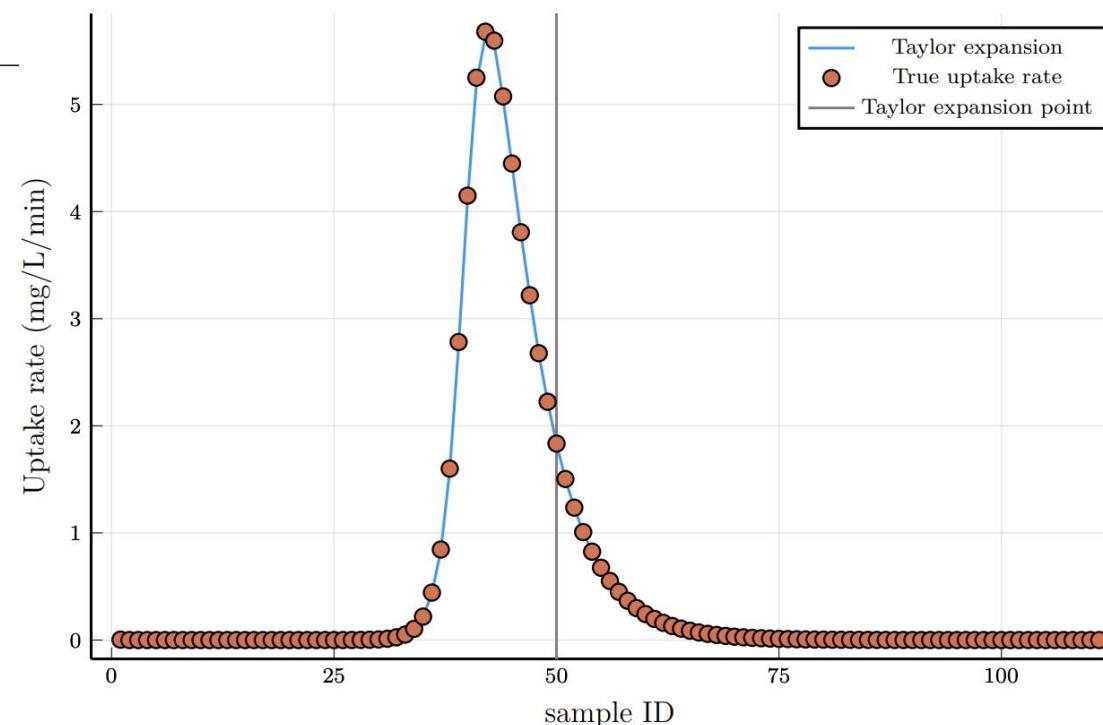
Table 5: Symbolic regression learned polynomials.

Isotherm	Kinetic	True kinetics	Learned kinetics
Langmuir	LDF	$0.22q^* - 0.22q$	$-0.535 - 0.225q + 0.234(q^*)$
Langmuir	improved LDF	$0.22(q^* + 0.2789q^*e^{\frac{-q}{2q^*}} - q)$	$-0.554 - 0.234q + 0.281(q^*)$
Langmuir	Vermeulen's	$0.22\frac{q^{*2}-q^2}{2.0q}$	$-0.6098 + 0.0122q + 0.263q^*$ $-0.00526qq^*$
Sips	LDF	$0.22q^* - 0.22q$	$0.198q^* - 0.200q$
Sips	improved LDF	$0.22(q^* + 0.2789q^*e^{\frac{-q}{2q^*}} - q)$	$0.277q^* - 0.241q$
Sips	Vermeulen's	$0.22\frac{q^{*2}-q^2}{2.0q}$	$-0.003557q^{*2} - 0.216q + 0.395q^*$

$$0.22(q^* + 0.2789q^*e^{\frac{-q}{2q^*}} - q)(49.23, 49.22) \approx 1.834 + 0.275q^* - 0.238q + \mathcal{O}(\|x^2\|)$$

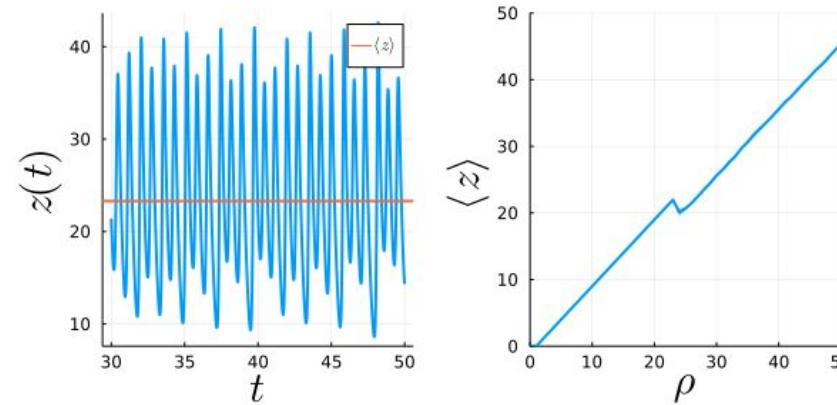
Santana, V. V., Costa, E., Rebello, C. M., Ribeiro, A. M., Rackauckas, C., & Nogueira, I. B. (2023). Efficient hybrid modeling and sorption model discovery for non-linear advection-diffusion-sorption systems: A systematic scientific machine learning approach. *Chemical Engineering Science*

Recovers equations with the same
2nd order Taylor expansion



**Can we generalize Differentiable
Simulation beyond continuous models?**

Differentiation of Chaotic Systems: Shadow Adjoints



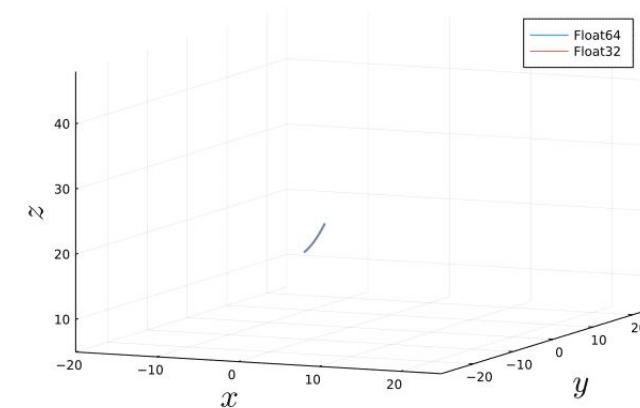
chaotic systems: trajectories diverge to $o(1)$ error ... but
shadowing lemma guarantees that the solution lies on
the attractor

$$\frac{d}{d\rho} \langle z \rangle_\infty \neq \lim_{T \rightarrow \infty} \frac{\partial}{\partial \rho} \langle z \rangle_T$$

- AD and finite differencing fails!

$$\left. \frac{d \langle z \rangle_\infty}{d \rho} \right|_{\rho=28} \approx -49899 \text{ (ForwardDiff)}$$

$$\left. \frac{d \langle z \rangle_\infty}{d \rho} \right|_{\rho=28} \approx 472 \text{ (Calculus)}$$



- Shadowing methods in DiffEqSensitivity.jl

$$\left. \frac{d \langle z \rangle_\infty}{d \rho} \right|_{\rho=28} \approx 1.028 \text{ (LSS/AdjointLSS)}$$

$$\left. \frac{d \langle z \rangle_\infty}{d \rho} \right|_{\rho=28} \approx 0.997 \text{ (NILSS)}$$

Differentiation of Chaotic Systems: Shadow Adjoints

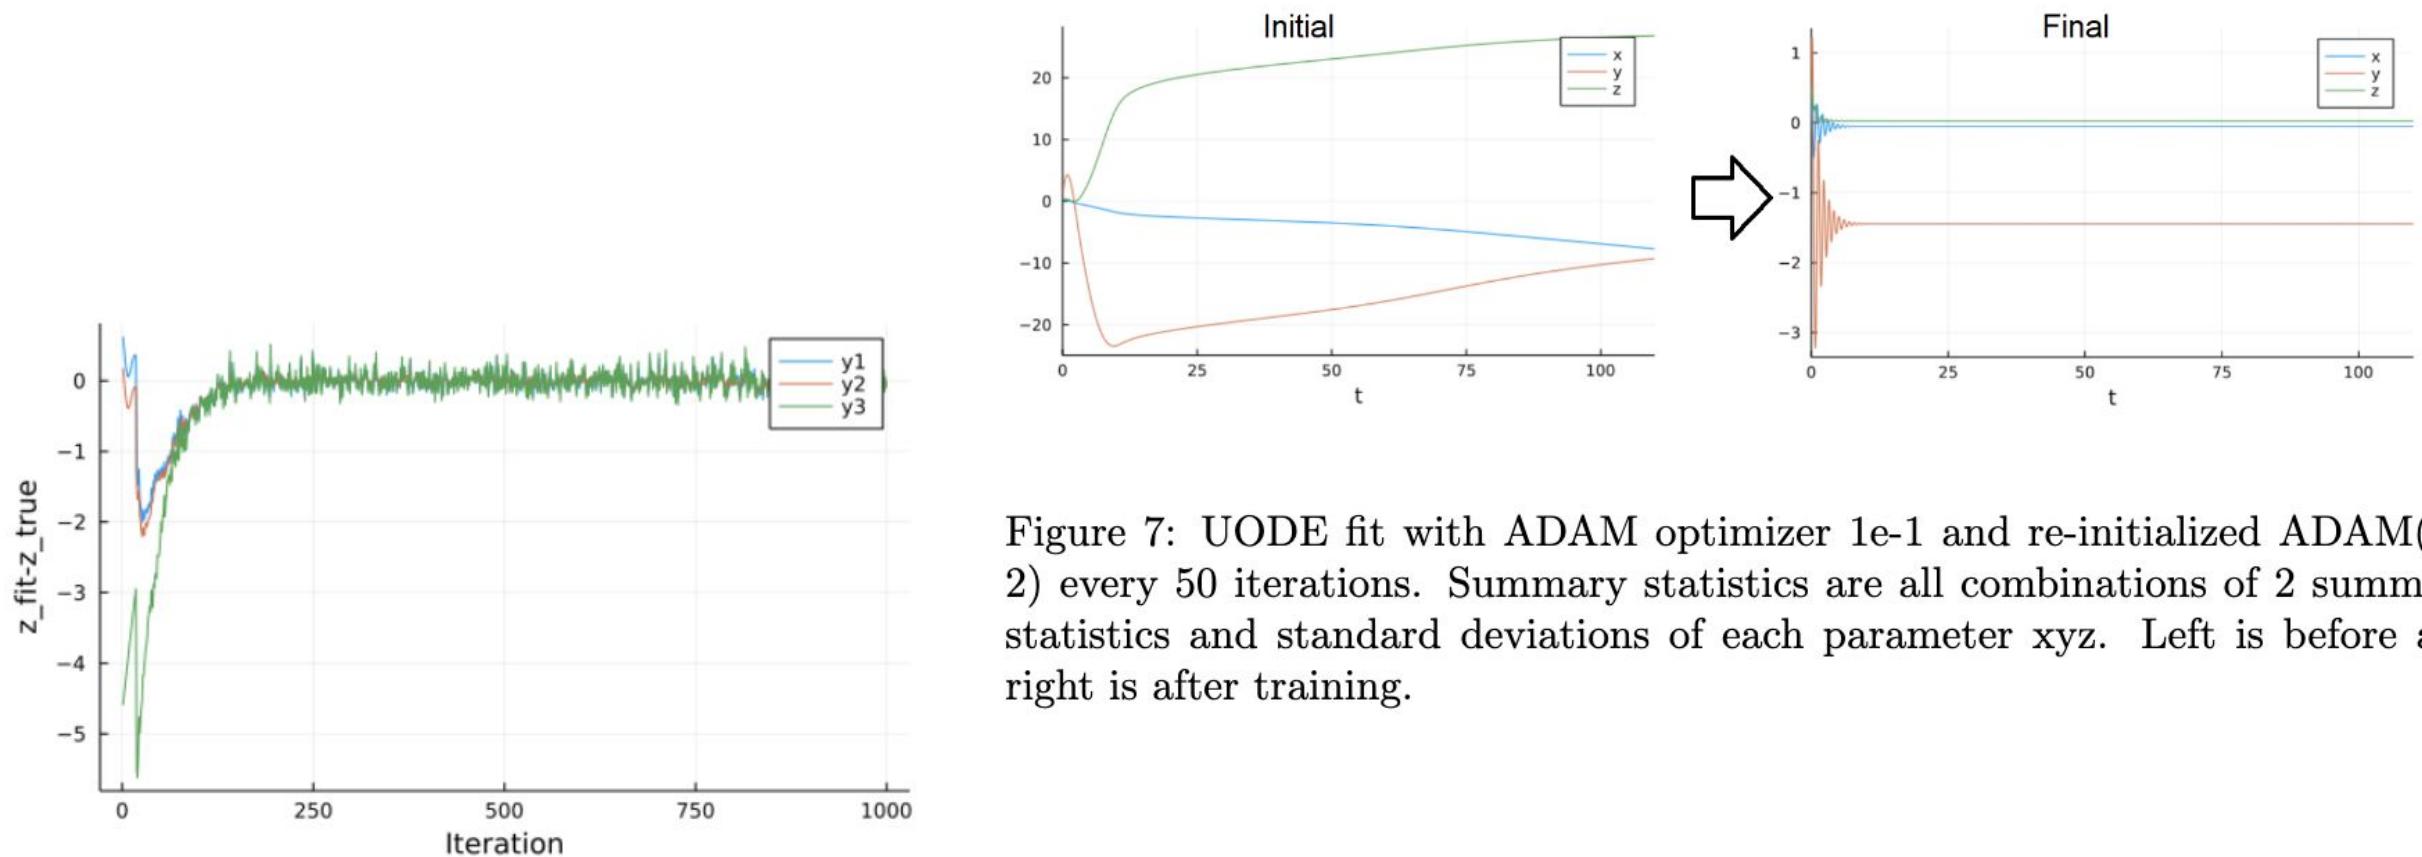


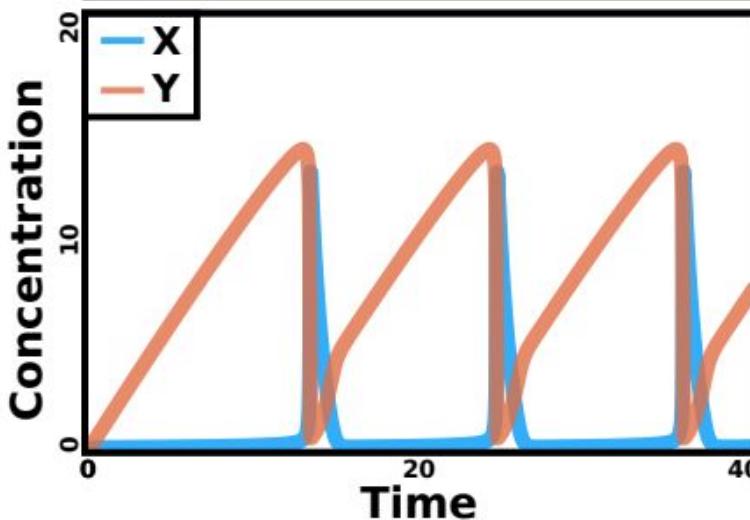
Figure 7: UODE fit with ADAM optimizer 1e-1 and re-initialized ADAM(1e-2) every 50 iterations. Summary statistics are all combinations of 2 summary statistics and standard deviations of each parameter xyz. Left is before and right is after training.

Figure 18: Summary statistic fitting corresponding to Figure 17. Fitting with all 3 Lorenz parameters free using 3 summary statistics ($\langle |x| \rangle$, $\langle |y| \rangle$, $\langle |z| \rangle$). ADAM optimizer learning rate is 1e-1.

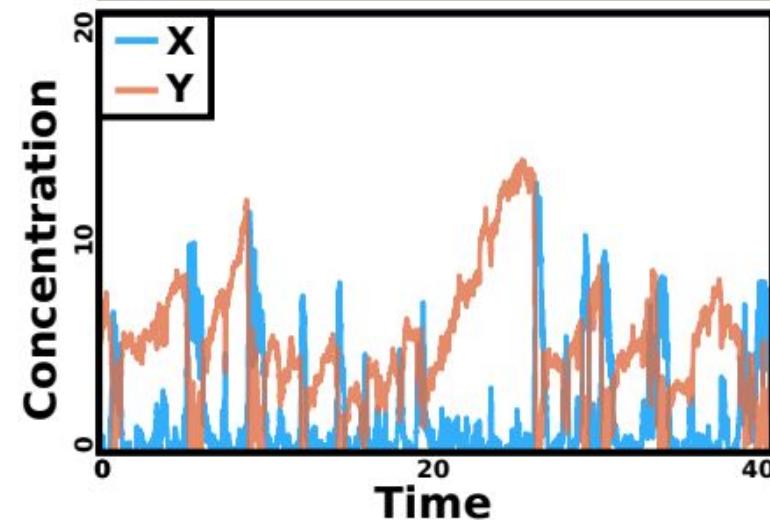
Poisson Jump Equations: Stochastic Chemical Reactions

A

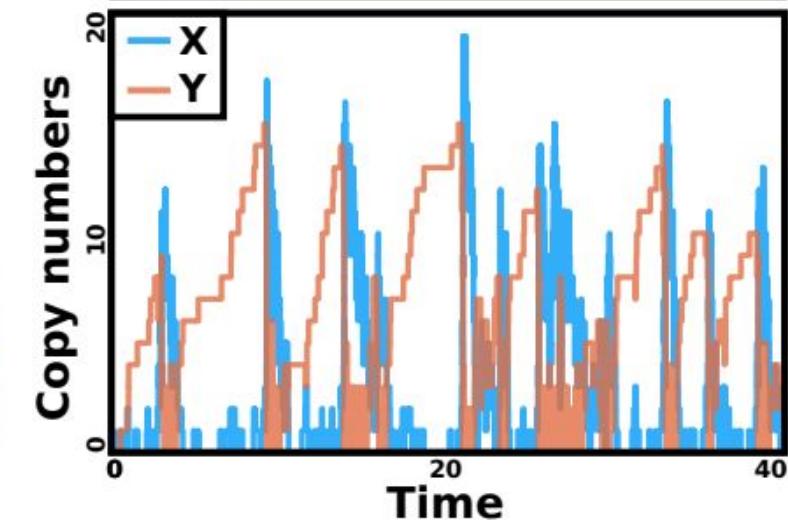
```
using OrdinaryDiffEq
u0 = [:X => 0., :Y => 0.]
tspan = (0., 60.)
p = [:A => 1.0, :B => 4.0]
oprob = ODEProblem(brusselator, u0, tspan, p)
sol = solve(oprob, Rosenbrock23())
plot(sol)
```

**B**

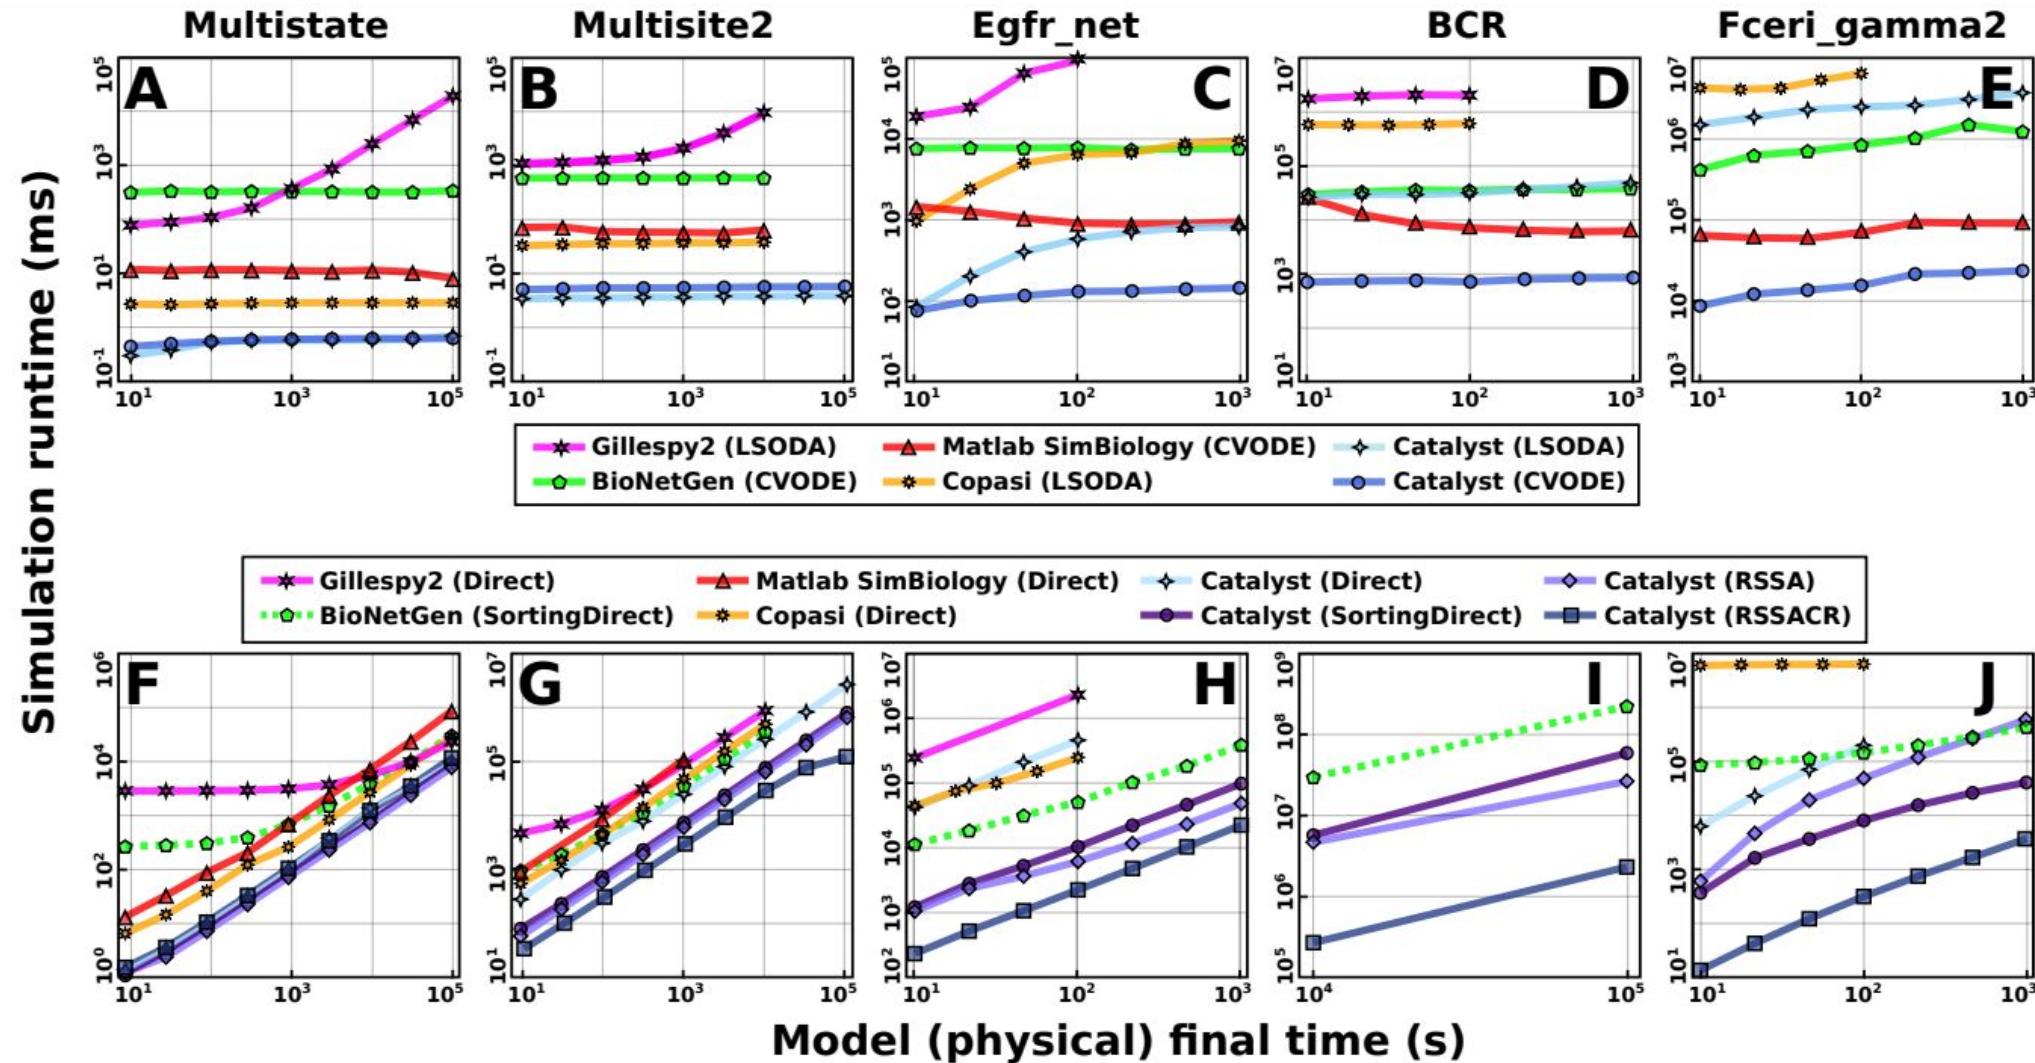
```
using StochasticDiffEq
u0 = [:X => 0., :Y => 5.]
tspan = (0., 60.)
p = [:A => 1.0, :B => 4.0]
sprob = SDEProblem(brusselator, u0, tspan, p)
sol = solve(sprob, ImplicitEM())
plot(sol)
```

**C**

```
using DiffEqJump
u0 = [:X => 0, :Y => 0]
tspan = (0., 60.)
p = [:A => 3.0, :B => 4.0]
dprob = DiscreteProblem(brusselator, u0, tspan, p)
jprob = JumpProblem(brusselator, dprob, Direct())
sol = solve(jprob, SSAStrong())
plot(sol)
```



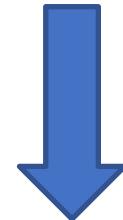
~100x performance improvements in Gillespie SSAs with Catalyst.jl



New Automatic Differentiation of Discrete Stochastic Spaces?

Traditional AD

$$f(p)$$



$$\frac{df(p)}{dp}$$

Discrete Stochastic AD

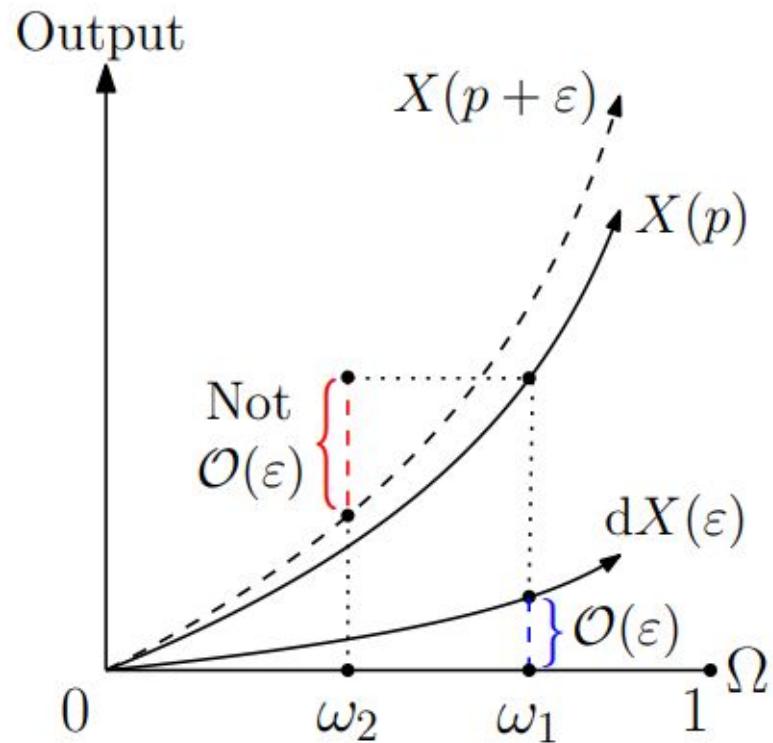
$$X(p)$$



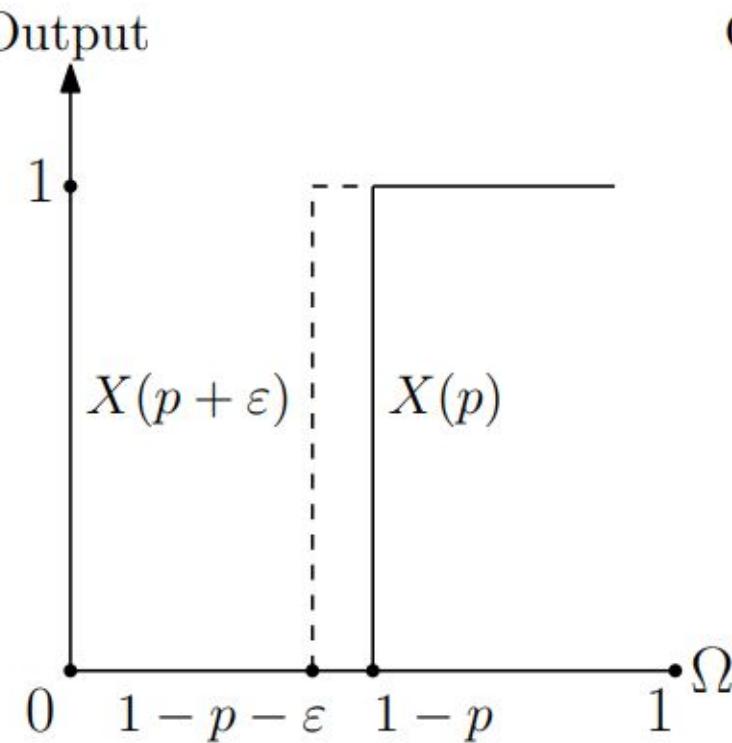
$$\mathbb{E}[\tilde{X}(p)] = \frac{d\mathbb{E}[X(p)]}{dp}$$

Automatic Differentiation of Programs with Discrete Randomness
Accepted to NeurIPS 2022!

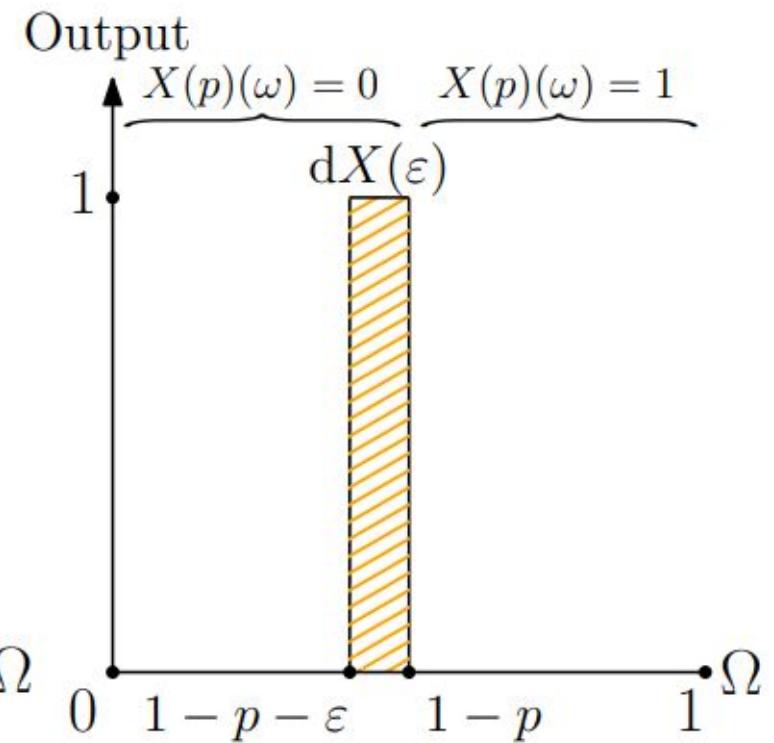
Infinitesimal Changes -> O(1) Changes



(a) $X(p) \sim \text{Exp}(p)$.

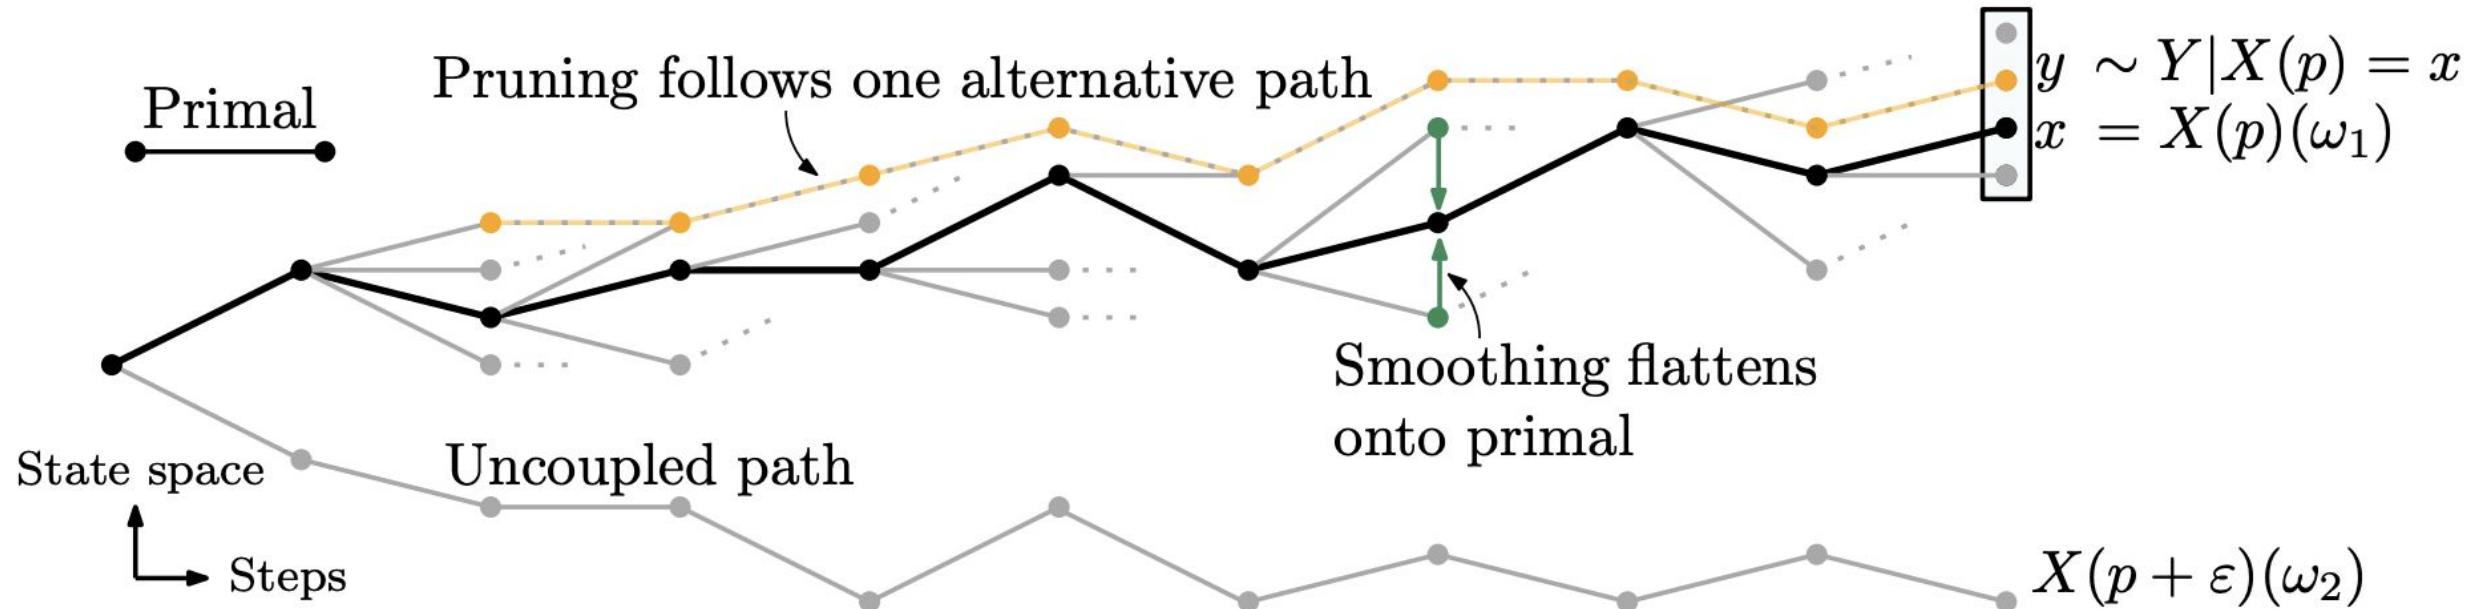


(b) $X(p) \sim \text{Ber}(p)$.



(c) $X(p) \sim \text{Ber}(p)$.

$O(1)$ changes with infinitesimal probability



StochasticAD.jl: Unbiased Stochastic Derivatives with Linear Cost

Perform AD
on
Agent-Based
models,
particle
filters,
chemical
reactions,
and more!

```
1 struct StochasticTriple
2     value # primal evaluation
3     δ # "infinitesimal" component
4     Δs # component of discrete change
5         # with "infinitesimal"
6         # probability
7 end
```

```
1 julia> using StochasticAD
2 julia> st = stochastic_triple(X, 0.6) # sample a single stochastic triple at p = 0.6
3 2.16 + 10.05ε + (0.64 with probability 12.5ε)
4 julia> derivative_contribution(st) # which produces a single derivative estimate...
5 16.79
6 julia> samples = [derivative_estimate(X, 0.6) for i in 1:1000] # take many estimates!
7 julia> println("d/dp of E[X(p)]: $(mean(samples)) ± $(std(samples) / sqrt(1000))")
8 d/dp of E[X(p)]: 19.39 ± 0.48
```

```
1 using Distributions
2 function X(p) p = 0.6 + ε
3     a = p^2 0.36 + 1.2ε
4     b = rand(Binomial(10, p))
5     6 + (1 with probability 10.0ε)
6     c = 2 * b + 3 * rand(Bernoulli(p))
7     12 + (3 with probability 12.5ε)
8     return a * c * rand(Normal(p, p^2))
9 end
```

StochasticAD.jl on Agent-Based Models

(a)

```
...
X = 0
for step in 1:n
    i = rand(Categorical(probs(X)))
    X += steps[i]
end
return f(X)
```

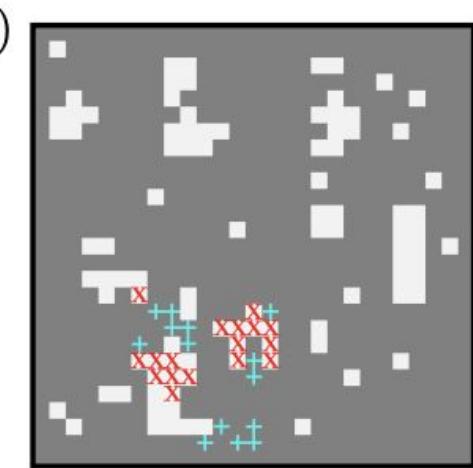
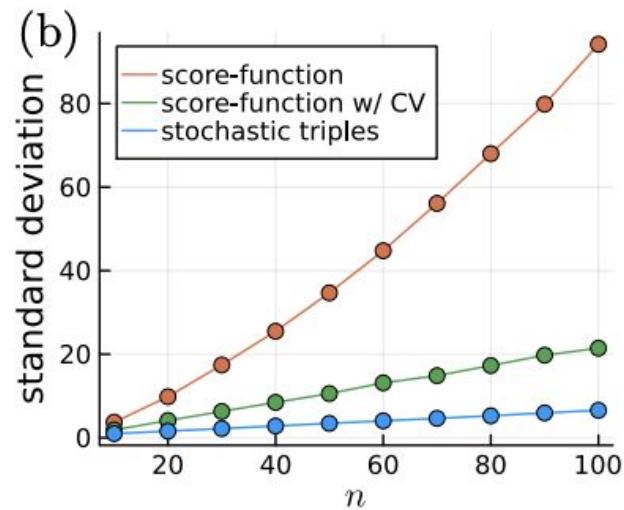


Figure 4: Automatic differentiation of discrete-time Markov processes. (a) Code snipped for a 1D random walk, which can be automatically differentiated by StochasticAD.jl; `probs(X)` gives the transition probabilities at value `X` and `steps[i]` gives the step size for the `i`th transition. (b) The variance of unbiased gradient estimates of the random walk program using stochastic triples and the score function, which is applied both without a control variate (CV) and with a pre-computed batch-average CV. (c) The final board for one run of the stochastic Game of Life with $N = 25$ and $T = 10$, where the “+” signs represent additional living cells (white) in the stochastic alternative path, and the “X” signs represent additional dead cells (grey).

But wait, why not PINNs?

**How does differentiable simulation compare
to PINNs?**

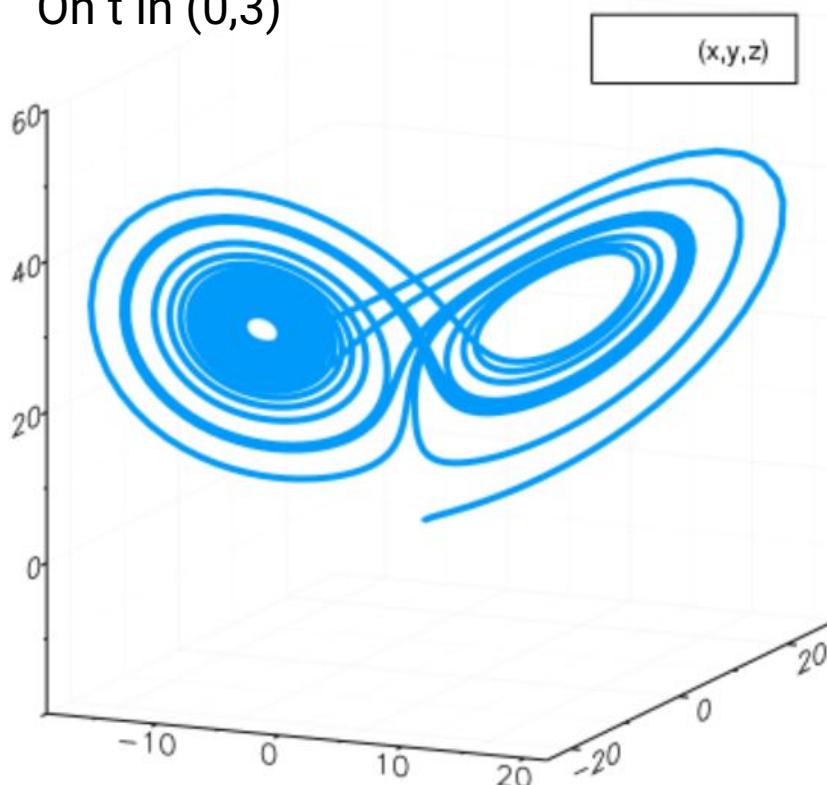


I wanna go fast

Keeping Neural Networks Small Keeps Speed For Inverse Problems

DeepXDE (TensorFlow Physics-Informed NN)

Problem: parameter estimation
of Lorenz equation from data
On t in (0,3)



```
Best model at step 57000:  
train loss: 5.91e-03  
test loss: 5.86e-03  
test metric: []
```

'train' took 362.351454 s

DiffEqFlux.jl (Julia UDEs)

```
opt = Opt(:LN_BOBYQA, 3)  
lower_bounds!(opt,[9.0,20.0,2.0])  
upper_bounds!(opt,[11.0,30.0,3.0])  
min_objective!(opt, obj_short.cost_function2)  
xtol_rel!(opt,1e-12)  
maxeval!(opt, 10000)  
@time (minf,minx,ret) = NLopt.optimize(opt,LocIniPar) # 0.1 seconds
```

```
0.032699 seconds (148.87 k allocations: 14.175 MiB)  
(2.7636309213683456e-18, [10.0, 28.0, 2.66], :XTOL_REACHED)
```

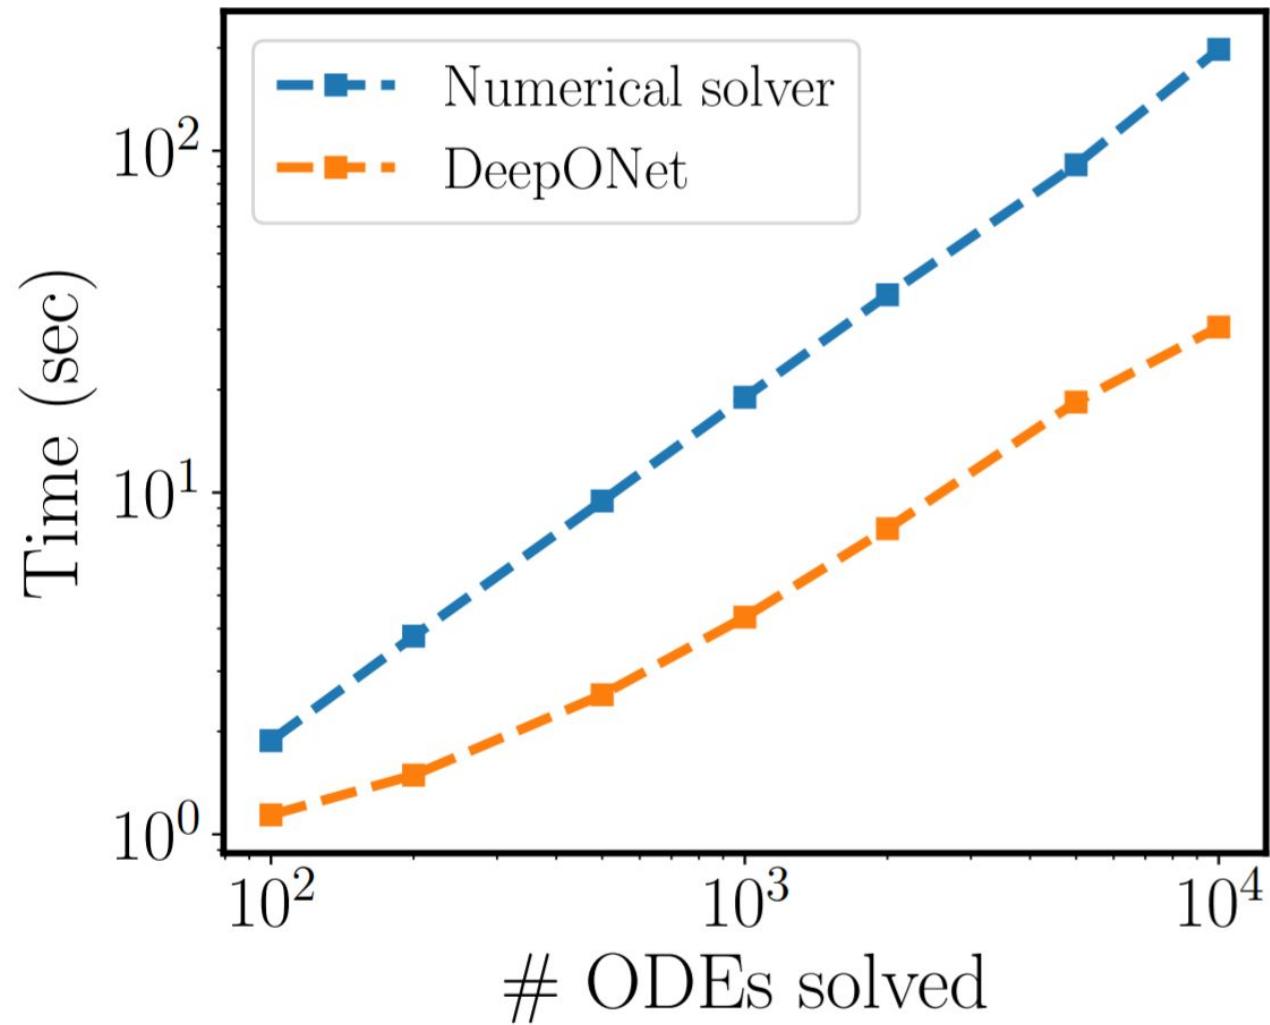
Note on Neural Networks “Outperforming” Classical Solvers

Long-time integration of parametric evolution equations with physics-informed DeepONets

Sifan Wang, Paris Perdikaris

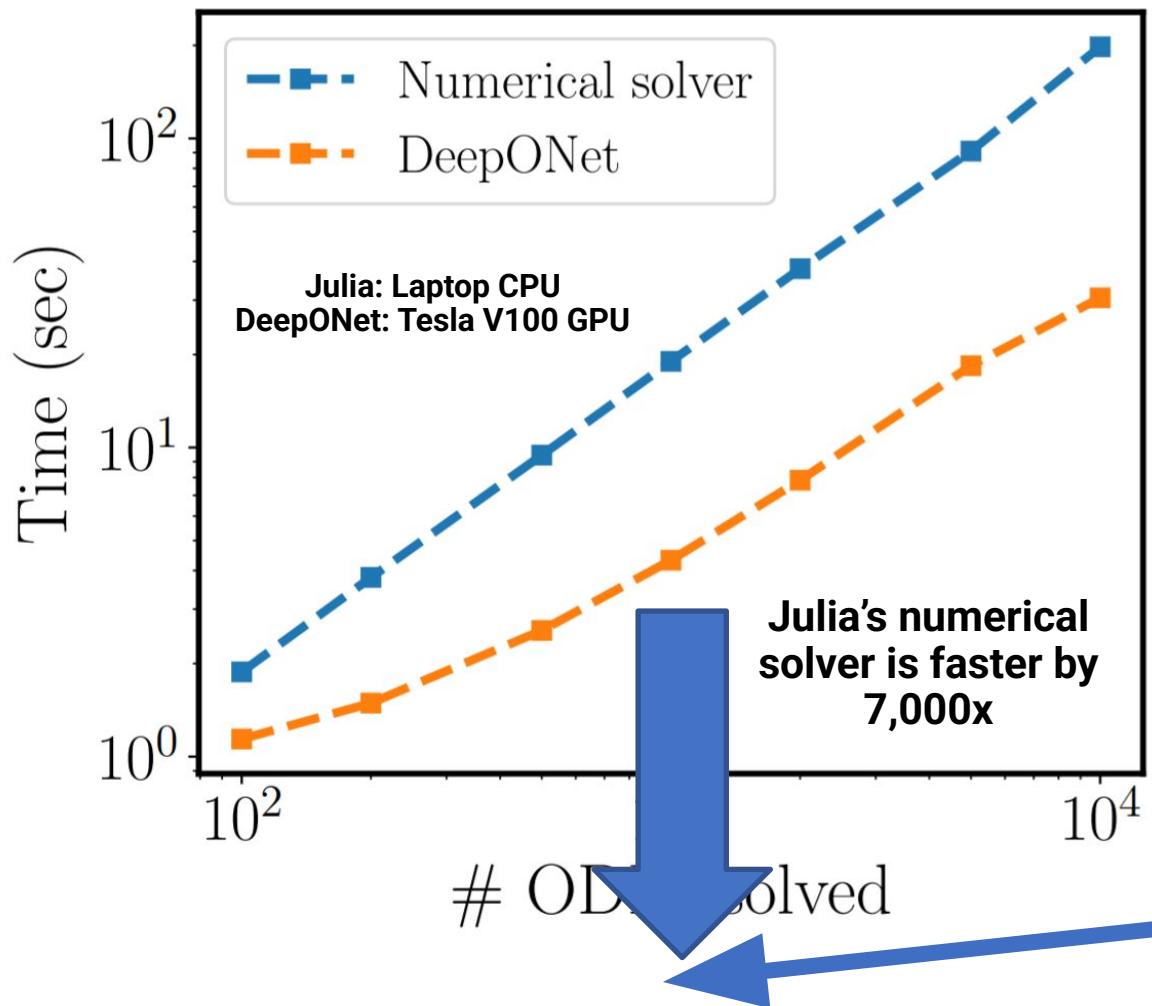
Ordinary and partial differential equations (ODEs/PDEs) play a paramount role in analyzing and simulating complex dynamic processes across all corners of science and engineering. In recent years machine learning tools are aspiring to introduce new effective ways of simulating PDEs, however existing approaches are not able to reliably return stable and accurate predictions across long temporal horizons. We aim to address this challenge by introducing an effective framework for learning infinite-dimensional operators that map random initial conditions to associated PDE solutions within a short time interval. Such latent operators can be parametrized by deep neural networks that are trained in an entirely self-supervised manner without requiring any paired input-output observations. Global long-time predictions across a range of initial conditions can be then obtained by iteratively evaluating the trained model using each prediction as the initial condition for the next evaluation step. This introduces a new approach to temporal domain decomposition that is shown to be effective in performing accurate long-time simulations for a wide range of parametric ODE and PDE systems, from wave propagation, to reaction-diffusion dynamics and stiff chemical kinetics, all at a fraction of the computational cost needed by classical numerical solvers.

Note on Neural Networks “Outperforming” Classical Solvers



Oh no, we're doomed!

Stay tuned...



```
using ModelingToolkit, OrdinaryDiffEq, StaticArrays

@variables t y₁(t) y₂(t) y₃(t)
@parameters k₁ k₂ k₃
D = Differential(t)
eqs = [D(y₁) ~ -k₁*y₁+k₃*y₂*y₃
       D(y₂) ~ k₁*y₁-k₂*y₂^2-k₃*y₂*y₃
       D(y₃) ~ k₂*y₂^2]

sys = ODESystem(eqs, t)
prob = ODEProblem{false}(sys, SA[y₁=>1f0, y₂=>0f0, y₃=>0f0], (0f0, 500f0),
                        SA[k₁=>4f-2, k₂=>3f7, k₃=>1f4], jac=true)

N = 1000
y₁s = rand(Float32, N)
y₂s = 1f-4 .* rand(Float32, N)
y₃s = rand(Float32, N)

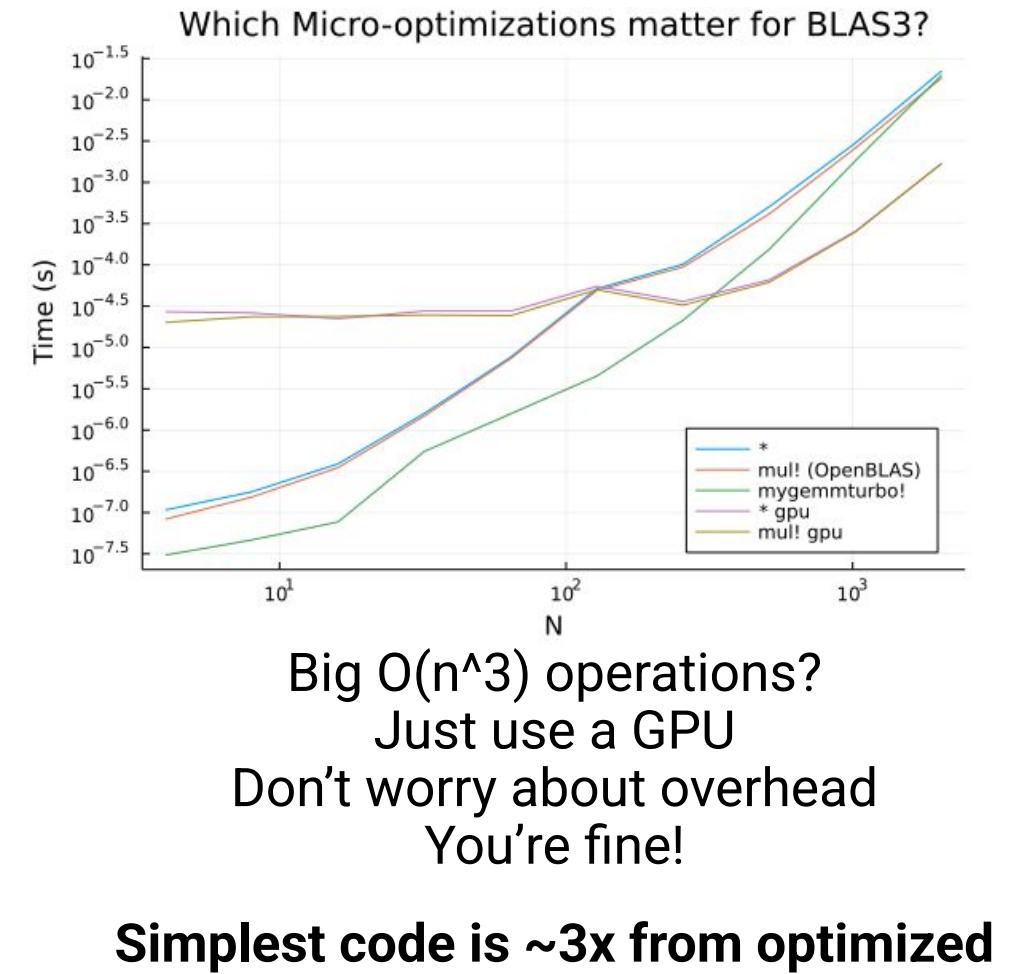
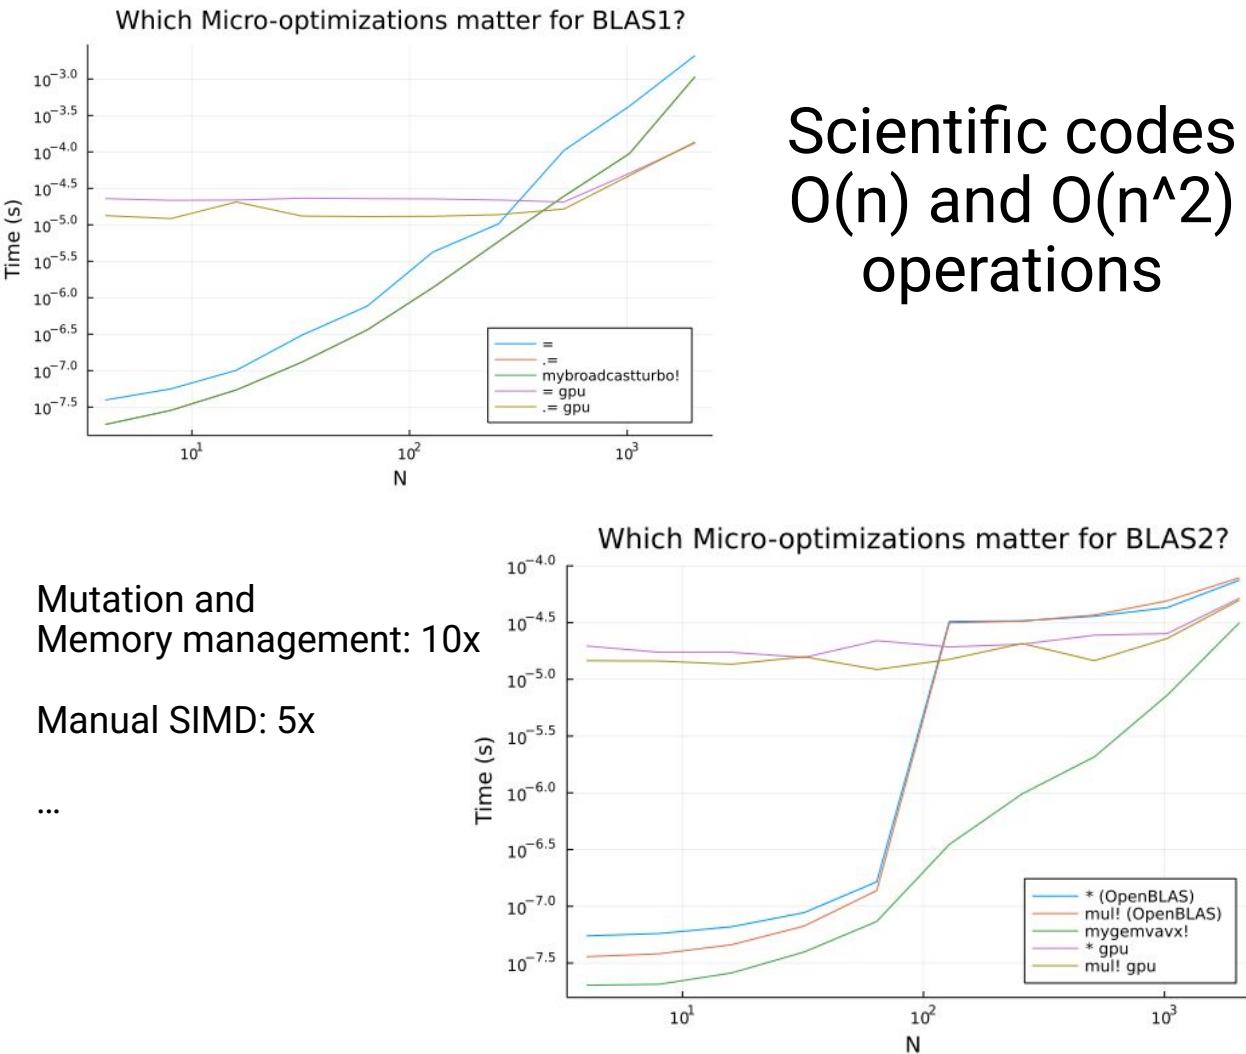
function prob_func(prob, i, repeat)
    remake(prob, p=SA[y₁s[i], y₂s[i], y₃s[i]])
end

monteprob = EnsembleProblem(prob, prob_func = prob_func, safetycopy=false)
solve(monteprob, Rodas5(), EnsembleThreads(), trajectories=1000)

@time solve(monteprob, Rodas5(), EnsembleThreads(), trajectories=1000)

#0.006486 seconds (172.26 k allocations: 16.740 MiB)
#0.006024 seconds (172.26 k allocations: 16.740 MiB)
#0.007074 seconds (172.26 k allocations: 16.740 MiB)
```

Code Optimization in Machine Learning vs Scientific Computing



SimpleChains + StaticArray Neural ODEs

```
sc = SimpleChain(  
    static(2),  
    Activation(x -> x.^3),  
    TurboDense{true}(tanh, static(50)),  
    TurboDense{true}(identity, static(2))  
)  
  
p_nn = SimpleChains.init_params(sc)  
  
f(u,p,t) = sc(u,p)
```

This function is plugged into an ODE solver and the L2 loss is calculated from the numerical solution and the NeuralODE output.

```
prob_nn = ODEProblem(f, u0, tspan)  
  
function predict_neuralode(p)  
    Array(solve(prob_nn,  
        Tsit5(); p=p, saveat=tsteps, sensealg=QuadratureAdjoint(autojacvec=Zygote  
        VJP())))  
end
```

**About a 5x improvement
~1000x in a nonlinear mixed effects context**

Caveat: Requires sufficiently small ODEs (<20)

And simulation processes are still improving rapidly.

New Parallelized Extrapolation Methods for Stiff ODEs: 2x-4x improvements

Elrod, Chris, Yingbo Ma, and Christopher Rackauckas.
"Parallelizing Explicit and Implicit Extrapolation
Methods for Ordinary Differential Equations." *arXiv*
preprint arXiv:2207.08135 (2022).

Accepted to HPEC 2022!

**Caveat: only good for 5-200
ODEs with no implicit
parallelism in f**

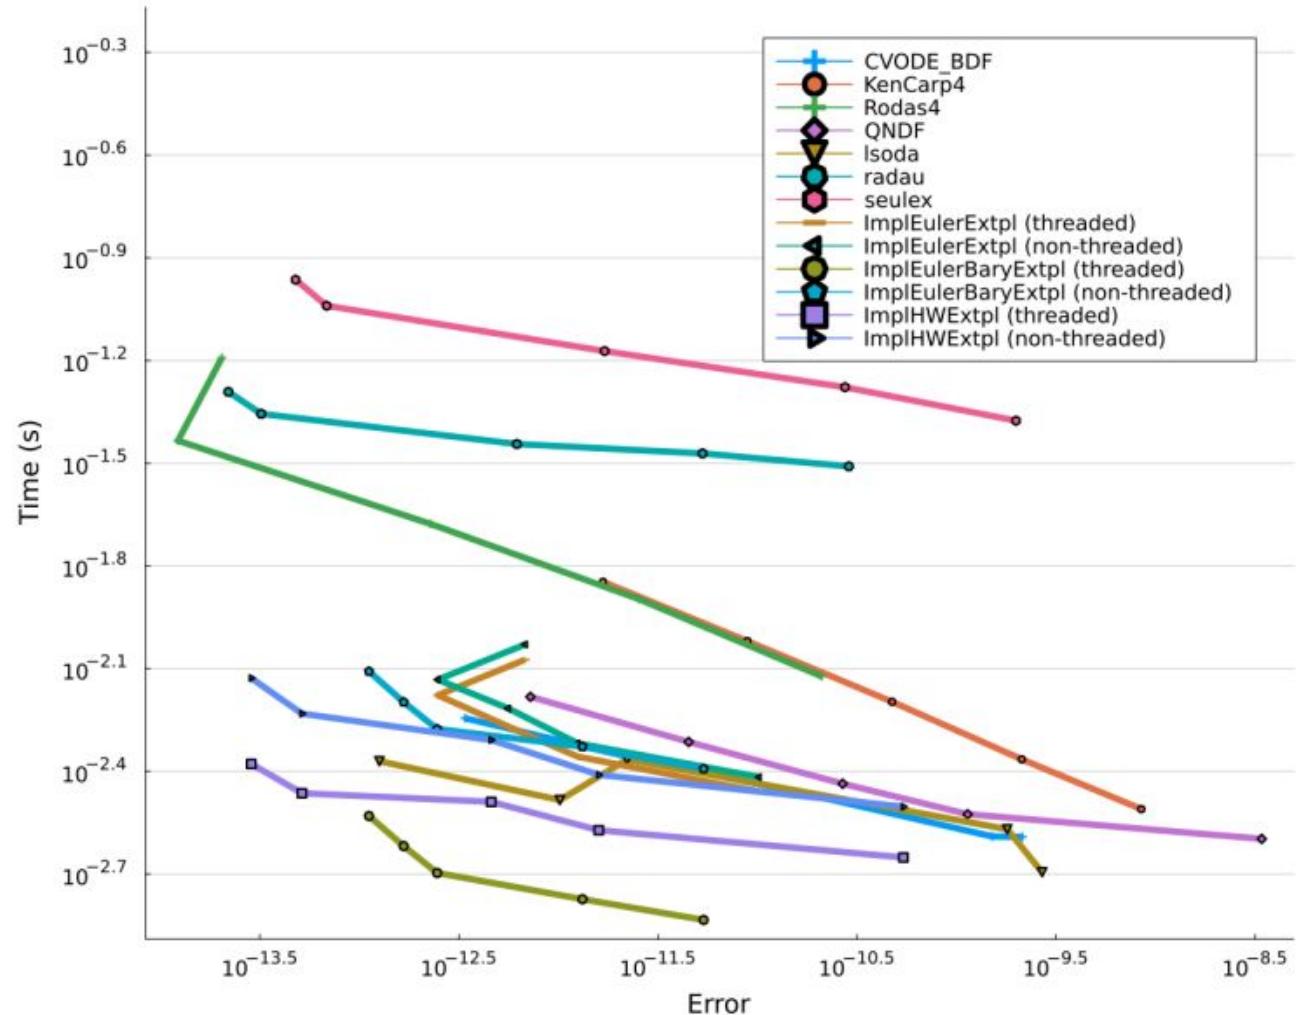


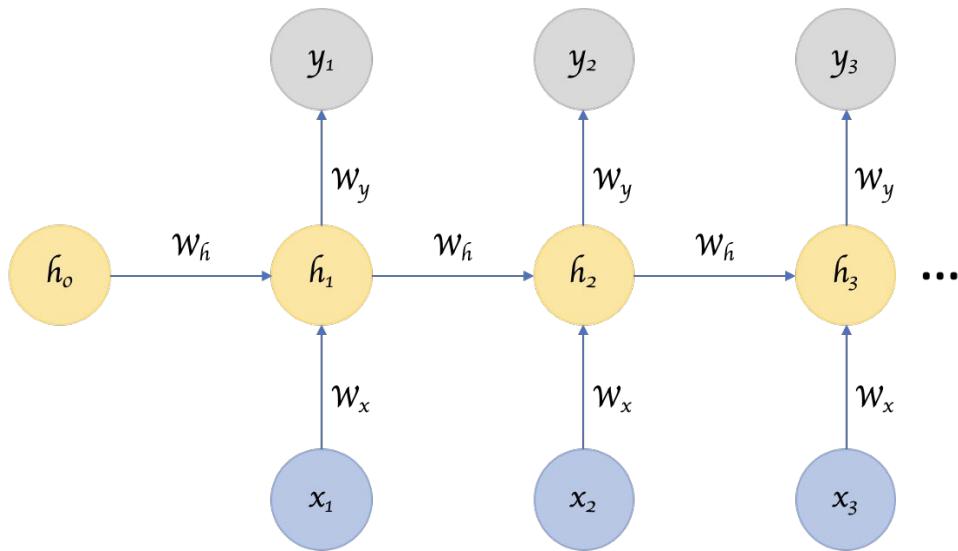
Fig. 6. Benchmark on QSP model with low tolerances [VI-B6].

Can we improve “standard machine learning” using these new differentiable tools?

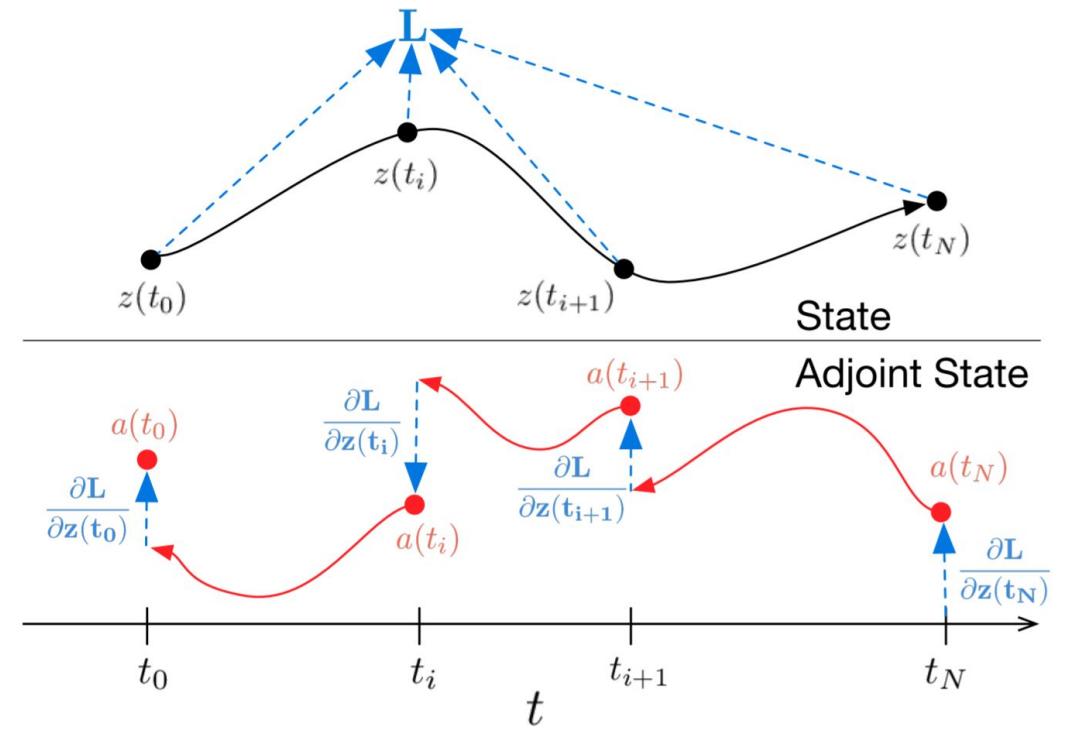
Using Neural ODEs as automated hyperparameter optimization

Neural ODEs as Adaptive Layer Methods

RNN: Euler's method with fixed number of layers

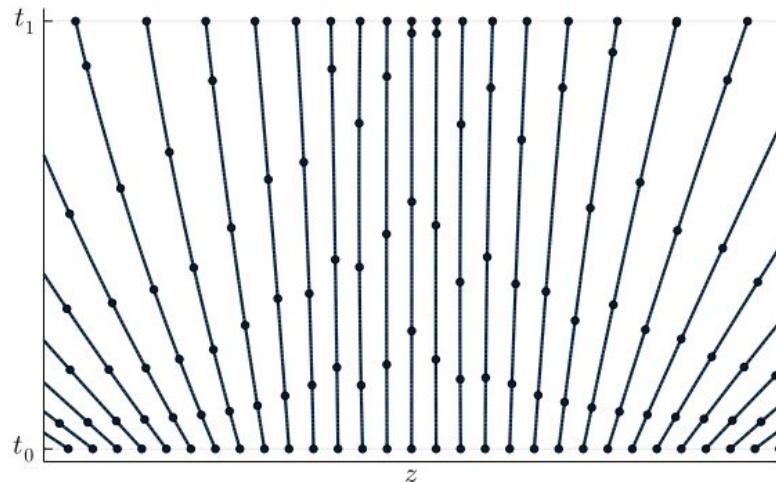
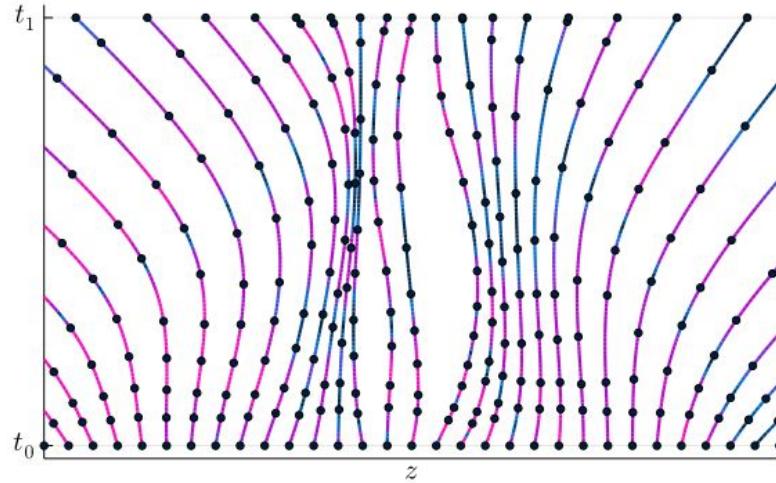


Neural ODE: Adaptive ODE solver chooses number of layers



Neural ODEs can be used on classical machine learning problems to automatically learn the required number of layers
ML Layer: Value is initial condition of ODE, output is solution of ODE at final time

If you only care about the end, why not learn the easiest dynamics you can?



Adaptive ODE solvers are correct to $K - 1$ derivatives and control error on the K th.

Use Taylor mode (high order) automatic differentiation to calculate:

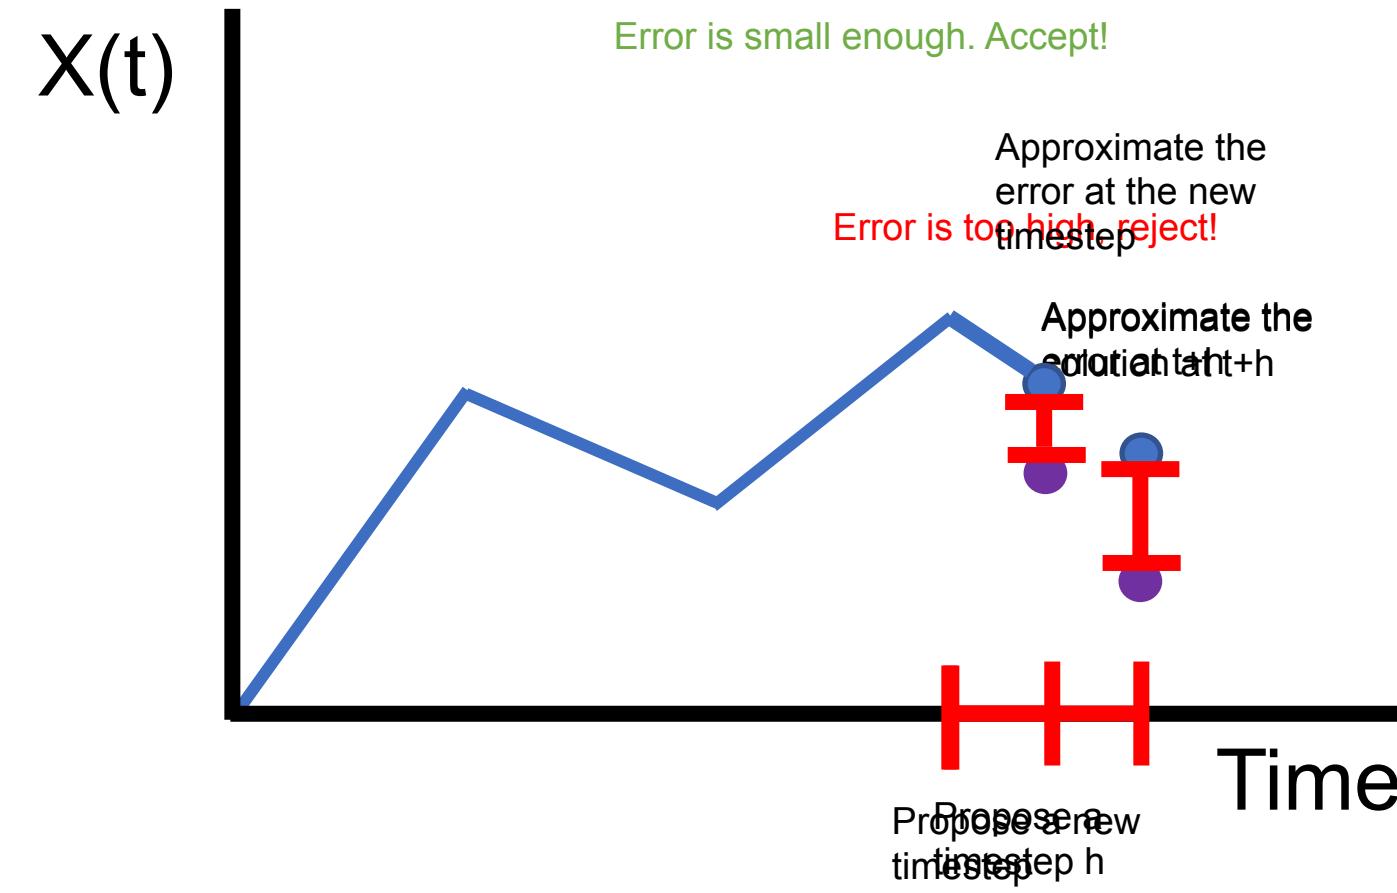
$$R_K(\theta) = \int_{t_0}^{t_1} \left\| \frac{d^K z(t)}{dt^K} \right\| dt$$

and regularize:

$$L_{reg}(\theta) = L(\theta) + R_K(\theta)$$

It does accelerate the learned dynamics, but training is expensive (10x slower!) because higher order automatic differentiation is exponentially expensive.

But Solvers “know” a lot about the equation!



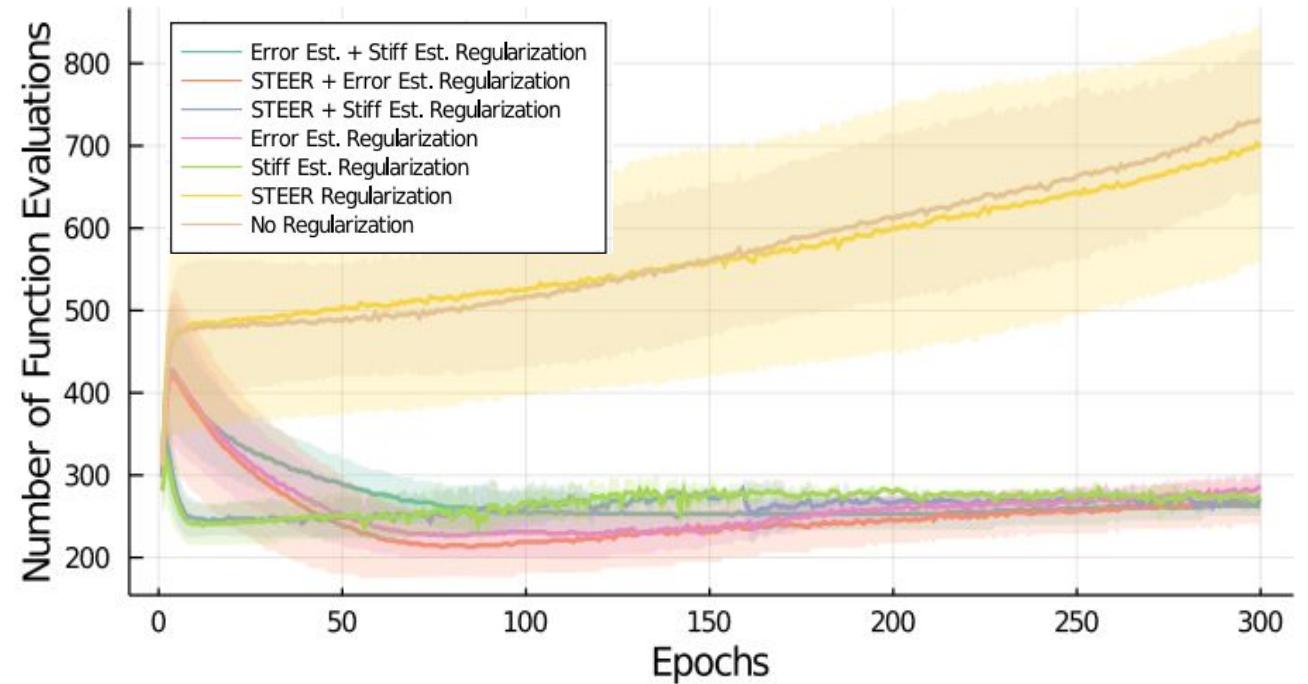
Adaptive ODE solvers already know “free” error estimates and stiffness estimators!

Idea: use the solver’s internal heuristics to regularize out “hard” dynamics

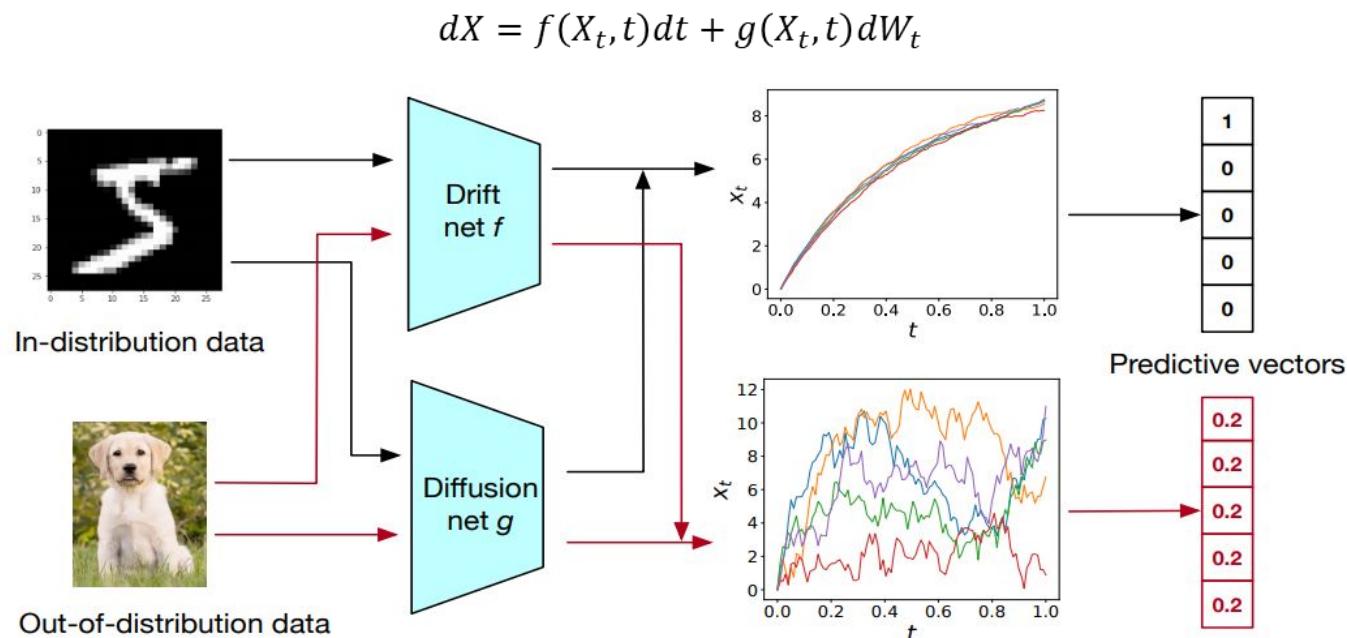
How to improve by an order of magnitude: use knowledge of numerical methods!

10x Neural ODE training vs previous regularization, 2x faster prediction time vs vanilla neural ODE

Method	Train Loss	Test Loss	Train Time (hr)	Prediction Time (s)
Vanilla NODE	3.48	3.55	1.75	0.53
TayNODE	4.21	4.21	12.3	0.22
SRNODE	3.52	3.58	0.87	0.20

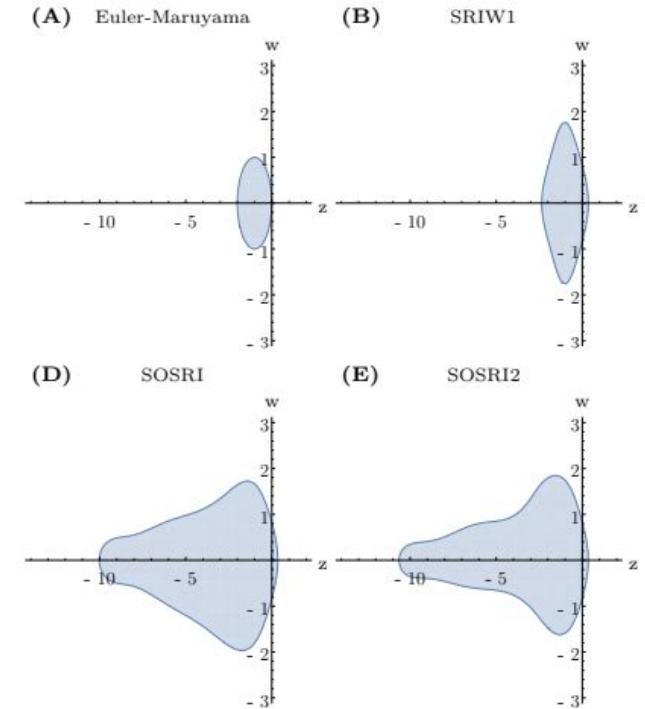


Neural SDEs improve generalization. Can we improve them?



Add noise and uncertainty quantification to continuous layer methods via stochastic differential equations

Liu, X., Si, S., Cao, Q., Kumar, S. and Hsieh, C.J., 2019. Neural sde: Stabilizing neural ode networks with stochastic noise. *arXiv preprint arXiv:1906.02355*.



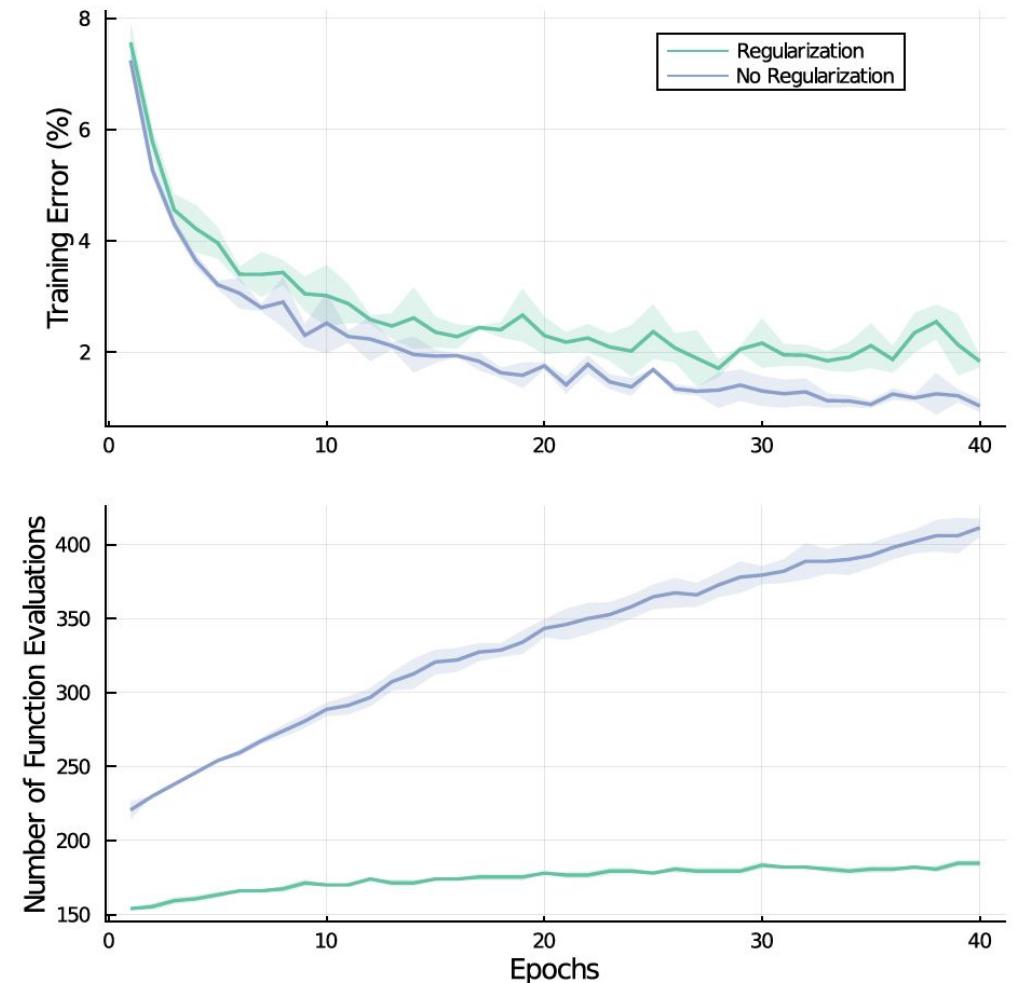
New improved stability SDE solvers with adaptivity and automatic stiffness detection

Rackauckas, C. and Nie, Q., 2020, September. Stability-optimized high order methods and stiffness detection for pathwise stiff stochastic differential equations. In 2020 IEEE High Performance Extreme Computing Conference (HPEC) (pp. 1-8). IEEE. (Quality Submission Award)

Major improvements to Neural SDEs on MNIST

Method	Train Accuracy (%)	Test Accuracy (%)	Train Time (hr)	Prediction Time (s)
Vanilla NSDE	98.97	96.95	6.32	15.07
RegNSDE	98.16	96.27	4.19	7.23

Double Neural SDE prediction speed!



Neural ODEs to Infinity are an “Infinite Layer” Network which can train fast!

DEQs are Steady State Problems

Fixed Point Iterations

Find z s.t.
$$z^* = \underbrace{f_\theta(z^*, x)}_{\text{Neural Network}}$$

Discrete Formulation

Continuous Formulation

Steady State Problem

Find $z(\infty)$ s.t.
$$\frac{dz}{dt} = \underbrace{f_\theta(z, t)}_{\text{Neural Network}} - z$$

DEQs are Steady State Problems

Steady State Problem

$$\text{Find } z(\infty) \text{ s.t. } \overbrace{\frac{dz}{dt}}^{\text{Neural Network}} = \underbrace{f_\theta(z, t)}_{\text{Neural Network}} - z$$

Adaptive ODE Solver with
Terminating Callback

$$z(\infty) = \boxed{\begin{aligned} z(T) &= \int_{t_0}^T (f_\theta(z(t), t) - z(t)) dt \\ \text{s.t. } f_\theta(z(T), T) &= z(T) \end{aligned}}$$

Infinite-Time Neural ODE

$$\boxed{z(\infty) = z(T) = \int_{t_0}^T (f_\theta(z(t), t) - z(t)) dt}$$

s.t. $f_\theta(z(T), T) = z(T)$

Continuous DEQ

$$z(T) = \int_{t_0}^T f_\theta(z(t), t) dt$$

Neural ODE

Won't Back Propagation be Expensive?

$$f_\theta(z(T), T) = z(T)$$

Implicit Function

$$\frac{dg_\theta(z^*)}{d\theta} = \frac{\partial g_\theta(z^*)}{\partial \theta} + \frac{\partial g_\theta(z^*)}{\partial z^*} \underbrace{\left(I - \frac{\partial f_\theta(z^*)}{\partial z^*} \right)^{-1}}_{\frac{\partial z^*}{\partial \theta}} \frac{\partial f_\theta(z^*)}{\partial \theta}$$

Only need to
backpropogate at the final
point!

Neural ODE Adjoints are Really Slow!

Model	Continuous	# of Params	Test Accuracy (%)	Training Time (s / batch)	Backward Pass (s / batch)	Prediction Time (s / batch)
Vanilla DEQ	✗	10.63M	88.913 ± 0.287	0.625 ± 0.165	0.111 ± 0.021	0.414 ± 0.222
	✓		89.367 ± 0.832		1.284 ± 0.011	0.739 ± 0.003
Skip DEQ	✗	11.19M	88.783 ± 0.178	0.588 ± 0.042	0.112 ± 0.006	0.314 ± 0.017
	✓		89.600 ± 0.947		0.697 ± 0.012	0.150 ± 0.013
Skip Reg. DEQ	✗	10.63M	88.773 ± 0.115	0.613 ± 0.048	0.109 ± 0.008	0.268 ± 0.031
	✓		90.107 ± 0.837		0.660 ± 0.019	0.125 ± 0.003
Neural ODE	✓	10.63M	89.047 ± 0.116	5.267 ± 0.078	4.569 ± 0.077	0.573 ± 0.010

Table 3: **CIFAR10 Classification with Large Neural Network:** Skip Reg. Continuous DEQ achieves the *highest test accuracy*. Continuous DEQs are faster than Neural ODEs during training by a factor of $4.1 \times - 7.98 \times$, with a speedup of $6.18 \times - 36.552 \times$ in the backward pass. However, we observe a prediction slowdown for Continuous DEQs of $1.4 \times - 2.36 \times$ against Discrete DEQs and $0.90 \times - 0.95 \times$ against Neural ODE.

Neural ODE Adjoint are Really Slow!

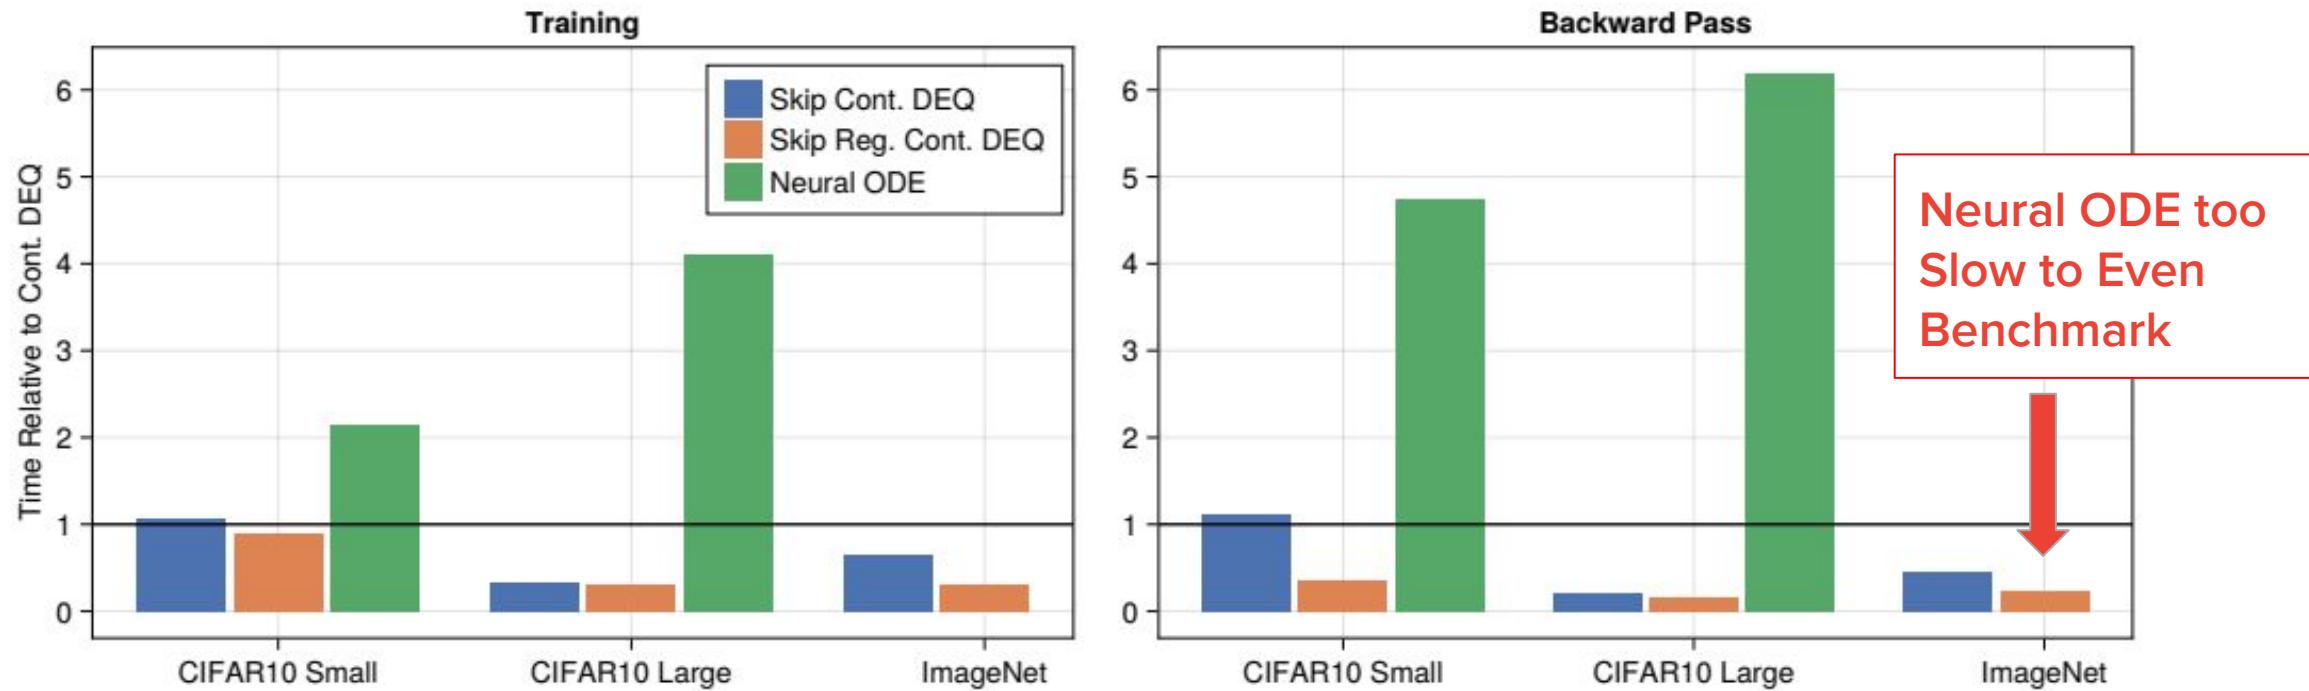


Figure 1: **Relative Training and Backward Pass Timings against Continuous DEQs (lower is better)**: In all scenarios, Neural ODEs take $4.7 - 6.182 \times$ more time in the backward pass compared to Vanilla Continuous DEQs. Whereas combining Skip (Reg.) with Continuous DEQs accelerates the backward pass by $2.8 - 5.9 \times$.

**Many other SciML Methods can be
rephrased as UDEs**

Solving Semilinear Parabolic PDEs (Nonlinear Black-Scholes) by Learning Structured Stochastic Differential Equations

- Semilinear Parabolic Form (Diffusion-Advection Equations, Hamilton-Jacobi-Bellman, Black-Scholes)

$$\begin{aligned} \frac{\partial u}{\partial t}(t, x) + \frac{1}{2} \text{Tr} \left(\sigma \sigma^T(t, x) (\text{Hess}_x u)(t, x) \right) + \nabla u(t, x) \cdot \mu(t, x) \\ + f(t, x, u(t, x), \sigma^T(t, x) \nabla u(t, x)) = 0 \end{aligned} \quad [1]$$

Then the solution of Eq. 1 satisfies the following BSDE (cf., e.g., refs. 8 and 9):

$$\begin{aligned} u(t, X_t) - u(0, X_0) \\ = - \int_0^t f(s, X_s, u(s, X_s), \sigma^T(s, X_s) \nabla u(s, X_s)) ds \\ + \int_0^t [\nabla u(s, X_s)]^T \sigma(s, X_s) dW_s. \end{aligned} \quad [3]$$

- Make $(\sigma^T \nabla u)(t, X)$ a neural network.
- Solve the resulting SDEs and learn $\sigma^T \nabla u$ via:

$$l(\theta) = \mathbb{E} \left[|g(X_{t_N}) - \hat{u}(\{X_{t_n}\}_{0 \leq n \leq N}, \{W_{t_n}\}_{0 \leq n \leq N})|^2 \right].$$

Simplified:

- Transform the PDE into a Backwards SDE: a stochastic boundary value problem. The unknown is a function!
- Learn the unknown function via neural network.
- Once learned, the PDE solution is known.

Solving Semilinear Parabolic PDEs (Nonlinear Black-Scholes) by Learning Structured Stochastic Differential Equations

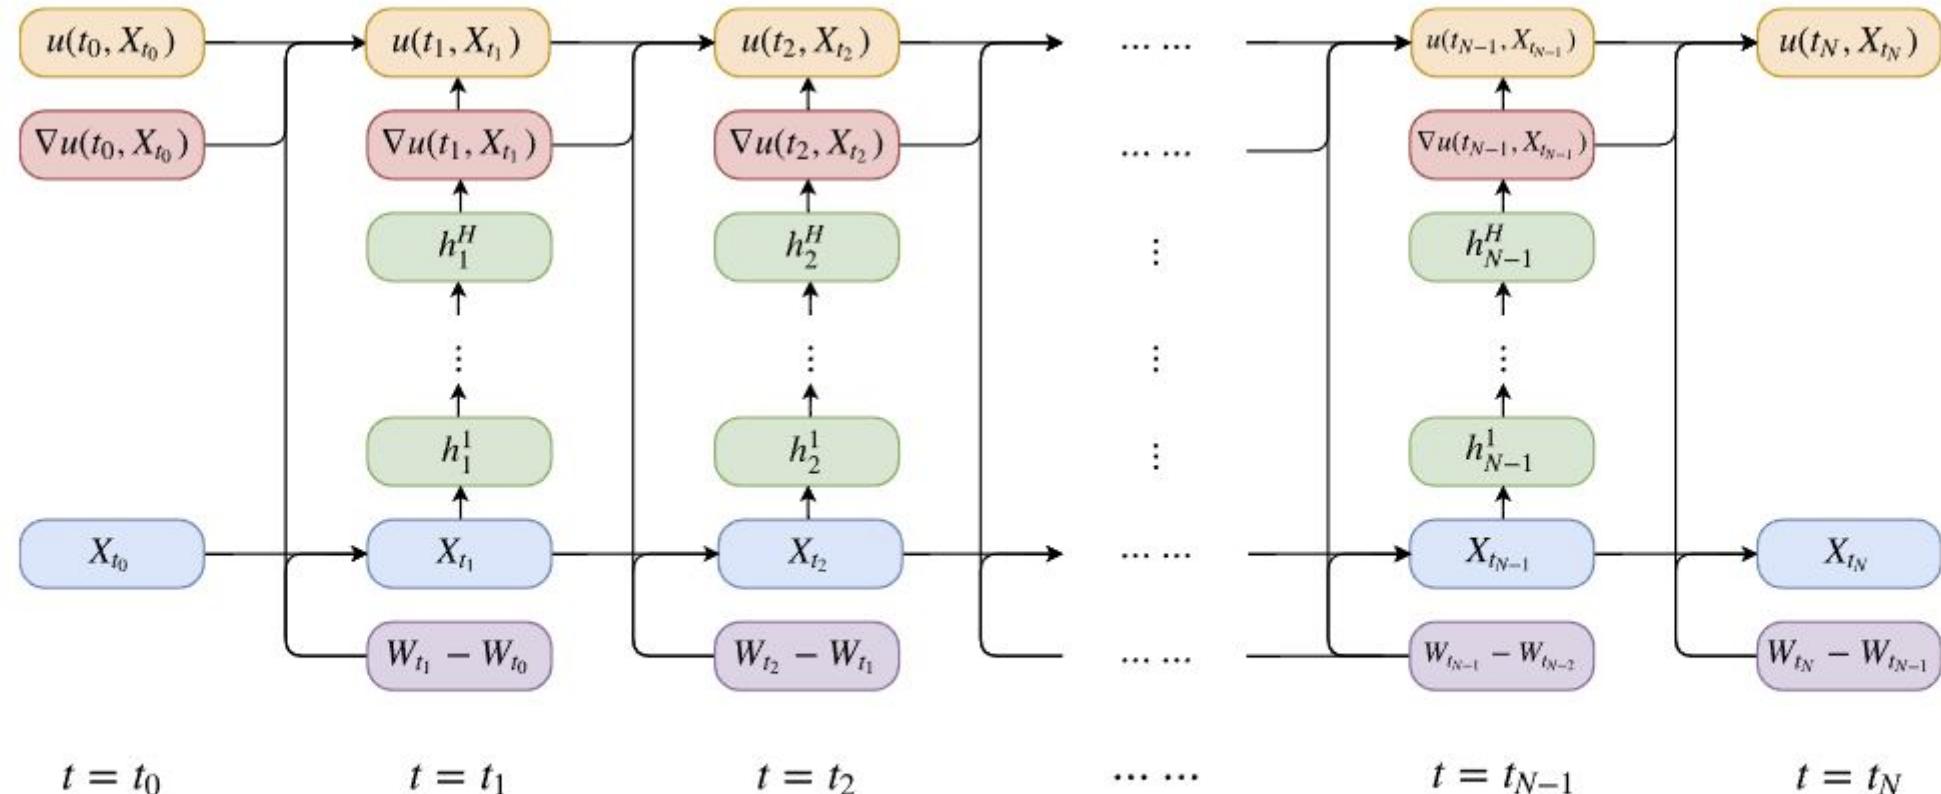
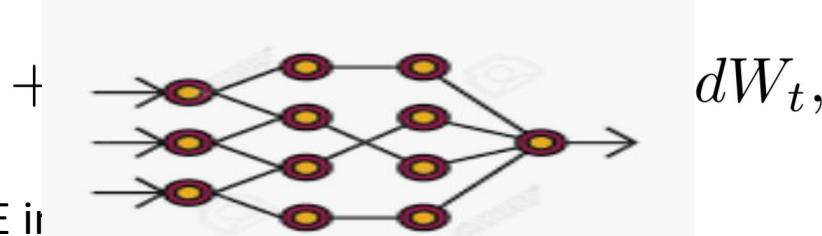


Fig. 4. Illustration of the network architecture for solving semilinear parabolic PDEs with H hidden layers for each subnetwork and N time intervals. The whole network has $(H+1)(N-1)$ layers in total that involve free parameters to be optimized simultaneously. Each column for $t = t_1, t_2, \dots, t_{N-1}$ corresponds to a subnetwork at time t . h_n^1, \dots, h_n^H are the intermediate neurons in the subnetwork at time $t = t_n$ for $n = 1, 2, \dots, N-1$.

Solving Semilinear Parabolic PDEs (Nonlinear Black-Scholes) by Learning Structured Stochastic Differential Equations

- This is equivalent to training the universal SDE:

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t,$$
$$dU_t = f(t, X_t, U_t, \sigma^T(t, X_t)\nabla u(t, X_t))dt$$



- Solves a 100 dimensional PDE if
- As a universal SDE, we can solve this with high order (less neural network evaluations), adaptivity, etc. methods using DiffEqFlux.jl. Thus we automatically extend the deep BSDE method to better discretizations by using this interpretation!

$$l(\theta) = \mathbb{E} \left[|g(X_{t_N}) - \hat{u}(\{X_{t_n}\}_{0 \leq n \leq N}, \{W_{t_n}\}_{0 \leq n \leq N})|^2 \right].$$



HighDimPDE.jl

Search docs

Home

◦ Algorithm overview

Getting started

Solver Algorithms

The MLP algorithm

The DeepSplitting algorithm

Feynman Kac formula

Home

Edit on GitHub

HighDimPDE.jl

HighDimPDE.jl is a Julia package to **solve Highly Dimensional non-linear, non-local PDEs** of the form

$$(\partial_t u)(t, x) = \int_{\Omega} f(t, x, \mathbf{x}, u(t, x), u(t, \mathbf{x}), (\nabla_x u)(t, x), (\nabla_x u)(t, \mathbf{x})) d\mathbf{x} \\ + \langle \mu(t, x), (\nabla_x u)(t, x) \rangle + \frac{1}{2} \text{Trace}(\sigma(t, x)[\sigma(t, x)]^*(\text{Hess}_x u)(t, x)).$$

where $u: [0, T] \times \Omega \rightarrow \mathbb{R}$, $\Omega \subseteq \mathbb{R}^d$ is subject to initial and boundary conditions, and where d is large.

HighDimPDE.jl implements solver algorithms that break down the curse of dimensionality, including

- the Deep Splitting scheme
- the Multi-Level Picard iterations scheme.

To make the most out of HighDimPDE.jl, we advise to first have a look at the

- documentation on the Feynman Kac formula,

as all solver algorithms heavily rely on it.

Algorithm overview

Features	DeepSplitting	MLP
Time discretization free	✗	✓
Mesh-free	✓	✓
Single point $x \in \mathbb{R}^d$ approximation	✓	✓
d -dimensional cube $[a, b]^d$ approximation	✓	✗
GPU	✓	✗
Gradient non-linearities	✓	✗

✓ : will be supported in the future